

# **COMP 201**

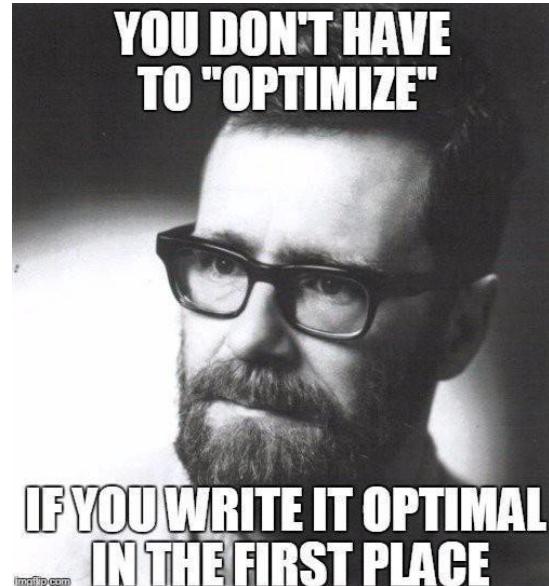
## Fall 2025



**Lab 9**  
Code Optimization

# What is Code Optimization?

- A program transformation technique to:
  - **improve** the intermediate code,
  - make it **consume fewer resources** (CPU, memory),
  - reduce the **size** of the code,
  - **speed up** execution.
- It is **not** optimizing an algorithm.
  - It is beyond our scope for now.



# Types of Code Optimization

- **Machine Independent Optimization**
- **Machine Dependent Optimization**

# Types of Code Optimization

## Machine Independent Optimization

- Improve the intermediate code to get a better target code.
- Does not involve any CPU registers, or absolute memory locations.

```
do {
    item = 10;
    value = value + item;
}
while(value<100);

// this code involves repeated
// assignment of 'item'. Instead:
```

# Types of Code Optimization

## Machine Independent Optimization

- Improve the intermediate code to get a better target code.
- Does not involve any CPU registers, or absolute memory locations.

```
do{
    item = 10;
    value = value + item;
}
while(value<100);

// this code involves repeated
// assignment of 'item'. Instead:

item = 10;
do{
    value = value + item;
}
while(value<100);
```

# Types of Code Optimization

## Machine Dependent Optimization

- **Goal:** Exploit hardware features.
- After the target code is generated,
  - optimization is done according to the target machine architecture.
- Involves use of special instructions,
- Allocating CPU registers,
- May have absolute memory references, rather than relative references etc.

# Compiler Optimizations

- **GCC** applies both machine independent and dependent optimizations.
- Without any optimization option, the compiler's goal is to **reduce the cost of compilation** and to make debugging produce the expected results.
- Turning on optimization flags makes the compiler attempt to **improve the performance** and/or decrease the **code size** at the expense of compilation time and possibly the **ability to debug** the program.
- Most optimizations are completely disabled at -O0 (or equivalently when -O level is not set), even if individual optimization flags are specified. Similarly, -Og suppresses many optimization passes.

# Compiler Optimizations

- **-O or -O1 option (Optimize).** the compiler tries to reduce code size and execution time, without performing any optimizations that take a great deal of compilation time.
- **-O2 (Optimize even more).** GCC performs nearly all supported optimizations that do not involve a space-speed tradeoff. As compared to -O, this option increases both compilation time and the performance of the generated code.
- **-O3 (Optimize yet more).** -O3 turns on all optimizations specified by -O2 and also more options such as loop unrolling and jamming.
- **-Os (Optimize for size).** -Os enables all -O2 optimizations except those that often increase code size.

# Machine Independent Techniques

- Techniques are various and vast
- Be careful about time you spend on optimization
- With practice, you can write your codes optimized in the first place and optimize it further after having a base code
- Profiling is an invaluable tool
- In practice you should reuse already existing optimized codes

# Inlining

- C functions can be recoded as **macros**
  - to obtain similar speedup on compilers with no inlining capability.
- This should be done after the code is completely debugged.
- No function call
  - Fewer instructions!

```
int foo(a, b)
{
    a = a - b;
    b++;
    a = a * b;
    return a;
}
```

Can be replaced by:

```
#define foo(a, b) (((a)-(b)) * ((b)+1))
```

# Avoid Pointer Dereference in Loop

- Pointer dereferencing creates lots of trouble in memory. It's better to assign it to some temporary variable and then use that temporary variable in the loop.

```
int a = 0;
int* iptr = &a;

for (int i = 1; i < 11; ++i) {
    *iptr = *iptr + i;
}

int a = 0;
int* iptr = &a;
// Dereferencing pointer outside loop & use temp var
int temp = *iptr;
for (int i = 1; i < 11; ++i) {
    temp = temp + i;
}
// Updating pointer using final value of temp
*iptr = temp;
```

# Avoid Pointer Dereference in Loop

- Pointer dereferencing creates lots of trouble in memory. It's better to assign it to some temporary variable and then use that temporary variable in the loop.

```
struct Warrior{  
    double damage; double HP;  
    ...  
};  
  
void main(){  
    struct Warrior* herol;  
    struct Warrior* enemies[n];  
    ...  
  
    for (int i = 0; i < n; ++i) {  
        enemies[i]->HP -= herol->damage;  
    }  
}
```

```
void main(){  
    struct Warrior* herol;  
    struct Warrior* enemies[n];  
    ...  
    double damage = herol->damage;  
    for (int i = 0; i < n; ++i) {  
        enemies[i]->HP -= damage;  
    }  
}
```

# Loop Unrolling

- `gcc -funroll-loops` will do this. But if you know that yours doesn't, you can modify your code to get the same effect.
- This way, the loop condition (`i < 100`) and the backward branch are executed only 21 times instead of 101.

```
for (i = 0; i < 100; i++) {  
    do_stuff(i);  
}  
}
```

Can be replaced by:

```
for (i = 0; i < 100; ) {  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
}  
}
```

# Loop Unrolling Caveat

- An unrolled loop is larger than the "rolled" version.
  - So, **may no longer fit** into the instruction cache (on the machines which have them).
  - This will make the unrolled version slower.
- In general, **function call** overhead is much **larger** than **loop-control** overhead.
  - So any savings from loop unrolling are insignificant in comparison to what you'd **achieve from inlining** in this case.

```
for (i = 0; i < 100; i++) {  
    do_stuff(i);  
}
```

Can be replaced by:

```
for (i = 0; i < 100; ) {  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
    do_stuff(i); i++;  
}
```

# Code Motion

- Avoid recomputing expressions that produce the same result each iteration.
- In particular, move loop-invariant code outside the loop.

```
void foo(double *a, double *b, long i, long n){  
    long j;  
    for (j = 0; j < n; j++)  
        a[n*i+j] = b[j];  
}
```

Can be replaced by:

```
void foo(double *a, double *b, long i, long n){  
    long j;  
  
    int ni = n*i;  
    for (j = 0; j < n; j++)  
        a[ni+j] = b[j];  
}
```

# Share Common Subexpressions

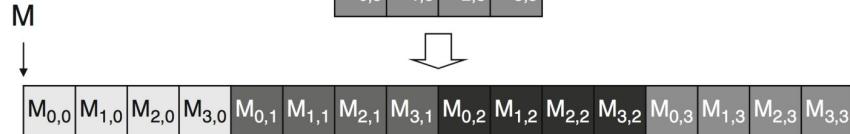
```
/* Sum neighbors of i,j */
up = val[(i-1)*n + j];
down = val[(i+1)*n + j];
left = val[i*n + j-1];
right = val[i*n + j+1];
sum = up + down + left + right;
```

3 multiplications  
→

```
leaq 1(%rsi), %rax # i+1
leaq -1(%rsi), %r8 # i-1
imulq %rcx, %rsi # i*n
imulq %rcx, %rax # (i+1)*n
imulq %rcx, %r8 # (i-1)*n
addq %rdx, %rsi # i*n+j
addq %rdx, %rax # (i+1)*n+j
addq %rdx, %r8 # (i-1)*n+j
```

- Especially problematic if this function is getting called inside a loop

M <sub>0,0</sub>	M <sub>1,0</sub>	M <sub>2,0</sub>	M <sub>3,0</sub>
M <sub>0,1</sub>	M <sub>1,1</sub>	M <sub>2,1</sub>	M <sub>3,1</sub>
M <sub>0,2</sub>	M <sub>1,2</sub>	M <sub>2,2</sub>	M <sub>3,2</sub>
M <sub>0,3</sub>	M <sub>1,3</sub>	M <sub>2,3</sub>	M <sub>3,3</sub>



# Share Common Subexpressions

```
/* Sum neighbors of i,j */
long inj = i*n + j;
up = val[inj - n];
down = val[inj + n];
left = val[inj - 1];
right = val[inj + 1];
sum = up + down + left +
      right;
```

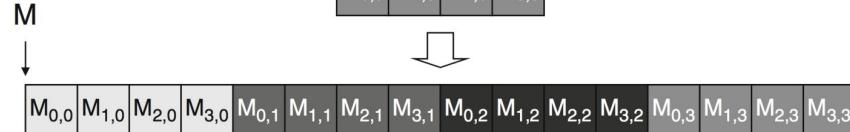
1 multiplication



```
imulq %rcx, %rsi # i*n
addq %rdx, %rsi # i*n+j
movq %rsi, %rax # i*n+j
subq %rcx, %rax # i*n+j-n
leaq (%rsi,%rcx), %rcx # i*n+j+n
```

- Reuse portions of expressions
- GCC will do this with `-O1`

M <sub>0,0</sub>	M <sub>1,0</sub>	M <sub>2,0</sub>	M <sub>3,0</sub>
M <sub>0,1</sub>	M <sub>1,1</sub>	M <sub>2,1</sub>	M <sub>3,1</sub>
M <sub>0,2</sub>	M <sub>1,2</sub>	M <sub>2,2</sub>	M <sub>3,2</sub>
M <sub>0,3</sub>	M <sub>1,3</sub>	M <sub>2,3</sub>	M <sub>3,3</sub>



# Reduction in Strength

- Replace costly operations with simpler ones
- E.g. replace Shift/add with multiply/divide
  - E.g. `x << 4; // is equivalent to x * 16;`

# Loop Jamming

- Combine adjacent loops which loop over the same range of the same variable.
- Incrementing and testing of **i** is done only half as often
- Provided the second loop does not access elements from future iterations
  - e.g **array[i+3]**.

```
for (i = 0; i < MAX; i++)      /* initialize 2d array to 0's */
    for (j = 0; j < MAX; j++)
        a[i][j] = 0.0;

for (i = 0; i < MAX; i++)      /* put 1's along the diagonal */
    a[i][i] = 1.0;
```

Can be replaced by:

```
for (i = 0; i < MAX; i++){
    for (j = 0; j < MAX; j++)
        a[i][j] = 0.0;      /* initialize 2d array to 0's */

    a[i][i] = 1.0;          /* put 1's along the diagonal */

}
```

# Loop Inversion

- Some architectures provide an efficient *decrement-and-branch-if-not-zero* instruction.
- Counting down to zero can reduce instruction count and branches
- Assuming the loop is insensitive to direction

```
for (i = 1; i < MAX; i++) {  
    ...  
}
```

Can be replaced by:

```
for (i = MAX; i--;) {  
    ...  
}
```

# Table Lookup

- Use a lookup table when the input domain is **small/bounded** and values are **reused** often
  - e.g. small factorial range
- Replaces computation with an array index (time ↔ memory tradeoff)
- Tables can be **static (compile-time)** or **initialized at startup** (runtime precompute)

```
long factorial(int i){  
    if (i == 0)  
        return 1;  
    else  
        return i * factorial(i - 1);  
}
```

Can be replaced by:

```
static long factorial_table[] =  
    {1, 1, 2, 6, 24, 120, 720 /* etc */};  
long factorial(int i){  
    return factorial_table[i];  
}
```

# Stack Usage

- Each function call consumes stack space for locals and parameters
- Large local arrays can easily exhaust stack space
- Large structs used as locals or passed by value have the same issue
- Prefer static, global, or heap allocation for large data
- Pass large objects by pointer instead of by value

# Recap: Struct Padding

- Generally when a struct is stored in RAM, it is padded to correspond to the word-size of the architecture of the CPU.
- Additional padding is provided for arrays to make the first bytes of each item in the array a multiple of the item size.

```
/* Assume 32 bit Architecture  
Sizeof foo = 12 bytes */
```

```
struct foo{  
    char c;  
    int x;  
    short s;  
};
```

0 c	1	2	3
4	5	6	7
8 s	9	10	11 padding

# Reduce Padding

- Field ordering affects struct size due to alignment and padding
- Group fields by alignment, placing the most strictly aligned types first
- Reordering fields can significantly reduce memory footprint
- **char** and **short** are commonly used for flags or mode values
- Multiple flags can be packed using bit-fields, at the cost of portability

```
/* sizeof = 64 bytes */  
/* sizeof = 48 bytes */  
  
struct foo {  
    float    a;  
    double   b;  
    float    c;  
    double   d;  
    short    e;  
    long     f;  
    short    g;  
    long     h;  
    char     i;  
    int      j;  
    char     k;  
    int      l;  
};  
};
```

# References

<https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

<http://www.cs.cmu.edu/afs/cs/academic/class/15213-s19/www/schedule.html>