COMP201
Computer Systems & Programming

Lecture #22 – More Cache Memories

KOÇ UNIVERSITY

Aykut Erdem // Koç University // Spring 2025

Recap

- Cache basics

- Principle of locality

- Cache memory organization and operation

Recap: Cache Memories

- Cache memories are small, fast SRAM-based memories managed automatically in hardware
  - Hold frequently accessed blocks of main memory
- CPU looks first for data in cache
- Typical system structure:
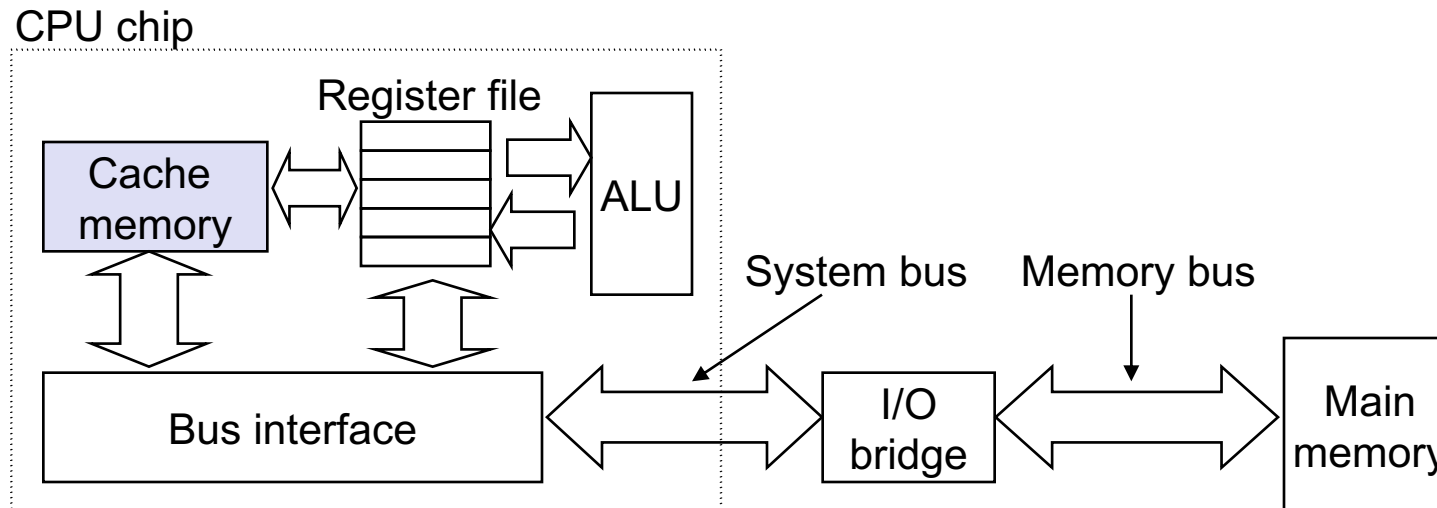
Recap: Why Caches Work

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently

- **Temporal locality:**
  – Recently referenced items are <u>likely</u> to be referenced again in the near future

- **Spatial locality:**
  – Items with nearby addresses <u>tend to</u> be referenced close together in time

Recap: Good Locality Example

- Does this function have good locality with respect to array `a`?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; j < M; j++)
        for (j = 0; i < N; i++)
            sum += a[i][j];
    return sum;
}
```

M = 3,
N = 4

a[0][0]	a[0][1]	a[0][2]	a[0][3]
a[1][0]	a[1][1]	a[1][2]	a[1][3]
a[2][0]	a[2][1]	a[2][2]	a[2][3]

Access Pattern:
stride = 1

1) a[0][0]
2) a[0][1]
3) a[0][2]
4) a[0][3]
5) a[1][0]
6) a[1][1]
7) a[1][2]
8) a[1][3]
9) a[2][0]
10) a[2][1]
11) a[2][2]
12) a[2][3]

Layout in Memory

a[0][0]	a[0][1]	a[0][2]	a[0][3]	a[1][0]	a[1][1]	a[1][2]	a[1][3]	a[2][0]	a[2][1]	a[2][2]	a[2][3]

Note: 76 is just one possible starting address of array a

76   92   108

Recap: Bad Locality Example

- Does this function have good locality with respect to array a?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

M = 3,
N = 4

a[0][0]	a[0][1]	a[0][2]	a[0][3]
a[1][0]	a[1][1]	a[1][2]	a[1][3]
a[2][0]	a[2][1]	a[2][2]	a[2][3]

Access Pattern:
stride = 4

1) a[0][0]
2) a[1][0]
3) a[2][0]
4) a[0][1]
5) a[1][1]
6) a[2][1]
7) a[0][2]
8) a[1][2]
9) a[2][2]
10) a[0][3]
11) a[1][3]
12) a[2][3]

Layout in Memory

a[0][0]	a[0][1]	a[0][2]	a[0][3]	a[1][0]	a[1][1]	a[1][2]	a[1][3]	a[2][0]	a[2][1]	a[2][2]	a[2][3]

Note: 76 is just one possible starting address of array a

76   92   108

Plan for Today

- Cache memory organization and operation
- Memory Mountain

Disclaimer: Slides for this lecture were borrowed from
—Randal E. Bryant and David R. O'Hallaroni's CMU 15-213 class
—Porter Jones' UW CSE 351 class

Lecture Plan

- Cache memory organization and operation
- Memory Mountain

Cache Organization

Cache

4	9	10	3

Smaller, faster, more expensive memory caches a subset of the blocks

10

Data is copied in block-sized transfer units

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Larger, slower, cheaper memory viewed as partitioned into "blocks"

Cache Performance Metrics

- Huge difference between a cache hit and a cache miss
  - Could be 100x speed difference between accessing cache and main memory (measured in *clock cycles*)

- Miss Rate (MR)
  - Fraction of memory references not found in cache (misses / accesses) = 1 - Hit Rate

- Hit Time (HT)
  - Time to deliver a block in the cache to the processor
    - Includes time to determine whether the block is in the cache

- Miss Penalty (MP)
  - Additional time required because of a miss

Can we have more than one cache?

- Why would we want to do that?
  - Avoid going to memory!

- Typical performance numbers:
  - Miss Rate
    - L1 MR = 3-10%
    - L2 MR = Quite small (*e.g.* < 1%), depending on parameters, etc.
  - Hit Time
    - L1 HT = 4 clock cycles
    - L2 HT = 10 clock cycles
  - Miss Penalty
    - P = 50-200 cycles for missing in L2 & going to main memory
    - Trend: increasing!

(1) Optimize L1 for high HT
(2) Optimize L2 for low MR

Cache Organization

- **Block Size (B):** unit of transfer between cache and main memory
  - Given in bytes and always a power of 2 (*e.g.* 64 bytes)
  - Blocks consist of adjacent bytes (differ in address by 1)
    - Spatial locality!

Cache Organization

- **Block Size (B):** unit of transfer between cache and main memory
  - Given in bytes and always a power of 2 (*e.g.* 64 bytes)
  - Blocks consist of adjacent bytes (differ in address by 1)
    - Spatial locality!

- Offset field
  - Low-order $\log_2(B) = b$ bits of address tell you which byte within a block
    - (address) mod 2^n = n lowest bits of address
  - (address) modulo (# of bytes in a block)

m-bit address:
(refers to byte in memory)

m - b bits	b bits
Block Number	Block Offset

Question

- If we have 6-bit addresses and block size B = 4 bytes, which block and byte does `0x15` refer to?

	Block Num	Block Offset
A.	1	1
B.	1	5
C.	5	1
D.	5	5
E.	We're lost…	

Recap:

slido

If we have 6-bit addresses and block size B = 4 bytes, which block and byte does 0x15 refer to?

ⓘ Start presenting to display the poll results on this slide.

17

Question

- If we have 6-bit addresses and block size B = 4 bytes, which block and byte does `0x15` refer to?

	Block Num	Block Offset
A.	1	1
B.	1	5
C.	**5**	**1**
D.	5	5
E.	We're lost…	

$$0x \quad 1 \quad 5$$

Address: 0b 0 1 0 1 0 1

Offset width = $\log_2(B) = \log_2(4) = 2$ bits

Cache Organization

- **Cache Size (C):** amount of *data* the cache can store
  - Cache can only hold so much data (subset of next level)
  - Given in bytes (C) or number of blocks (C/B)
  - <u>Example</u>: C = 32 KiB = 512 blocks if using 64-byte blocks

- Where should data go in the cache?
  - We need a mapping from memory addresses to specific locations in the cache to make checking the cache for an address **fast**

- What is a data structure that provides fast lookup?
  - Hash table!

# Review: Hash Tables for Fast Lookup

Insert:

5

27

34

102

119

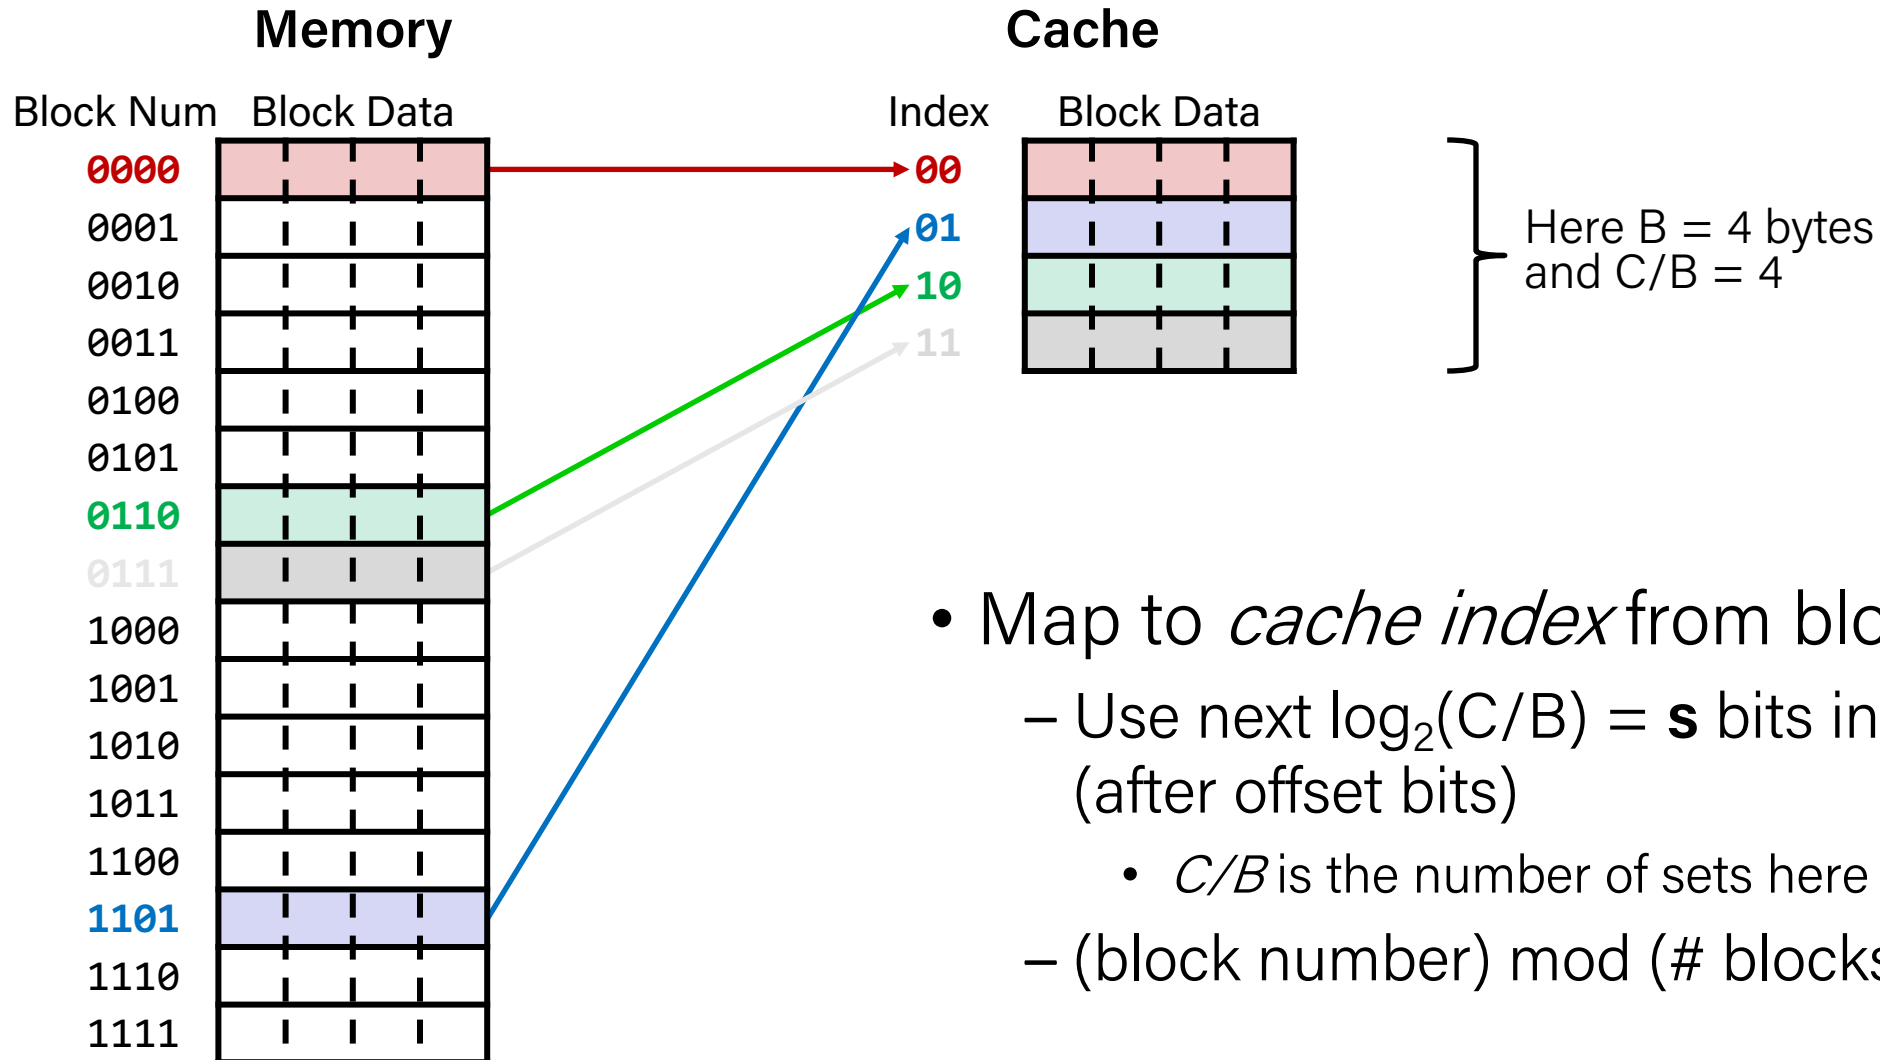Apply hash function to map
data to "buckets"

hash(27) % N (10) = 7

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	

Review: Hash Tables for Fast Lookup

Insert:

5

27

34

102

119

0
1
2
3
4
5
6
7
8
9

Apply hash function to map data to "buckets"

hash(27) % N (10) = 7

Place Data in Cache by Hashing Address

Memory

Cache

Block Num   Block Data

0000
0001
0010
0011
0100
0101
0110
0111
1000
1001
1010
1011
1100
1101
1110
1111

Index   Block Data

00
01
10
11

Here B = 4 bytes
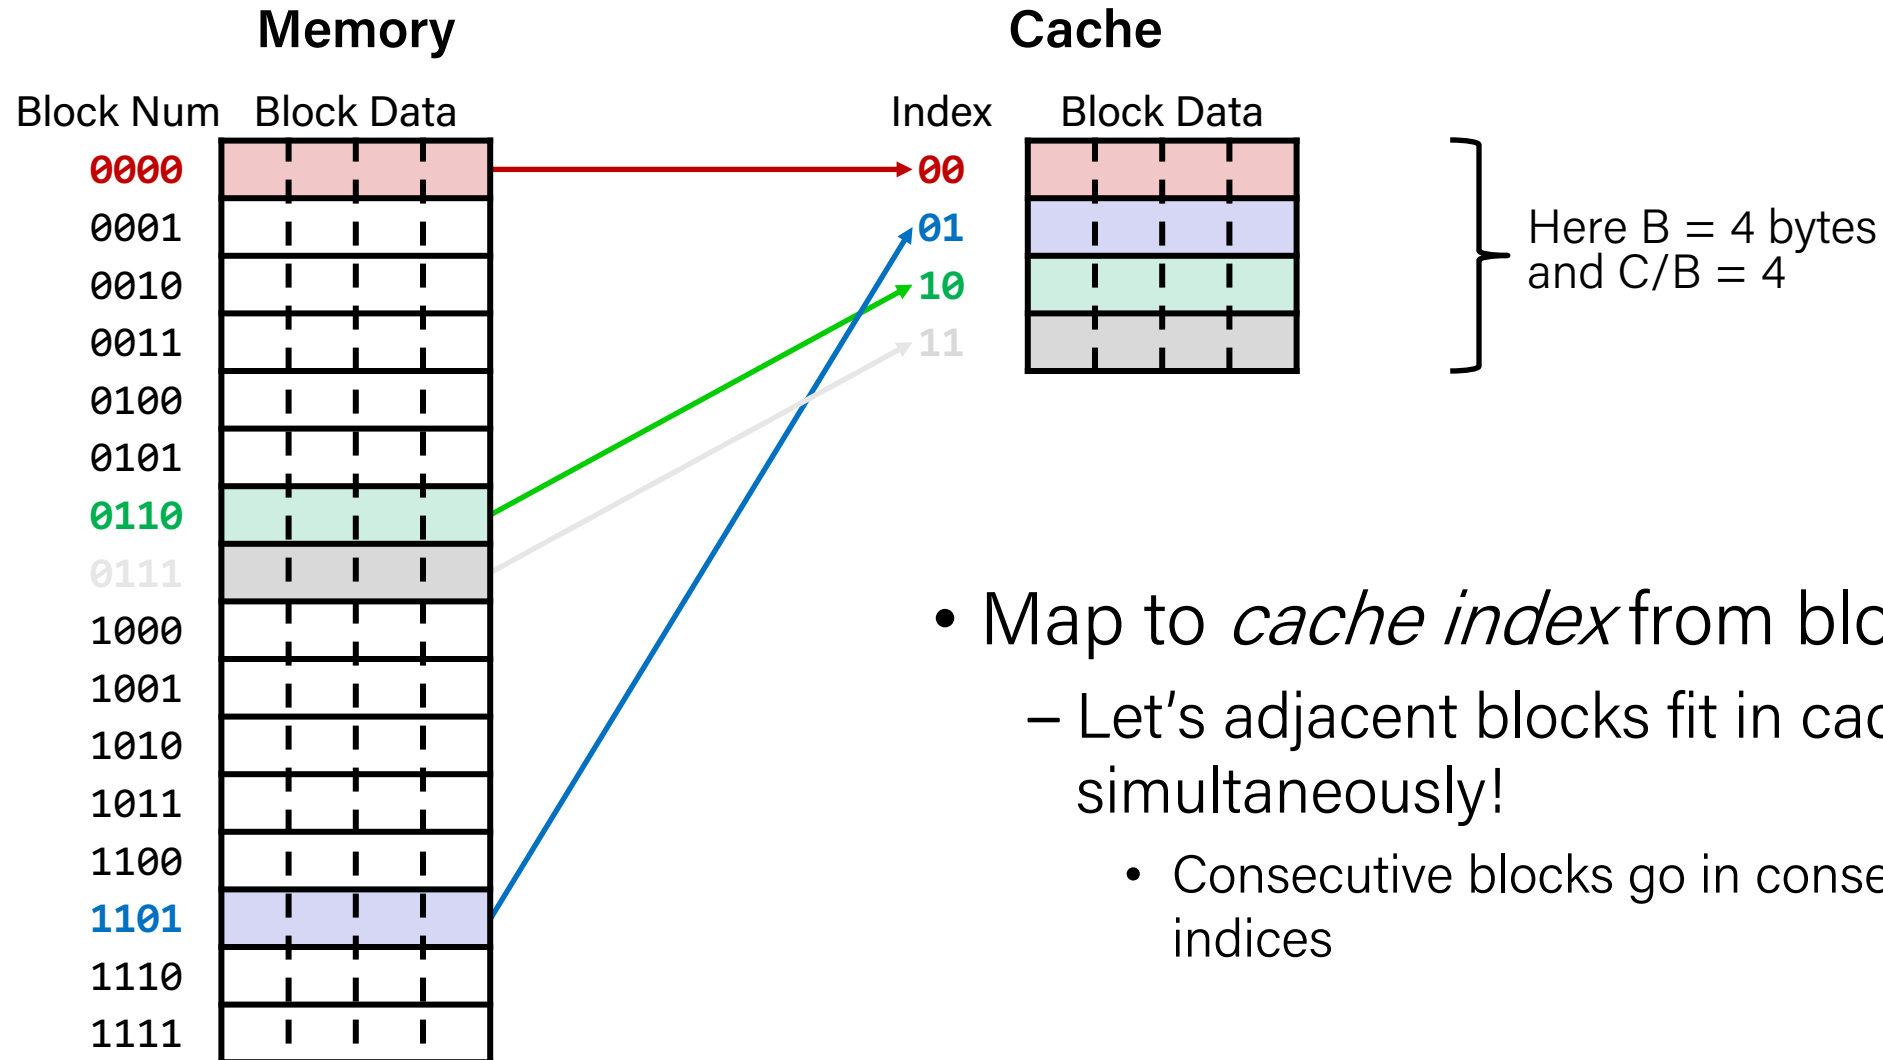and C/B = 4

- Map to *cache index* from block number
  - Use next $\log_2(C/B) = $ **s** bits in the address (after offset bits)
    - *C/B* is the number of sets here
  - (block number) mod (# blocks in cache)

Place Data in Cache by Hashing Address

Memory

Block Num   Block Data

0000	
0001	
0010	
0011	
0100	
0101	
0110	
0111	
1000	
1001	
1010	
1011	
1100	
1101	
1110	
1111	

Cache

Index   Block Data

00	
01	
10	
11	

Here B = 4 bytes
and C/B = 4

- Map to *cache index* from block number
  - Let's adjacent blocks fit in cache simultaneously!
    - Consecutive blocks go in consecutive cache indices

Practice Question

- 6-bit addresses, block size B = 4 bytes, and our cache holds S = 4 blocks.

- A request for address **0x2A** results in a cache miss. Which set index does this block get loaded into and which 3 other addresses are loaded along with it?

Practice Question

- 6-bit addresses, block size B = 4 bytes, and our cache holds S = 4 blocks.

  C = S $\times$ B = 16 bytes   b = $\log_2(4)$ = 2 bits   s = $\log_2(4)$ = 2 bits

- A request for address **0x2A** results in a cache miss. Which set index does this block get loaded into and which 3 other addresses are loaded along with it?

  0x 2   A
  Address: 0b 1 0 | 1 0 | 1 0
                        index offset
           block number

  addresses w/block number 1010
 0b101000 = 0x28
  0b101001 = 0x29
  0b101010 = 0x2A
  0b101011   = 0x2B

  These are loaded into cache!

# Place Data in Cache by Hashing Address

Memory

Cache

Block Num   Block Data

Index   Block Data

Block Num		Index	
0000		00	
0001		01	
0010		10	
0011		11	
0100			
0101			
0110			
0111			
1000			
1001			
1010			
1011			
1100			
1101			
1110			
1111			

Here B = 4 bytes
and C/B = 4

- Collusion!
  - This might confuse the cache later when we access the data
  - Solution?

Tags Differentiate Blocks in Same Index

- Tag = rest of address bits
  - **t** bits = **m** − **s** − **b**
  - Check this during a cache lookup

Checking for a Requested Address

- CPU sends address request for chunk of data
  - Address and requested data are not the same thing!
    - Analogy: your friend ≠ their phone number

- TIO address breakdown:

m-bit address:

Tag (**t**)	Index (**s**)	Offset (**b**)

Block Number

- Index field tells you where to look in cache
- Tag field lets you check that data is the block you want
- Offset field selects specified start byte within block

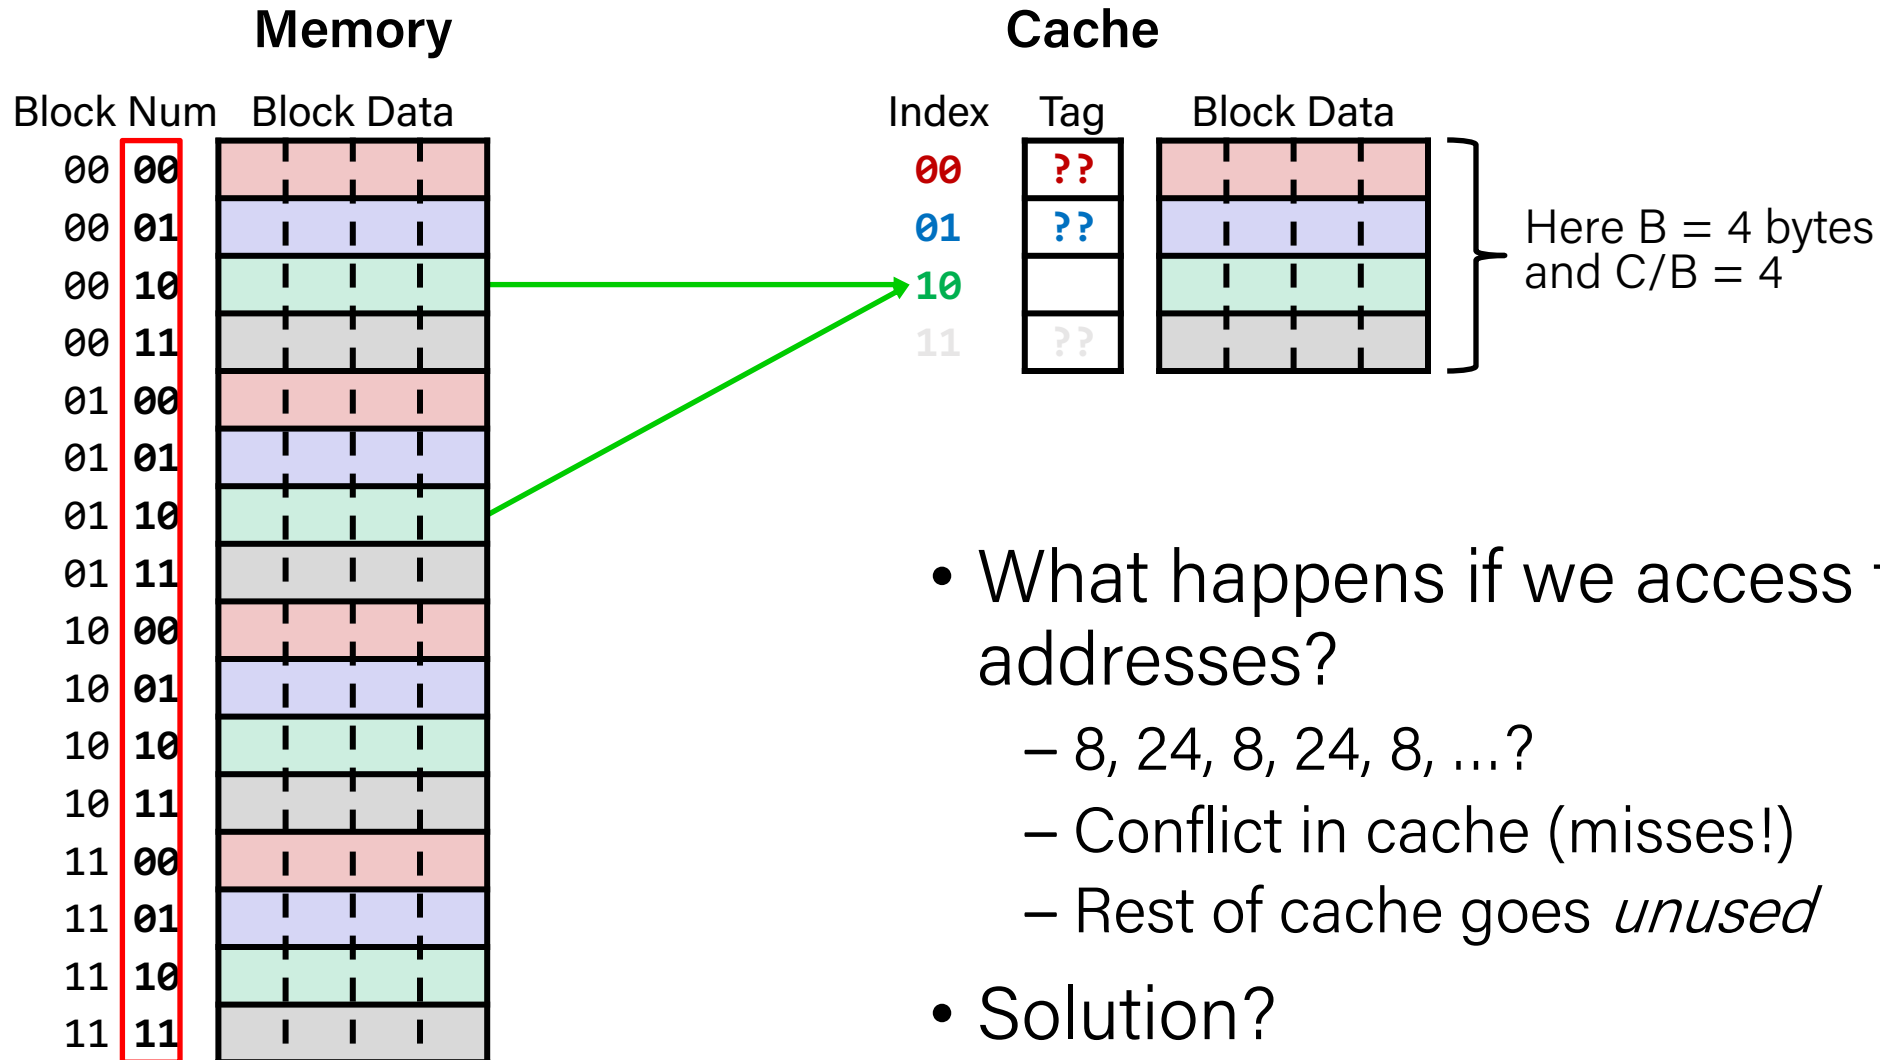- **Note: t** and **s** sizes will change based on hash function

Checking for a Requested Address Example

- Using 8-bit addresses.
- Cache Params: block size (B) = 4 bytes, cache size (C) = 32 bytes (which means number of sets is C/B = 8 sets).
  - Offset bits (b) = $\log_2$(B) = 2 bits
  - Index bits (s) = $\log_2$(number of sets) = 3 bits
  - Tag bits (t) = Rest of the bits in the address = 8 – 2 – 3 = 3 bits

m-bit address:

Tag (**t**)	Index (**s**)	Offset (**b**)

Block Number

- What are the fields for address 0xBA?
  - Tag bits (unique id for block):
  - Index bits (cache set block maps to):
  - Offset bits (byte offset within block):

Checking for a Requested Address Example

- Using 8-bit addresses.

- Cache Params: block size (B) = 4 bytes, cache size (C) = 32 bytes (which means number of sets is C/B = 8 sets).
    - Offset bits (b) = $\log_2$(B) = 2 bits
    - Index bits (s) = $\log_2$(number of sets) = 3 bits
    - Tag bits (t) = Rest of the bits in the address = 8 – 2 – 3 = 3 bits

m-bit address:

Tag (**t**)	Index (**s**)	Offset (**b**)

Block Number

- What are the fields for address 0xBA?
    - Tag bits (unique id for block):   0x5   <u>101 110 10</u>
    - Index bits (cache set block maps to):   0x6   5   6   2
   - Offset bits (byte offset within block): 0x2

30

Direct-Mapped Cache Problem

Memory

Block Num   Block Data

00	**00**
00	**01**
00	**10**
00	**11**
01	**00**
01	**01**
01	**10**
01	**11**
10	**00**
10	**01**
10	**10**
10	**11**
11	**00**
11	**01**
11	**10**
11	**11**

Cache

Index   Tag   Block Data

00	**??**
01	**??**
10	
11	**??**

Here B = 4 bytes
and C/B = 4

- What happens if we access the following addresses?
  – 8, 24, 8, 24, 8, …?
  – Conflict in cache (misses!)
  – Rest of cache goes *unused*
- Solution?

Associativity

- What if we could store data in any place in the cache?
  - More complicated hardware = more power consumed, slower

- So we *combine* the two ideas:
  - Each address maps to exactly one set
  - Each set can store block in more than one way

1-way:
8 sets,
1 block each
direct-mapped

2-way:
4 sets,
2 blocks each

4-way:
2 sets,
4 blocks each

8-way:
1 set,
8 blocks
fully associative

Cache Puzzle

- Based on the following behavior, which of the following block sizes is <u>NOT</u> possible for our cache?
  - Cache starts *empty*, also known as a **cold cache**
  - Access (addr: hit/miss) stream:
    - (14: miss), (15: hit), (16: miss)

  A. 4 bytes

  B. 8 bytes

  C. 16 bytes

  D. 32 bytes

  E. We're lost...

33

slido

Please download and install the Slido app on all computers you use

Based on the following behavior, which of the following block size is NOT possible for our cache?

ⓘ Start presenting to display the poll results on this slide.

Cache Puzzle

- Based on the following behavior, which of the following block size is <u>NOT</u> possible for our cache?
  - Cache starts *empty*, also known as a **cold cache**
  - Access (addr: hit/miss) stream:
    - (14: miss), (15: hit), (16: miss)

<u>hit:</u> block is already in cache!
<u>miss:</u> block is not in cache, pulls block from memory and puts it in cache

① Pulls block /w 14 into cache   **②** 15 is in the same block at 14   **③** 16 is not in block w/ 14 and 15

A. 4 bytes

B. 8 bytes

C. 16 bytes

D. 32 bytes

E. We're lost…

|	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Mem																	
K = 4															✗	✓	✗
K = 8															✗	✓	✗
K = 16															✗	✓	✗
K = 32															✗	✓	✓

Cache Organization

- **Associativity (E)**: # of ways for each set
    - Such a cache is called an "*E-way set associative cache*"
    - We now index into cache *sets*, of which there are S = C/B/E
    - Use lowest $\log_2(C/B/E) = s$ bits of block address
        - <u>Direct-mapped</u>:   E = 1, so $s = \log_2(C/B)$ as we saw previously
        - <u>Fully associative</u>:   E = C/B, so $s = 0$ bits
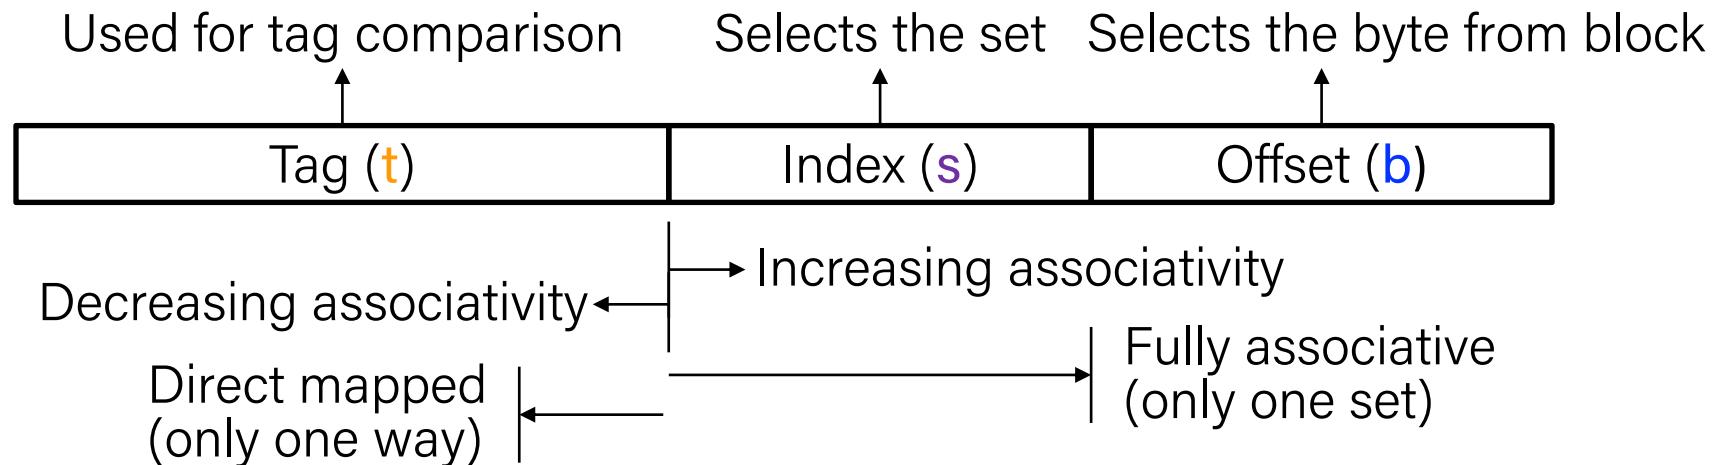
Used for tag comparison   Selects the set   Selects the byte from block

Tag (t)	Index (s)	Offset (b)

Decreasing associativity ← | → Increasing associativity

Direct mapped
(only one way)

Fully associative
(only one set)

Example Placement

- Where would data from address `0x1833` be placed?
  - Binary: `0b 0001 1000 0011 0011`

$t = m - s - b$ $s = \log2(C/B/E)$ $b = \log_2(B)$

m-bit address:

Tag(t)	Index (s)	Offset (b)

$s = ?$
Direct-mapped

Set	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

$s = ?$
2-way set associative

Set	Tag	Data
0		
1		
2		
3		

$s = ?$
4-way set associative

Set	Tag	Data
0		
1		

Example Placement

- Where would data from address `0x1833` be placed?
  - Binary: `0b 0001 1000 0011 0011`

  <small>E = 4
E = 2
E = 1</small>

$t = m - s - b$ $s = \log2(C/B/E)$ $b = \log_2(B)$

m-bit address:

Tag(t)	Index (s)	Offset (b)

$s = \log_2(8)=3$ bits
Direct-mapped

Set	Tag	Data
0		
1		
2		
3		✓
4		
5		
6		
7		

$s = \log_2(8/2)=2$ bits
2-way set associative

Set	Tag	Data
0		
1		
2		
3		✓ ✓

$s = \log_2(8/4)=1$ bit
4-way set associative

Set	Tag	Data
0		
1		✓ ✓ ✓ ✓

38

Block Placement

- *Any* empty block in the correct set may be used to store block

- If there are no empty blocks, which one should we replace?
  - No choice for direct-mapped caches
  - Caches typically use something close to **least recently used (LRU)** (hardware usually implements "*not most recently used*")

Direct-mapped

Set	Tag	Data
0		
1		
2		
3		✓
4		
5		
6		
7		

2-way set associative

Set	Tag	Data
0		
1		
2		
3		✓
		✓

4-way set associative

Set	Tag	Data
0		
1		✓
		✓
		✓
		✓

Question

- We have a cache of size 2 KB with block size of 128 bytes. If our cache has 2 sets, what is its associativity?

    A. 2

    B. 4

    C. 8

    D. 16

- If addresses are 16 bits wide, how wide is the Tag field?

Question

- We have a cache of size 2 KB with block size of 128 bytes. If our cache has 2 sets, what is its associativity?

$(S = 2)$

A. 2

B. 4

C. 8

D. 16

num blocks $= C / K = 2^{11}/2^7 = 2^4 = 16$ blocks

blocks per set $= E = 16 / 2 = 8$

- If addresses are 16 bits wide, how wide is the Tag field? $= 16 - 7 - 1 = 8$

General Cache Organization (S, E, B)

$E = 2^e$ lines per set

set

line

$S = 2^s$ sets

Cache size:
$C = S \times E \times B$ data bytes

v | tag | 0 | 1 | 2 | ...... | B-1

valid bit

$B = 2^b$ bytes per cache block (the data)
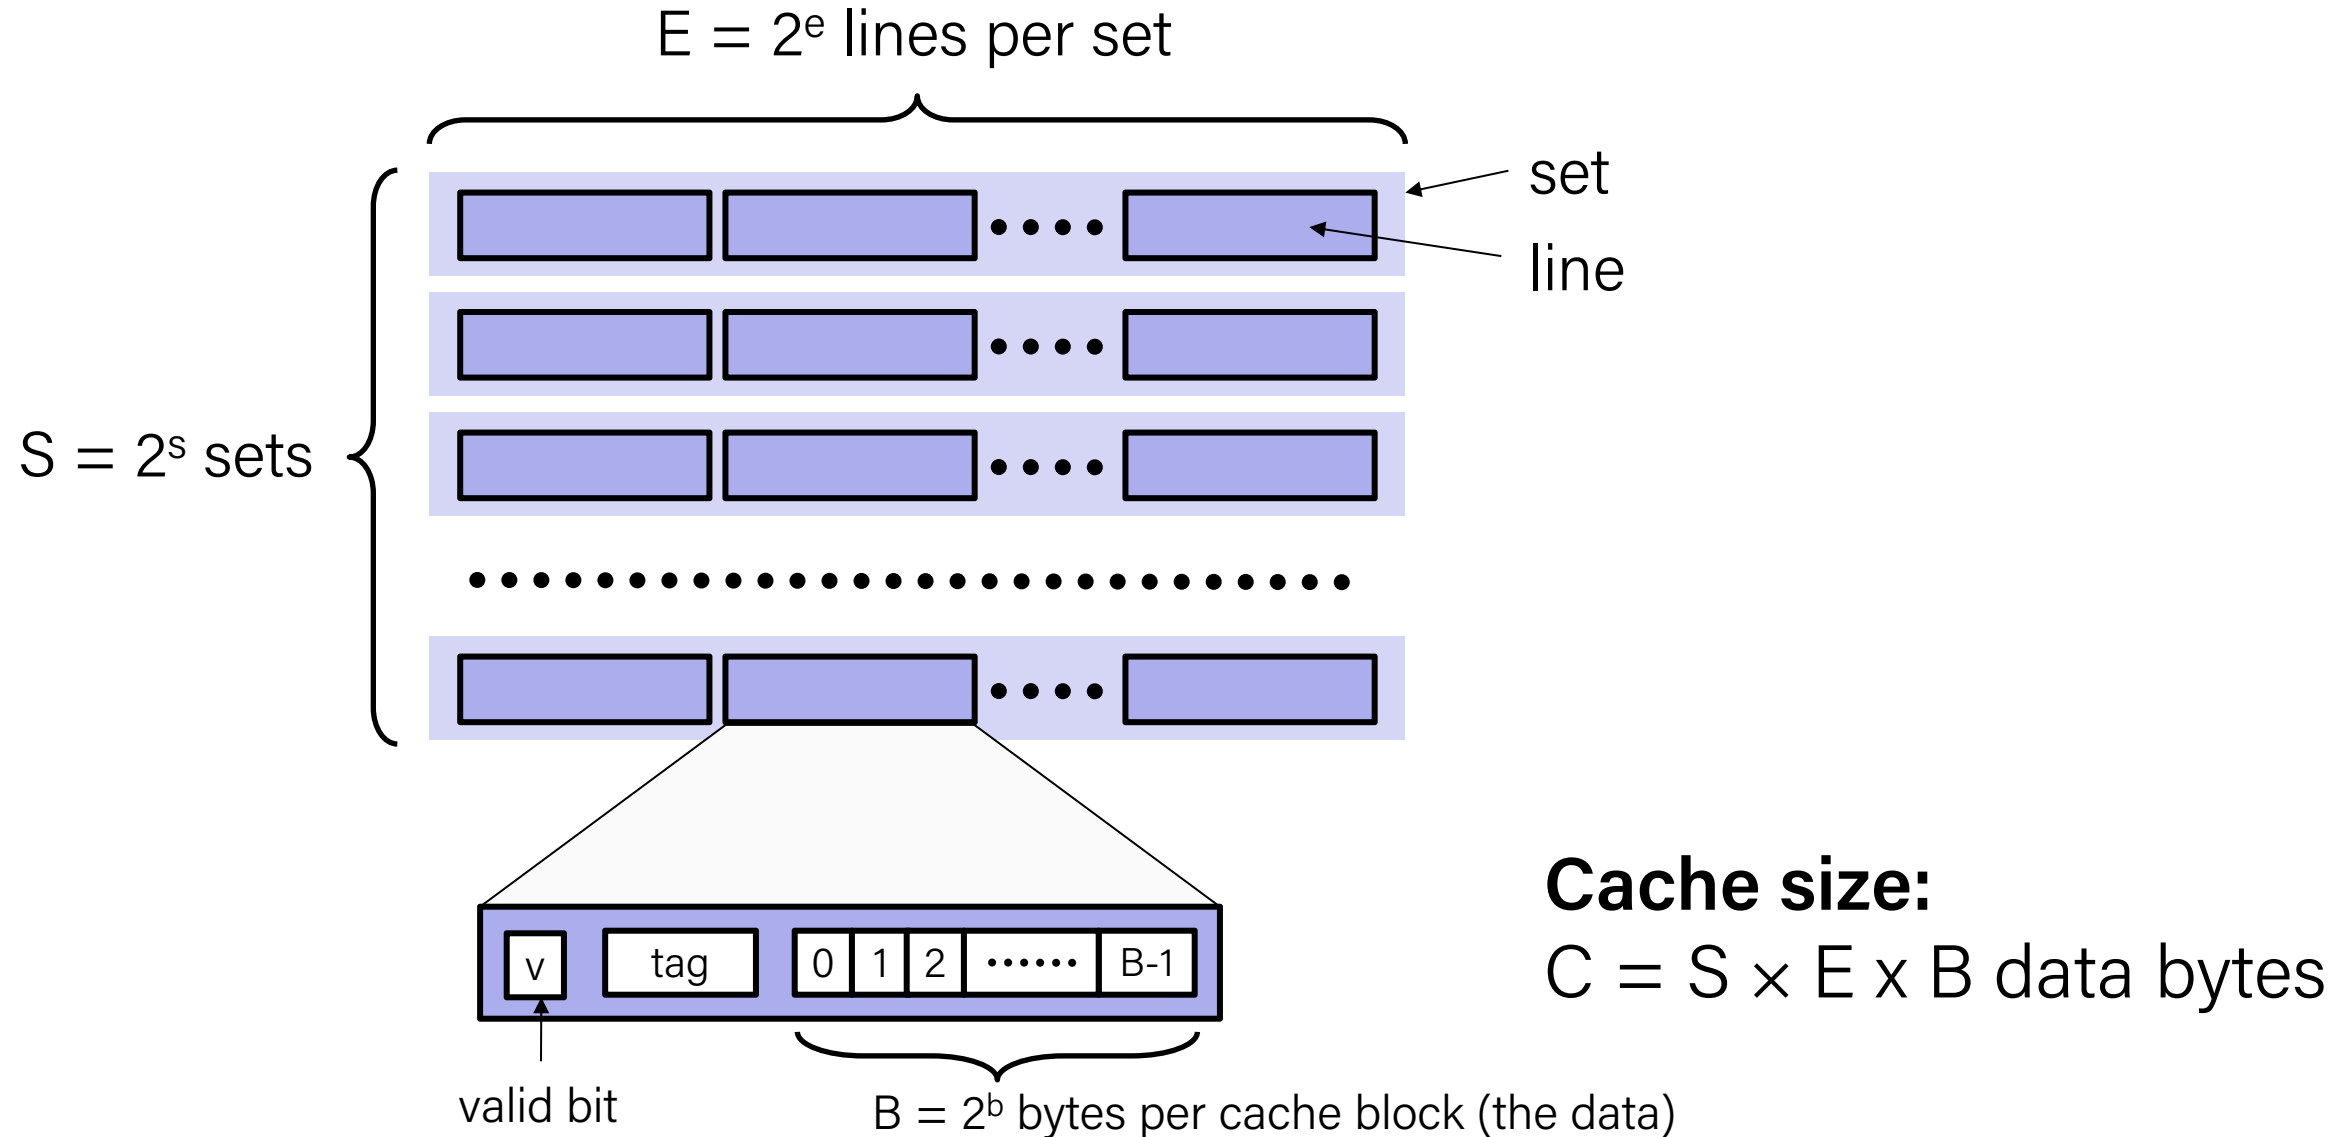
Cache Read
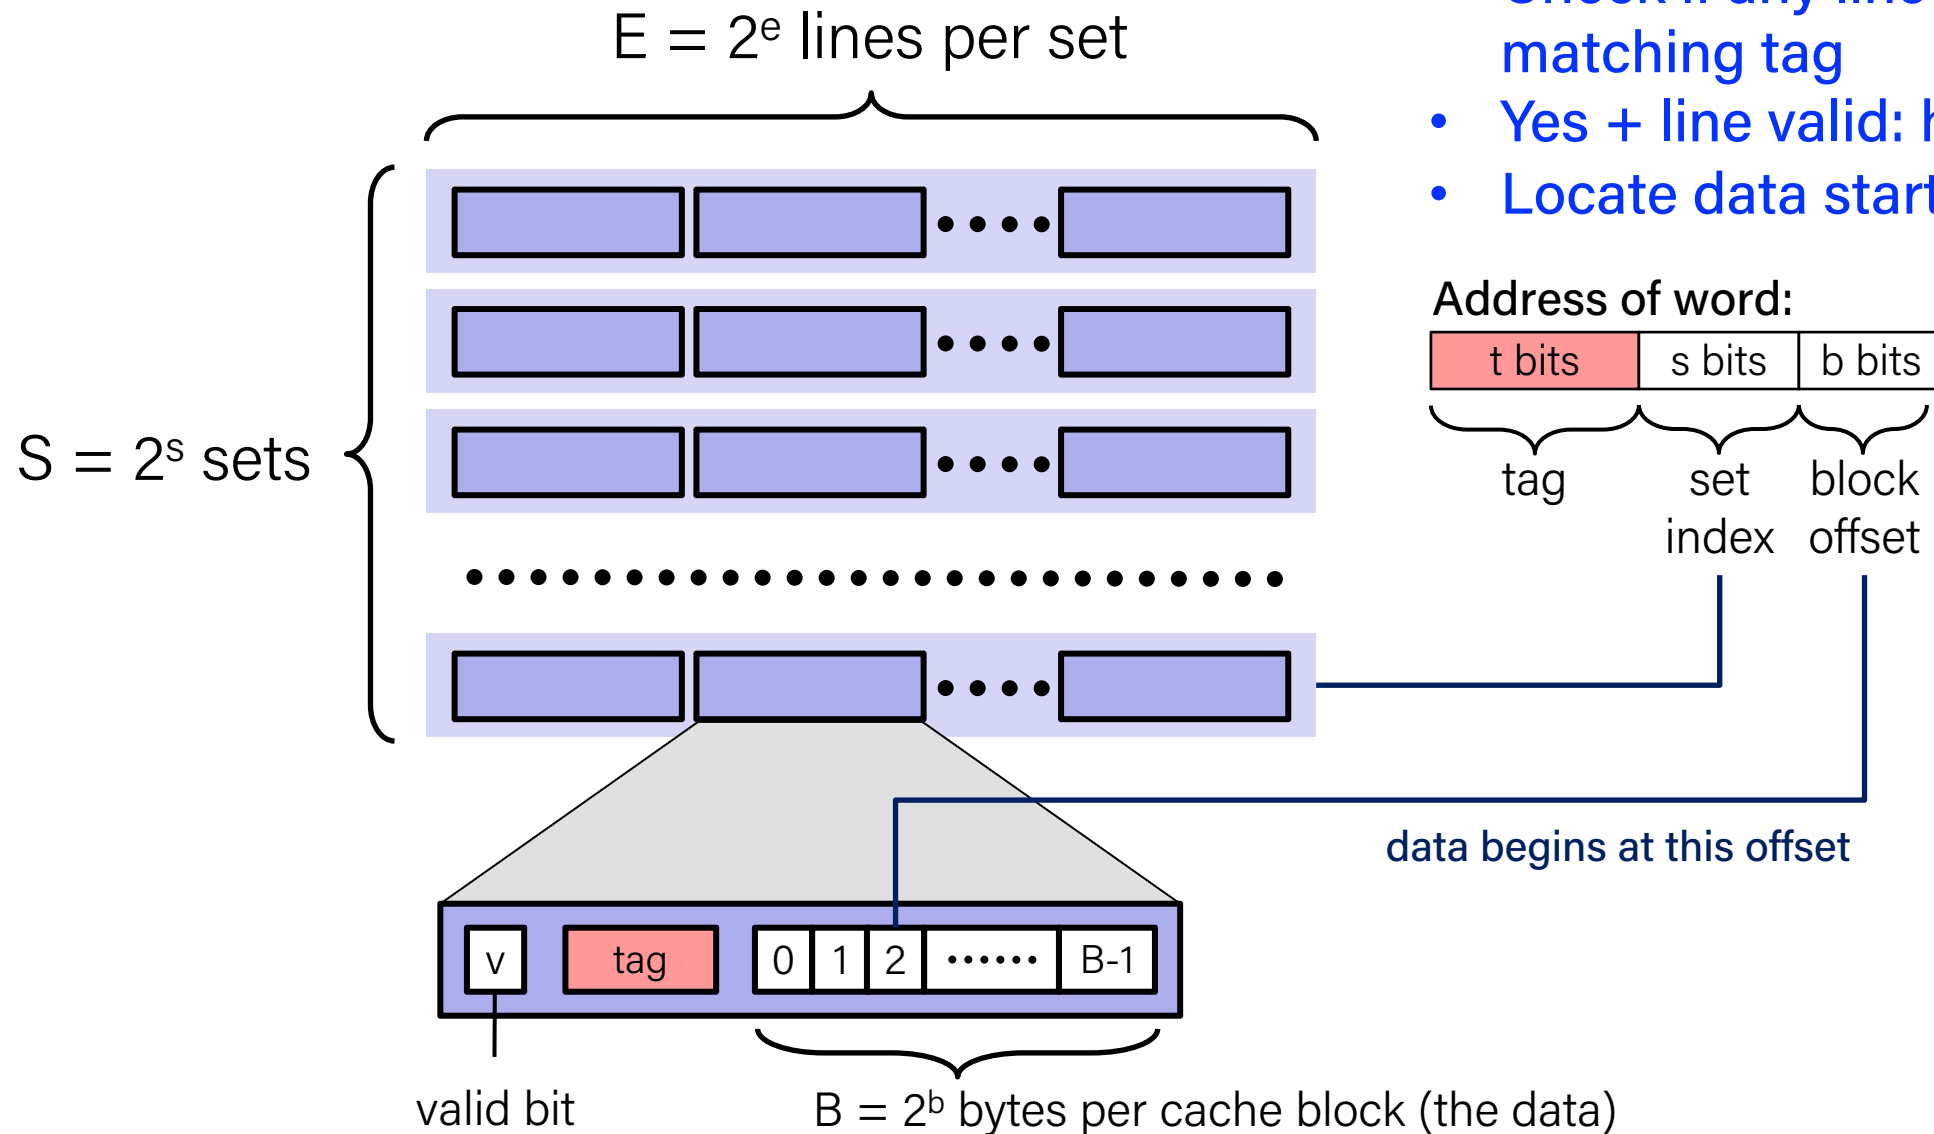
- **Locate set**
- **Check if any line in set has matching tag**
- **Yes + line valid: hit**
- **Locate data starting at offset**

$E = 2^e$ lines per set

$S = 2^s$ sets

Address of word:

t bits	s bits	b bits

tag   set index   block offset

data begins at this offset

v	tag	0	1	2	$\cdots\cdots$	B-1

valid bit

$B = 2^b$ bytes per cache block (the data)

43

Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set
Assume: cache block size 8 bytes

$S = 2^s$ sets

Address of `int`:

| t bits | 0...01 | 100 |

find set

Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set
Assume: cache block size 8 bytes

valid? +match: assume yes = hit

Address of int:

| t bits | 0...01 | 100 |

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

block offset

Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set
Assume: cache block size 8 bytes

valid?   +match: assume yes = hit

Address of `int`:

| t bits | 0...01 | 100 |

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

block offset

`int` (4 Bytes) is here
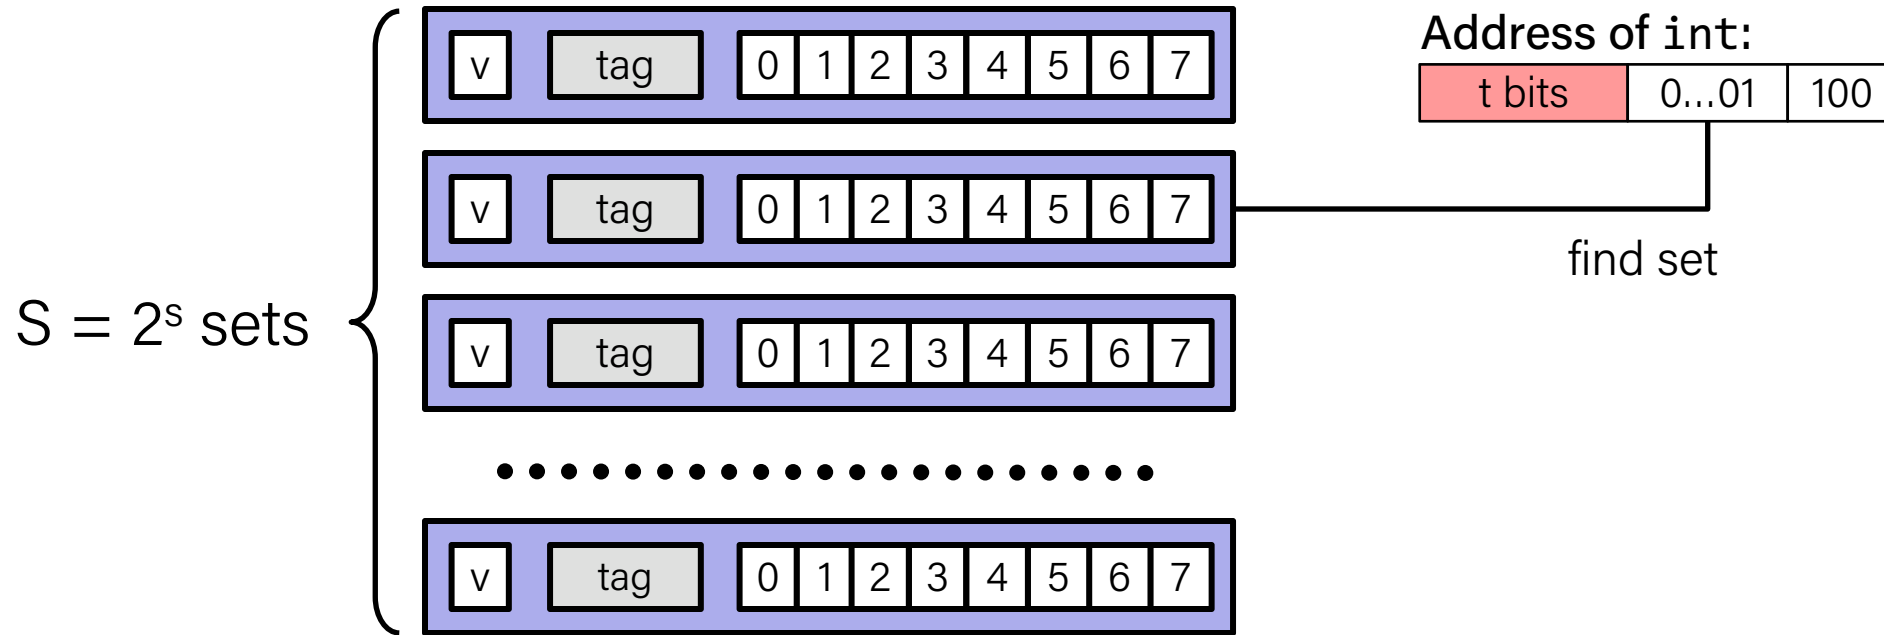
If tag doesn't match: old line is evicted and replaced

Direct-Mapped Cache Simulation

t=1	s=2	b=1
x	xx	x

M=16 bytes (4-bit addresses), B=2 bytes/block,
S=4 sets, E=1 Blocks/set

Address trace (reads, one byte per read):

0	$[\underline{000}0_2]$,	miss
1	$[\underline{000}1_2]$,	hit
7	$[\underline{011}1_2]$,	miss
8	$[\underline{100}0_2]$,	miss
0	$[\underline{000}0_2]$	miss

	v	Tag	Block
Set 0	1	0	M[0-1]
Set 1			
Set 2			
Set 3	1	0	M[6-7]

E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set
Assume: cache block size 8 bytes

Address of `short int`:

| t bits | 0...01 | 100 |

find set

E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set
Assume: cache block size 8 bytes

Address of `short int:`

| t bits | 0...01 | 100 |

compare both

valid? + match: yes = hit

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

block offset

E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set
Assume: cache block size 8 bytes

Address of `short int`:

t bits	0...01	100

compare both

valid? + match: yes = hit

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

| v | tag | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

block offset

`short int` (2 Bytes) is here
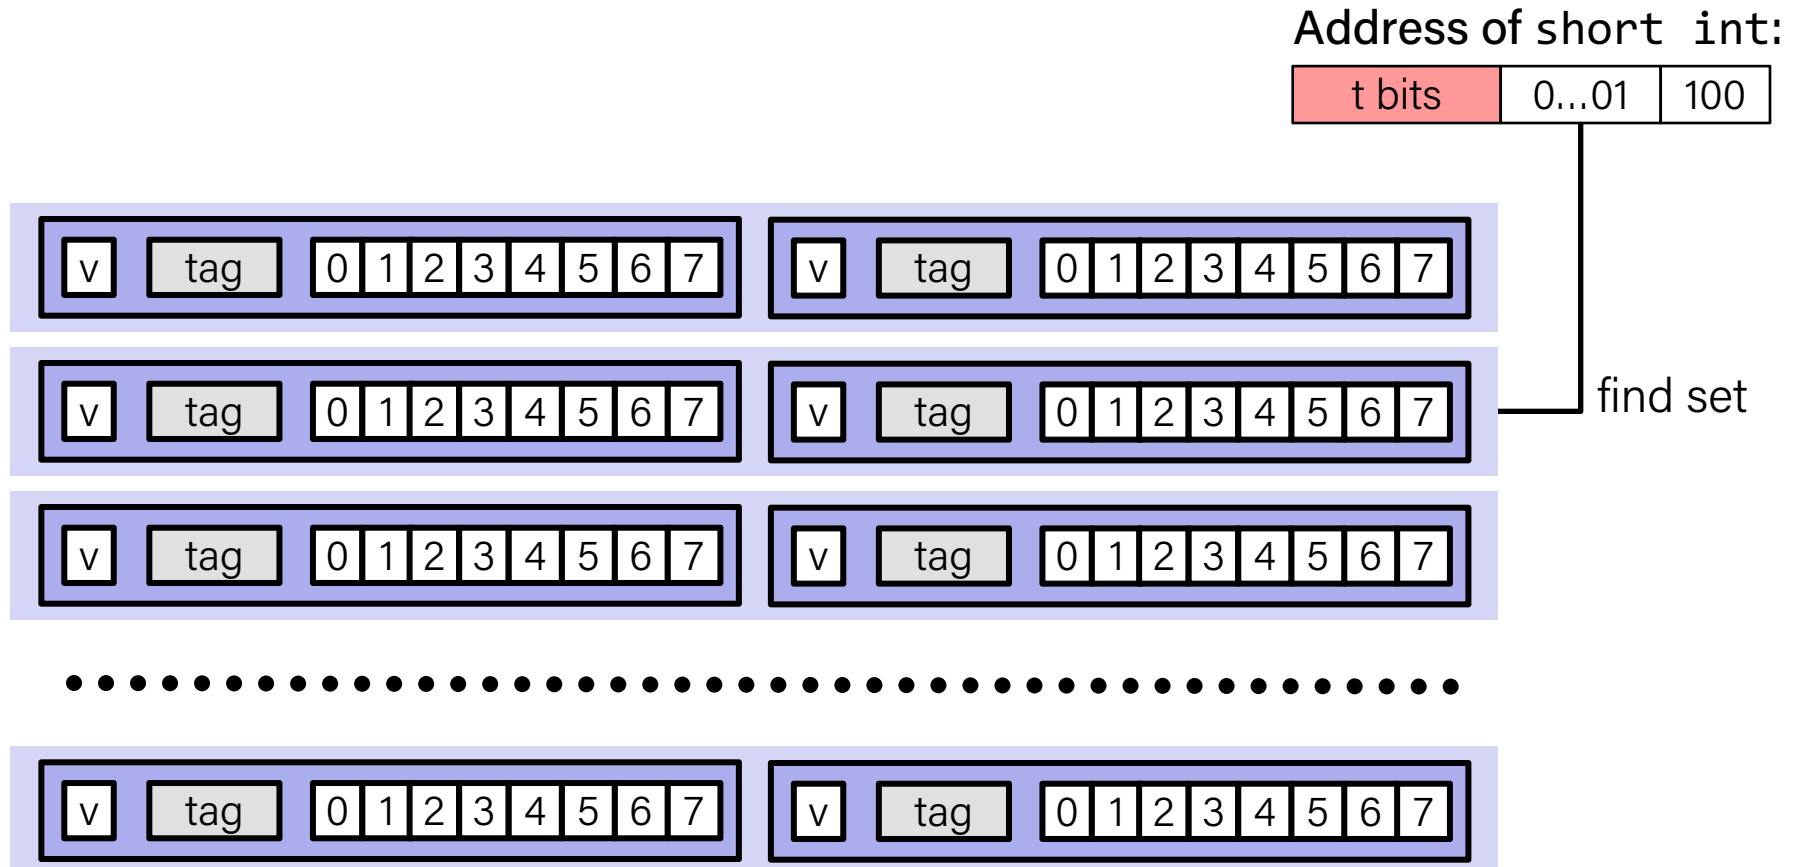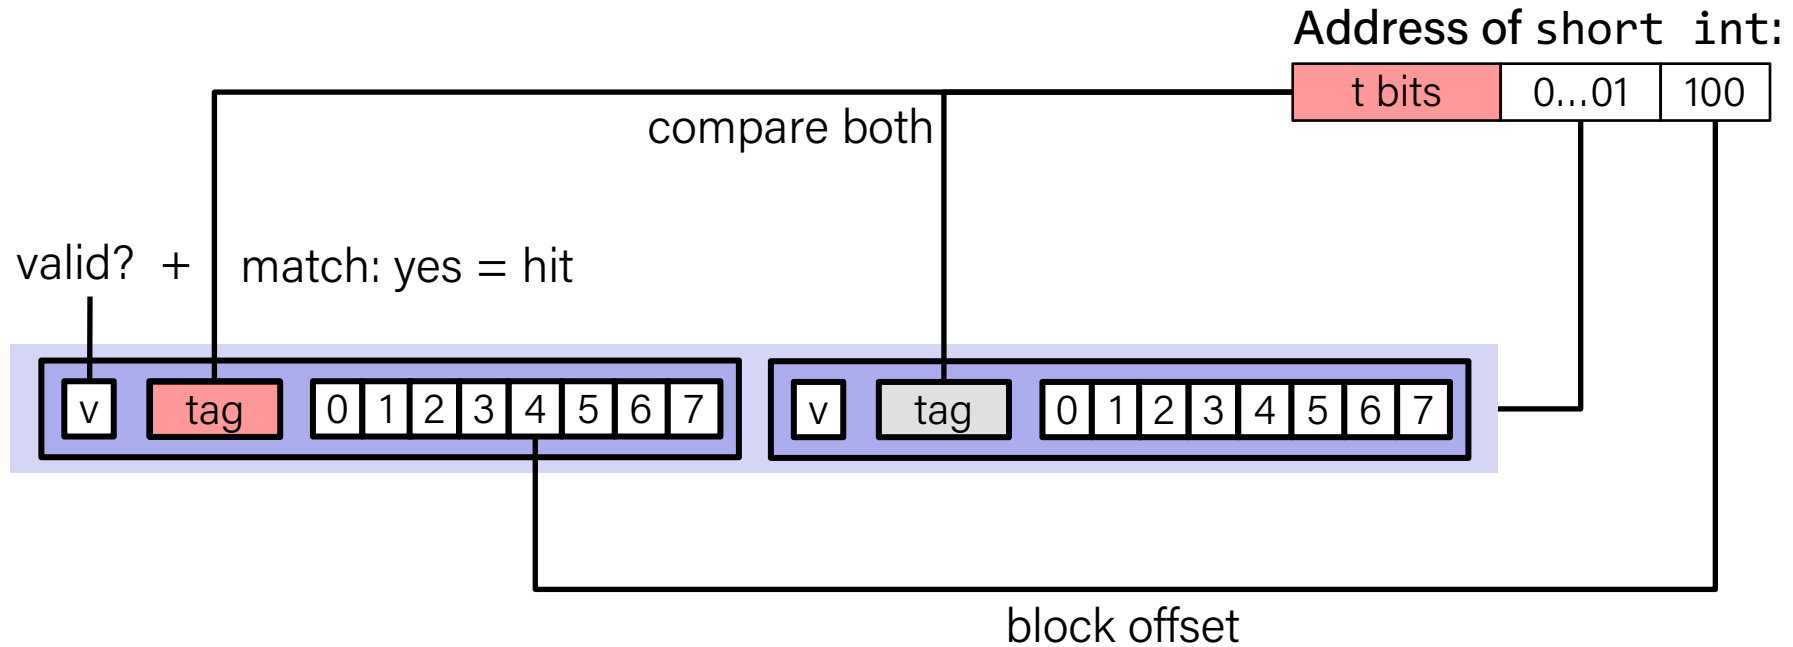
No match:
- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

2-Way Set Associative Cache Simulation

t=2 s=1 b=1

XX	X	X

M=16 byte addresses, B=2 bytes/block,
S=2 sets, E=2 blocks/set

Address trace (reads, one byte per read):

0	[0000_2],	miss
1	[0001_2],	hit
7	[0111_2],	miss
8	[1000_2],	miss
0	[0000_2]	hit

	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

What about writes?

- **Multiple copies of data exist:**
  - L1, L2, L3, Main Memory, Disk

- **What to do on a write-hit?**
  - Write-through (write immediately to memory)
  - Write-back (defer write to memory until replacement of line)
    - Need a dirty bit (line different from memory or not)

- **What to do on a write-miss?**
  - Write-allocate (load into cache, update line in cache)
    - Good if more writes to the location follow
  - No-write-allocate (writes straight to memory, does not load into cache)

- **Typical**
  - Write-through + No-write-allocate
  - **Write-back + Write-allocate**

Write-back, Write Allocate Example

Note: While unrealistic, this example assumes that all requests have offset 0 and are for a block's worth of data.

Cache:

Valid	Dirty	Tag	Block Contents
1	0	G	0xBEEF

There is only one set in this tiny cache, so the tag is the entire block number!

Memory:

Block Num

F	0xCAFE
G	0xBEEF

Write-back, Write Allocate Example

1) `mov $0xFACE, (F)`

Write Miss!

Cache:

Valid	Dirty	Tag	Block Contents
1	0	G	0xBEEF

Step 1: Bring `F` into cache

Memory:

Block Num

F	0xCAFE
G	0xBEEF

54

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`
   Write Miss

Valid Dirty   Tag   Block Contents

Cache: | 1 | 0 | F | 0xCAFE |

Step 1: Bring `F` into cache

Memory:

Block Num

F   0xCAFE

G   0xBEEF

Step 2: Write `0xFACE` to cache only and set the dirty bit

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`

Write Miss

Cache:

| Valid	Dirty	Tag	Block Contents
1	1	F	0xFACE

Memory:

Block Num

⋮	
F	0xCAFE
⋮	
G	0xBEEF
⋮	

Step 1: Bring `F` into cache

Step 2: Write `0xFACE` to cache only and set the dirty bit

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`
      Write Miss

(2) `mov $0xFEED, (F)`
      Write Hit!

Cache:

Valid	Dirty	Tag	Block Contents
1	1	F	0xFACE

<u>Step</u>: Write `0xFEED` to cache only (and set the dirty bit)

Memory:

Block Num

F	0xCAFE

G	0xBEEF

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`
   Write Miss

(2) `mov $0xFEED, (F)`
   Write Hit!

Cache:

Valid	Dirty	Tag	Block Contents
1	1	F	0xFEED

Memory:

Block Num

F	0xCAFE
G	0xBEEF

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`
     Write Miss

(2) `mov $0xFEED, (F)`
     Write Hit!

(3) `mov (G), %ax`
     Read Miss!

Cache:

	Valid	Dirty	Tag	Block Contents
	1	1	F	0xFEED

Step 1: Write `F` back to memory since it is dirty

Memory:

Block Num	
F	0xCAFE
G	0xBEEF

Write-back, Write Allocate Example

(1) `mov $0xFACE, (F)`
Write Miss

(2) `mov $0xFEED, (F)`
Write Hit!

(3) `mov (G), %ax`
Read Miss!

Cache:

	Valid	Dirty	Tag	Block Contents
	1	0	G	0xBEEF

Memory:

Block Num

F	0xFEED
G	0xBEEF

Step 1: Write `F` back to memory since it is dirty

Step 2: Bring `G` into the cache so that we can copy it into `%ax`

Cache Simulator

Polling Question

- Which of the following cache statements is FALSE?

  A. We can reduce compulsory misses by decreasing our block size

  B. We can reduce conflict misses by increasing associativity

  C. A write-back cache will save time for code with good temporal locality on writes

  D. A write-through cache will always match data with the memory hierarchy level below it

  E. We're lost...

Polling Question

- Which of the following cache statements is FALSE?

    A. We can reduce compulsory misses by decreasing our block size
    smaller block size pulls fewer bytes into cache on a miss

    B. We can reduce conflict misses by increasing associativity
    more options to place blocks before evictions occur

    C. A write-back cache will save time for code with good temporal locality on writes
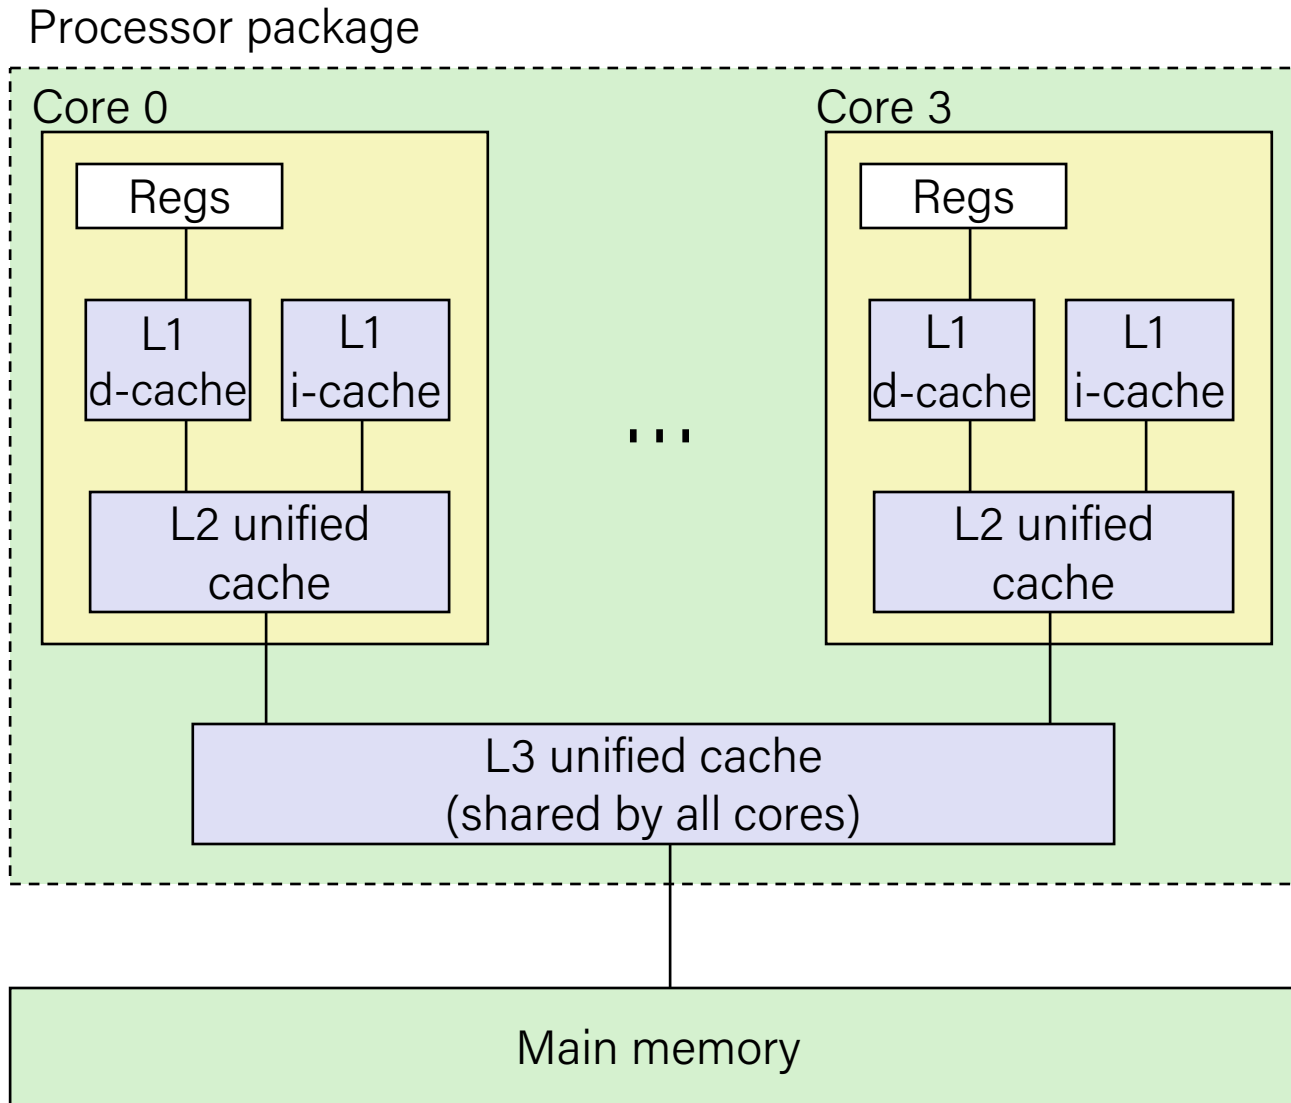   yes, its main goal is data consistency

    D. A write-through cache will always match data with the memory hierarchy level below it
    frequently-used blocks rarely get evicted, so fewer write-backs

    E. We're lost…

Intel Core i7 Cache Hierarchy

L1 i-cache and d-cache:
    32 KB, 8-way,
    Access: 4 cycles

L2 unified cache:
    256 KB, 8-way,
    Access: 10 cycles

L3 unified cache:
    8 MB, 16-way,
    Access: 40-75 cycles

Block size: 64 bytes for all caches.

Lecture Plan

- Cache memory organization and operation

- The memory mountain

Writing Cache Friendly Code

- Make the common case go fast
  - Focus on the inner loops of the core functions

- Minimize the misses in the inner loops
  - Repeated references to variables are good (temporal locality)
  - Stride-1 reference patterns are good (spatial locality)

Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories

The Memory Mountain

- **Read throughput** (read bandwidth)
  - Number of bytes read from memory per second (MB/s)

- **Memory mountain:** Measured read throughput as a function of spatial and temporal locality.
  - Compact way to characterize memory system performance.

Memory Mountain Test Function

```c
long data[MAXELEMS];  /* Global array to traverse */

/* test – Iterate over first "elems" elements of
 *        array "data" with stride of "stride", using
 *        using 4x4 loop unrolling.
 */
int test(int elems, int stride) {
    long i, sx2=stride*2, sx3=stride*3, sx4=stride*4;
    long acc0 = 0, acc1 = 0, acc2 = 0, acc3 = 0;
    long length = elems, limit = length – sx4;

    /* Combine 4 elements at a time */
    for (i = 0; i < limit; i += sx4) {
        acc0 = acc0 + data[i];
        acc1 = acc1 + data[i+stride];
        acc2 = acc2 + data[i+sx2];
        acc3 = acc3 + data[i+sx3];
    }

    /* Finish any remaining elements */
    for (; i < length; i++) {
        acc0 = acc0 + data[i];
    }
    return ((acc0 + acc1) + (acc2 + acc3));
}
```

Call `test()` with many combinations of `elems` and `stride`.

For each elems and stride:

1. Call test() once to warm up the caches.

2. Call test() again and measure the read throughput(MB/s)

mountain/mountain.c

The Memory Mountain

Core i7 Haswell
2.1 GHz
32 KB L1 d-cache
256 KB L2 cache
8 MB L3 cache
64 B block size

69

Recap

- Cache memories can have significant performance impact

- You can write your programs to exploit this!

  - Focus on the inner loops, where bulk of computations and memory accesses occur.

  - Try to maximize spatial locality by reading data objects with sequentially with stride 1.

  - Try to maximize temporal locality by using a data object as often as possible once it's read from memory.

Next time: *Optimization*