

COMP201

Computer Systems & Programming

Lecture #18 – Introduction to x86-64 Assembly



KOÇ
UNIVERSITY

Aykut Erdem // Koç University // Fall 2020

Good news, everyone!

- Course Evaluations are now available via Koç Üniversitesi Mobil (deadline extended to November 23)



Recap

1. **Bits and Bytes** - *How can a computer represent numbers?*
2. **Chars and C-Strings** - *How can a computer represent and manipulate more complex data like text?*
3. **Pointers, Stack and Heap** – *How can we effectively manage all types of memory in our programs?*
4. **Generics** - *How can we use our knowledge of memory and data representation to write code that works with any data type?*

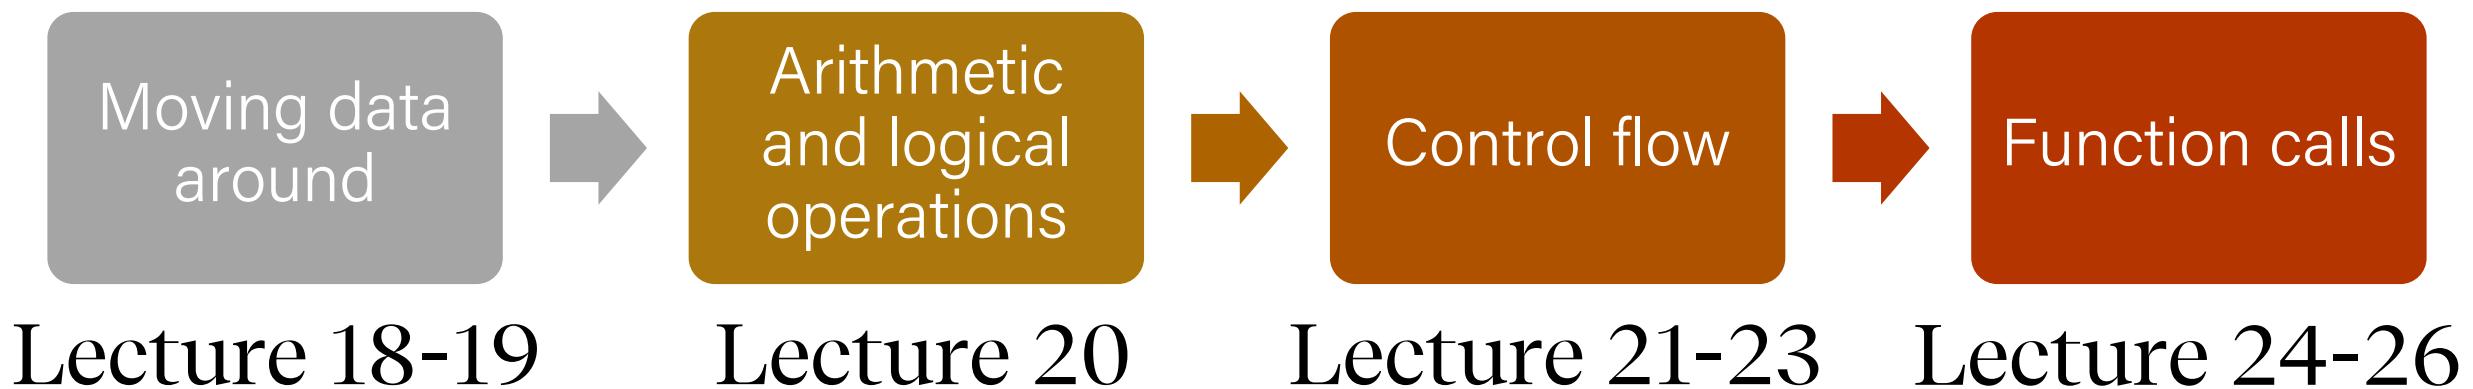
Course Overview

1. **Bits and Bytes** - *How can a computer represent numbers?*
2. **Chars and C-Strings** - *How can a computer represent and manipulate more complex data like text?*
3. **Pointers, Stack and Heap** – *How can we effectively manage all types of memory in our programs?*
4. **Generics** - *How can we use our knowledge of memory and data representation to write code that works with any data type?*

5. **Assembly** - *How does a computer interpret and execute C programs?*
6. **Heap Allocators** - *How do core memory-allocation operations like malloc and free work?*
7. **The Memory Hierarchy** – *How to improve the performance of application programs by improving their temporal and spatial locality?*
8. **Code Optimization** – *How write C code so that a compiler can then generate efficient machine code?*
9. **Linking** – *How static and dynamic linking work?*

COMP201 Topic 6: How does a computer interpret and execute C programs?

Learning Assembly



Learning Goals

- Learn what assembly language is and why it is important
- Become familiar with the format of human-readable assembly and x86
- Learn the **mov** instruction and how data moves around at the assembly level

Plan for Today

- **Overview:** GCC and Assembly
- **Demo:** Looking at an executable
- Registers and The Assembly Level of Abstraction
- The **mov** instruction

Disclaimer: Slides for this lecture were borrowed from
—Nick Troccoli's Stanford CS107 class

Lecture Plan

- **Overview:** GCC and Assembly
- **Demo:** Looking at an executable
- Registers and The Assembly Level of Abstraction
- The `mov` instruction

Bits all the way down

Data representation so far

- Integer (`unsigned int`, 2's complement `signed int`)
- Floating Points (IEEE single (`float`) and double (`double`) precision)
- `char` (ASCII)
- Address (`unsigned long`)
- Aggregates (arrays, `structs`)

The code itself is binary too!

- Instructions (machine encoding)

GCC

- **GCC** is the compiler that converts your human-readable code into machine-readable instructions.
- C, and other languages, are high-level abstractions we use to write code efficiently. But computers don't really understand things like data structures, variable types, etc. Compilers are the translator!
- Pure machine code is 1s and 0s – everything is bits, even your programs! But we can read it in a human-readable form called **assembly**. (Engineers used to write code in assembly before C).
- There may be multiple assembly instructions needed to encode a single C instruction.
- We're going to go behind the curtain to see what the assembly code for our programs looks like.

Lecture Plan

- Overview: GCC and Assembly
- **Demo:** Looking at an executable
- Registers and The Assembly Level of Abstraction
- The `mov` instruction

Demo: Looking at an Executable (objdump -d)



Our First Assembly

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;  
    for (int i = 0; i < nelems; i++) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

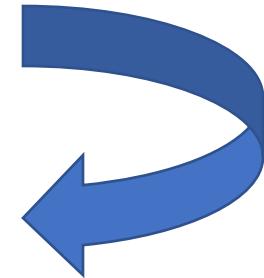
What does this look like in assembly?

Our First Assembly

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;  
    for (int i = 0; i < nelems; i++) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

00000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00	mov	\$0x0,%edx
4005bb:	b8 00 00 00 00	mov	\$0x0,%eax
4005c0:	eb 09	jmp	4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq	%edx,%rcx
4005c5:	03 04 8f	add	(%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add	\$0x1,%edx
4005cb:	39 f2	cmp	%esi,%edx
4005cd:	7c f3	jl	4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz	retq



make
objdump -d sum

Our First Assembly

00000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00	mov \$0x0,%edx
4005bb:	b8 00 00 00 00	mov \$0x0,%eax
4005c0:	eb 09	jmp 4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq %edx,%rcx
4005c5:	03 04 8f	add (%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add \$0x1,%edx
4005cb:	39 f2	cmp %esi,%edx
4005cd:	7c f3	jl 4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz retq

Our First Assembly

00000000004005b6 <sum_array>:

4005b6: b8 00 00 00 00

4005b7: b8 00 00 00 00

This is the name of the function (same as C) and the memory address where the code for this function starts.

4005c8: 83 c2 01

4005cb: 39 f2

4005cd: 7c f3

4005cf: f3 c3

mov \$0x0,%edx
\$0x0,%eax
4005cb <sum_array+0x15>
vs1q %edx,%rcx
(%rdi,%rcx,4),%eax
\$0x1,%edx
%esi,%edx
jl 4005c2 <sum_array+0xc>
repz retq

Our First Assembly

00000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00	mov \$0x0,%edx
4005bb:	b8 00 00 00 00	mov \$0x0,%eax
4005c0:	eb 09	jmp 4005cb <sum_array+0x15>
4005c2:	40 00 00 00 00	add \$0x40,%eax
4005c5:	00 00 00 00 00	add \$0x0,%eax
4005c8:	89 d0	mov %eax,%eax
4005cb:	39 72	cmpl %eax,%eax
4005cd:	7c f3	j1 4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz retq

These are the memory addresses where each of the instructions live. Sequential instructions are sequential in memory.

Our First Assembly

00000000004005b6 <sum_array>:

```
4005b6:    ba 00 00 00 00  
4005bb:    b8 00 00 00 00  
4005c0:    eb 09
```

This is the assembly code:
“human-readable” versions of
each machine code instruction.

```
4005cd:    39 12  
4005cd:    7c f3  
4005cf:    f3 c3
```

```
mov    $0x0,%edx  
mov    $0x0,%eax  
jmp    4005cb <sum_array+0x15>  
movslq %edx,%rcx  
add    (%rdi,%rcx,4),%eax  
add    $0x1,%edx  
cmp    %esi,%edx  
jl    4005c2 <sum_array+0xc>  
repz  retq
```

Our First Assembly

0000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00
4005bb:	b8 00 00 00 00
4005c0:	eb 09
4005c2:	48 63 ca
4005c5:	03 04 8f
4005c8:	83 c2 01
4005cb:	39 f2
4005cd:	7c f3
4005cf:	f3 c3

mov \$0x0,%edx
add %dx,%dx

This is the machine code: raw hexadecimal instructions, representing binary as read by the computer. Different instructions may be different byte lengths.

repz retq

Our First Assembly

0000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00	mov \$0x0,%edx
4005bb:	b8 00 00 00 00	mov \$0x0,%eax
4005c0:	eb 09	jmp 4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq %edx,%rcx
4005c5:	03 04 8f	add (%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add \$0x1,%edx
4005cb:	39 f2	cmp %esi,%edx
4005cd:	7c f3	jl 4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz retq

Our First Assembly

0000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00 00	mov	\$0x0,%edx
4005bb:	b8 00 00 00 00 00	mov	\$0x0,%eax
4005c0:	eb 09	jmp	4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq	%edx,%rcx
4005c5:	03 04 8f	add	(%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add	\$0x1,%edx
4005cb:	39 f2	cpx	%esi,%edx
4005cd:	7c f3	jl	4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz	retq



Each instruction has an operation name ("opcode").

Our First Assembly

0000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00 00	mov \$0x0,%edx
4005bb:	b8 00 00 00 00 00	mov \$0x0,%eax
4005c0:	eb 09	jmp 4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq %edx,%rcx
4005c5:	03 04 8f	add (%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add \$0x1,%edx
4005cb:	39 f2	cmp %esi,%edx
4005cd:	7c f3	jl 4005c2 <sum_array+0xc>
4005cf:	f3 c3	

Each instruction can also have arguments ("operands").

Our First Assembly

0000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00 00	mov	\$0x0,%edx
4005bb:	b8 00 00 00 00 00	mov	\$0x0,%eax
4005c0:	eb 09	jmp	4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq	%edx,%rcx
4005c5:	03 04 8f	add	(%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add	\$0x1,%edx
4005cb:	39 f2	cmp	%esi,%edx
4005cd:	7c f3	jl	4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz	retq

\$[number] means a constant value, or “immediate” (e.g. 1 here).

Our First Assembly

00000000004005b6 <sum_array>:

4005b6:	ba 00 00 00 00 00	mov	\$0x0,%edx
4005bb:	b8 00 00 00 00 00	mov	\$0x0,%eax
4005c0:	eb 09	jmp	4005cb <sum_array+0x15>
4005c2:	48 63 ca	movslq	%edx,%rcx
4005c5:	03 04 8f	add	(%rdi,%rcx,4),%eax
4005c8:	83 c2 01	add	\$0x1,%edx
4005cb:	39 f2	cmp	%esi,%edx
4005cd:	7c f3	jl	4005c2 <sum_array+0xc>
4005cf:	f3 c3	repz	retq



%[name] means a register, a storage location on the CPU (e.g. edx here).

Lecture Plan

- Overview: GCC and Assembly
- Demo: Looking at an executable
- Registers and The Assembly Level of Abstraction
- The `mov` instruction

Assembly Abstraction

- C abstracts away the low-level details of machine code. It lets us work using variables, variable types, and other higher-level abstractions.
- C and other languages let us write code that works on most machines.
- Assembly code is just bytes! No variable types, no type checking, etc.
- Assembly/machine code is processor-specific.
- What is the level of abstraction for assembly code?

Registers



Registers



%rax



%rsi



%r8



%r12



%rbx



%rdi



%r9



%r13



%rcx



%rbp



%r10



%r14



%rdx



%rsp



%r11



%r15

Registers

What is a register?

A register is a fast read/write memory slot right on the CPU that can hold variable values.

Registers are **not** located in memory.

Registers

- A **register** is a 64-bit space inside the processor.
- There are 16 registers available, each with a unique name.
- Registers are like “scratch paper” for the processor. Data being calculated or manipulated is moved to registers first. Operations are performed on registers.
- Registers also hold parameters and return values for functions.
- Registers are extremely *fast* memory!
- Processor instructions consist mostly of moving data into/out of registers and performing arithmetic on them. This is the level of logic your program must be in to execute!

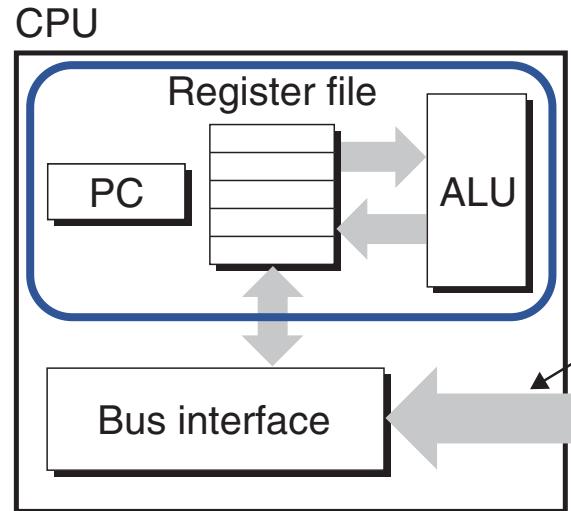
Machine-Level Code

Assembly instructions manipulate these registers. For example:

- One instruction adds two numbers in registers
- One instruction transfers data from a register to memory
- One instruction transfers data from memory to a register

Computer architecture

registers accessed by name
ALU is main workhorse of CPU



System bus

I/O bridge

Memory bus

Main memory

"hello, world\n"
hello code

memory needed for program execution (stack, heap, etc.) accessed by address

I/O bus

Disk controller

Disk

hello executable stored on disk

disk/server stores program when not executing

USB controller

Graphics adapter

Mouse Keyboard

Display

Expansion slots for other devices such as network adapters

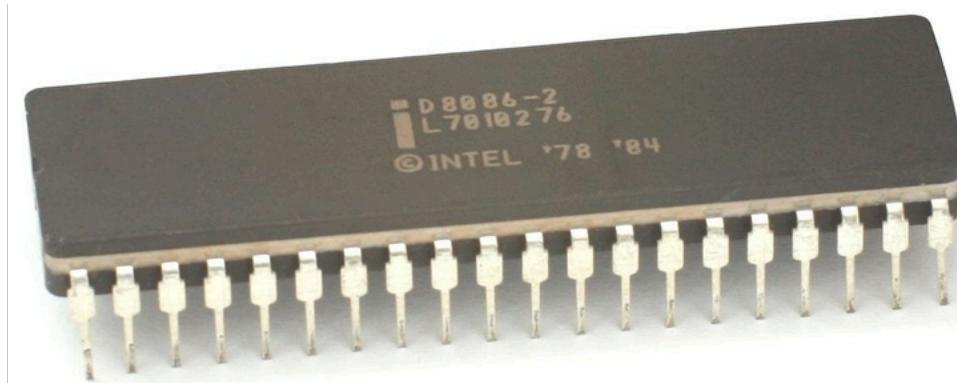
GCC And Assembly

- GCC compiles your program – it lays out memory on the stack and heap and generates assembly instructions to access and do calculations on those memory locations.
- Here's what the "assembly-level abstraction" of C code might look like:

C	Assembly Abstraction
int sum = x + y;	<ol style="list-style-type: none">1) Copy <i>x</i> into register 12) Copy <i>y</i> into register 23) Add register 2 to register 14) Write register 1 to memory for <i>sum</i>

Assembly

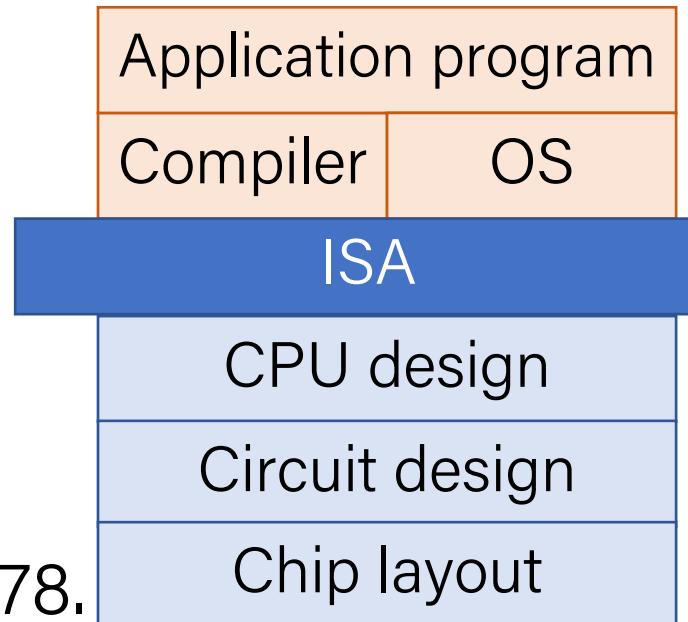
- We are going to learn the **x86-64** instruction set architecture. This instruction set is used by Intel and AMD processors.
- There are many other instruction sets: ARM, MIPS, etc.



Instruction set architecture (ISA)

A contract between program/compiler and hardware:

- Defines operations that the processor (CPU) can execute
- Data read/write/transfer operations
- Control mechanisms



Intel originally designed their instruction set back in 1978.

- Legacy support is a huge issue for x86-64
- Originally 16-bit processor, then 32 bit, now 64 bit.
These design choices dictated the register sizes
(and even register/instruction names).



Lecture Plan

- **Overview:** GCC and Assembly
- **Demo:** Looking at an executable
- Registers and The Assembly Level of Abstraction
- The **mov** instruction

mov

The **mov** instruction copies bytes from one place to another; it is similar to the assignment operator (=) in C.

mov **src,dst**

The **src** and **dst** can each be one of:

- Immediate (constant value, like a number) (*only src*) **\$0x104**
 - Register **%rbx**
 - Memory Location
(*at most one of src, dst*) Direct address **0x6005c0**

Operand Forms: Immediate

mov \$0x104, _____



*Copy the value 0x104
into some
destination.*

Operand Forms: Registers

mov

%rbx, _____

Copy the value in register %rbx into some destination.

mov

_____ **,%rbx**

Copy the value from some source into register %rbx.

Operand Forms: Absolute Addresses

mov

0x104, _____

*Copy the value at
address 0x104 into
some destination.*

mov

_____ , 0x104

*Copy the value from
some source into the
memory at address
0x104.*

Practice #1: Operand Forms

What are the results of the following move instructions (executed separately)? For this problem, assume the value 5 is stored at address 0x42, and the value 8 is stored in %rbx.

1. **mov \$0x42,%rax** Move 0x42 into %rax

2. **mov 0x42,%rax** Move 5 into %rax

3. **mov %rbx,0x55** Move 8 to the address %rax

Operand Forms: Indirect

mov

(%rbx), _____

Copy the value at the address stored in register %rbx into some destination.

mov

_____ , **(%rbx)**

Copy the value from some source into the memory at the address stored in register %rbx.

Operand Forms: Base + Displacement

mov

0x10(%rax), _____

*Copy the value at the address (**0x10 plus** what is stored in register %rax) into some destination.*

mov

_____, **0x10(%rax)**

*Copy the value from some source into the memory at the address (**0x10 plus** what is stored in register %rax).*

Operand Forms: Indexed

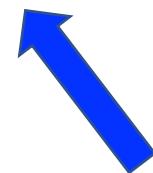
mov

(%rax,%rdx), _____



mov

_____,(%rax,%rdx)



Copy the value at the address which is (the sum of the values in registers %rax and %rdx) into some destination.

Copy the value from some source into the memory at the address which is (the sum of the values in registers %rax and %rdx).

Operand Forms: Indexed

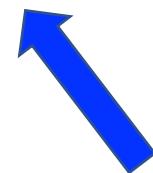
mov

0x10(%rax,%rdx), _____



mov

_____,**0x10(%rax,%rdx)**



Copy the value at the address which is (the sum of 0x10 plus the values in registers %rax and %rdx) into some destination.

Practice #2: Operand Forms

What are the results of the following move instructions (executed separately)? For this problem, assume

the value $0x11$ is stored at address $0x10C$,

the value $0xAB$ is stored at address $0x104$,

$0x100$ is stored in register $\%rax$ and $0x3$ is stored in $\%rdx$.

- | | | |
|---------------|---------------------------|---------------------------------------|
| 1. mov | \$0x42, (%rax) | Move $0x42$ to memory address $0x100$ |
| 2. mov | 4(%rax), %rcx | Move $0xAB$ into $\%rcx$ |
| 3. mov | 9(%rax,%rdx), %rcx | Move $0x11$ into $\%rcx$ |

$Imm(r_b, r_i)$ is equivalent to address $Imm + R[r_b] + R[r_i]$

Displacement: positive or negative constant (if missing, = 0)

Base: register (if missing, = 0)

Index: register (if missing, = 0)

Operand Forms: Scaled Indexed

mov (,%rdx,4), _____

mov _____, (,%rdx,4)

Copy the value at the address which is
(4 times the value in register %rdx)
into some destination.

The scaling factor
(e.g. 4 here) must
be hardcoded to
be either 1, 2, 4 or
8.

Copy the value from some source into the
memory at the address which is (4 times
the value in register %rdx).

Operand Forms: Scaled Indexed

mov

$0x4(,%rdx,4),\underline{\hspace{1cm}}$

*Copy the value at the address which is
(4 times the value in register %rdx, **plus**
0x4), into some destination.*

mov

$\underline{\hspace{1cm}},0x4(,%rdx,4)$

*Copy the value from some source into the
memory at the address which is (4 times the
value in register %rdx, **plus 0x4**).*

Operand Forms: Scaled Indexed

mov

(%rax,%rdx,2),_____

mov

_____,(%rax,%rdx,2)

Copy the value at the address which is (the value in register %rax plus 2 times the value in register %rdx) into some destination.

Copy the value from some source into the memory at the address which is (the value in register %rax plus 2 times the value in register %rdx).

Operand Forms: Scaled Indexed

mov

0x4(%rax,%rdx,2),_____

Copy the value at the address which is (0x4 plus the value in register %rax plus 2 times the value in register %rdx) into some destination.

mov

_____,0x4(%rax,%rdx,2)

Copy the value from some source into the memory at the address which is (0x4 plus the value in register %rax plus 2 times the value in register %rdx).

Most General Operand Form

Imm(r_b, r_i, s)

is equivalent to...

Imm + R[r_b] + R[r_i]*s

Most General Operand Form

$\text{Imm}(r_b, r_i, s)$ is equivalent to
address $\text{Imm} + R[r_b] + R[r_i]*s$

Displacement:

pos/neg constant
(if missing, = 0)

Index: register
(if missing, = 0)

Base: register (if
missing, = 0)

Scale must be
1,2,4, or 8
(if missing, = 1)

Operand Forms

Type	Form	Operand Value	Name
Immediate	\$Imm	Imm	Immediate
Register	r_a	$R[r_a]$	Register
Memory	Imm	$M[Imm]$	Absolute
Memory	(r_a)	$M[R[r_a]]$	Indirect
Memory	$Imm(r_b)$	$M[Imm + R[r_b]]$	Base + displacement
Memory	(r_b, r_i)	$M[R[r_b] + R[r_i]]$	Indexed
Memory	$Imm(r_b, r_i)$	$M[Imm + R[r_b] + R[r_i]]$	Indexed
Memory	$(, r_i, s)$	$M[R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(, r_i, s)$	$M[Imm + R[r_i] \cdot s]$	Scaled indexed
Memory	(r_b, r_i, s)	$M[R[r_b] + R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$	Scaled indexed

Figure 3.3 from the book: “Operand forms. Operands can denote immediate (constant) values, register values, or values from memory. The scaling factor s must be either 1, 2, 4, or 8.”

Practice #3: Operand Forms

What are the results of the following move instructions (executed separately)? For this problem, assume

the value *0x1* is stored in register %rcx,
the value *0x100* is stored in register %rax,
the value *0x3* is stored in register %rdx, and
the value *0x11* is stored at address *0x10C*.

1. **mov \$0x42,0xfc(,%rcx,4)**

Move 0x42 to memory address 0x100

2. **mov (%rax,%rdx,4),%rbx**

Move 0x11 into %rbx

Imm(r_b , r_i , s) is equivalent to
address **Imm + R[r_b] + R[r_i]*s**

Displacement	Base	Index	Scale
			(1,2,4,8)

Goals of indirect addressing: C

Why are there so many
forms of indirect addressing?

We see these indirect addressing
paradigms in C as well!

Extra Practice

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
long arr[5];
```

```
...
```

```
long num = ____???____;
```

```
// %rdi stores arr, %rcx stores 3, and %rax stores num
mov (%rdi, %rcx, 8),%rax
```

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
long arr[5];
```

```
...
```

```
long num = arr[3];
```

```
// %rdi stores arr, %rcx stores 3, and %rax stores num
mov (%rdi,%rcx,8),%rax
```

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
int x = ...  
int *ptr = malloc(...);  
      ???      = x;
```

```
// %ecx stores x, %rax stores ptr  
mov %ecx,(%rax)
```

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
int x = ...  
int *ptr = malloc(...);  
*ptr = x;
```

```
// %ecx stores x, %rax stores ptr  
mov %ecx,(%rax)
```

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
char str[5];
```

```
...
```

```
_____ = 'c';
```

```
// %rcx stores str, %rdx stores 2
mov $0x63,(%rcx,%rdx,1)
```

Extra Practice

slido

Event code:
73165

Fill in the blank to complete the code that generated the assembly below.

```
char str[5];
```

```
...
```

```
str[2] = 'c';
```

```
// %rcx stores str, %rdx stores 2
mov $0x63,(%rcx,%rdx,1)
```

Our First Assembly

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;  
    for (int i = 0; i < nelems; i++) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

We're 1/4th of the way to understanding assembly!
What looks understandable right now?

Some notes:

- Registers store addresses and values
- `mov src, dst` **copies** value into dst
- `sizeof(int)` is 4
- Instructions executed sequentially

0000000004005b6 <sum_array>:

```
4005b6: ba 00 00 00 00  
4005bb: b8 00 00 00 00  
4005c0: eb 09  
4005c2: 48 63 ca  
4005c5: 03 04 8f  
4005c8: 83 c2 01
```

```
mov    $0x0,%edx  
mov    $0x0,%eax  
jmp    4005cb <sum_array+0x15>  
movslq %edx,%rcx  
add    (%rdi,%rcx,4),%eax  
add    $0x1,%edx  
cmp    %esi,%edx  
jl    4005c2 <sum_array+0xc>  
repz  retq
```

We'll come back to this example in future lectures!



Additional Reading

The screenshot shows a web browser window displaying the Ars Technica website. The URL in the address bar is <https://arstechnica.com/gadgets/2002/03/an-introduction-to-64-bit-computing-and-x86-64/>. The main content is an article titled "An Introduction to 64-bit Computing and x86-64" by Jon Stokes, published on 3/12/2002 at 1:36 AM. The article discusses AMD's plans to extend the x86 ISA to 64 bits. It includes a sidebar with social sharing icons (110 shares, Facebook, Twitter) and a video thumbnail for "ARS VIDEO: 30 People Play Mario Kart 8 From Newbies to Pros". The video thumbnail shows a person wearing a Mario Kart hat.

An Introduction to 64-bit Computing and x86-64

When I first heard that AMD had plans to extend the x86 ISA to 64 bits, I ...

JON STOKES - 3/12/2002, 1:36 AM

110

f

Twitter

When I first heard that AMD had plans to extend the x86 ISA to 64 bits, I thought it was a terrible idea. Though x86 is the world's most successful ISA, it's also the world's most widely disparaged. Programmers, analysts, architecture buffs, and enthusiasts often see x86 as a leaden albatross around the neck of the entire computing industry, and like the Mariner's albatross we were all hoping it would just fall off at some point and slip quietly into the deep. But in spite of such hopes, I really knew better. In fact, I've argued elsewhere that x86 isn't going away anytime soon, and it no longer makes any real sense to gripe about its quirks from a performance perspective. I won't recap that argument here, but I can sum it up briefly.

Most of us would probably assent to the following statement: "there's a huge global market for mainstream business and consumer software, and the overwhelming majority of that software just so happens to use the x86 ISA." This statement is true, as far as it goes, but framing x86's role in the software industry this way misses an important point. In my article "[The Future of x86 and the Concept of the ISA](#)," I argue that a statement like the following would provide a more relevant assessment of the true state of the software industry: "There's a huge global market for mainstream business and consumer *x86 software*, and several smaller markets for software written to other ISAs." All discussions of the desktop prospects of widely ported operating systems (i.e. Linux) or of the possible effects of greater open source market penetration aside, this statement should still ring true to anyone who's acquainted with the present realities of the installed base of IT and consumer software.

If we think realistically about most of the world's commercial software not as "software" in the abstract but as x86 binary code, then it becomes apparent that improvements to the x86 ISA

ARS VIDEO

30 People Play Mario Kart 8 From Newbies to Pros

They usually don't have very many tailings.

<https://arstechnica.com/gadgets/2002/03/an-introduction-to-64-bit-computing-and-x86-64/>

Recap

- **Overview:** GCC and Assembly
- **Demo:** Looking at an executable
- Registers and The Assembly Level of Abstraction
- The **mov** instruction

Next time: diving deeper into assembly