

COMP541

DEEP LEARNING

Lecture #9 – Graph Neural Networks

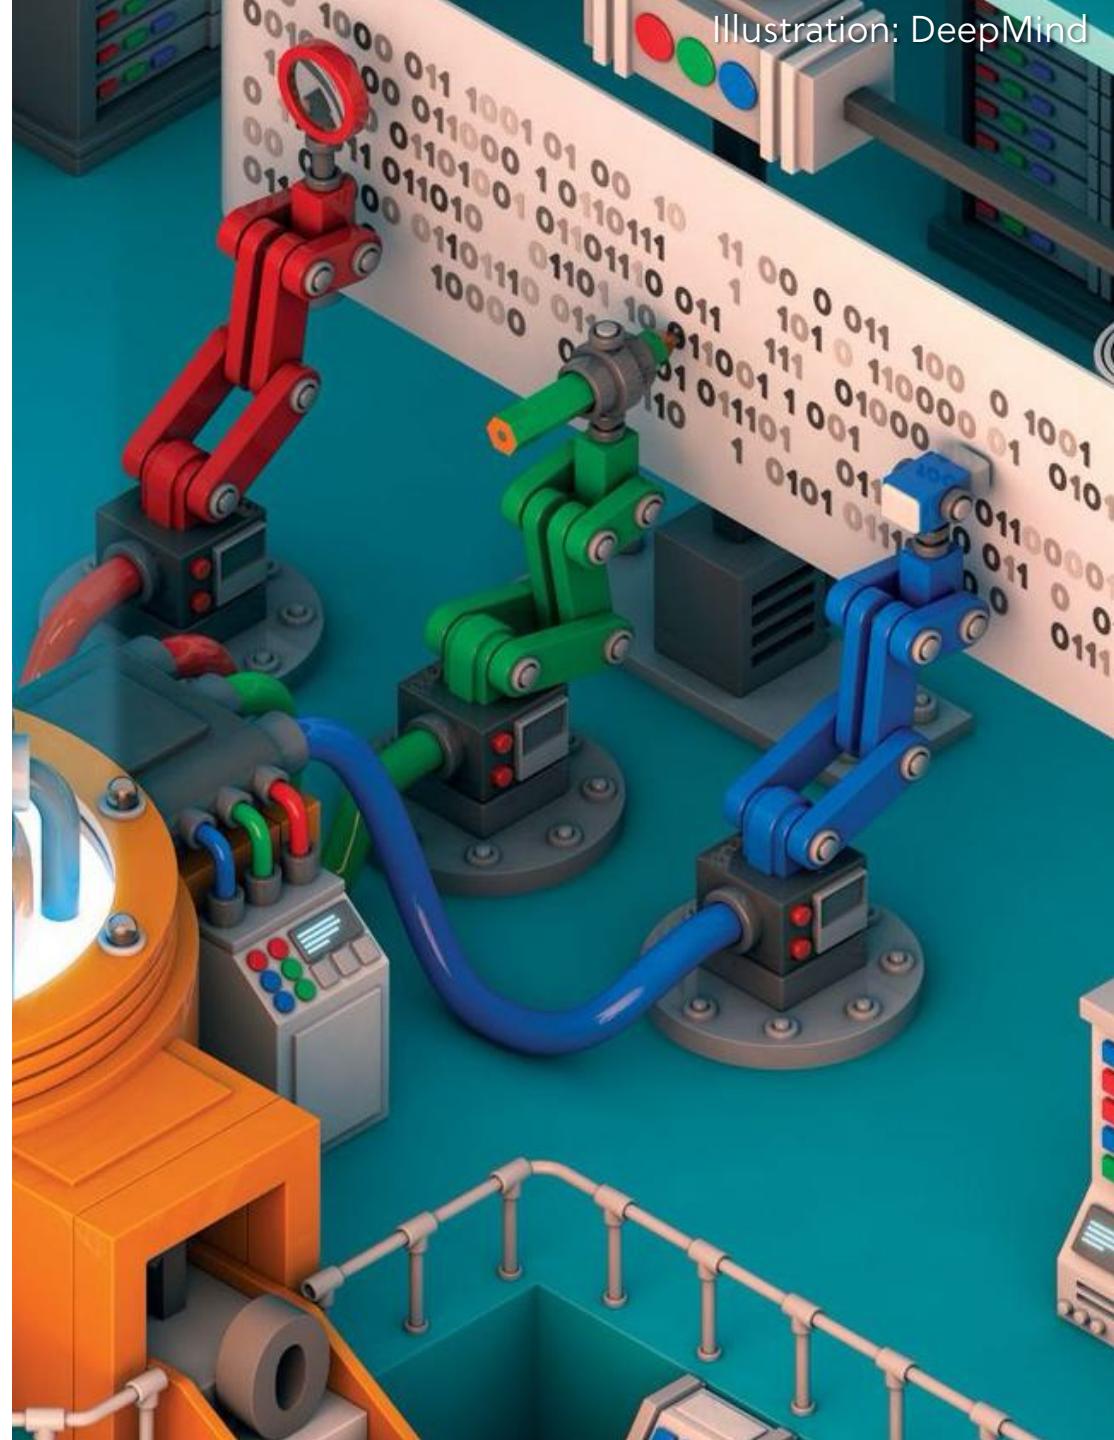


KOÇ
UNIVERSITY

Aykut Erdem // Koç University // Fall 2025

Previously on COMP541

- content-based attention
- location-based attention
- soft vs. hard attention
- case study: Show, Attend and Tell
- self-attention
- case study: Transformer networks



Lecture overview

- graph structured data
- graph neural nets (GNNs)
- GNNs for “classical” network problems
- **Disclaimer:** Much of the material and slides for this lecture were borrowed from
 - Yujia Li and Oriol Vinyals' tutorial on Graph Nets
 - Thomas Kipf's talk on structured deep models: deep Learning on graphs and beyond
 - Minji Yoon's CMU 10707 slides

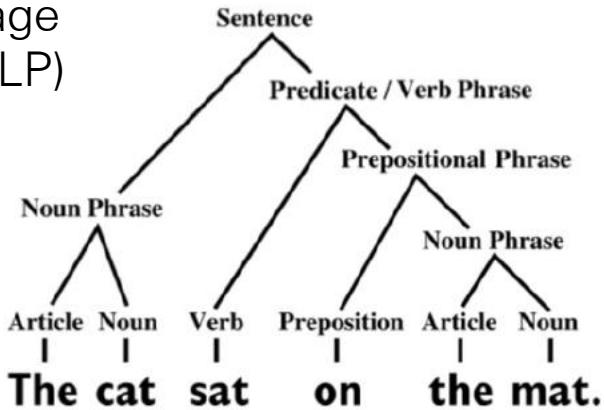
Deep Learning



Speech data

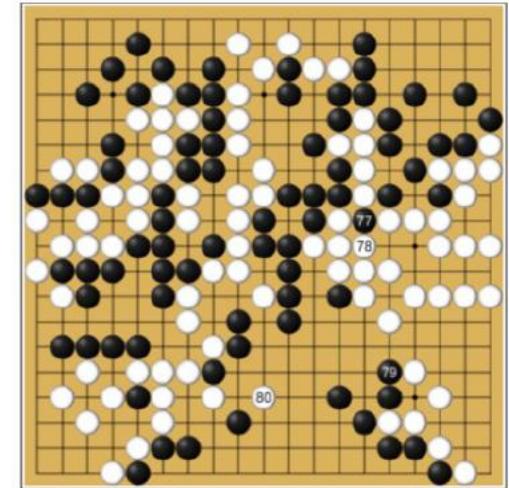


Natural language processing (NLP)



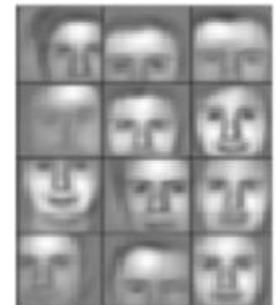
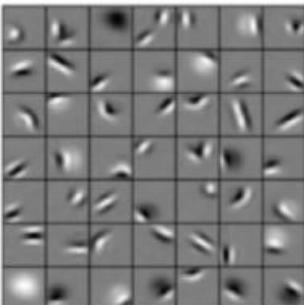
...

Grid games



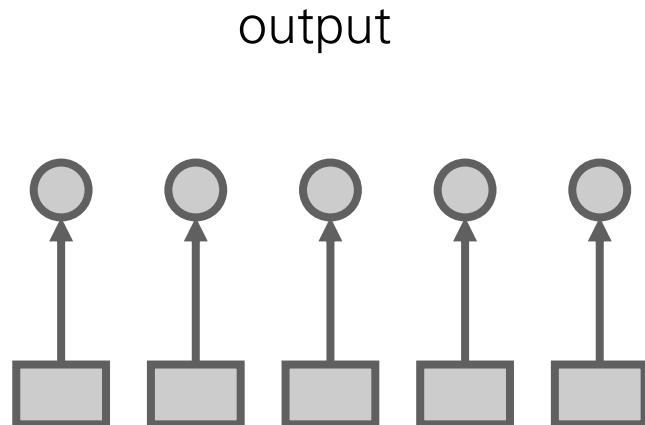
Deep neural nets that exploit:

- translation equivariance (weight sharing)
- hierarchical compositionality

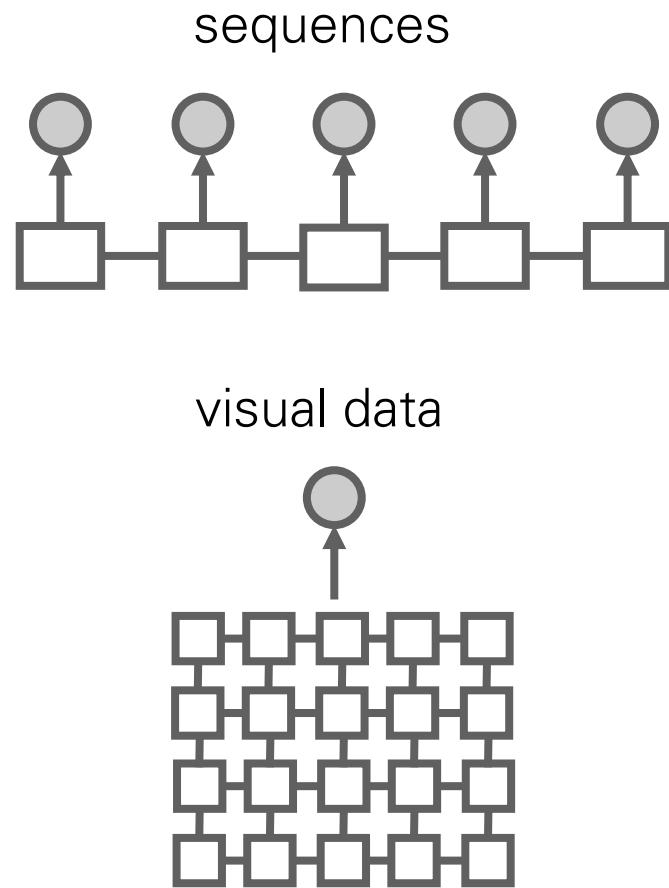


Modeling Structured Data

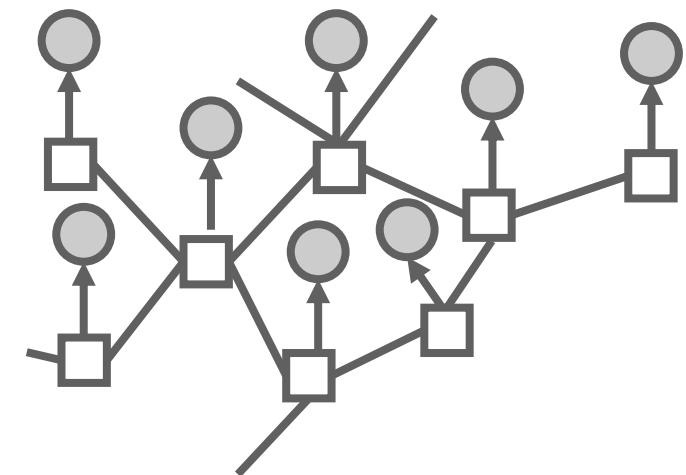
Unstructured Data



Data with Rigid Structure

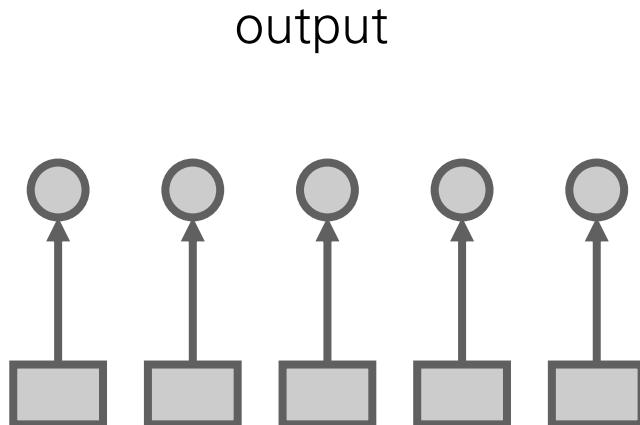


Graph Structured Data

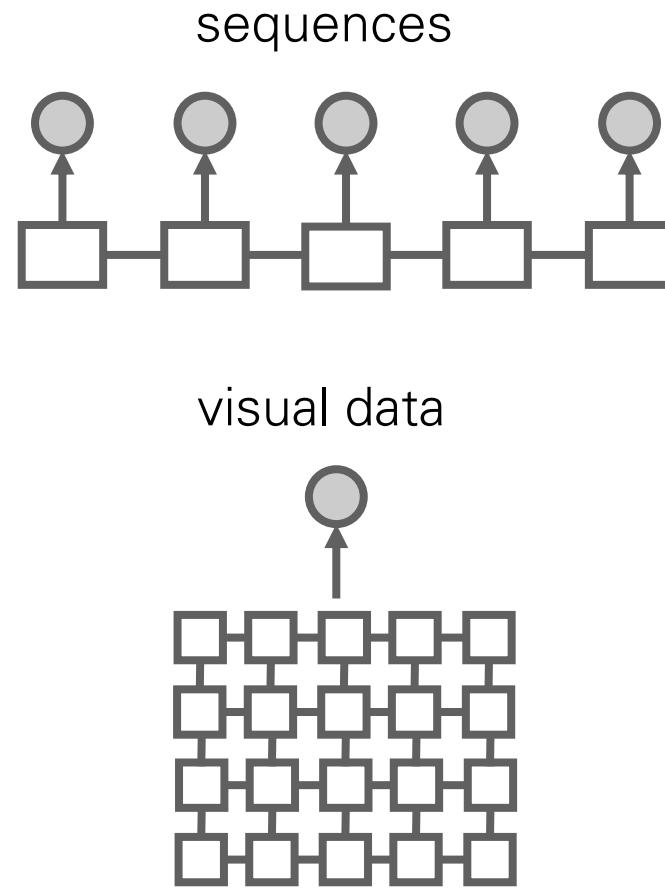


Modeling Structured Data

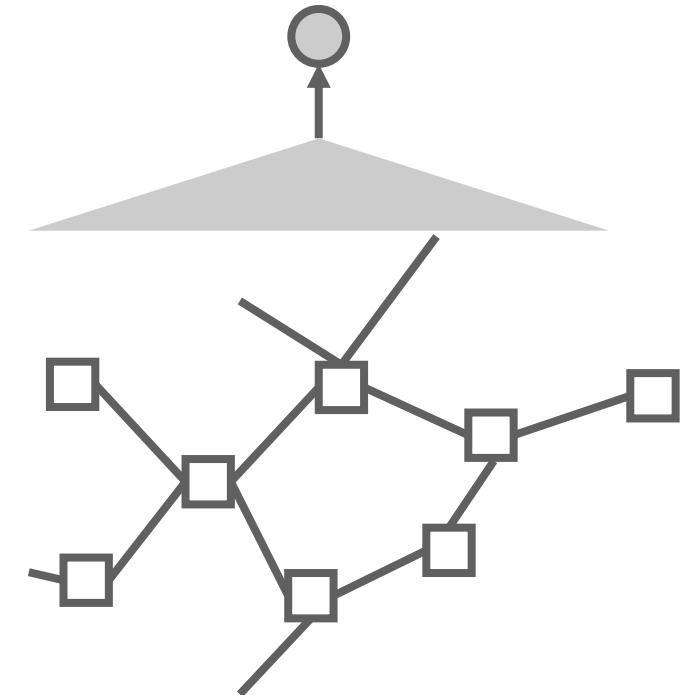
Unstructured Data



Data with Rigid Structure



Graph Structured Data

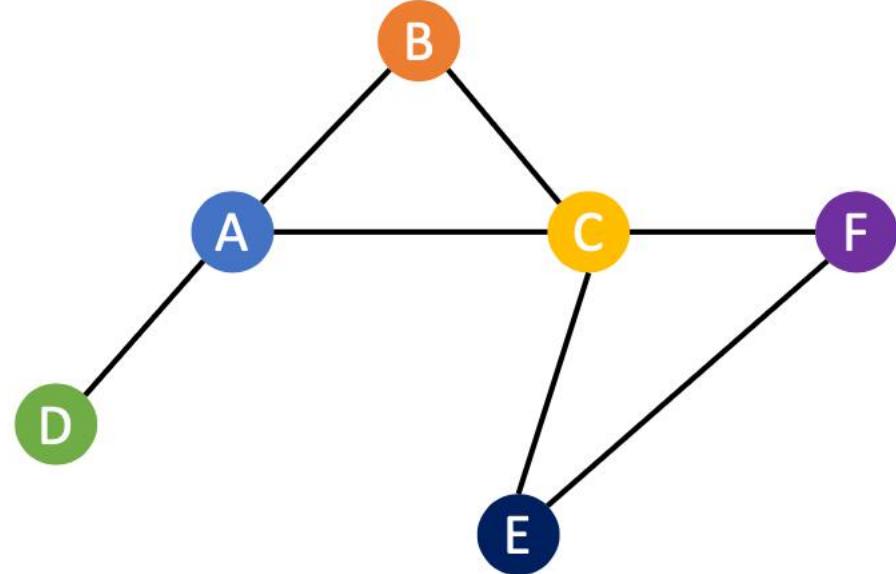


What is a graph

- A graph is composed of
 - Nodes (also called vertices)
 - Edges connecting a pair of nodes

presented in an **adjacency matrix**

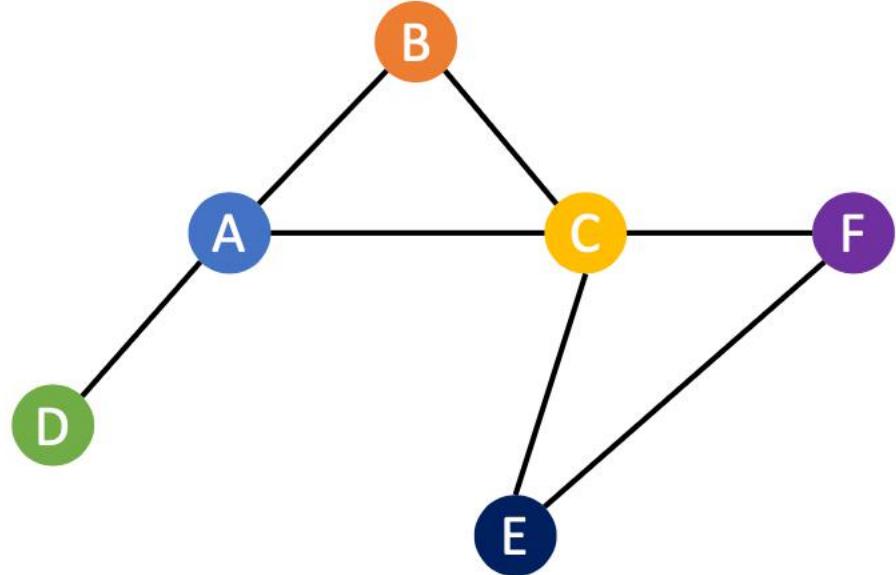
	A	B	C	D	E	F
A		1	1	1		
B	1		1			
C	1	1			1	1
D	1					
E		1				1
F		1		1		



What is a graph

- A graph is composed of
 - Nodes (also called vertices)
 - Edges connecting a pair of nodespresented in an **adjacency matrix**
- Nodes can have **feature vectors**

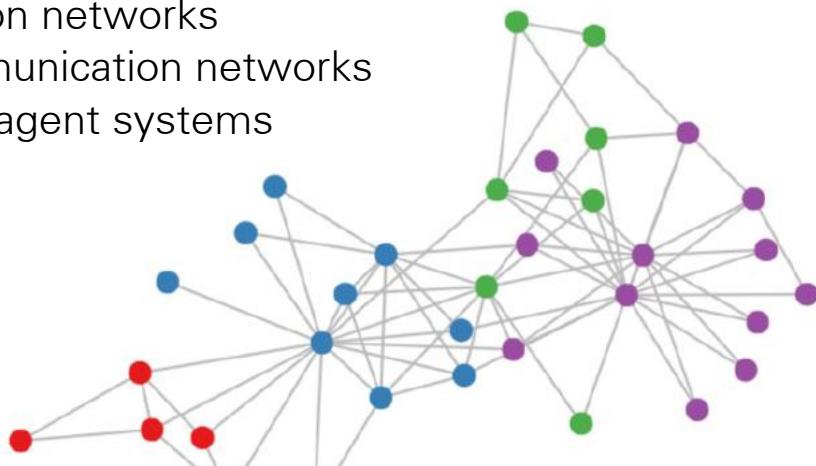
A	X_A
B	X_B
C	X_C
D	X_D
E	X_E
F	X_F



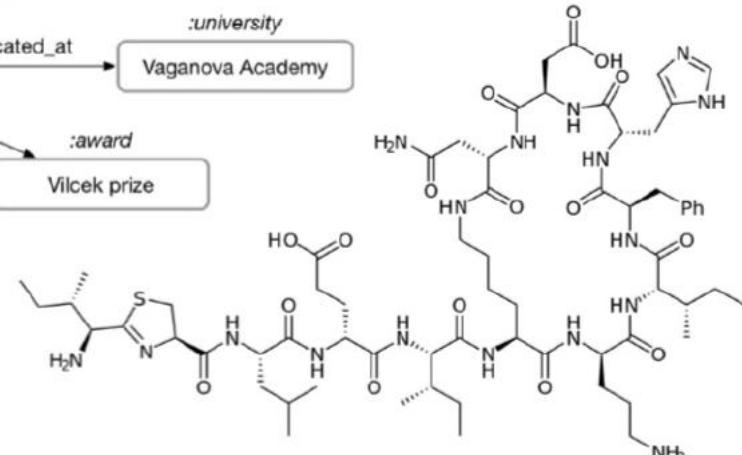
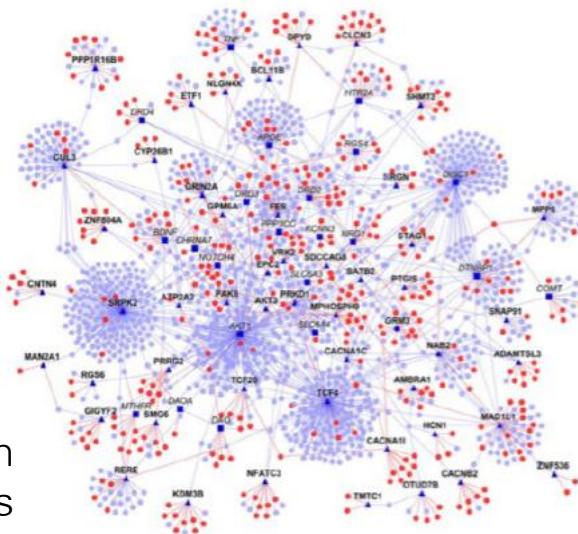
Graph structured data

- A lot of real-world data does not “live” on grids

- Social networks
- Citation networks
- Communication networks
- Multi-agent systems



Protein interaction networks



Molecules



Road maps

Standard deep learning architectures like CNNs and RNNs don't work here!

Graph Neural Networks have a large impact on...

DeepMind > Blog > Traffic prediction with advanced Graph Neural Networks



BLOG POST
RESEARCH 03 SEP 2020

Traffic prediction with advanced Graph Neural Networks

**Food Discovery with Uber Eats:
Using Graph Learning to Power
Recommendations**

Ankit Jain, Isaac Liu, Ankur Sarda, and Piero Molino

0

December 4, 2019



Pinterest Engineering

Aug 15, 2018 · 8 min read

PinSage: A new graph convolutional neural network for web-scale recommender systems

Ruining He | Pinterest engineer, Pinterest Labs

Web image search gets better with graph neural networks

A new approach to image search uses images returned by traditional search methods as nodes in a graph neural network through which similarity signals are passed in cross-modal retrieval.

amazon | science

PUBLICATION



FOLLOW

➤ P-Companion: A principled framework for diversified complementary product recommendation

By Junheng Hao, Tong Zhao, Jin Li, Xin Luna Dong, Christos Faloutsos, Yizhou Sun, Wei Wang

2020

Graph Neural Networks have a large impact on...

GCN-RL Circuit Designer: Transferable Transistor Sizing with Graph Neural Networks and Reinforcement Learning

Hanrui Wang¹, Kuan Wang¹, Jiacheng Yang¹, Linxiao Shen², Nan Sun², Hae-Seung Lee¹, Song Han¹

¹Massachusetts Institute of Technology

²UT Austin



The next big thing: the use of graph neural networks to discover particles

September 24, 2020 | Zack Savitsky



Machine learning algorithms can beat the world's hardest video games in minutes and solve complex equations faster than the collective efforts of generations of physicists. But the conventional algorithms still struggle to pick out stop signs on a busy street.

Object identification continues to hamper the field of machine learning — especially when the pictures are multidimensional and complicated, like the ones particle detectors take of collisions in high-energy physics experiments. However, a new class of neural networks is helping these models boost their pattern recognition abilities, and the technology may soon be implemented in particle physics experiments to optimize data analysis.

npj | computational materials

Explore content ▾ About the journal ▾ Publish with us ▾

nature > npj computational materials > articles > article

Article | Open Access | Published: 03 June 2021

Benchmarking graph neural networks for materials chemistry

Victor Fung , Jiaxin Zhang, Eric Juarez & Bobby G. Sumpter

npj Computational Materials 7, Article number: 84 (2021) | Cite this article

7807 Accesses | 7 Citations | 41 Altmetric | Metrics

nature

View all journals | Search Login

Explore content ▾ About the journal ▾ Publish with us ▾

nature > articles > article

Article | Published: 09 June 2021

A graph placement methodology for fast chip design

Azalia Mirhoseini , Anna Goldie , Mustafa Yazgan, Joe Wenjie Jiang, Ebrahim Songhori, Shen Wang, Young-Joon Lee, Eric Johnson, Omkar Pathak, Azade Naji, Jiwoo Pak, Andy Tong, Kavya Srinivasa, William Hang, Emre Tuncer, Quoc V. Le, James Laudon, Richard Ho, Roger Carpenter & Jeff Dean

Graph Neural Networks have a large impact on...

nature

Explore content ▾ About the journal ▾ Publish with us ▾

Subscribe

nature > news > article

NEWS | 01 December 2021

DeepMind's AI helps untangle the mathematics of knots

The machine-learning techniques could benefit other areas of maths that involve large data sets.

Patterns

Opinion

Neural algorithmic reasoning

Petar Veličković^{1,*} and Charles Blundell¹

¹DeepMind, London, Greater London, UK

*Correspondence: petarv@google.com

<https://doi.org/10.1016/j.patter.2021.100273>

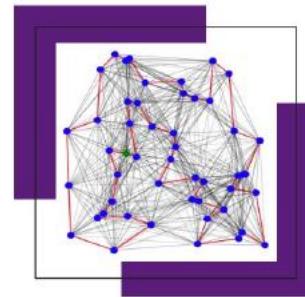
We present neural algorithmic reasoning—the art of building neural networks that are able to execute algorithmic computation—and provide our opinion on its transformative potential for running classical algorithms on inputs previously considered inaccessible to them.



institute for pure & applied mathematics

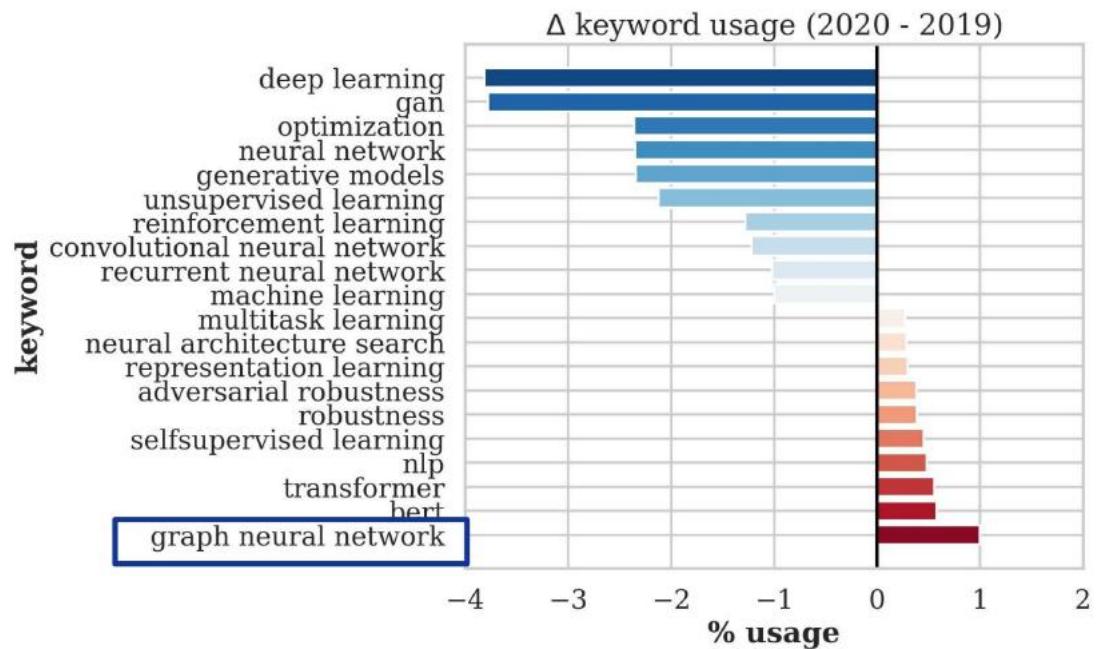
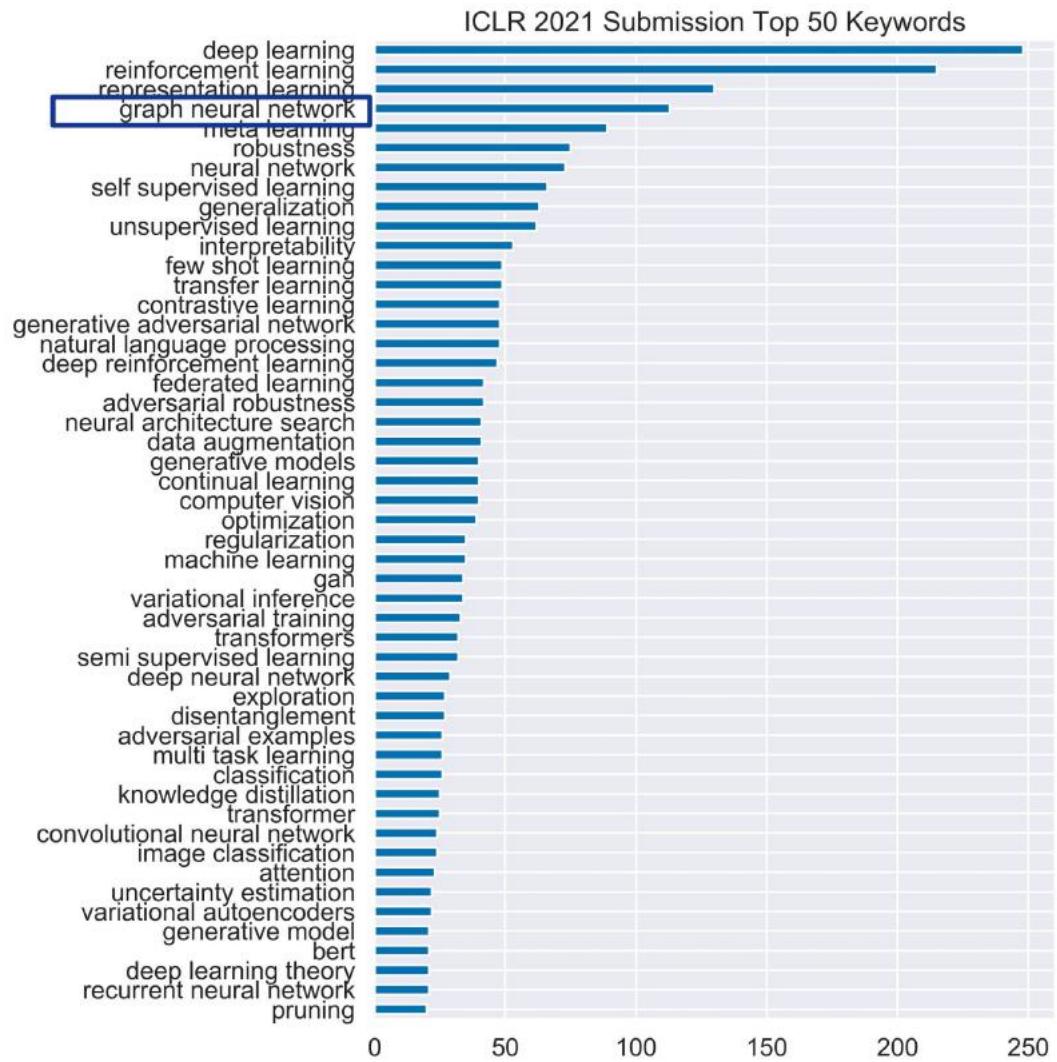
Deep Learning and Combinatorial Optimization

February 22 - 25, 2021



CellPress
OPEN ACCESS

A very hot research topic



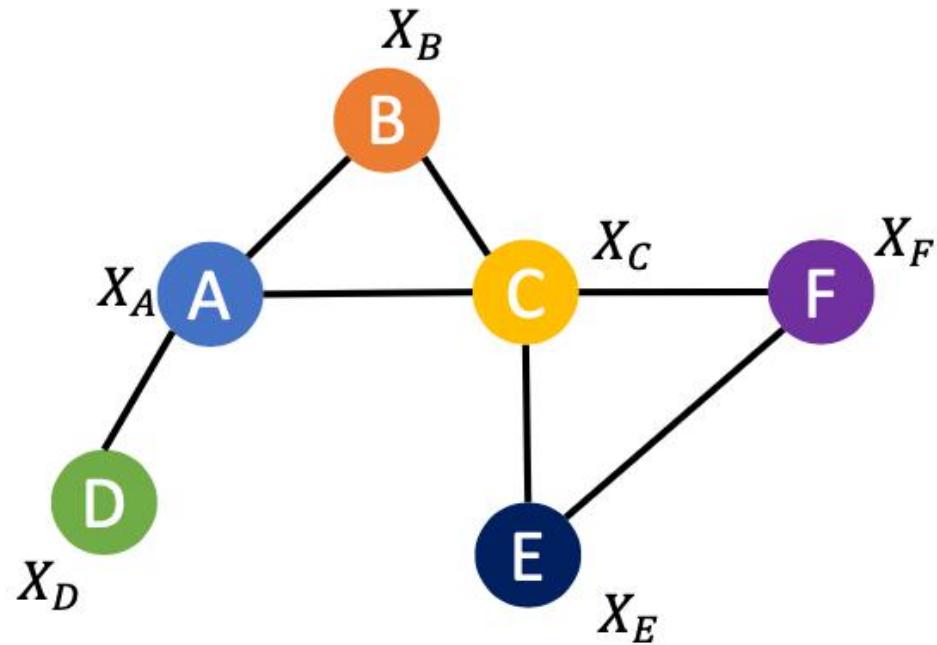
Recipe for a good model for graphs

- Handle different types of graph prediction problems
Requires: **Representations for graphs, nodes and edges**
- Handle graphs of varying sizes and structure
Requires: **A parametrization independent of graph size and structure**
- Handle arbitrary node ordering
Requires: **A model invariant to node permutations**
- Utilize graph structure
Requires: **A mechanism to communicate information on graphs**

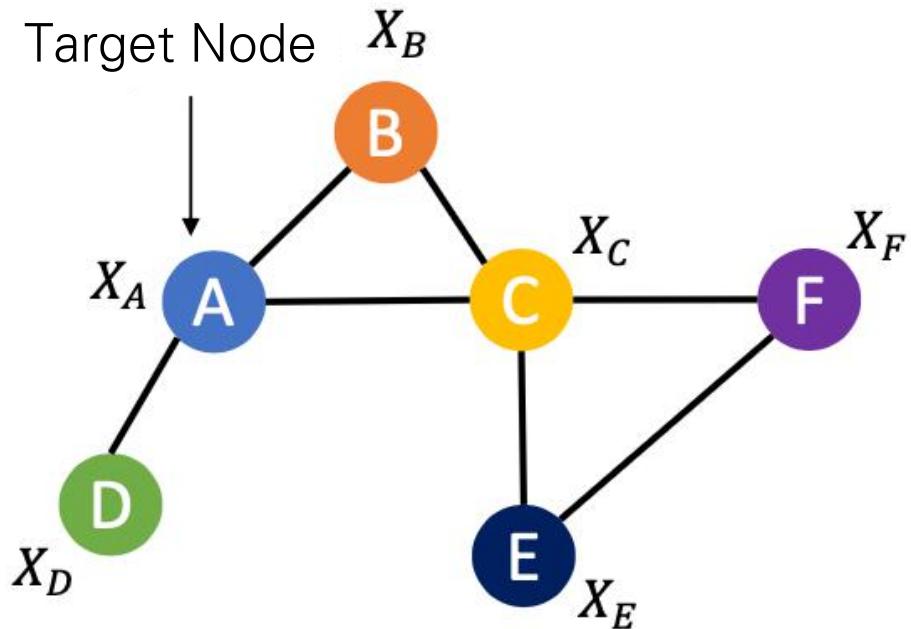
What is Graph Neural Network?

Problem definition

- Given
 - A graph
 - Node attributes
 - (part of nodes are labeled)
- Find
 - Node embeddings
- Predict
 - Labels for the remaining nodes

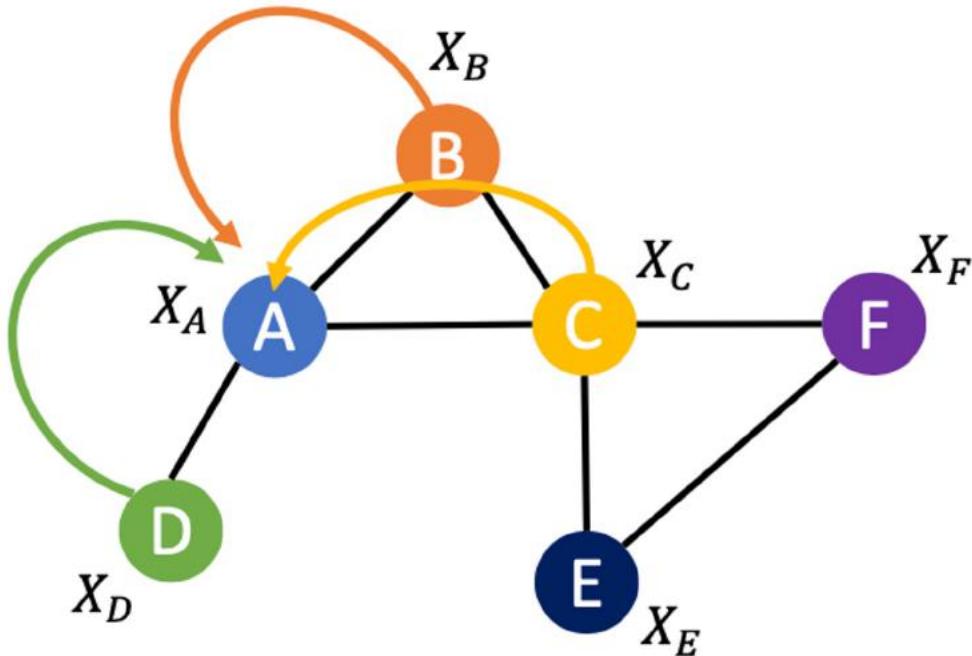


Graph Neural Networks (GNNs)



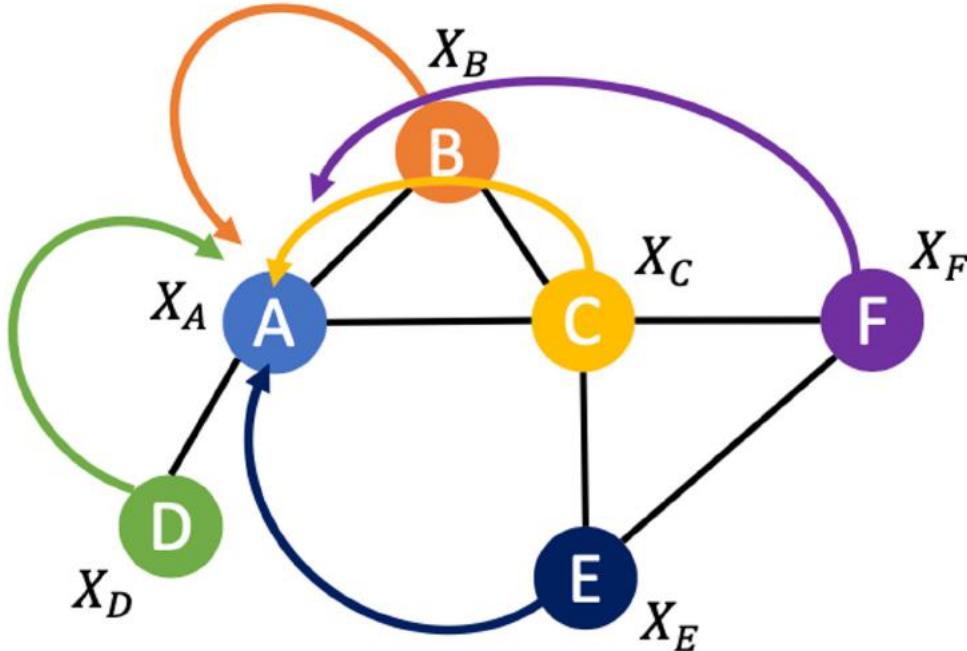
“Homophily: connected nodes are related/informative/similar”

Graph Neural Networks (GNNs)



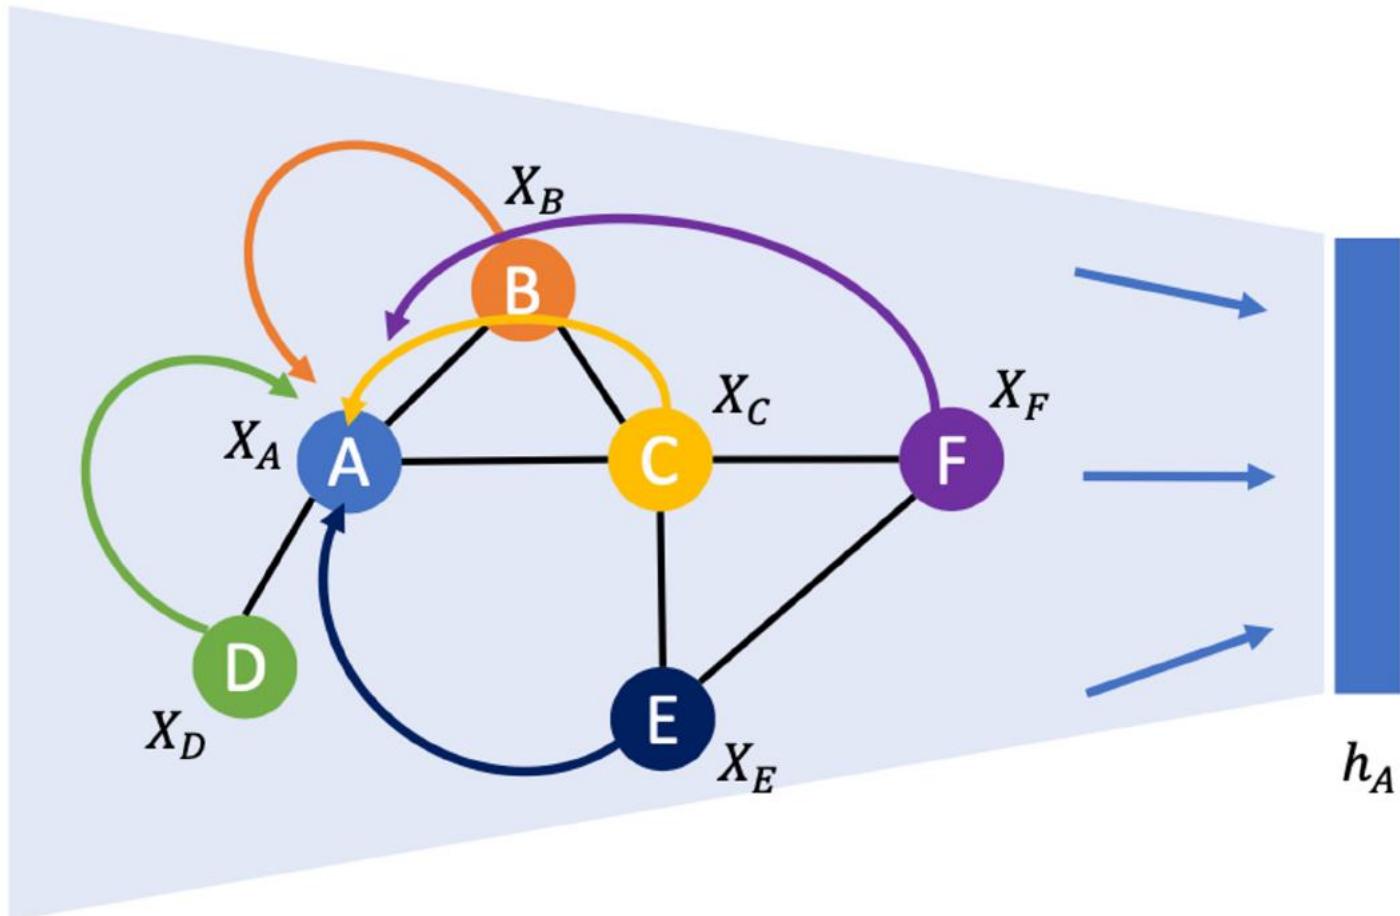
“Homophily: connected nodes are related/informative/similar”

Graph Neural Networks (GNNs)

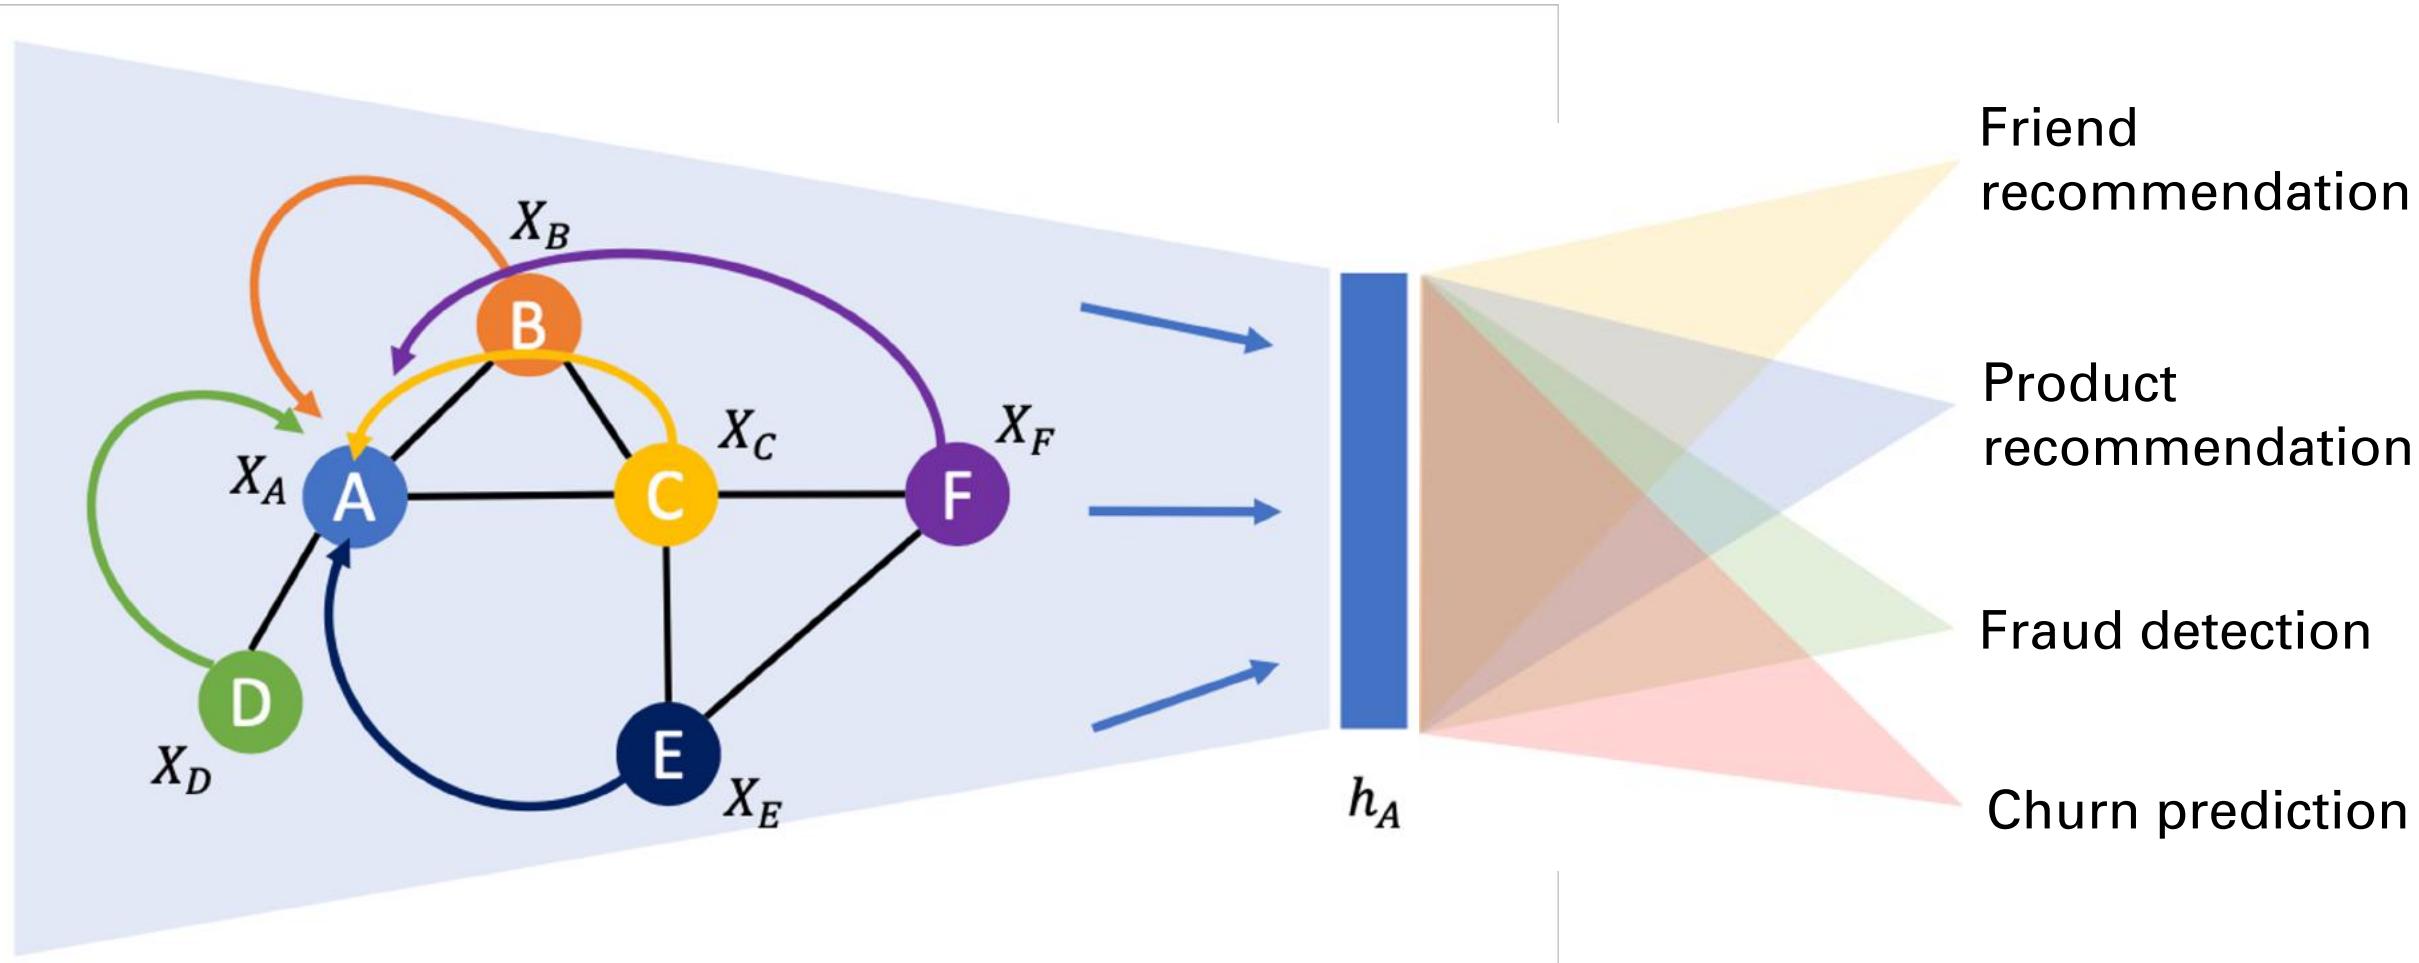


“Homophily: connected nodes are related/informative/similar”

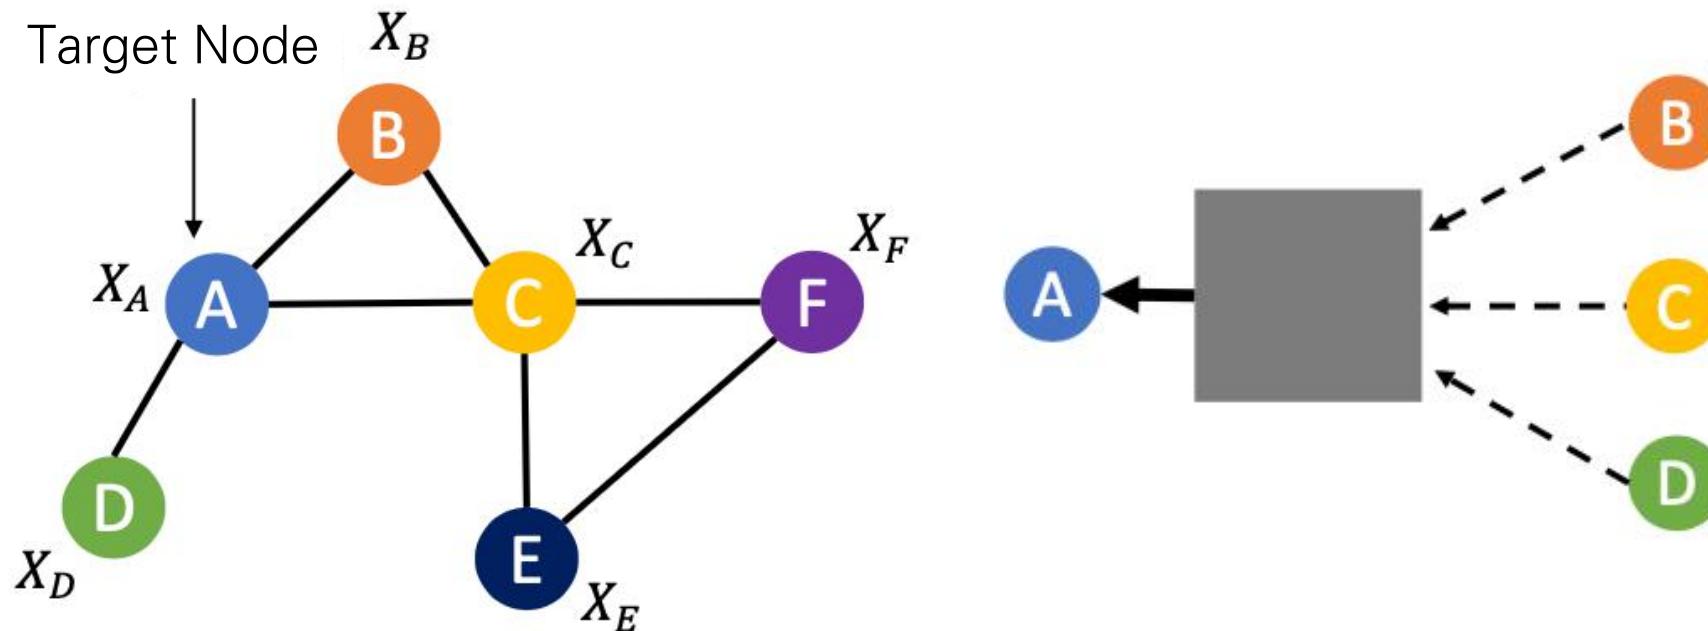
Graph Neural Networks (GNNs)



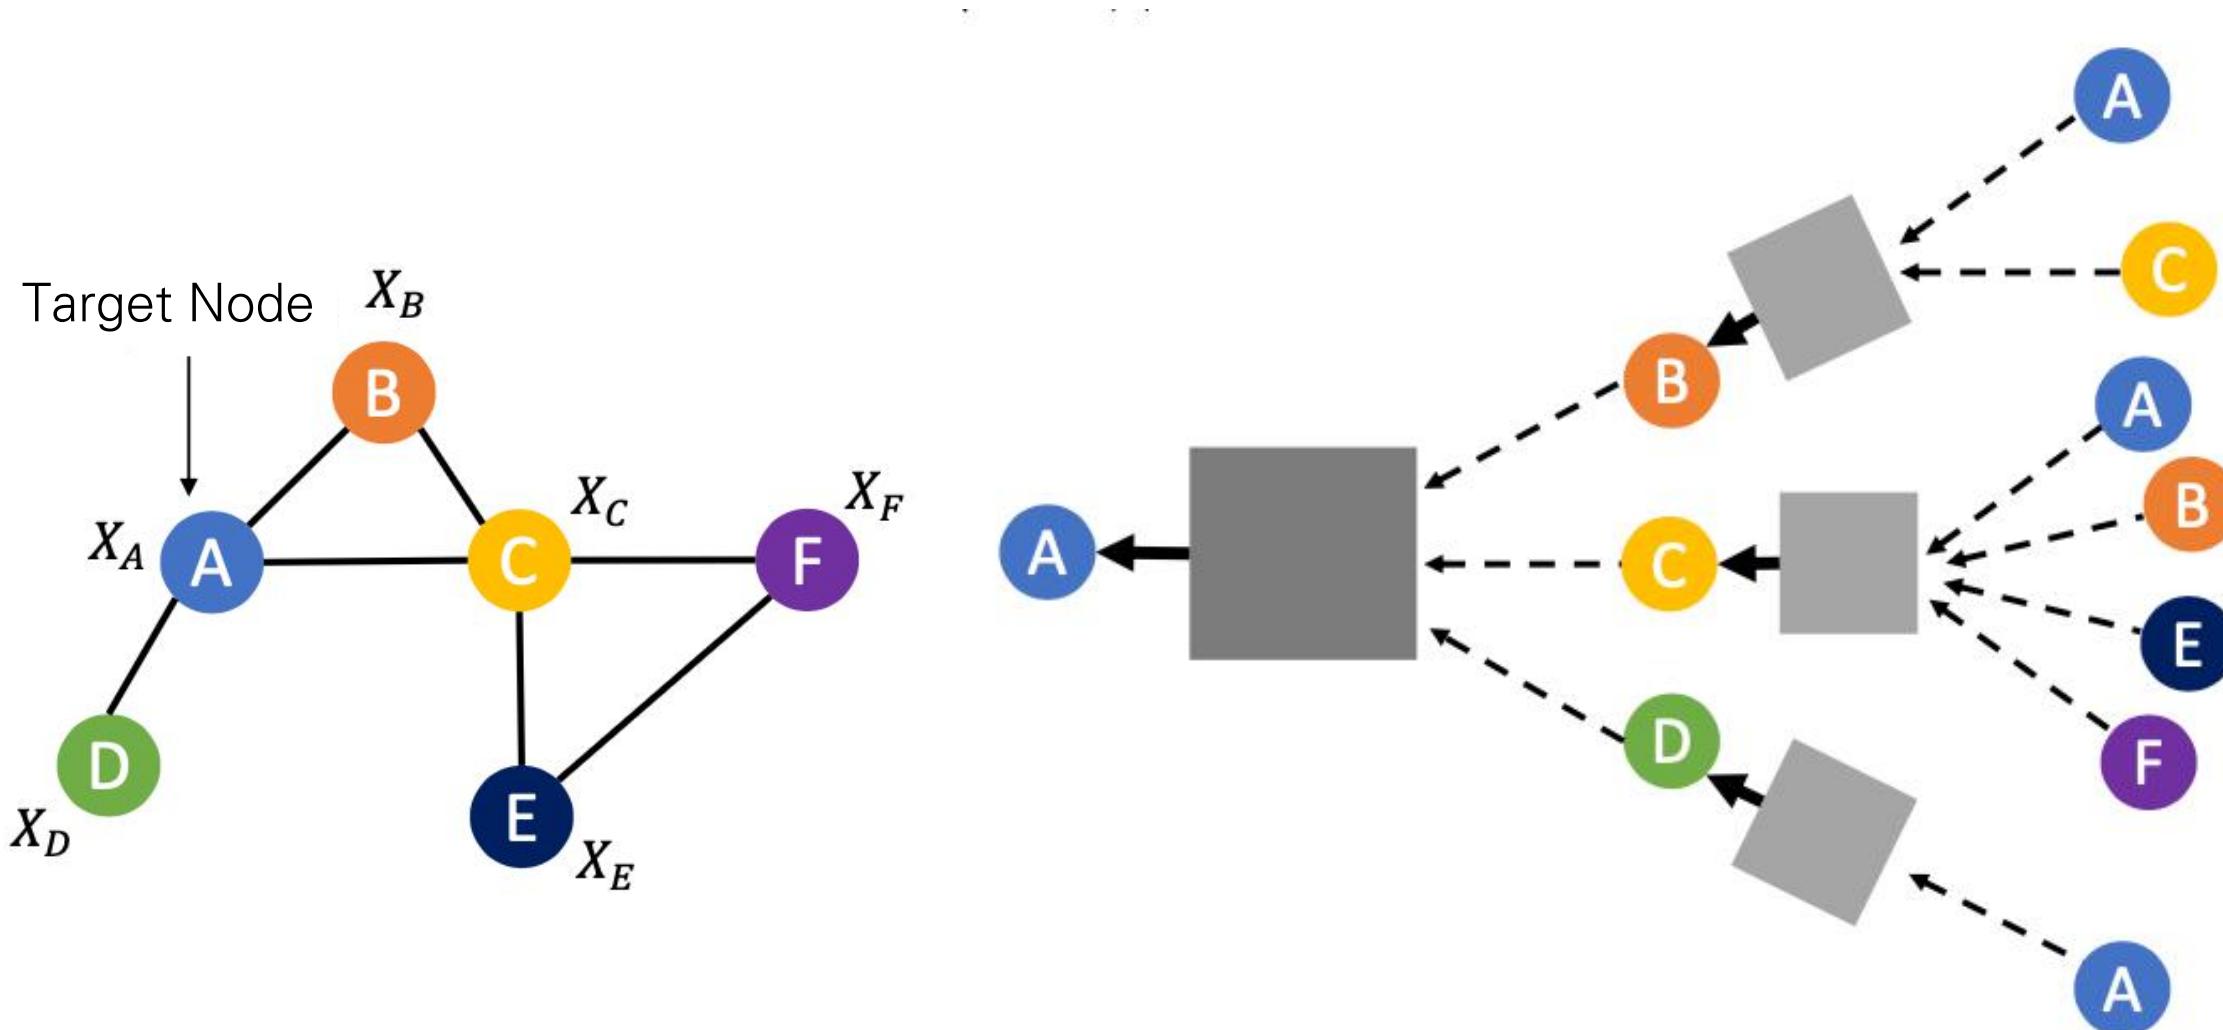
Graph Neural Networks (GNNs)



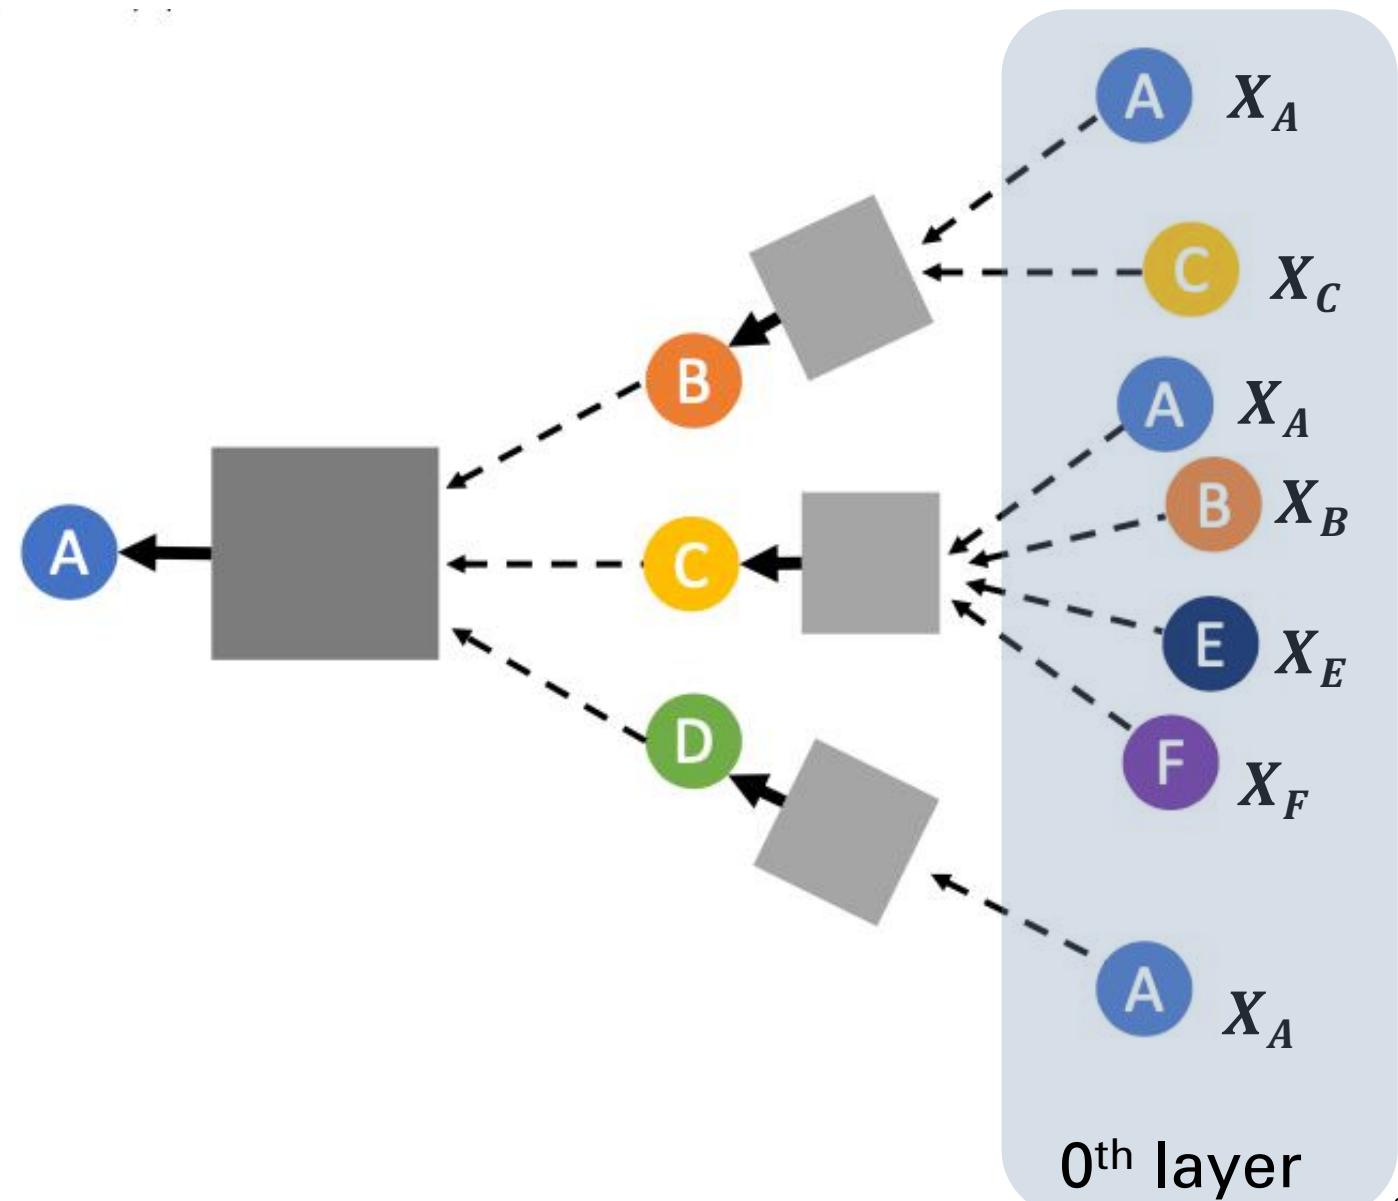
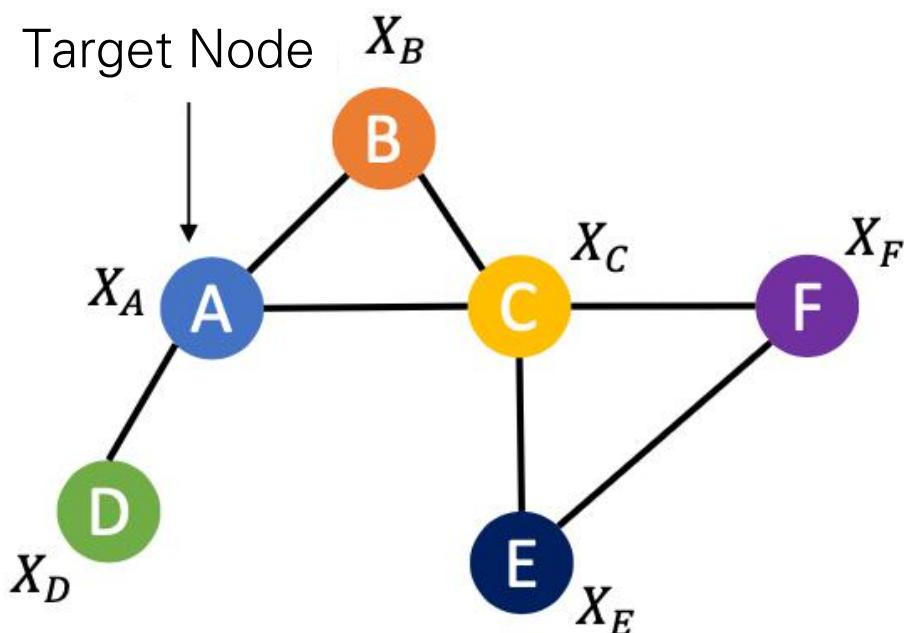
Graph Neural Networks (GNNs)



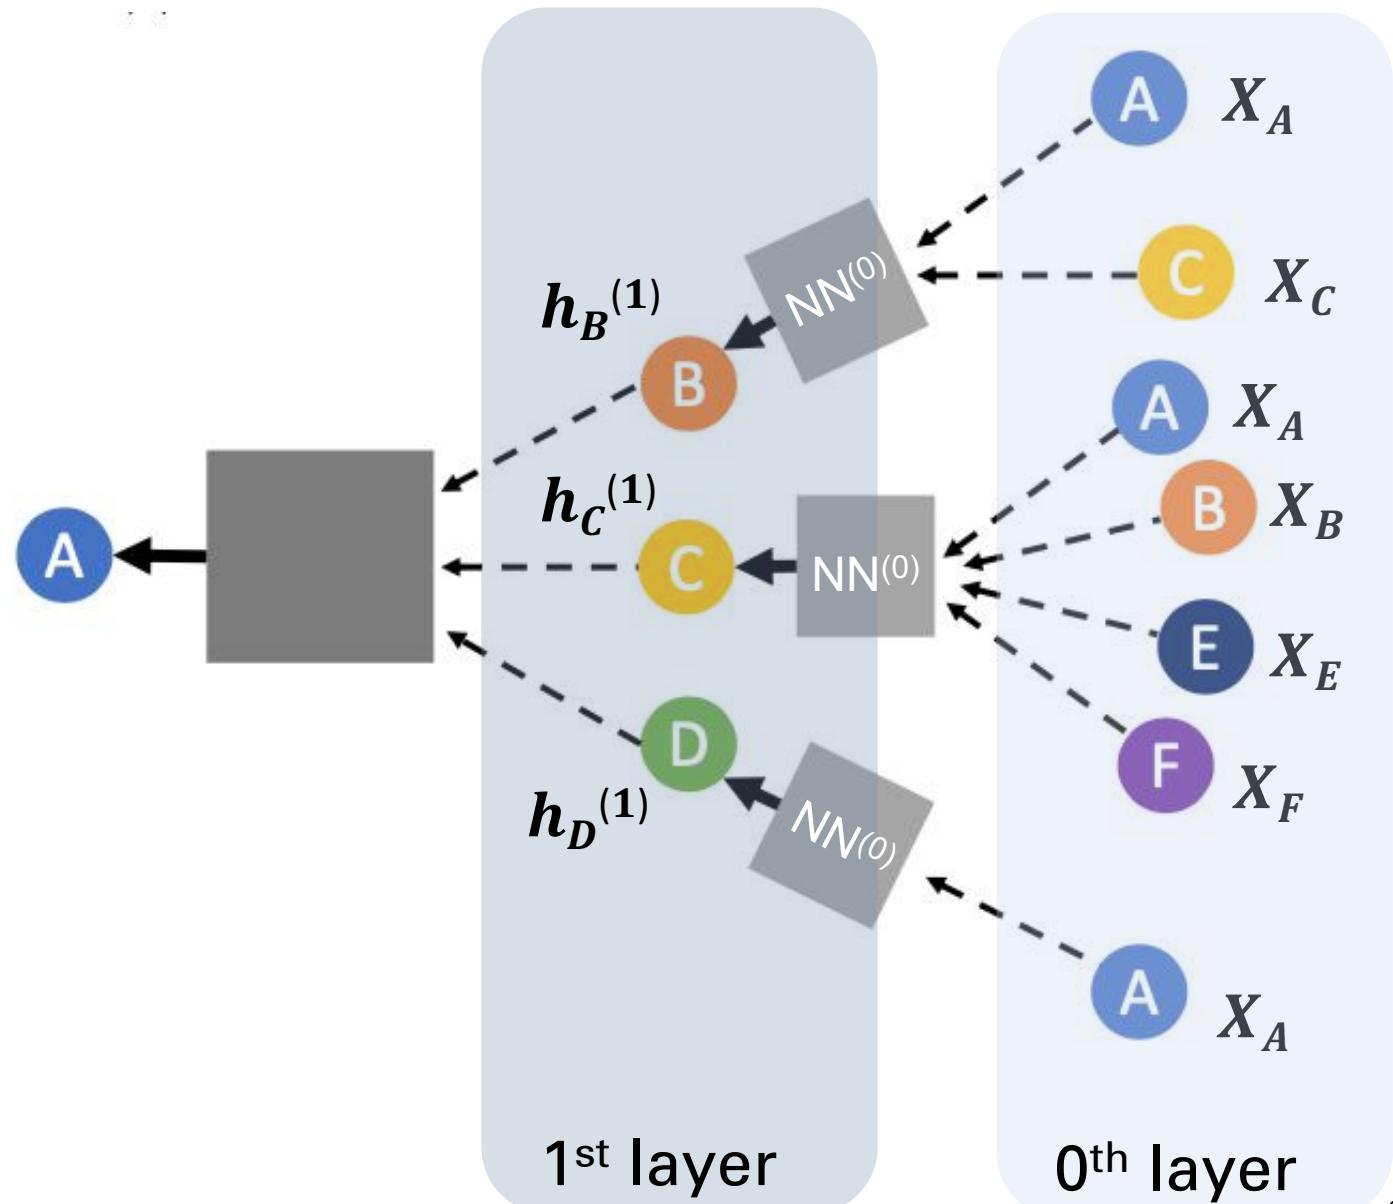
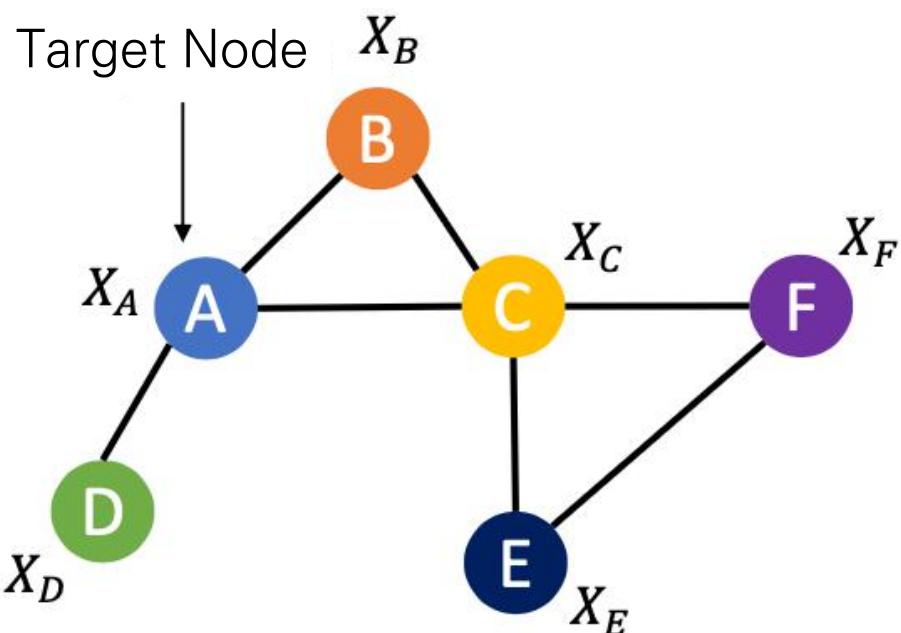
Graph Neural Networks (GNNs)



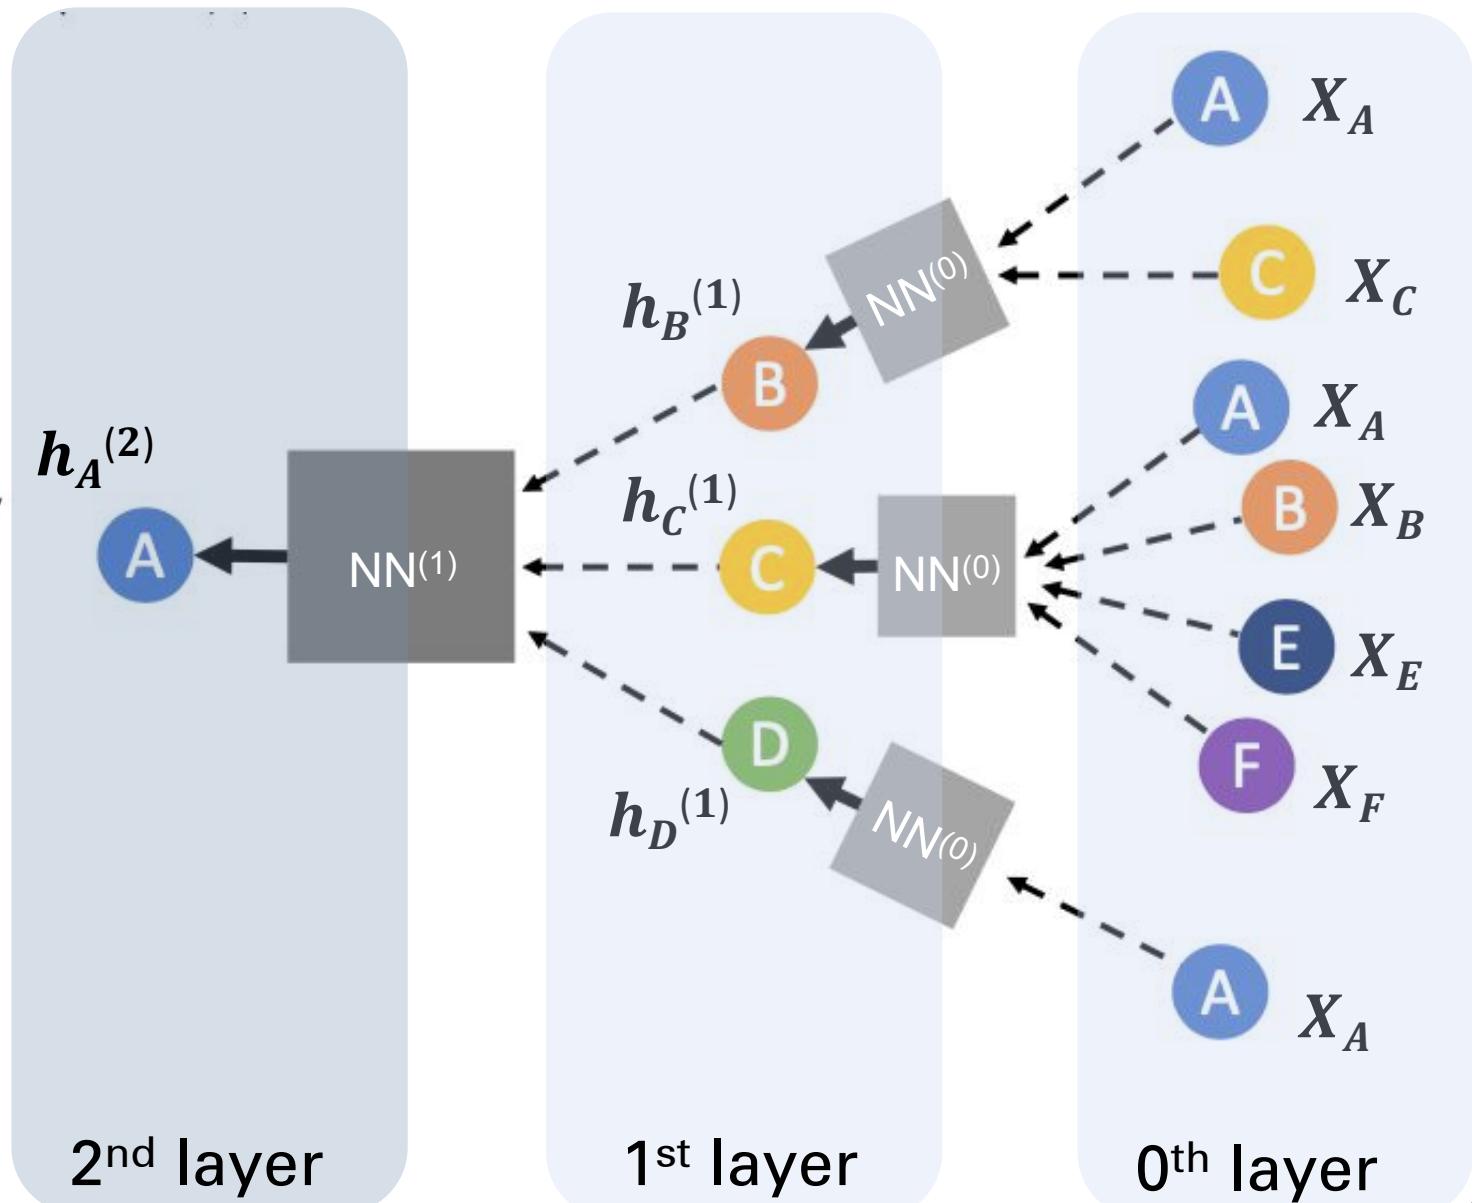
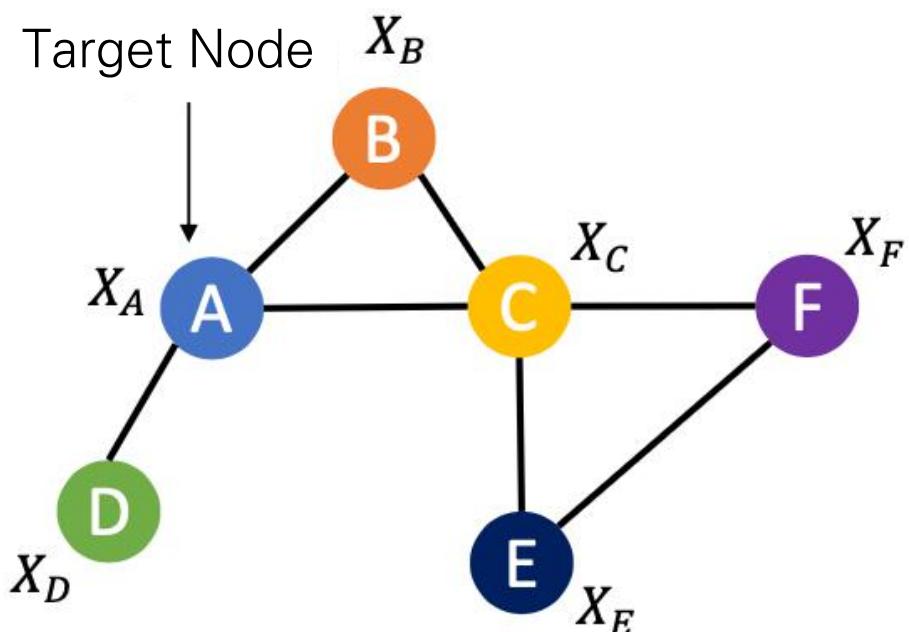
Graph Neural Networks (GNNs)



Graph Neural Networks (GNNs)



Graph Neural Networks (GNNs)



Graph Neural Networks (GNNs)

1. Aggregate messages from neighbors

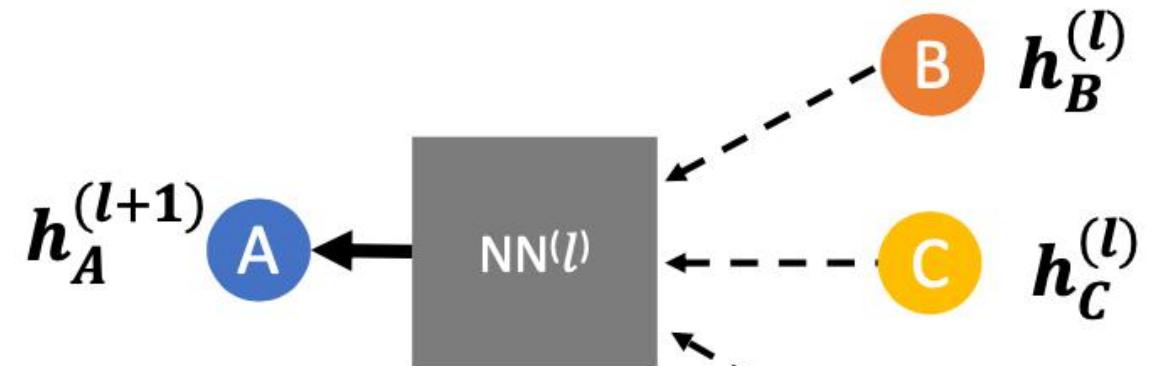
$h_v^{(l)}$: node embedding of v at l -th layer

$\mathcal{N}(v)$: neighboring nodes of v

$f^{(l)}$: aggregation function at l -th layer

$m_v^{(l)}$: message vector of v at l -th layer

$$\begin{aligned}m_A^{(l)} &= f^{(l)} \left(h_A^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(A)\} \right) \\&= f^{(l)} \left(h_A^{(l)}, h_B^{(l)}, h_C^{(l)}, h_D^{(l)} \right)\end{aligned}$$



Neighbors of node A

$$\mathcal{N}(A) = \{B, C, D\}$$

Graph Neural Networks (GNNs)

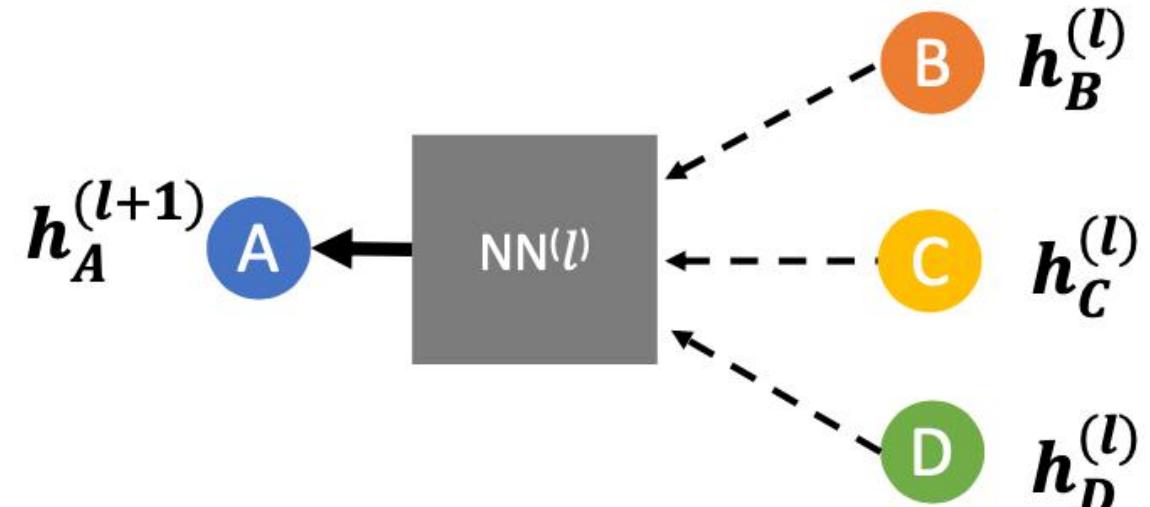
1. Aggregate messages from neighbors

$$\begin{aligned} m_A^{(l)} &= \mathbf{f}^{(l)} \left(h_A^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(A)\} \right) \\ &= \mathbf{f}^{(l)} \left(h_A^{(l)}, h_B^{(l)} h_C^{(l)} h_D^{(l)} \right) \end{aligned}$$

2. Transform messages

$\mathbf{g}^{(l)}$: transformation function at l -th layer

$$h_A^{(l+1)} = \mathbf{g}^{(l)}(m_A^{(l)})$$



Neighbors of node A

$$\mathcal{N}(A) = \{B, C, D\}$$

Graph Neural Networks (GNNs)

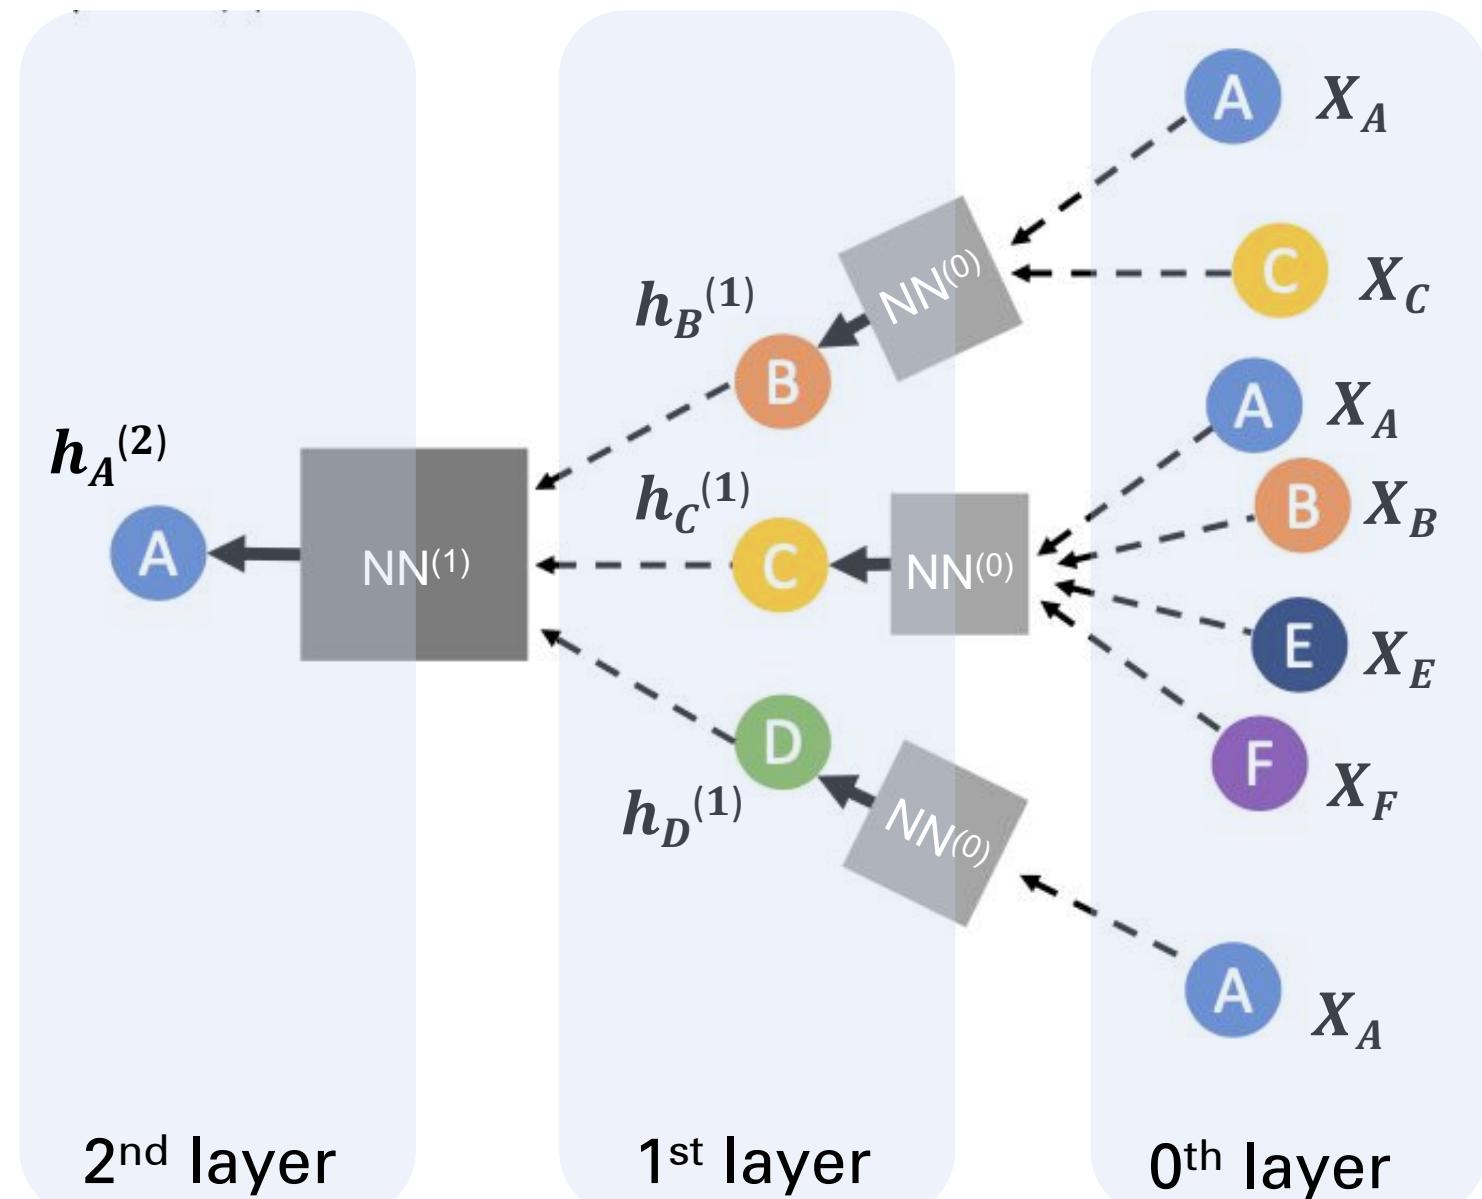
In each layer l ,
for each target node v :

1. Aggregate messages

$$m_v^{(l)} = \mathbf{f}^{(l)} \left(h_v^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(v)\} \right)$$

2. Transform messages

$$h_v^{(l+1)} = \mathbf{g}^{(l)}(m_v^{(l)})$$



Graph Neural Networks (GNNs)

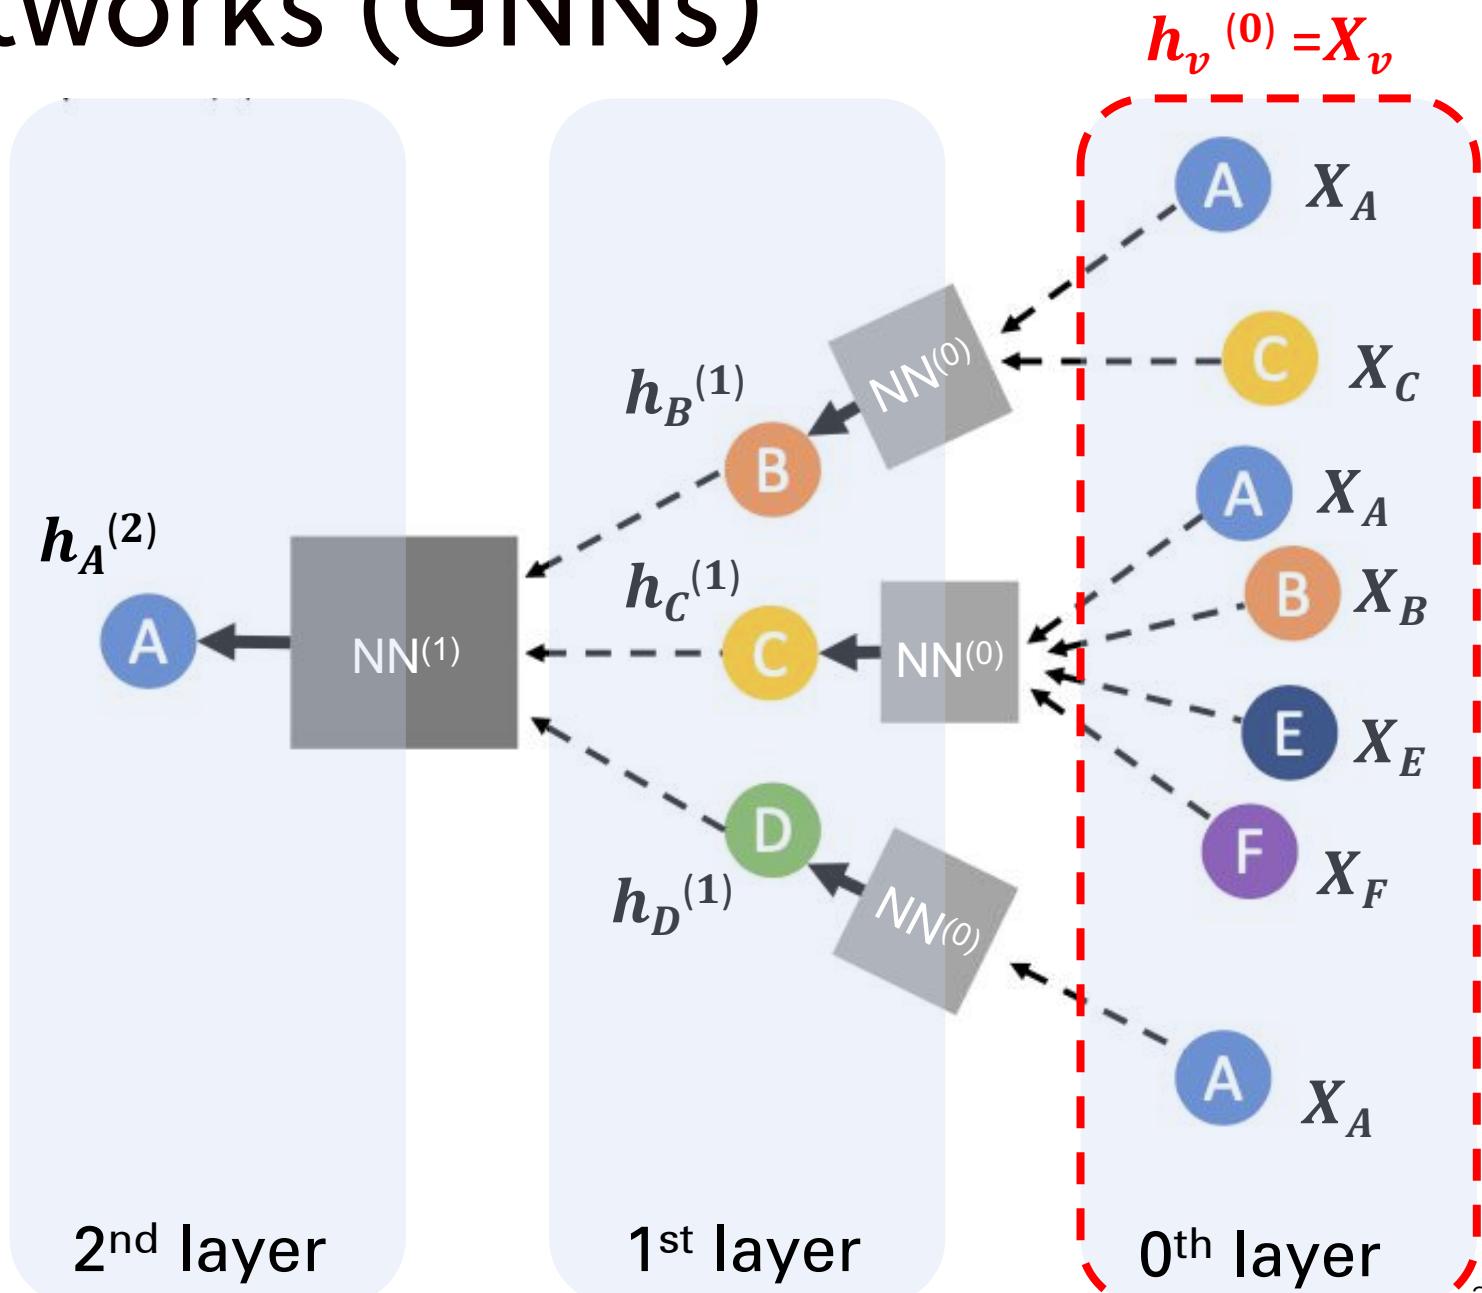
In each layer l ,
for each target node v :

1. Aggregate messages

$$m_v^{(l)} = f^{(l)} \left(h_v^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(v)\} \right)$$

2. Transform messages

$$h_v^{(l+1)} = g^{(l)}(m_v^{(l)})$$



Graph Neural Networks (GNNs)

In each layer l ,
for each target node v :

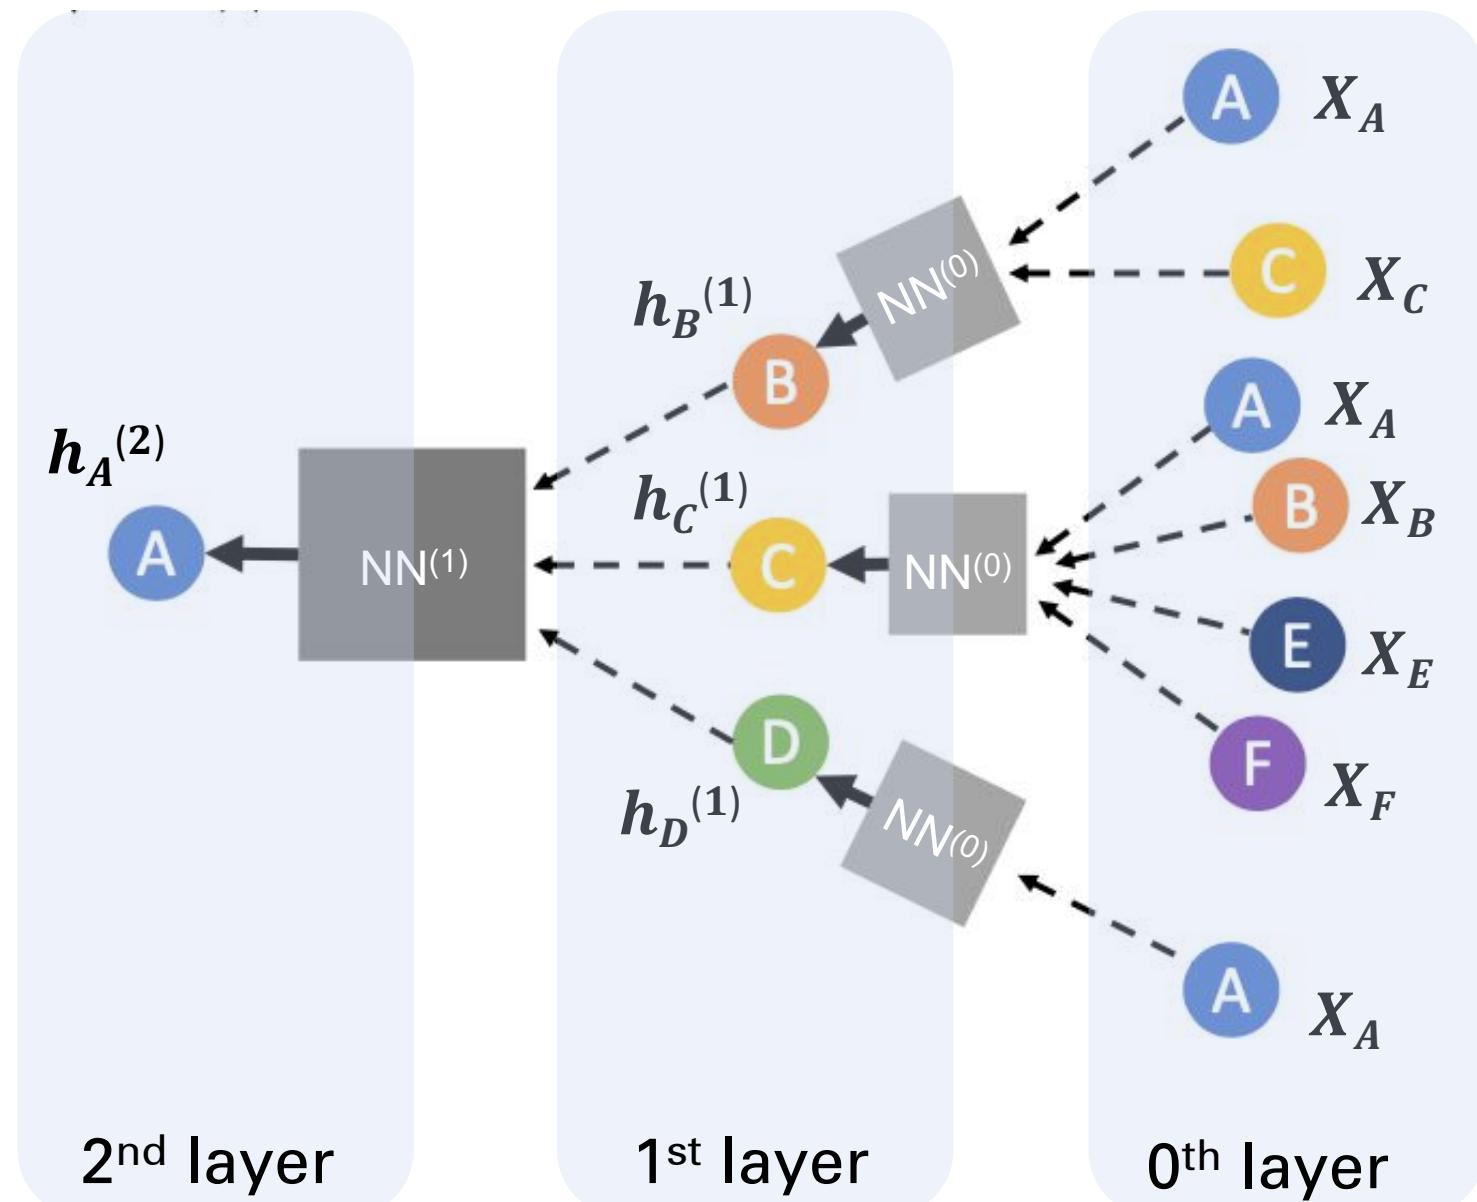
1. Aggregate messages

$$m_v^{(l)} = \boxed{f^{(l)}}\left(h_v^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(v)\}\right)$$

2. Transform messages

$$h_v^{(l+1)} = \boxed{g^{(l)}}(m_v^{(l)})$$

GNN models mostly differ
in how these functions are
defined..



Graph Neural Networks (GNNs)

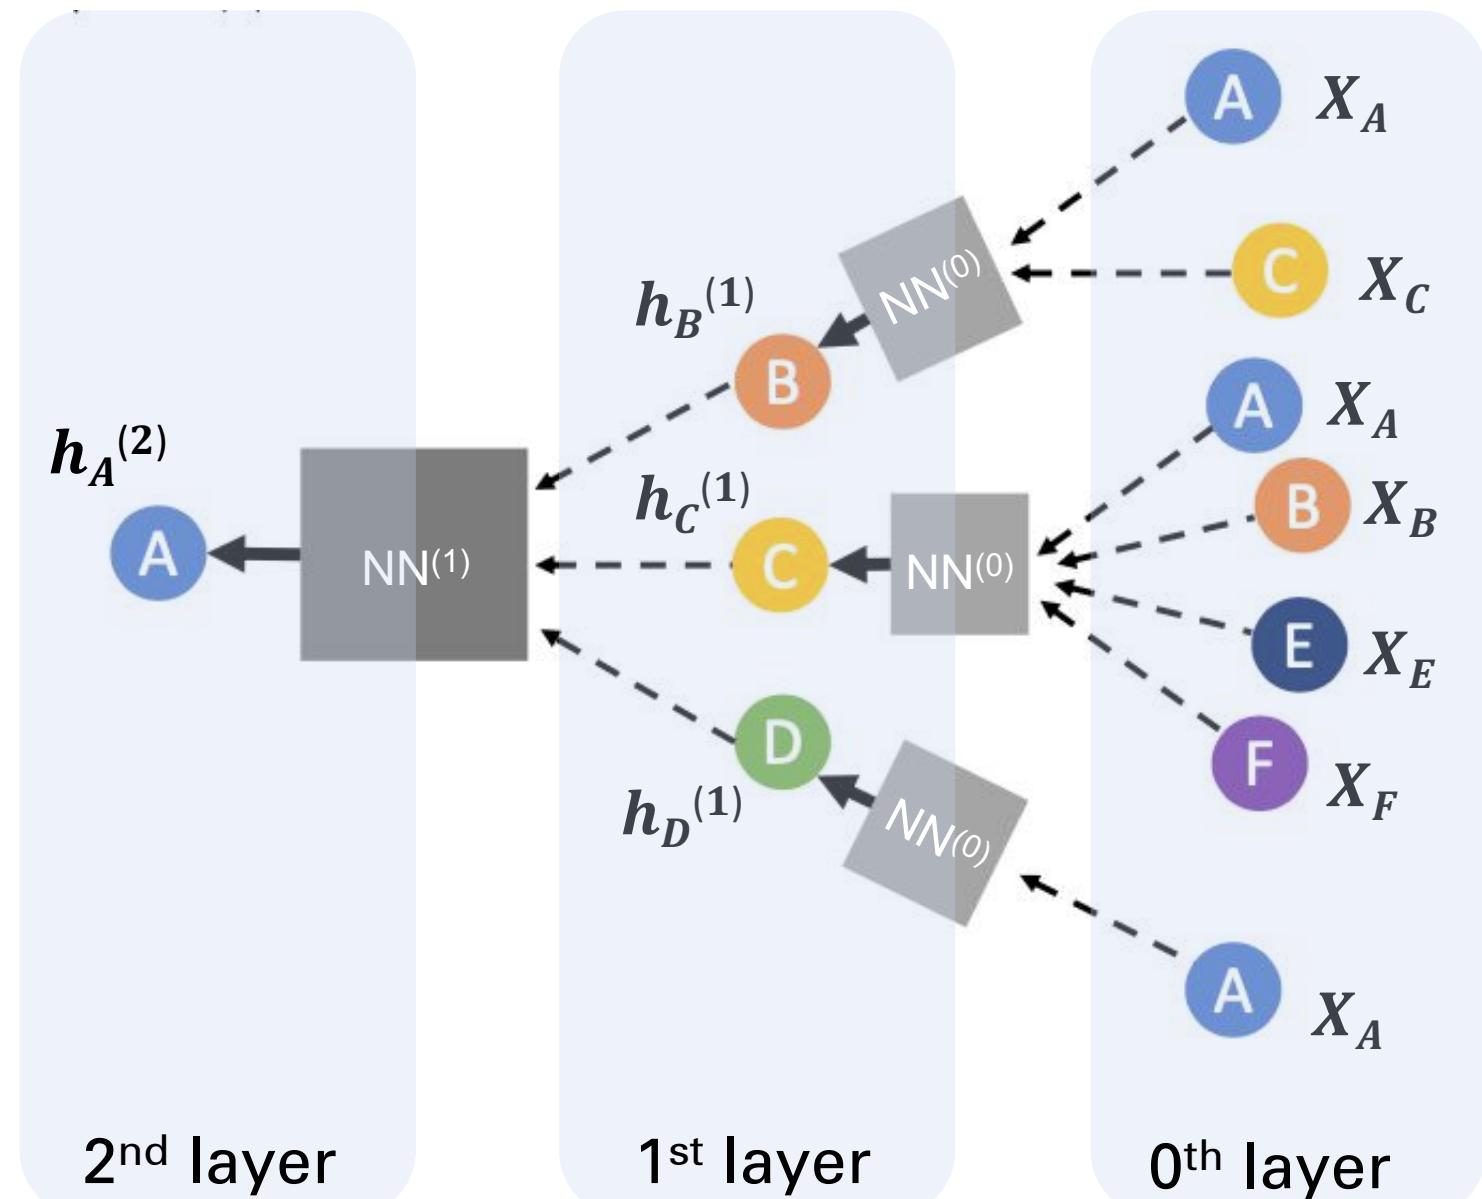
In each layer l ,
for each target node v :

1. Aggregate messages

$$m_v^{(l)} = \mathbf{f}^{(l)} \left(h_v^{(l)}, \{h_u^{(l)} : u \in \mathcal{N}(v)\} \right)$$

2. Transform messages

$$h_v^{(l+1)} = \mathbf{g}^{(l)}(m_v^{(l)})$$



Graph Neural Networks (GNNs)

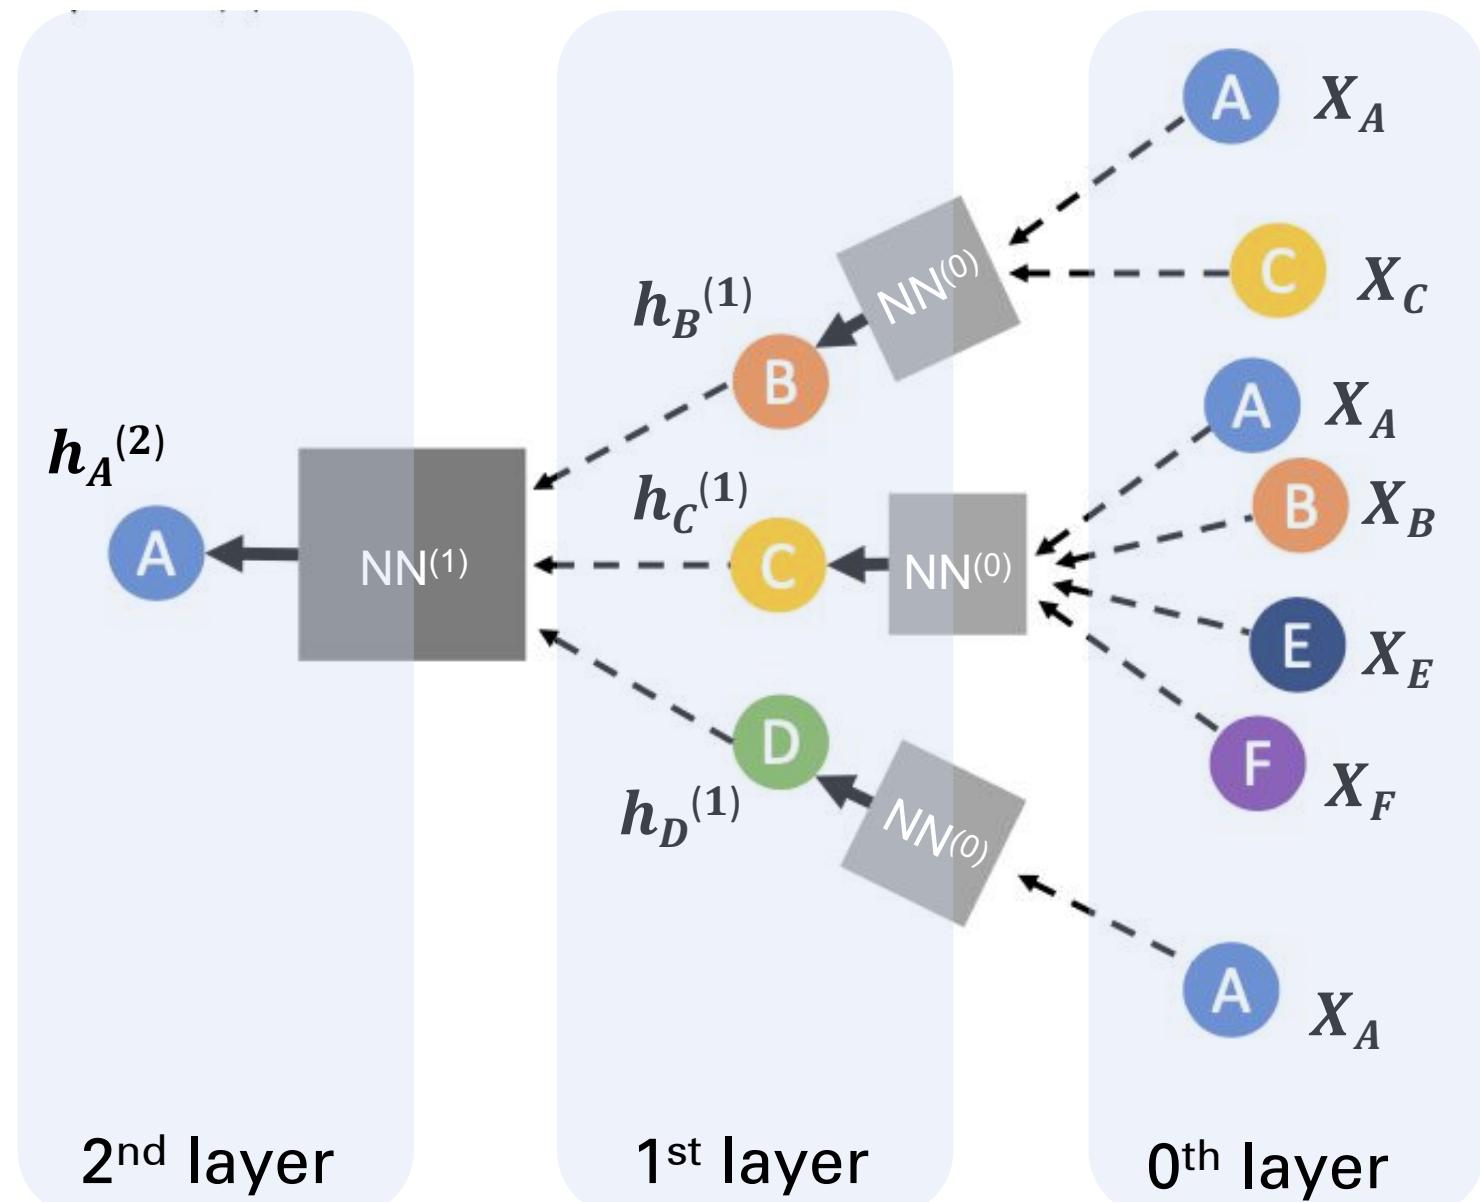
Graph Convolutional Networks^[1]

1. Aggregate messages

$$m_v^{(l)} = \frac{1}{|\mathcal{N}(v) + 1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l)}$$

2. Transform messages

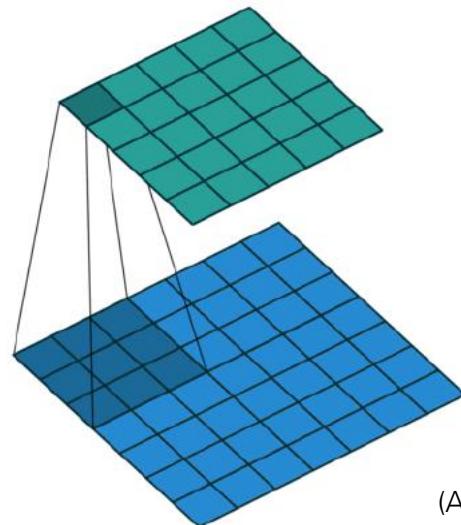
$$h_v^{(l+1)} = \sigma(\mathbf{W}^{(l)} \circ m_v^{(l)})$$



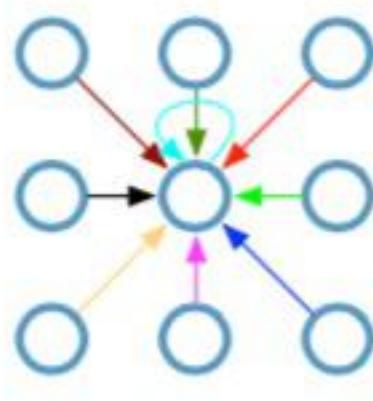
[1] Kipf, Thomas N., et al. "Semi-supervised classification with graph convolutional networks."

Recap: Convolutional neural networks (on grids)

Single CNN layer
with 3x3 filter:

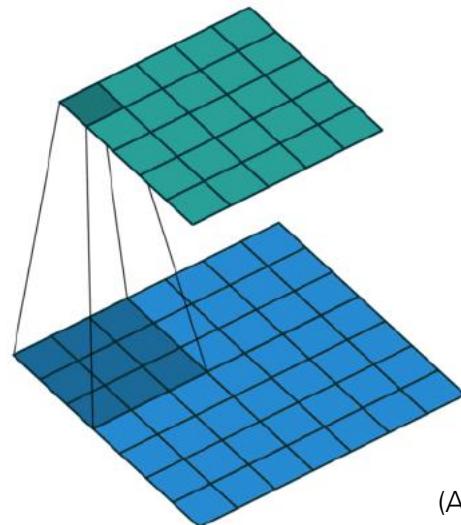


(Animation by
Vincent Dumoulin)

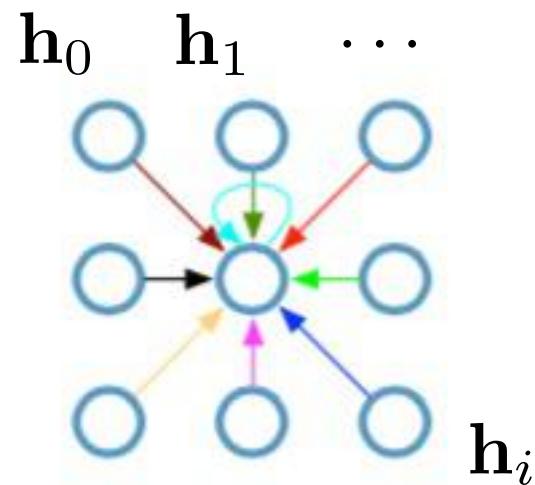


Recap: Convolutional neural networks (on grids)

Single CNN layer
with 3x3 filter:

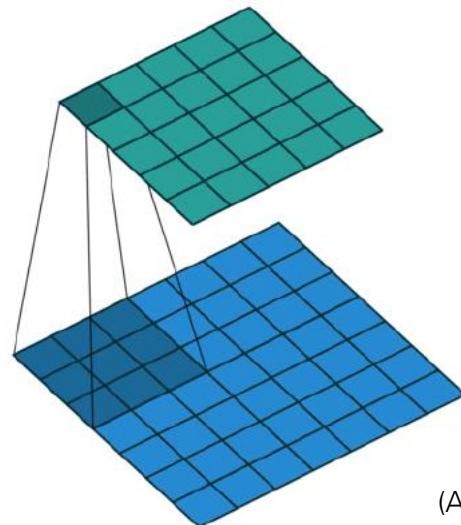


(Animation by
Vincent Dumoulin)

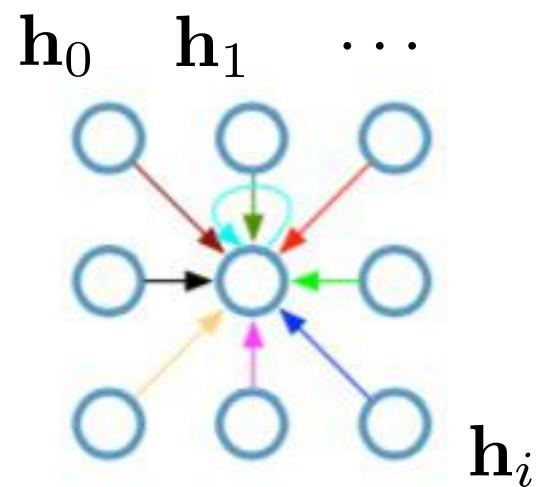


Recap: Convolutional neural networks (on grids)

Single CNN layer
with 3x3 filter:



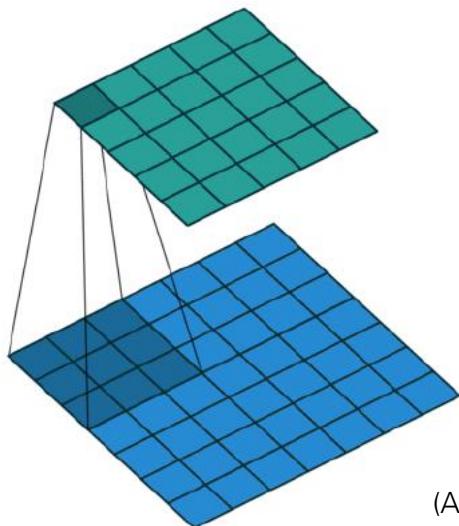
(Animation by
Vincent Dumoulin)



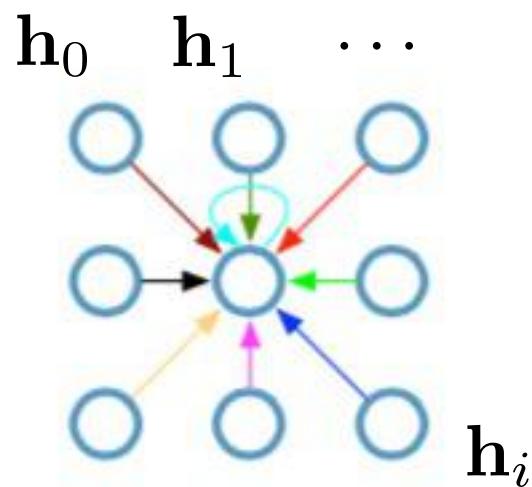
$\mathbf{h}_i \in \mathbb{R}^F$ are (hidden layer) activations of a pixel/node

Recap: Convolutional neural networks (on grids)

Single CNN layer
with 3x3 filter:



(Animation by
Vincent Dumoulin)



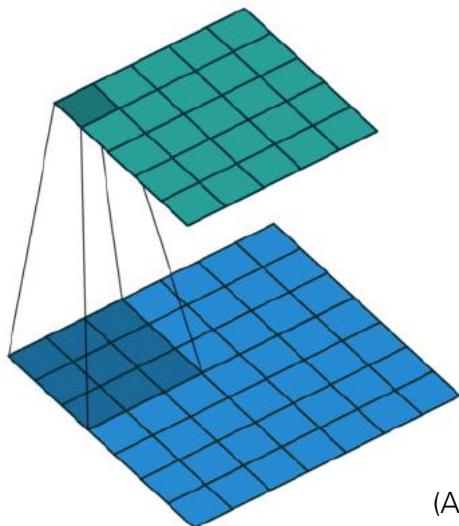
$\mathbf{h}_i \in \mathbb{R}^F$ are (hidden layer) activations of a pixel/node

Update for a single pixel:

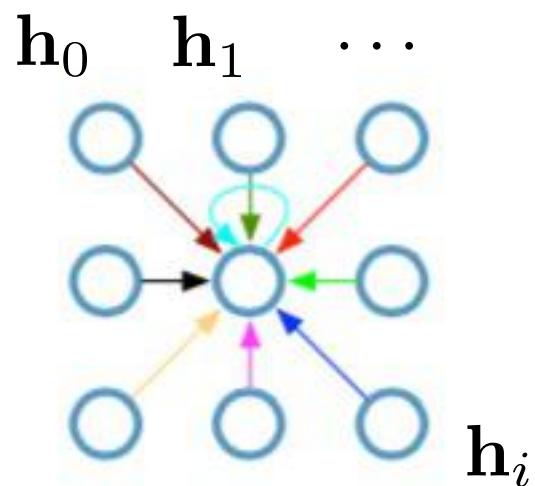
- Transform messages individually $\mathbf{W}_i \mathbf{h}_i$
- Add everything up $\sum_i \mathbf{W}_i \mathbf{h}_i$

Recap: Convolutional neural networks (on grids)

Single CNN layer
with 3x3 filter:



(Animation by
Vincent Dumoulin)



$\mathbf{h}_i \in \mathbb{R}^F$ are (hidden layer) activations of a pixel/node

Update for a single pixel:

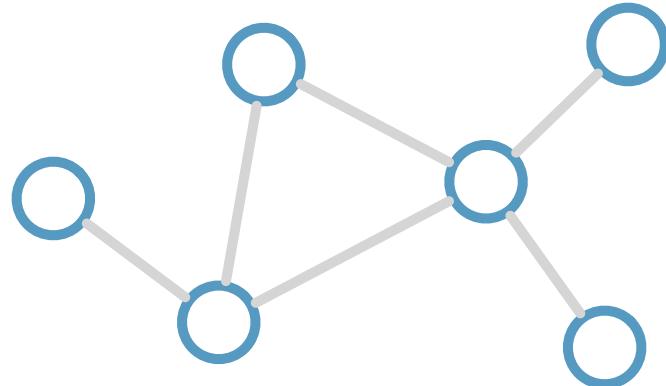
- Transform messages individually $\mathbf{W}_i \mathbf{h}_i$
- Add everything up $\sum_i \mathbf{W}_i \mathbf{h}_i$

Full update:

$$\mathbf{h}_4^{(l+1)} = \sigma \left(\mathbf{W}_0^{(l)} \mathbf{h}_0^{(l)} + \mathbf{W}_1^{(l)} \mathbf{h}_1^{(l)} + \cdots + \mathbf{W}_8^{(l)} \mathbf{h}_8^{(l)} \right)$$

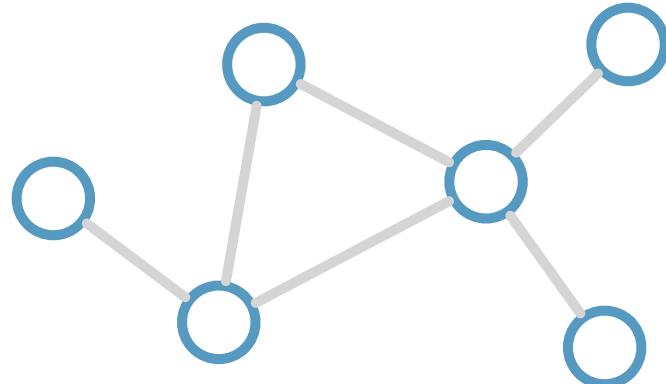
Graph Convolutional Networks (GCNs)

Consider this
undirected graph:

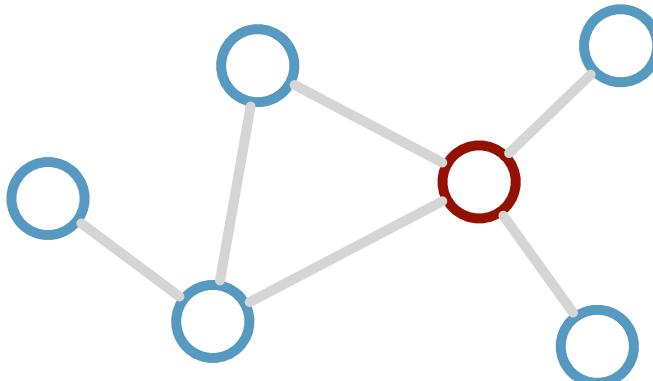


Graph Convolutional Networks (GCNs)

Consider this
undirected graph:

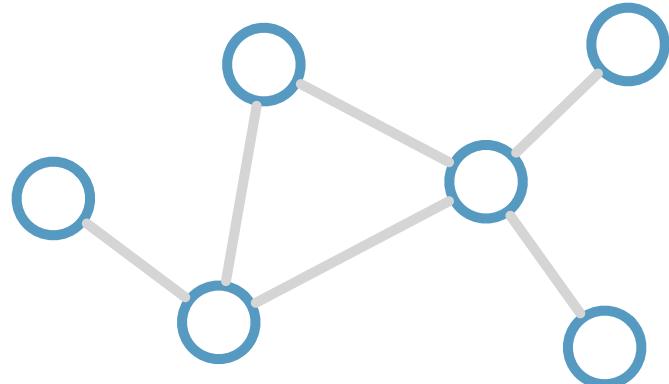


Consider update
for node in red:

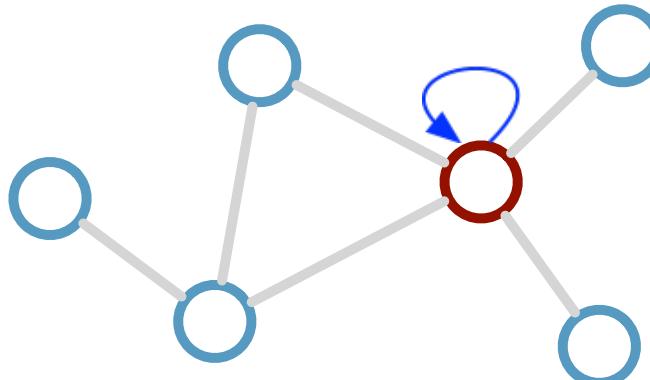


Graph Convolutional Networks (GCNs)

Consider this
undirected graph:

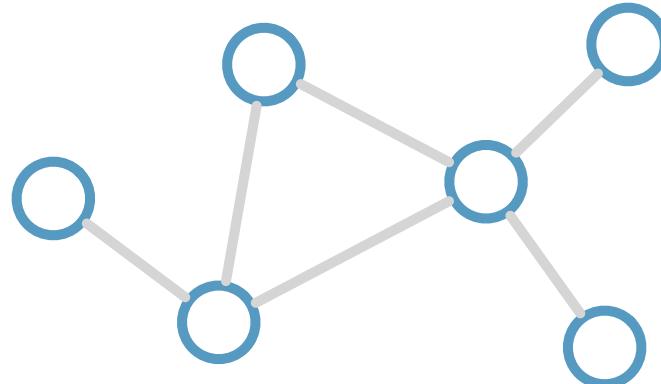


Consider update
for node in red:

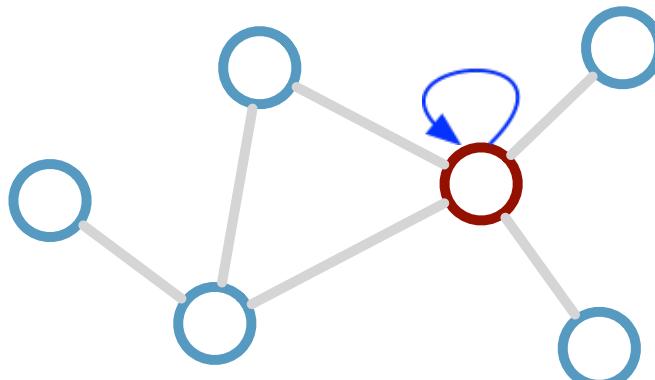


Graph Convolutional Networks (GCNs)

Consider this
undirected graph:



Consider update
for node in red:

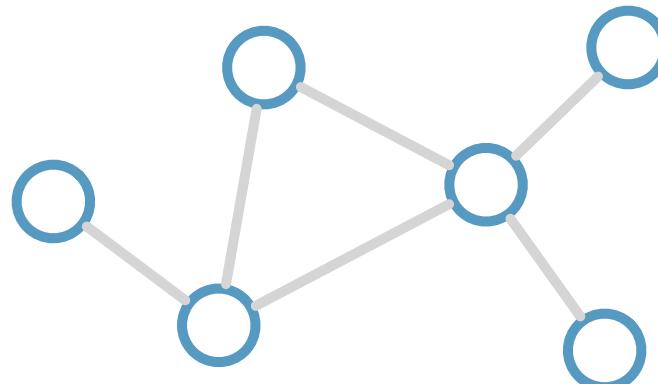


Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

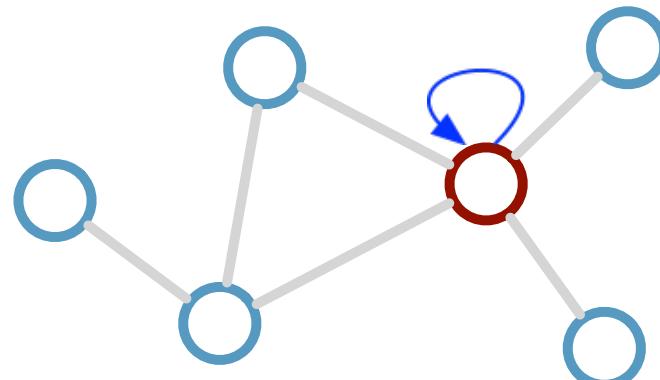
\mathcal{N}_i : neighbor indices c_{ij} : norm. constant (fixed/trainable)

Graph Convolutional Networks (GCNs)

Consider this undirected graph:



Consider update for node in red:



Desirable properties:

- Weight sharing over all locations
- Invariance to permutations
- Linear complexity $O(E)$
- Applicable both in transductive and inductive settings

Update rule:

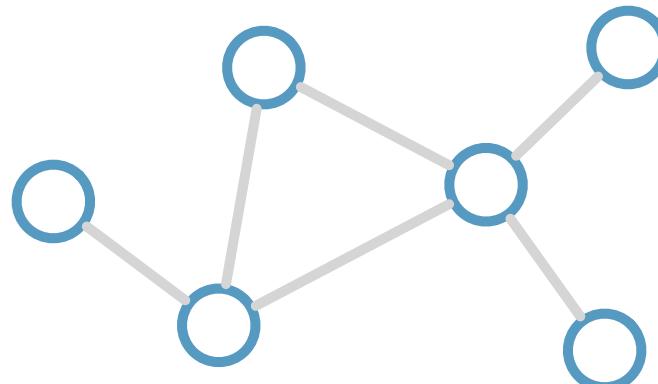
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

\mathcal{N}_i : neighbor indices

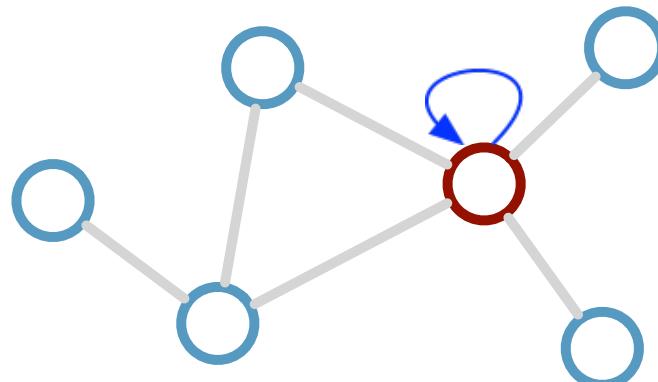
c_{ij} : norm. constant (fixed/trainable)

Graph Convolutional Networks (GCNs)

Consider this undirected graph:



Consider update for node in red:



Update rule:

$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

\mathcal{N}_i : neighbor indices

c_{ij} : norm. constant (fixed/trainable)

Desirable properties:

- Weight sharing over all locations
- Invariance to permutations
- Linear complexity $O(E)$
- Applicable both in transductive and inductive settings

Limitations:

- Requires gating mechanism / residual connections for depth
- Only indirect support for edge features

Graph Neural Networks (GNNs)

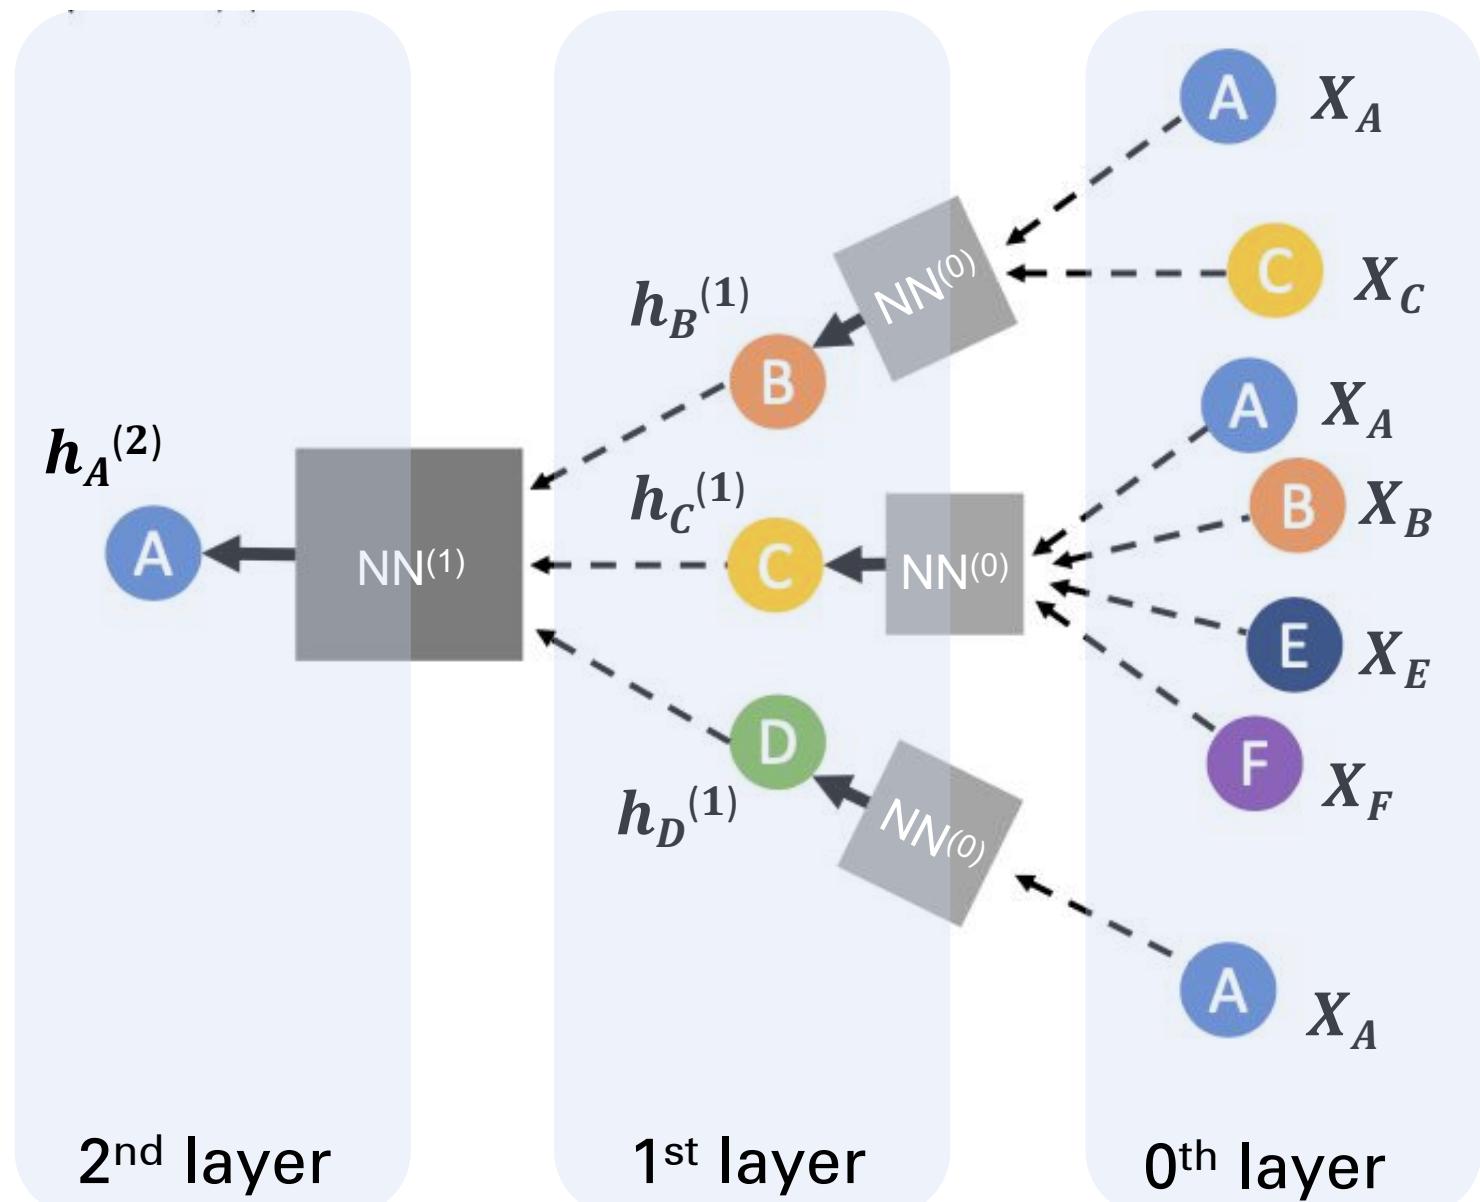
Graph Convolutional Networks^[1]

1. Aggregate messages

$$m_v^{(l)} = \frac{1}{|\mathcal{N}(v) + 1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l)}$$

2. Transform messages

$$h_v^{(l+1)} = \sigma(\mathbf{W}^{(l)} \circ m_v^{(l)})$$



[1] Kipf, Thomas N., et al. "Semi-supervised classification with graph convolutional networks."

Graph Neural Networks (GNNs)

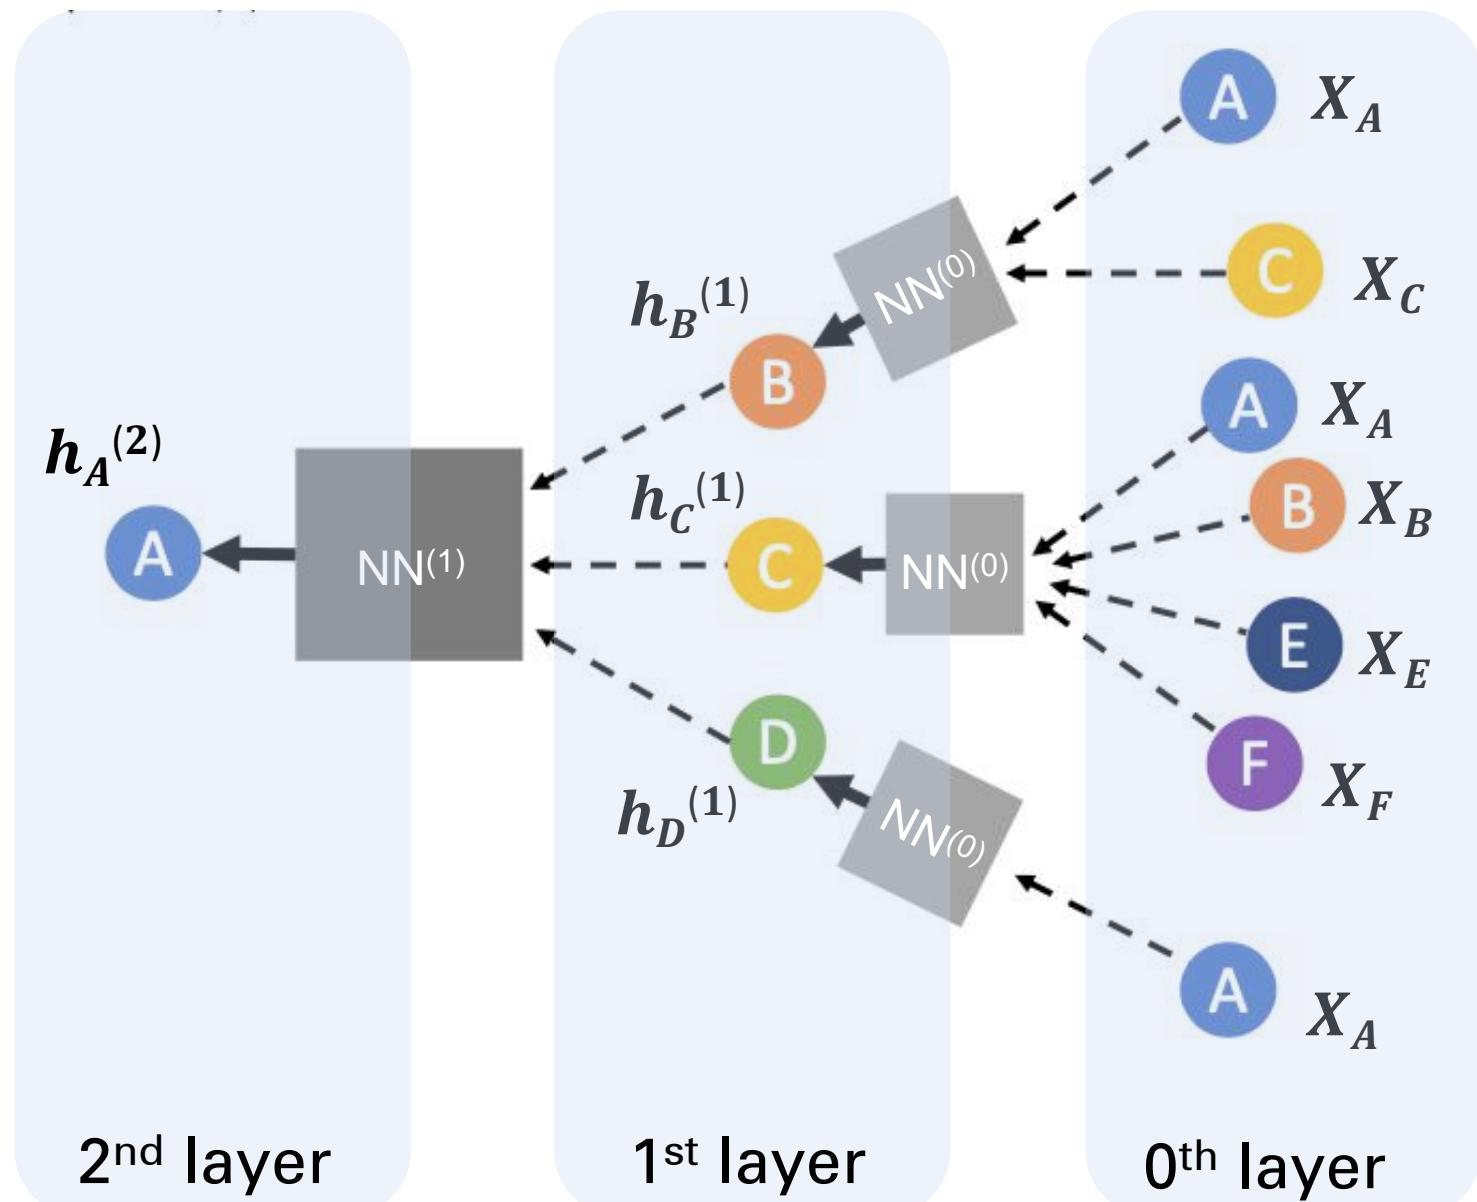
Graph Isomorphism
Networks^[2]

1. Aggregate messages

$$m_v^{(l)} = \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l)}$$

2. Transform messages

$$h_v^{(l+1)} = \sigma(\mathbf{W}^{(l)} \circ m_v^{(l)})$$



[2] Xu, Keyulu, et al. "How powerful are graph neural networks?".

Graph Neural Networks (GNNs)

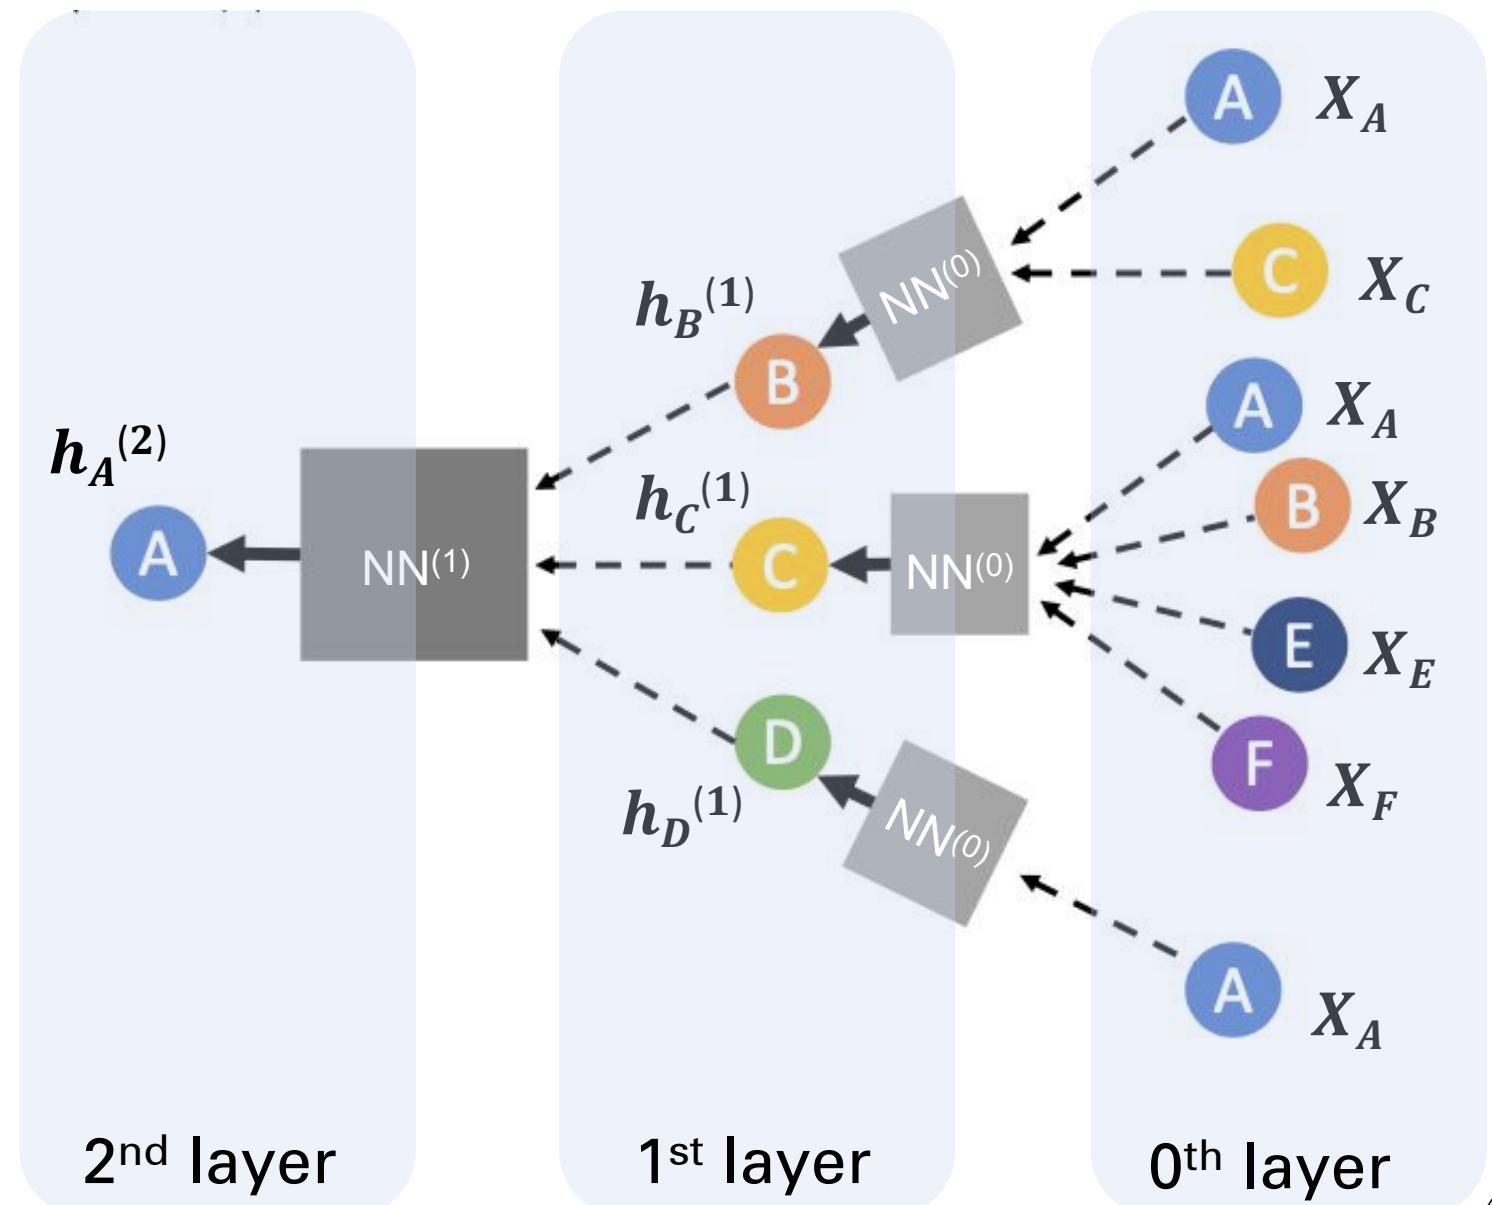
Simplified Graph
Convolutional Networks^[2]

1. Aggregate messages

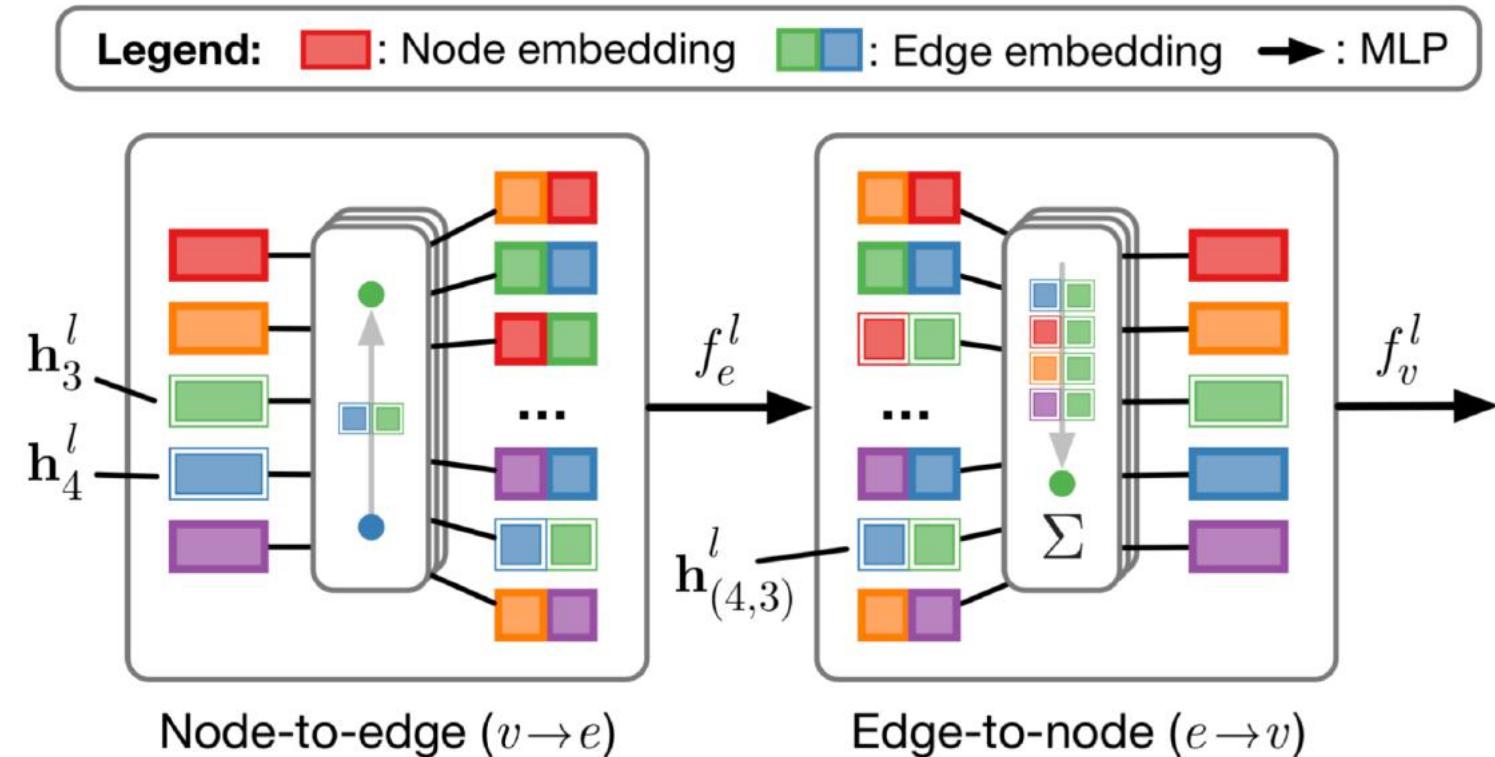
$$m_v^{(l)} = \frac{1}{|\mathcal{N}(v) + 1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l)}$$

2. Transform messages

$$h_v^{(l+1)} = \mathbf{W}^{(l)} \circ m_v^{(l)}$$



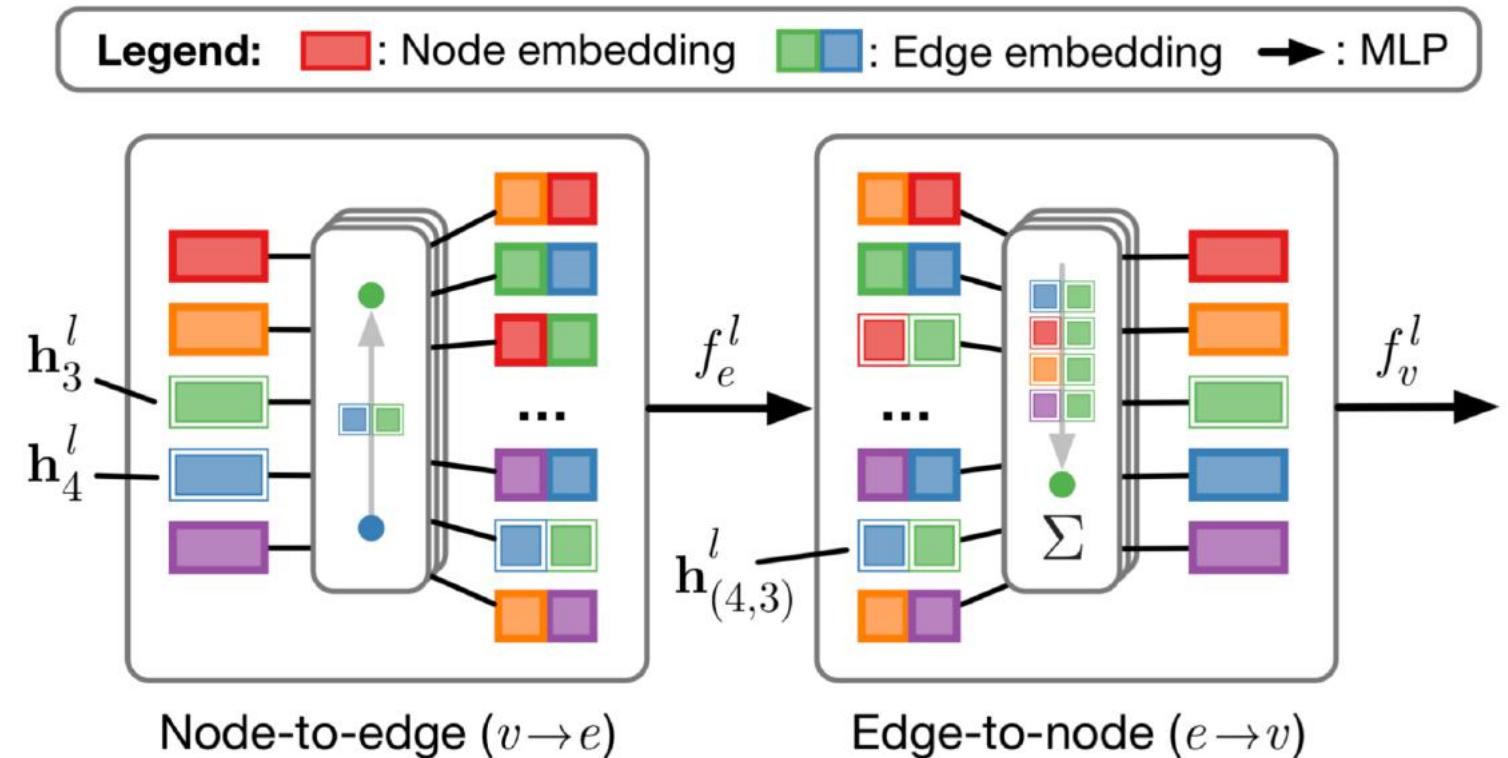
GCNs with edge embeddings



Formally: $v \rightarrow e : \mathbf{h}_{(i,j)}^l = f_e^l ([\mathbf{h}_i^l, \mathbf{h}_j^l, \mathbf{x}_{(i,j)}])$

$e \rightarrow v : \mathbf{h}_j^{l+1} = f_v^l \left(\left[\sum_{i \in \mathcal{N}_j} \mathbf{h}_{(i,j)}^l, \mathbf{x}_j \right] \right)$

GCNs with edge embeddings



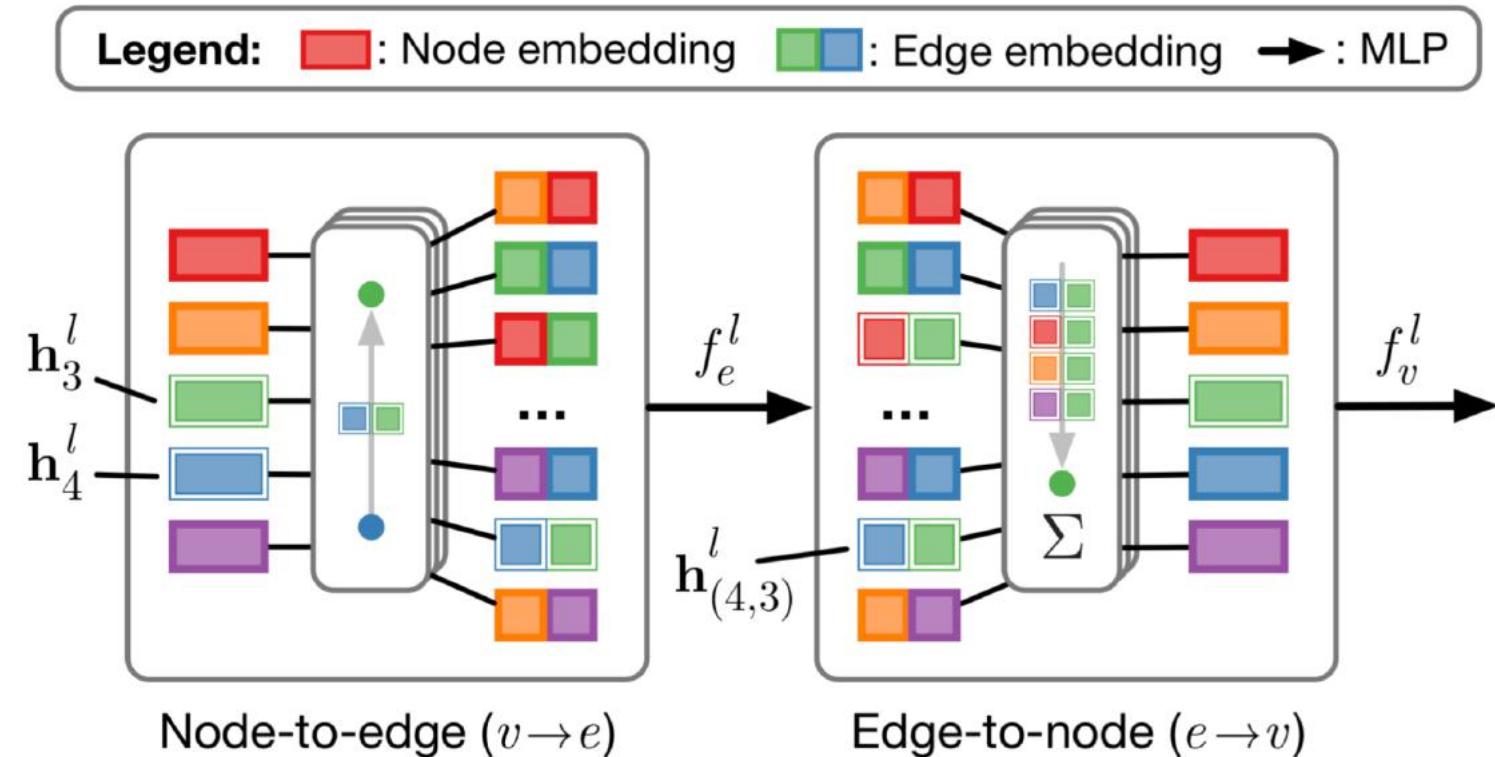
Formally: $v \rightarrow e : \mathbf{h}_{(i,j)}^l = f_e^l ([\mathbf{h}_i^l, \mathbf{h}_j^l, \mathbf{x}_{(i,j)}])$

$e \rightarrow v : \mathbf{h}_j^{l+1} = f_v^l \left(\left[\sum_{i \in \mathcal{N}_j} \mathbf{h}_{(i,j)}^l, \mathbf{x}_j \right] \right)$

Pros:

- Supports edge features
- More expressive than GCN
- As general as it gets (?)
- Supports sparse matrix ops

GCNs with edge embeddings



Formally: $v \rightarrow e : \mathbf{h}_{(i,j)}^l = f_e^l ([\mathbf{h}_i^l, \mathbf{h}_j^l, \mathbf{x}_{(i,j)}])$

$e \rightarrow v : \mathbf{h}_j^{l+1} = f_v^l \left(\left[\sum_{i \in \mathcal{N}_j} \mathbf{h}_{(i,j)}^l, \mathbf{x}_j \right] \right)$

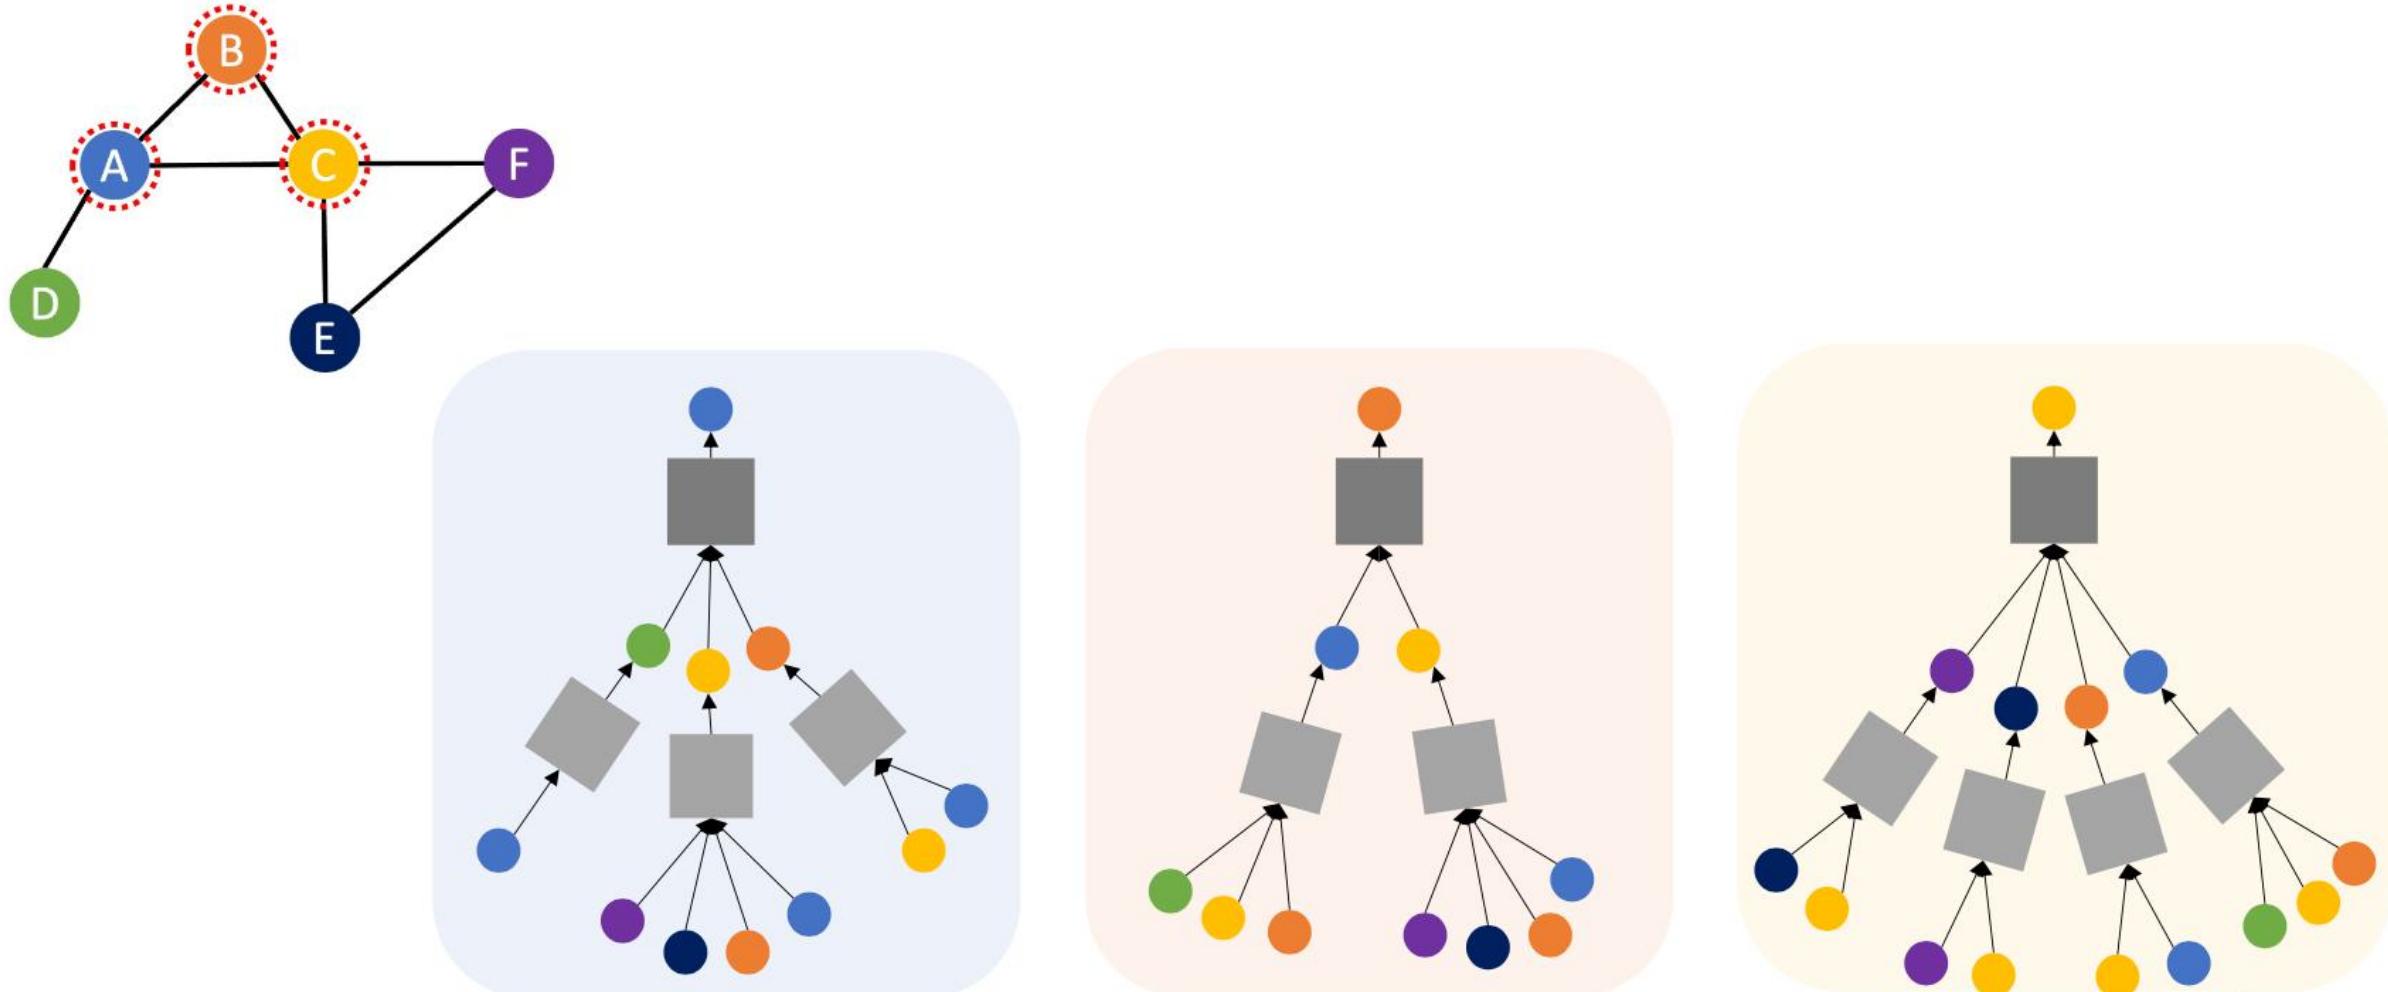
Pros:

- Supports edge features
- More expressive than GCN
- As general as it gets (?)
- Supports sparse matrix ops

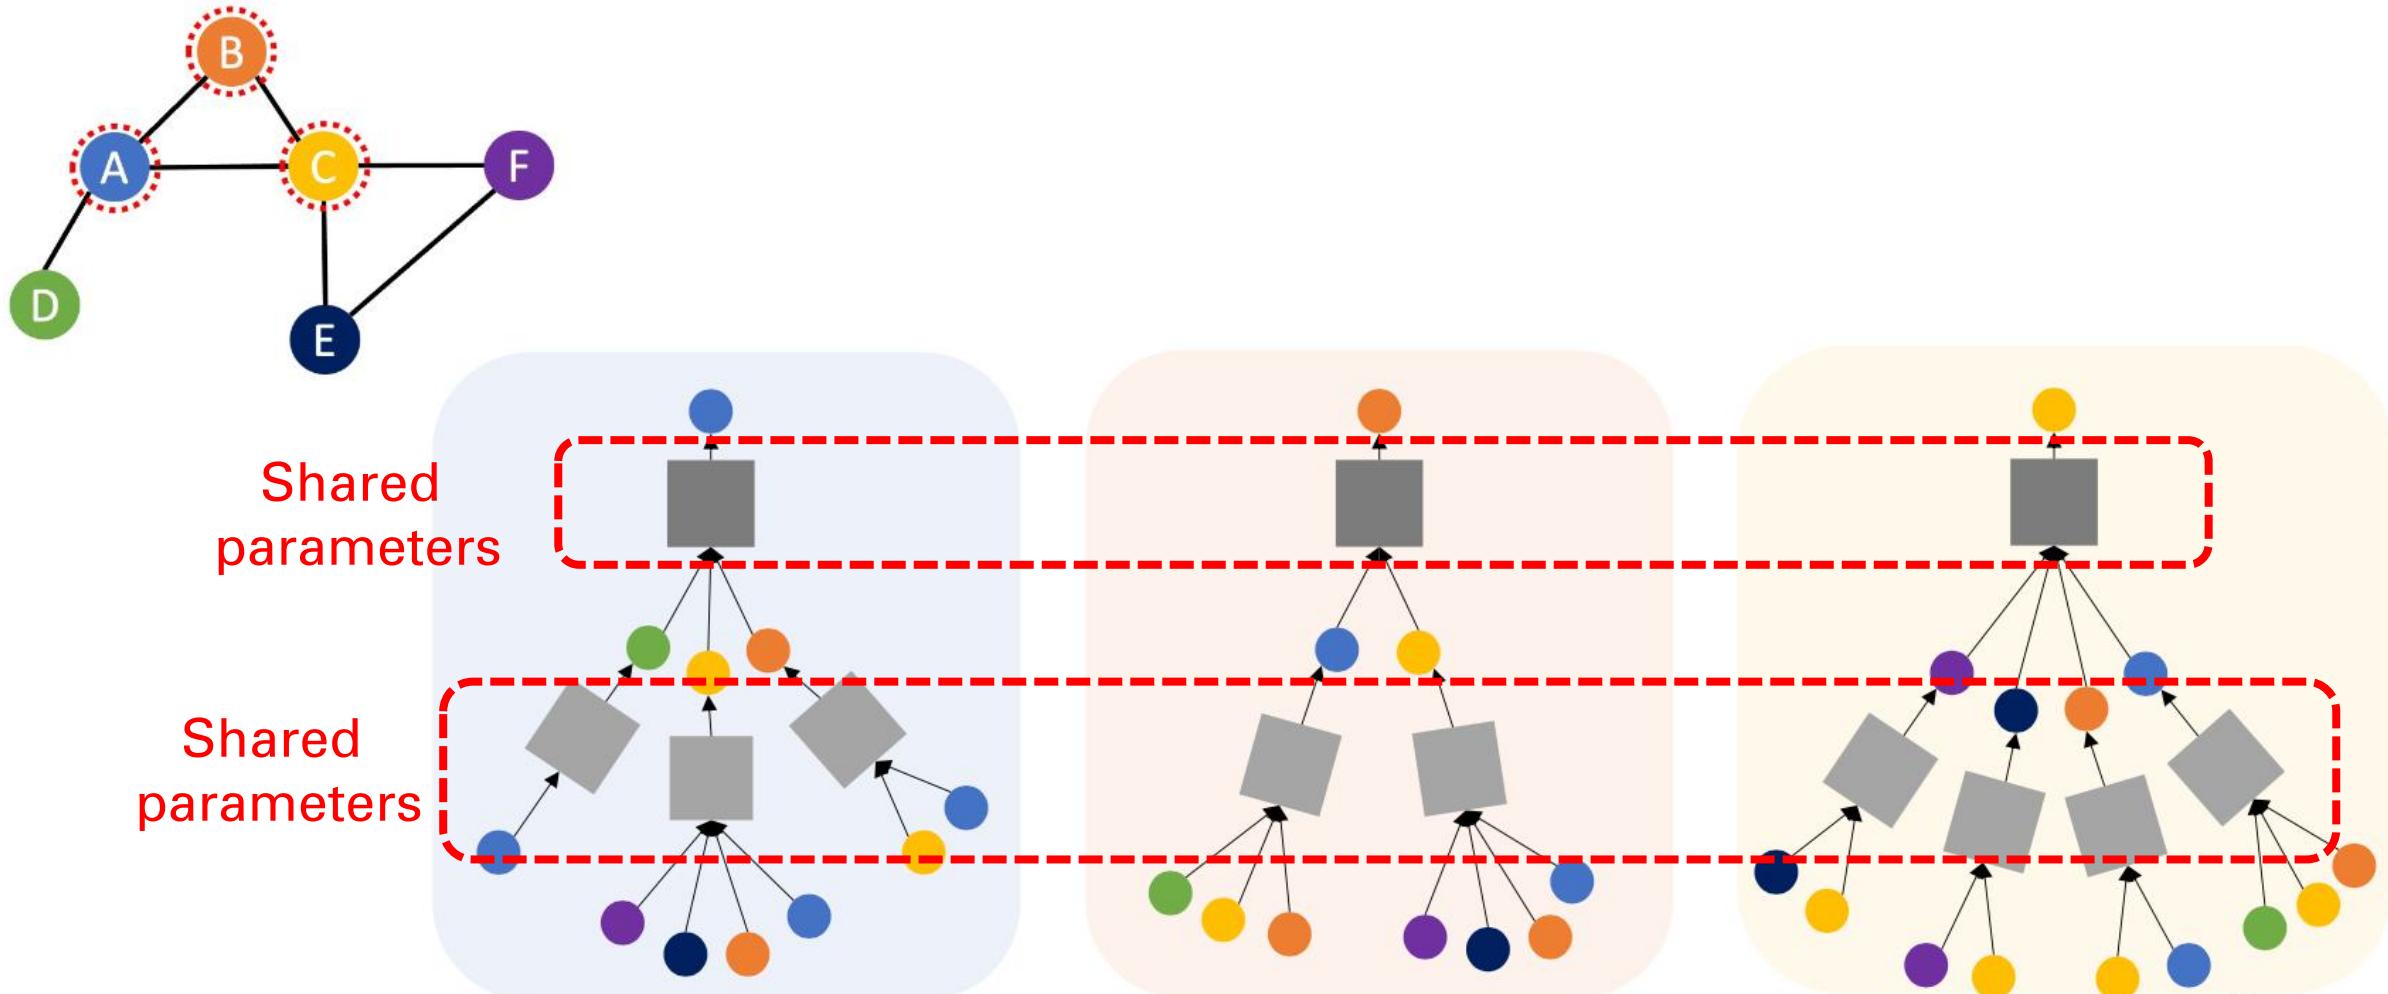
Cons:

- Need to store intermediate edge-based activations
- Difficult to implement with subsampling
- In practice limited to small graphs

Computation graphs

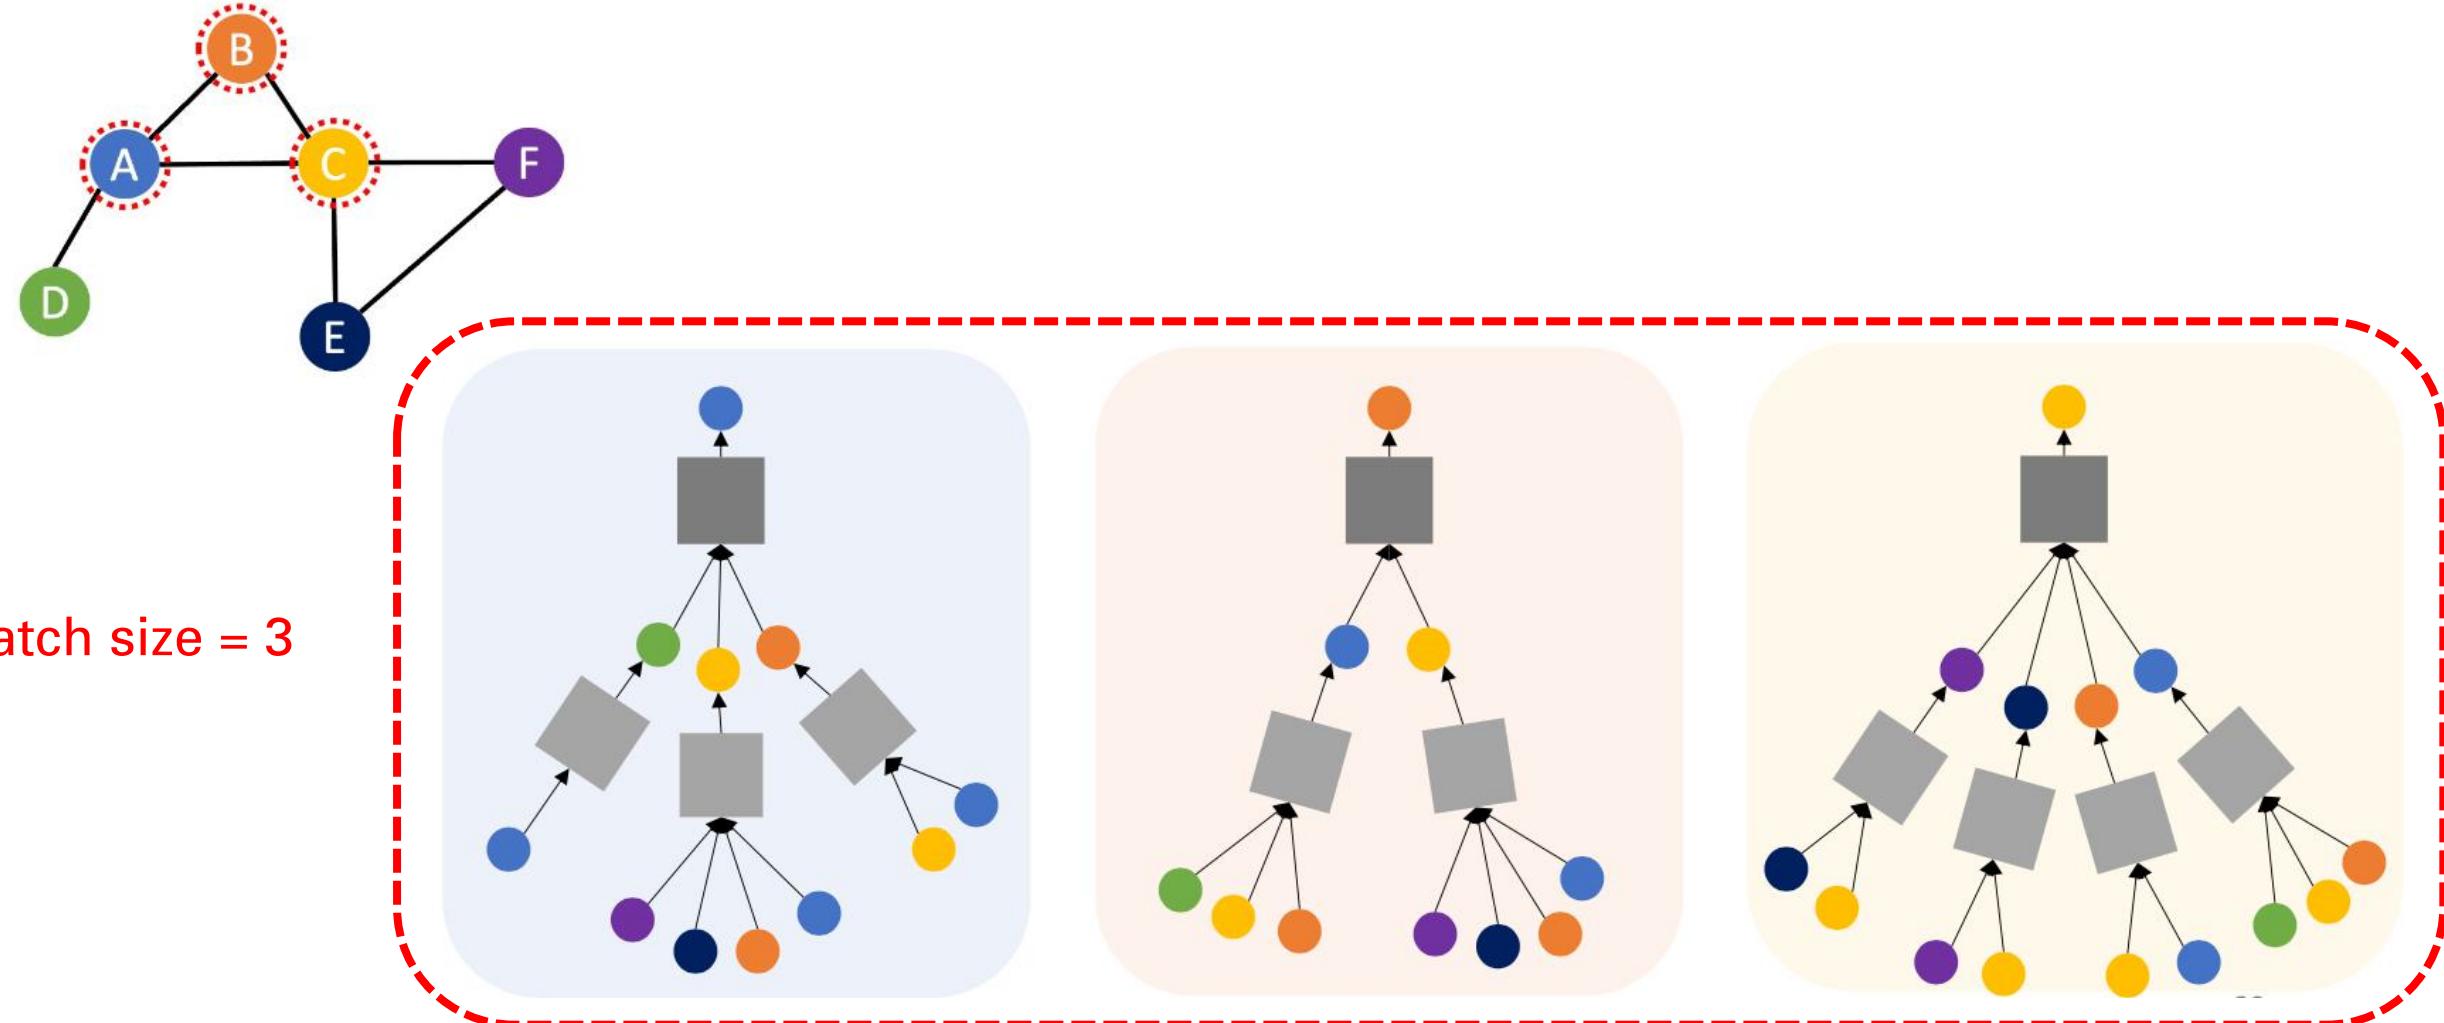


Computation graphs

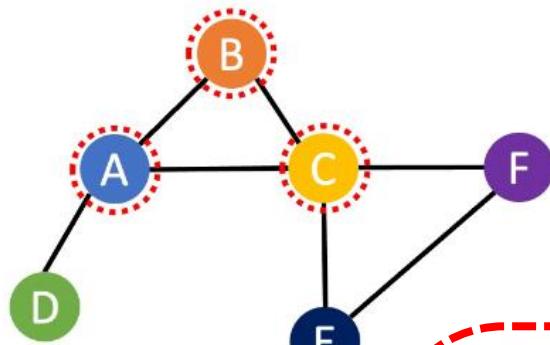


Batch execution

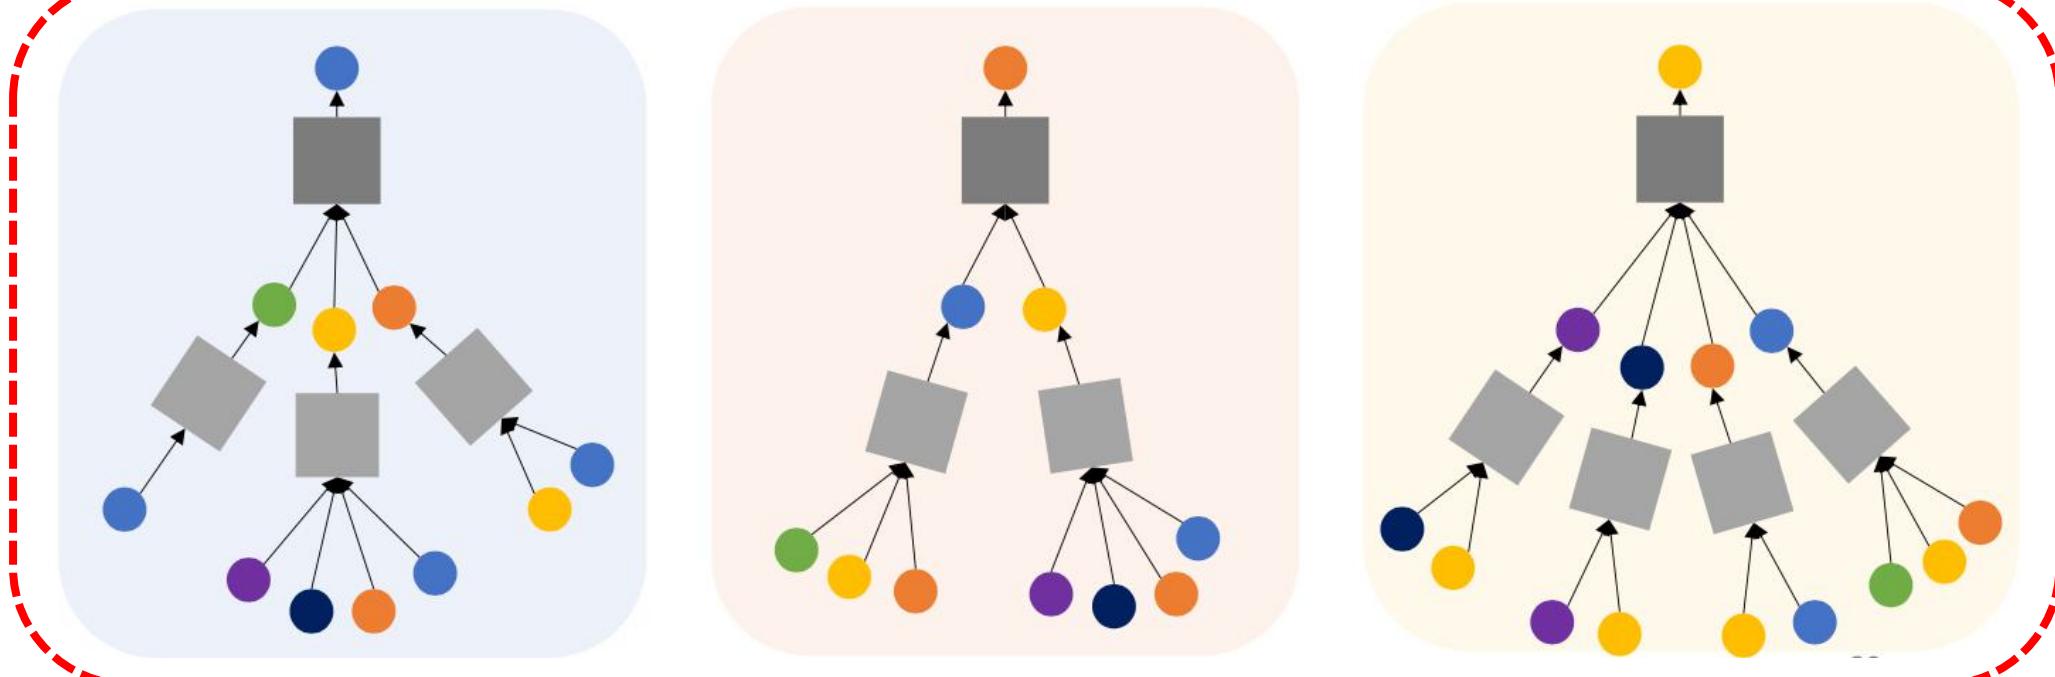
$$h_v^{(l)} = \sigma(\mathbf{W}^{(l)} \circ (\frac{1}{|\mathcal{N}(v)+1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l-1)}))$$



Batch execution



Batch size = 3

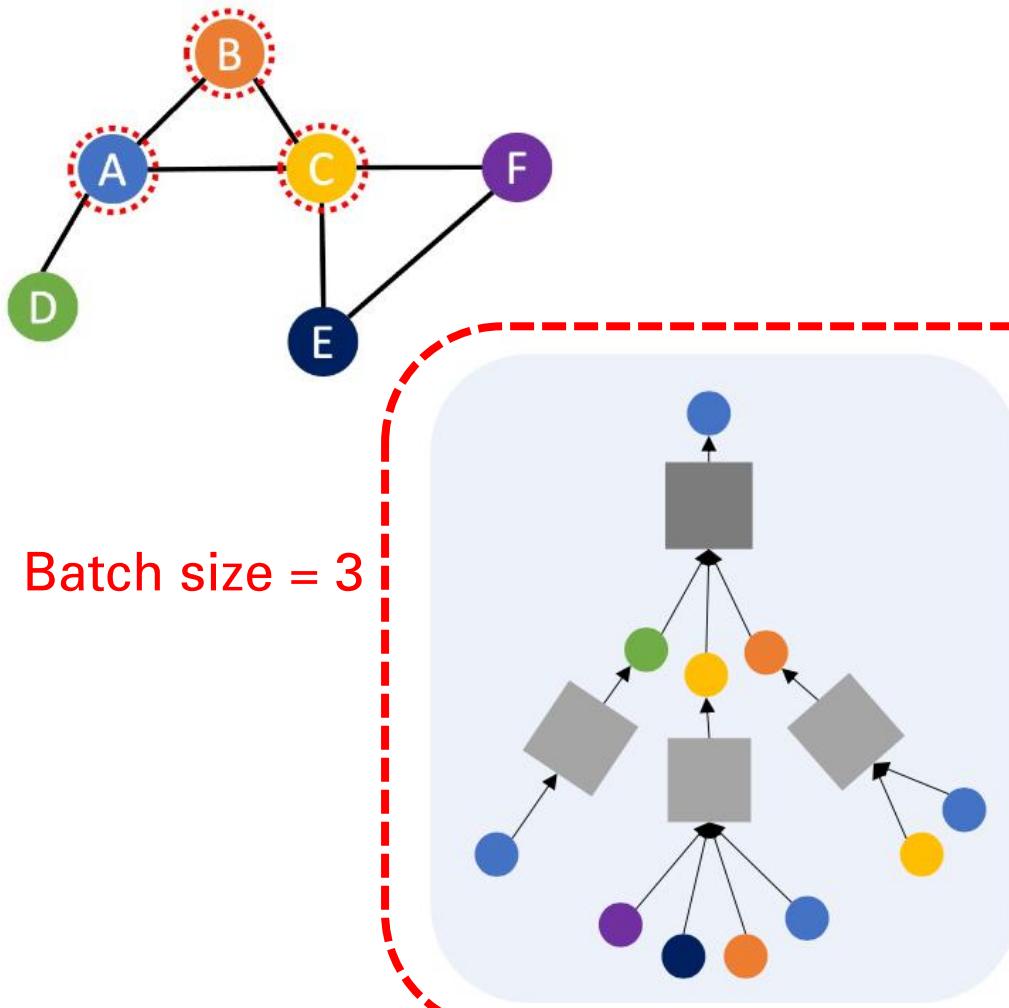


$$h_v^{(l)} = \sigma(\mathbf{W}^{(l)} \circ (\frac{1}{|\mathcal{N}(v)+1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l-1)}))$$

$$\mathbf{H}^{(l)} = \sigma((\widetilde{\mathbf{A} + \mathbf{I}}) \mathbf{H}^{(l-1)} \mathbf{W}^{(l)})$$

Node embedding matrix
(row-normalized) Adjacency matrix

Batch execution

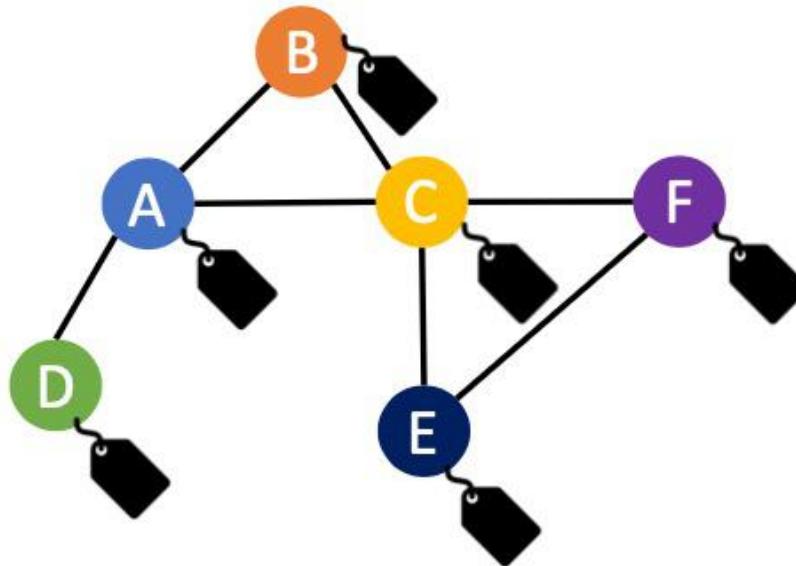


$$h_v^{(l)} = \sigma(\mathbf{W}^{(l)} \circ (\frac{1}{|\mathcal{N}(v)+1|} \sum_{u \in \mathcal{N}(v) \cup \{v\}} h_u^{(l-1)}))$$
$$\mathbf{H}^{(l)} = \sigma((\widetilde{\mathbf{A} + \mathbf{I}}) \mathbf{H}^{(l-1)} \mathbf{W}^{(l)})$$

Fixed **Trainable**

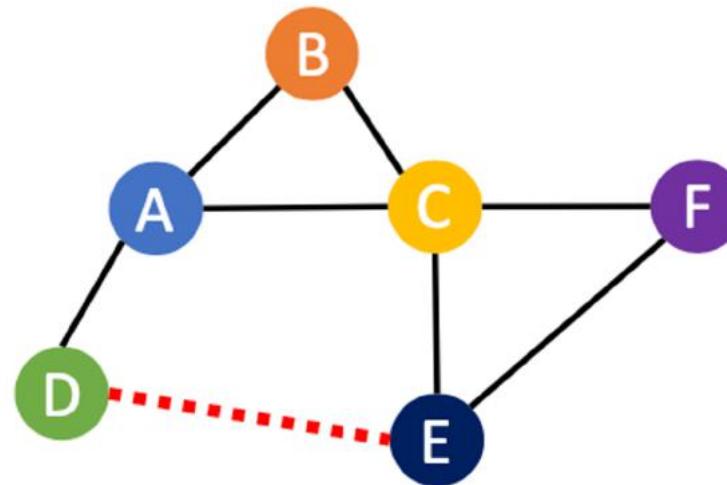
Downstream tasks

- Node-level prediction



Downstream tasks

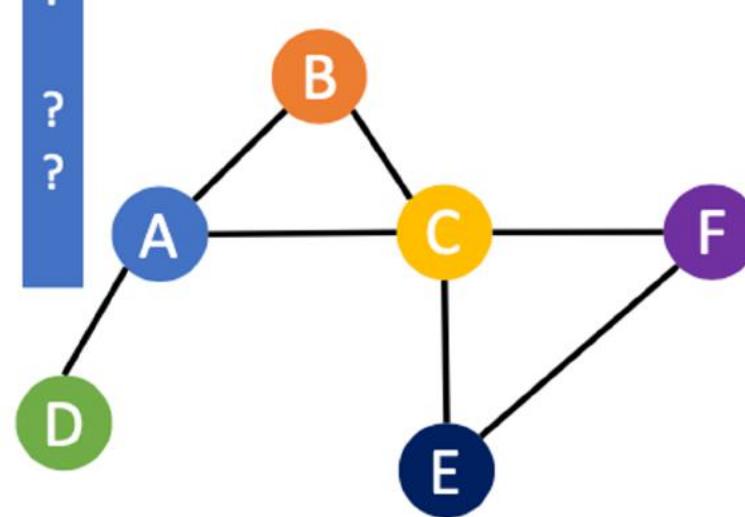
- Node-level prediction
- Edge-level prediction



D and E are related enough
to be connected?

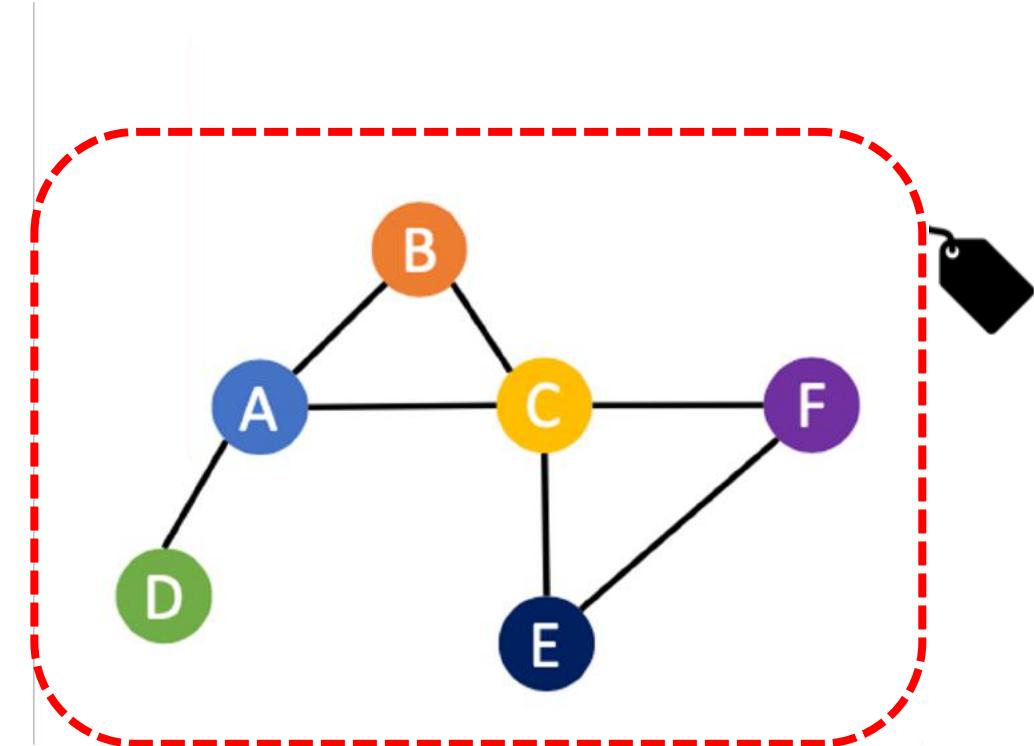
Downstream tasks

- Node-level prediction
- Edge-level prediction
- Attribute-level prediction



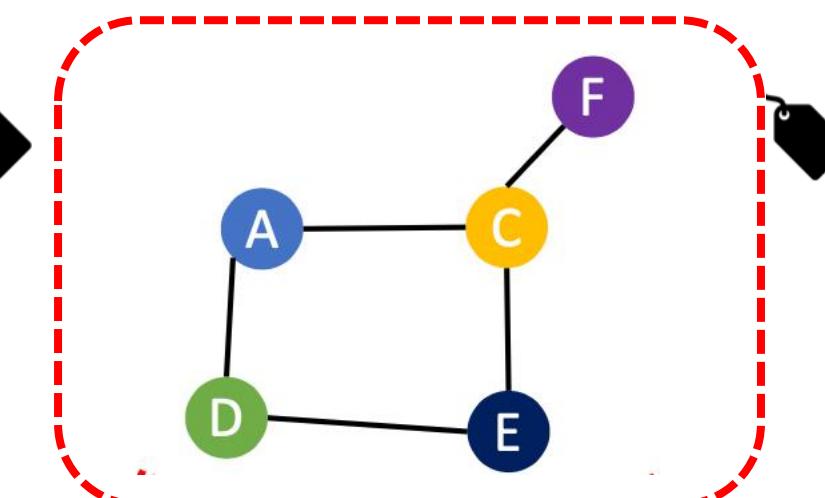
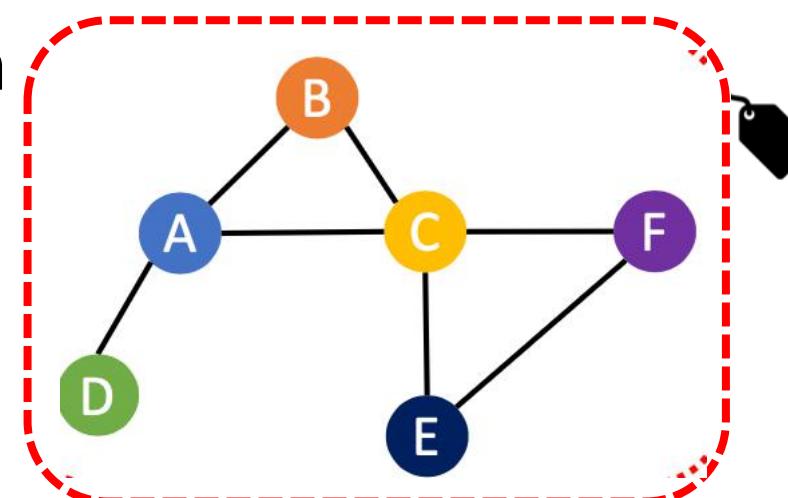
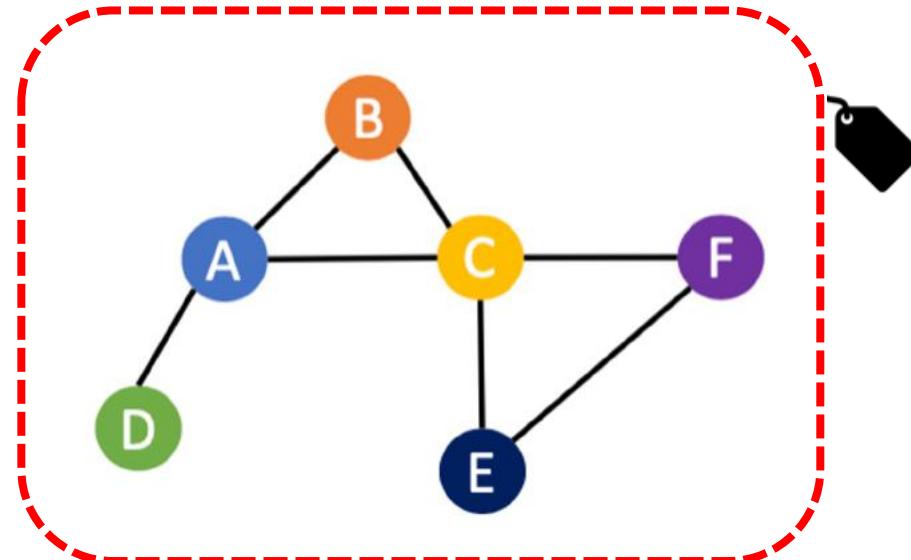
Downstream tasks

- Node-level prediction
- Edge-level prediction
- Attribute-level prediction
- Graph-level prediction



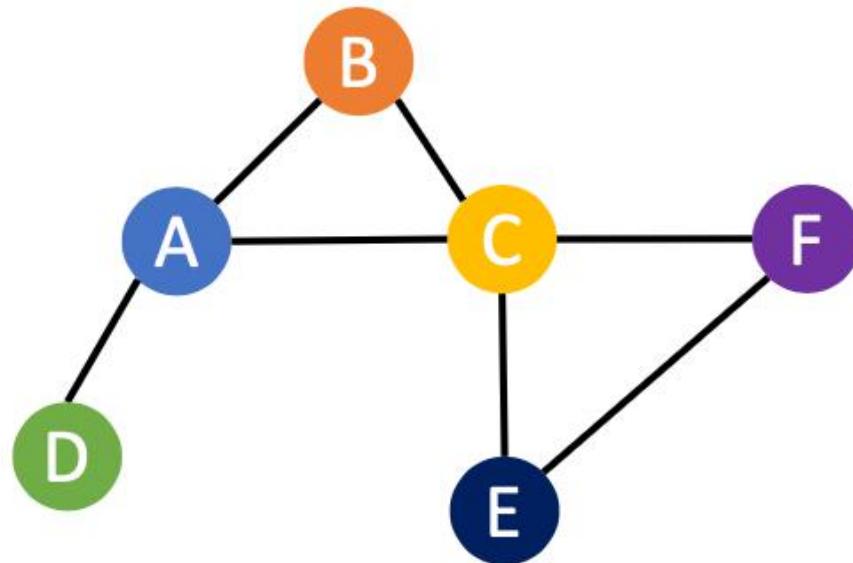
Downstream tasks

- Node-level prediction
- Edge-level prediction
- Attribute-level prediction
- Graph-level prediction

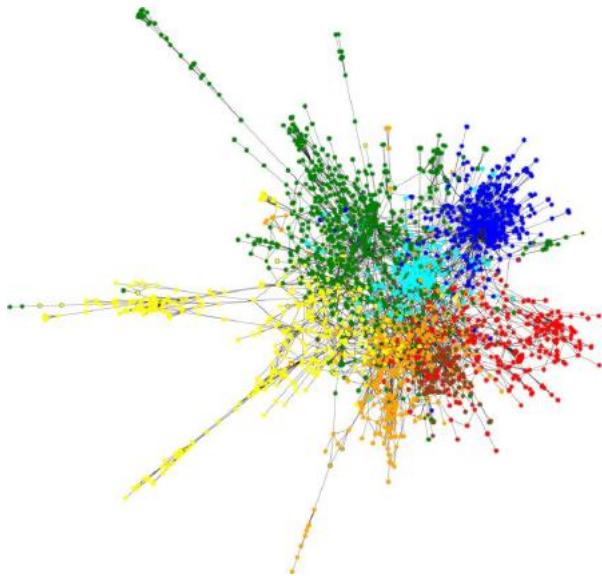
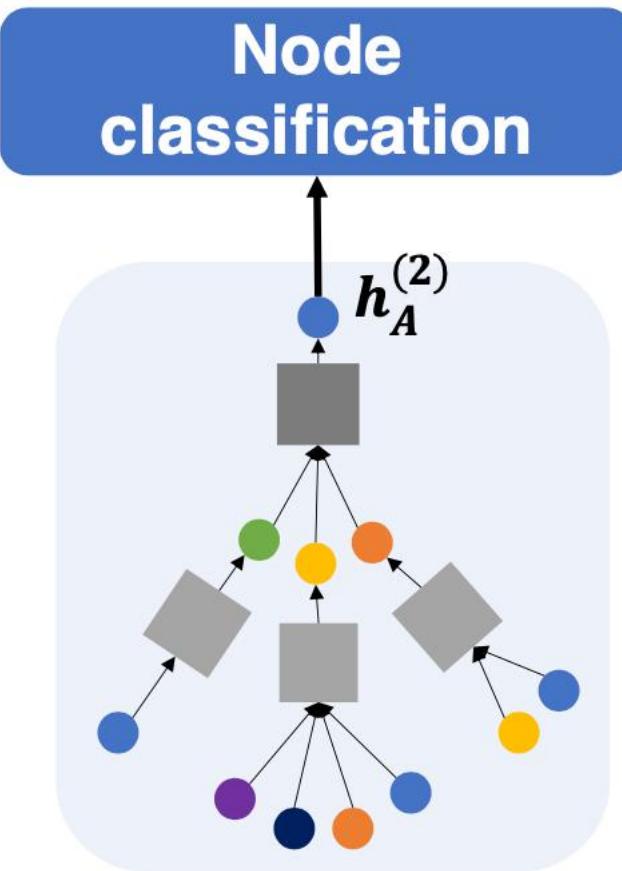


Downstream tasks

- Node-level prediction
- Edge-level prediction
- Attribute-level prediction
- Graph-level prediction

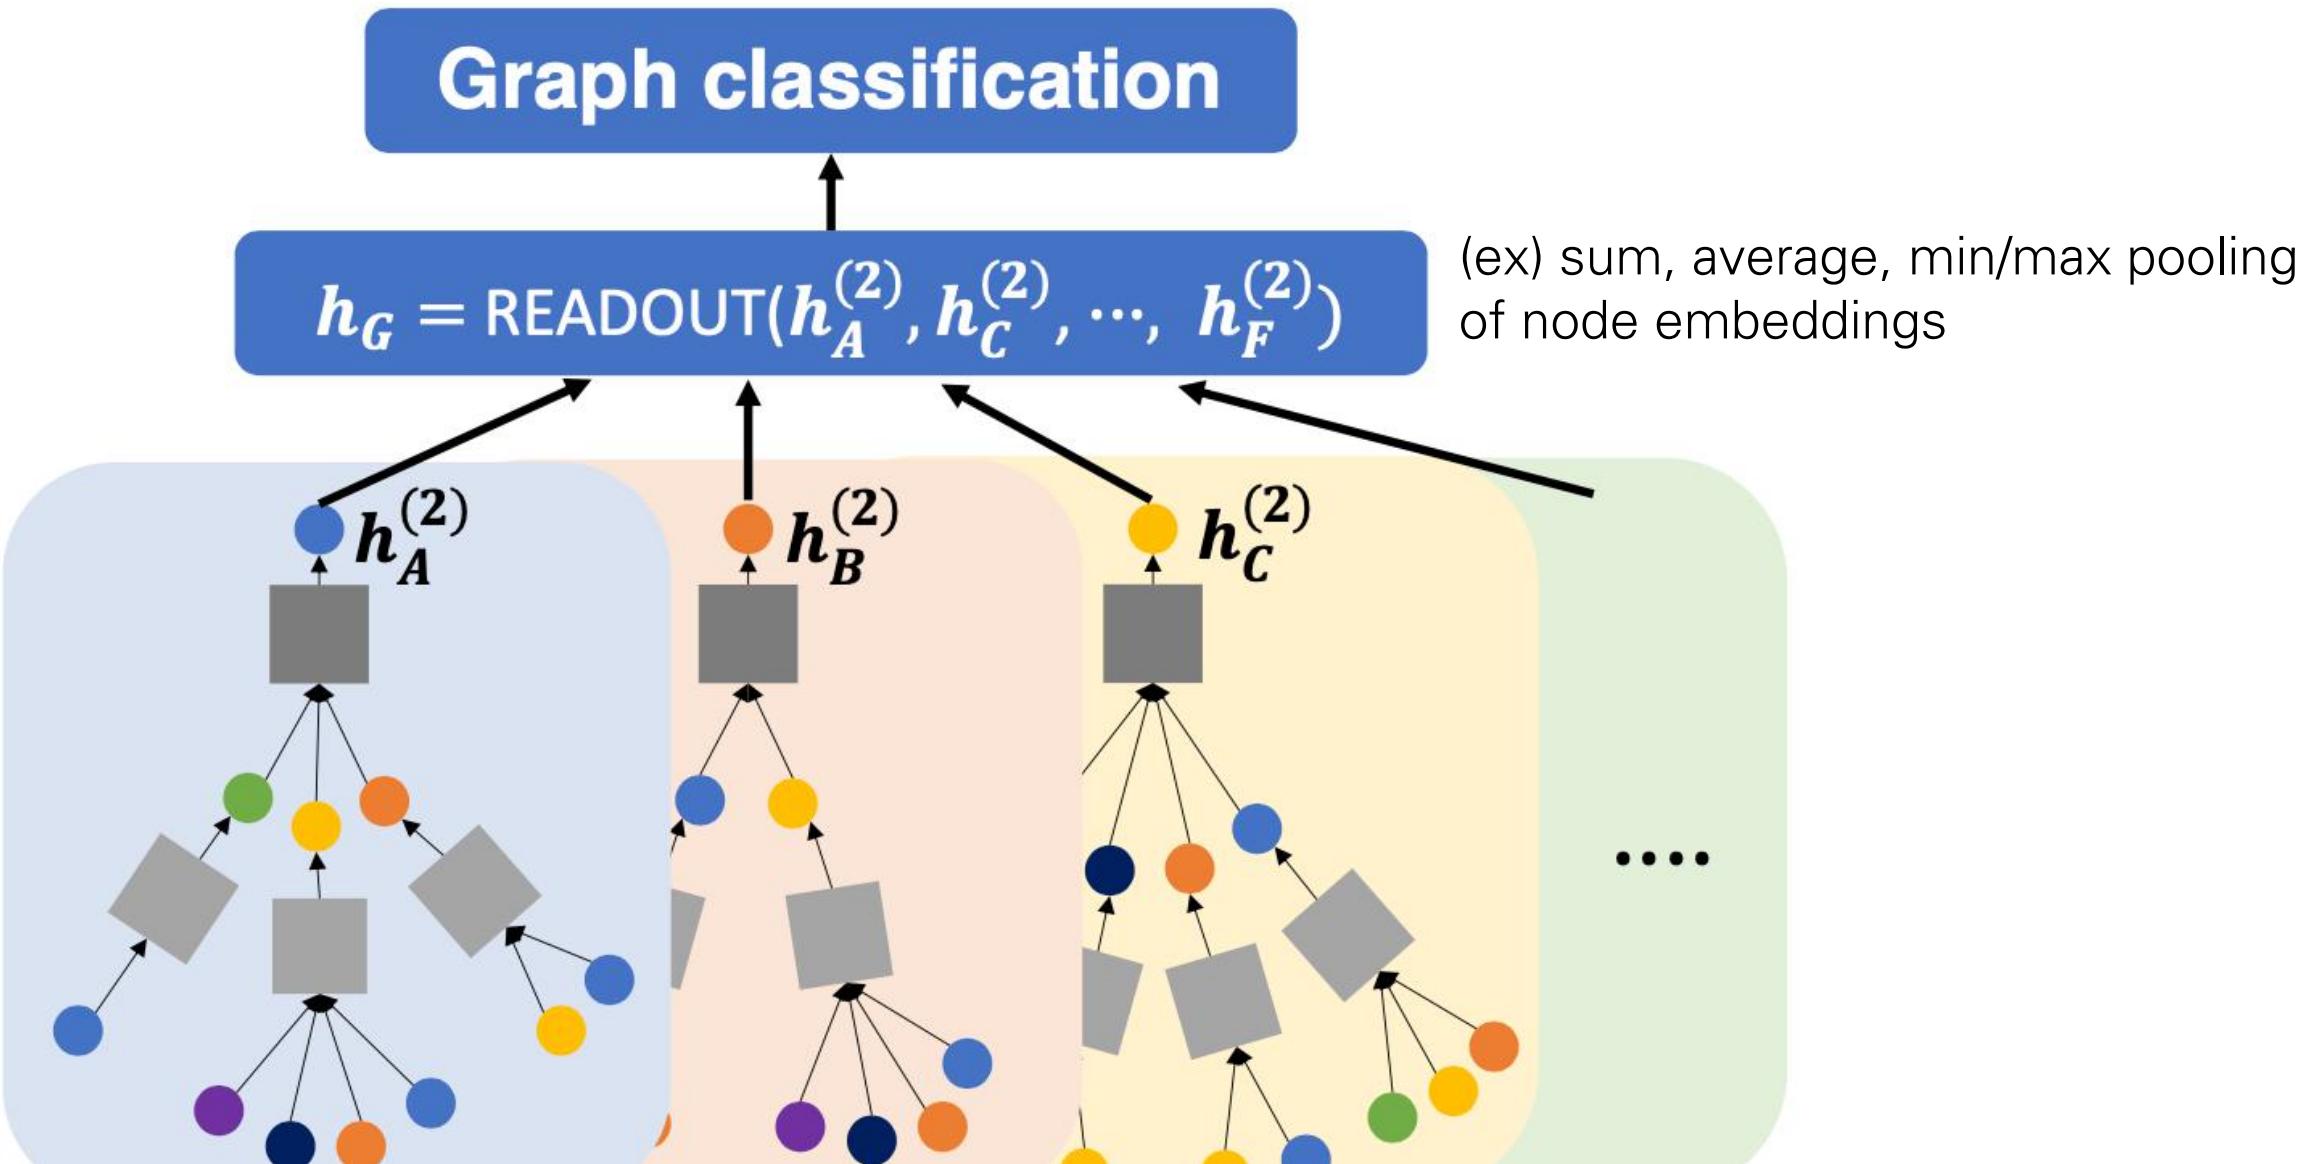


Node-level prediction tasks

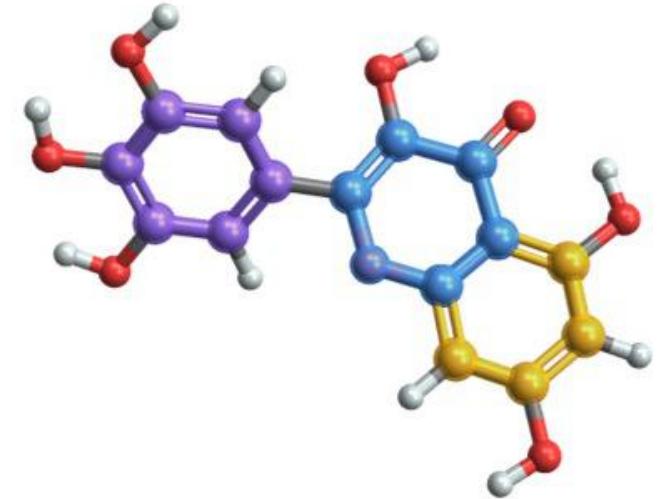
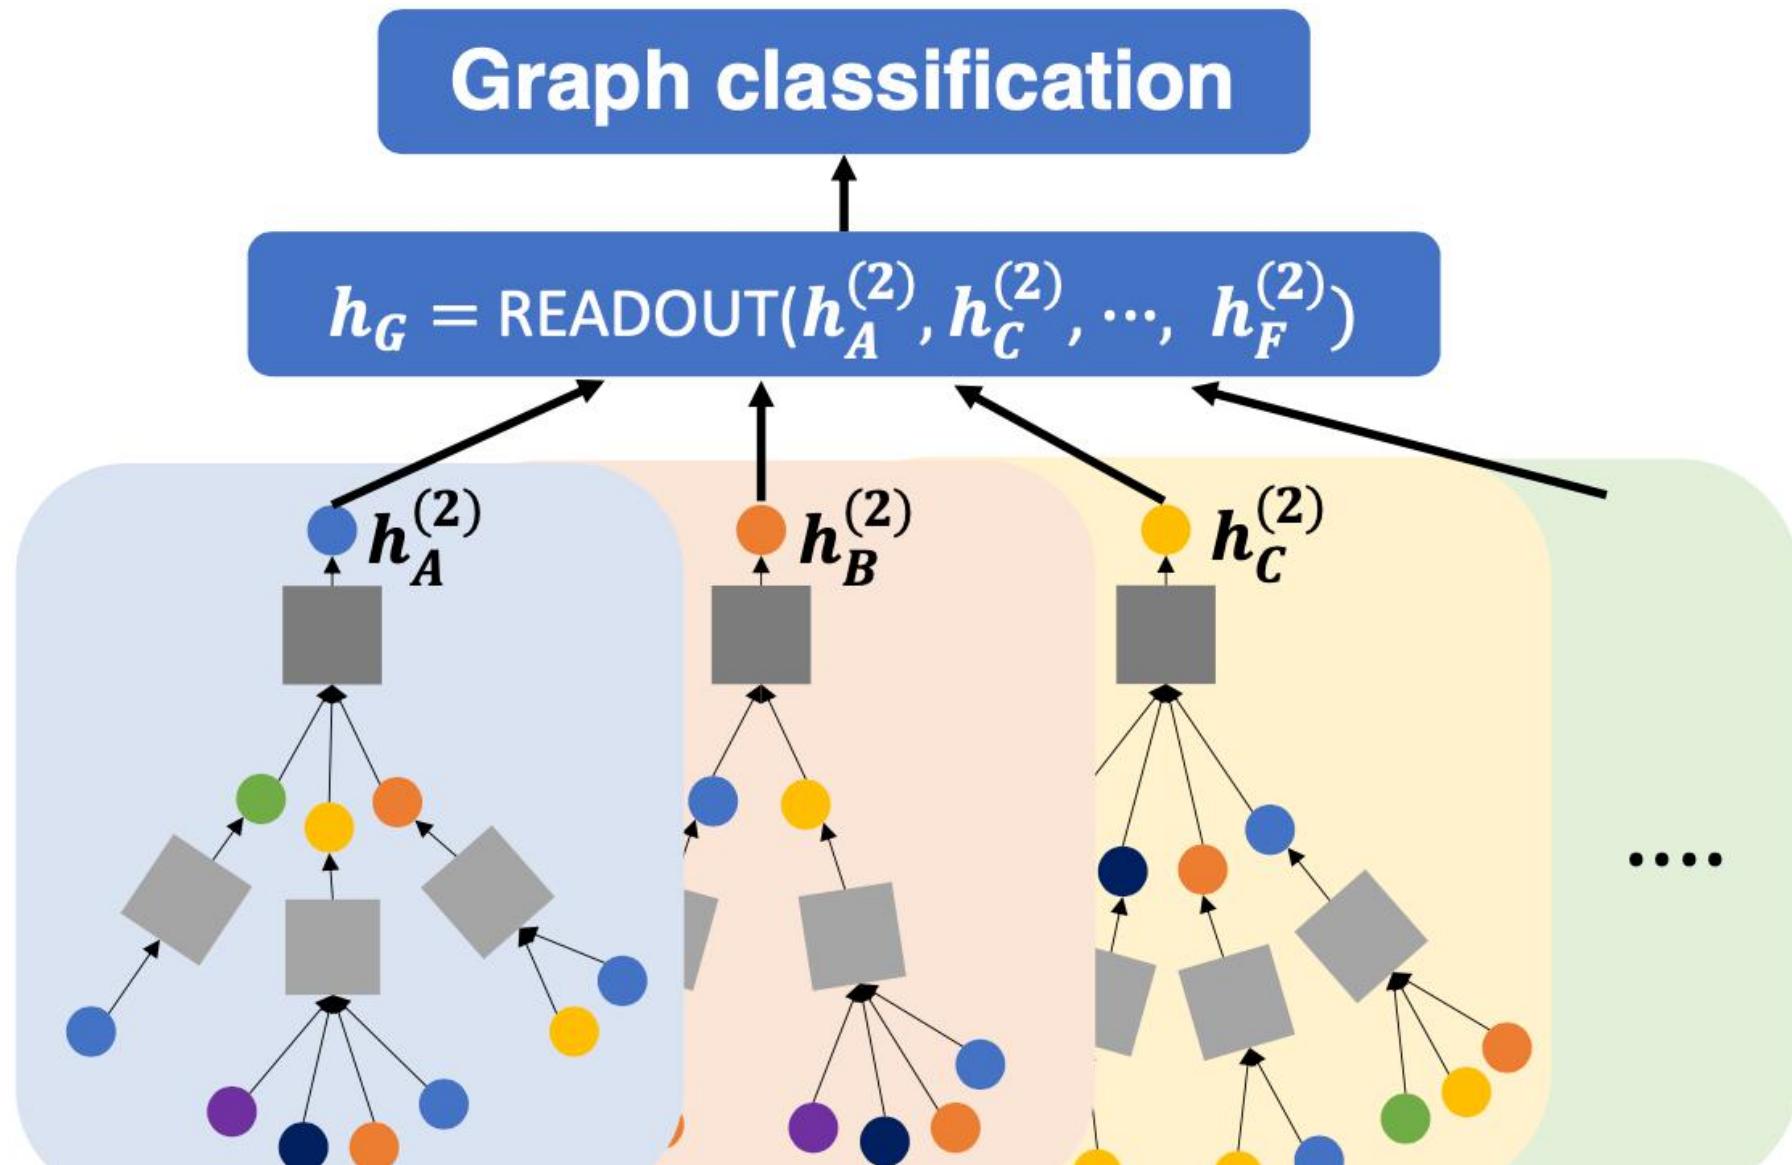


- Classify papers into topics on citation networks
- Cluster posts into subgroups on Reddit networks
- Classify products into categories on Amazon co-purchase graphs

Graph-level prediction tasks



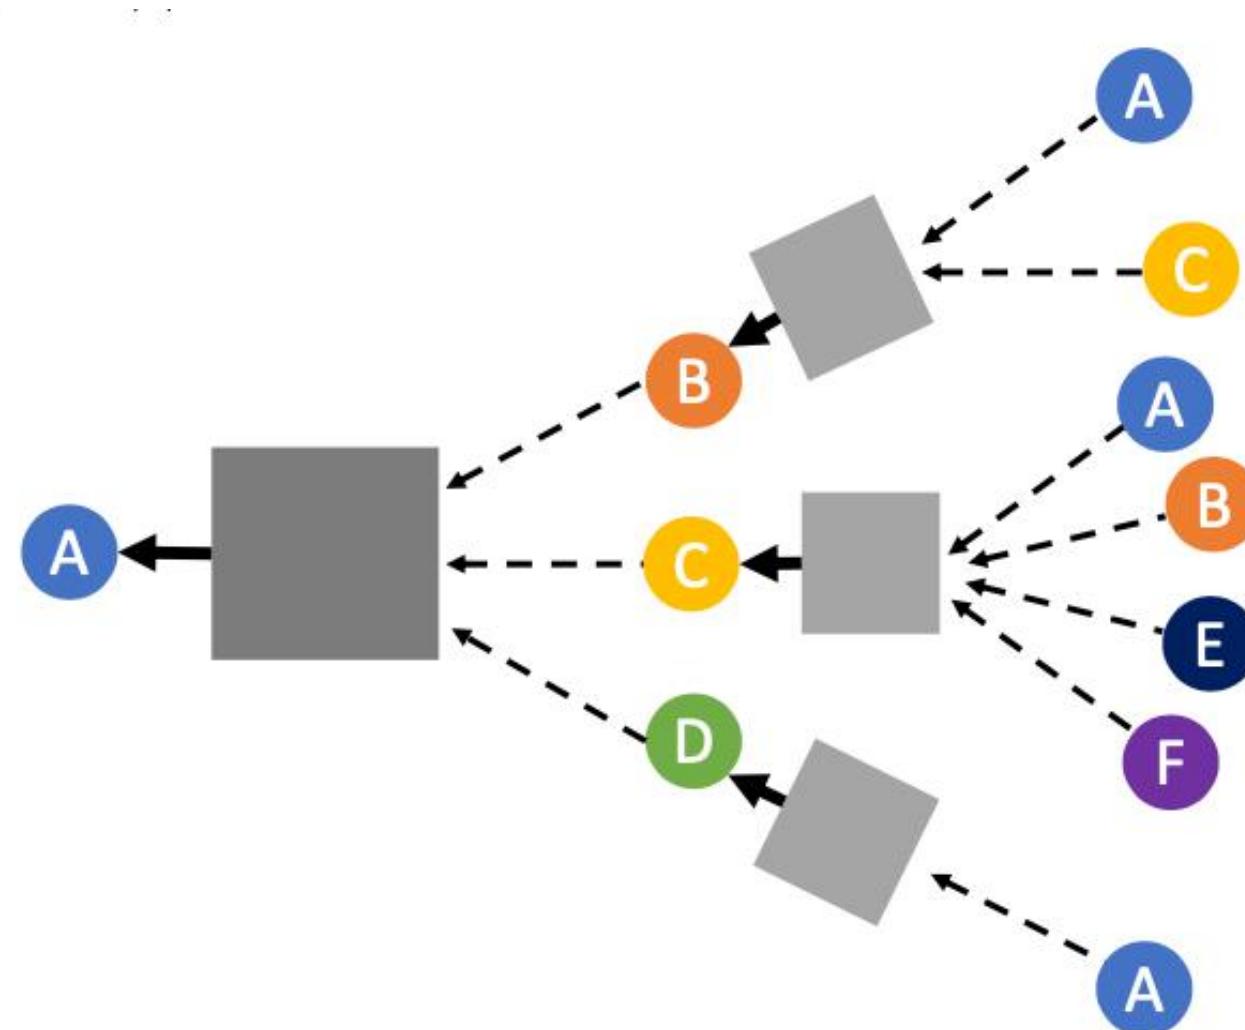
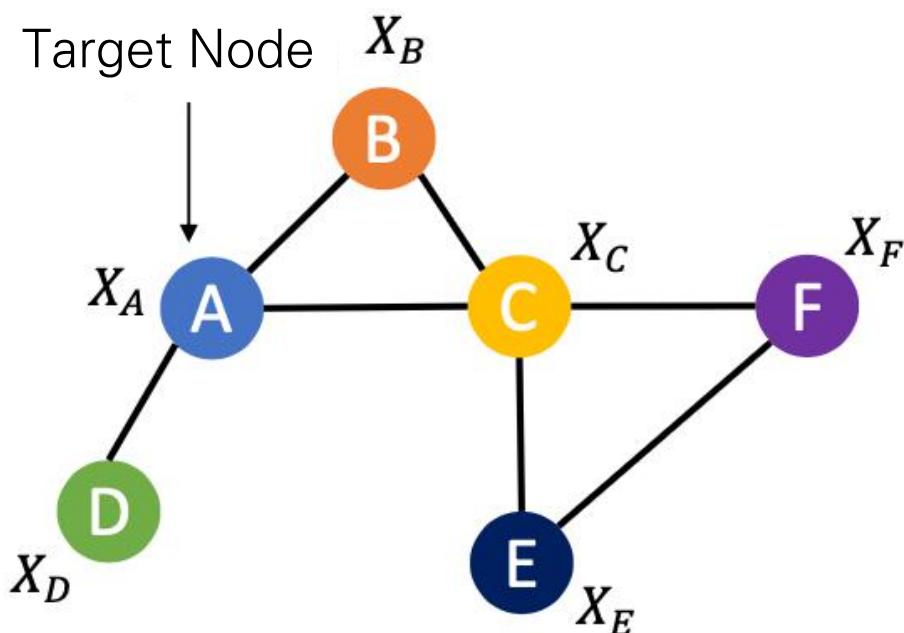
Graph-level prediction tasks



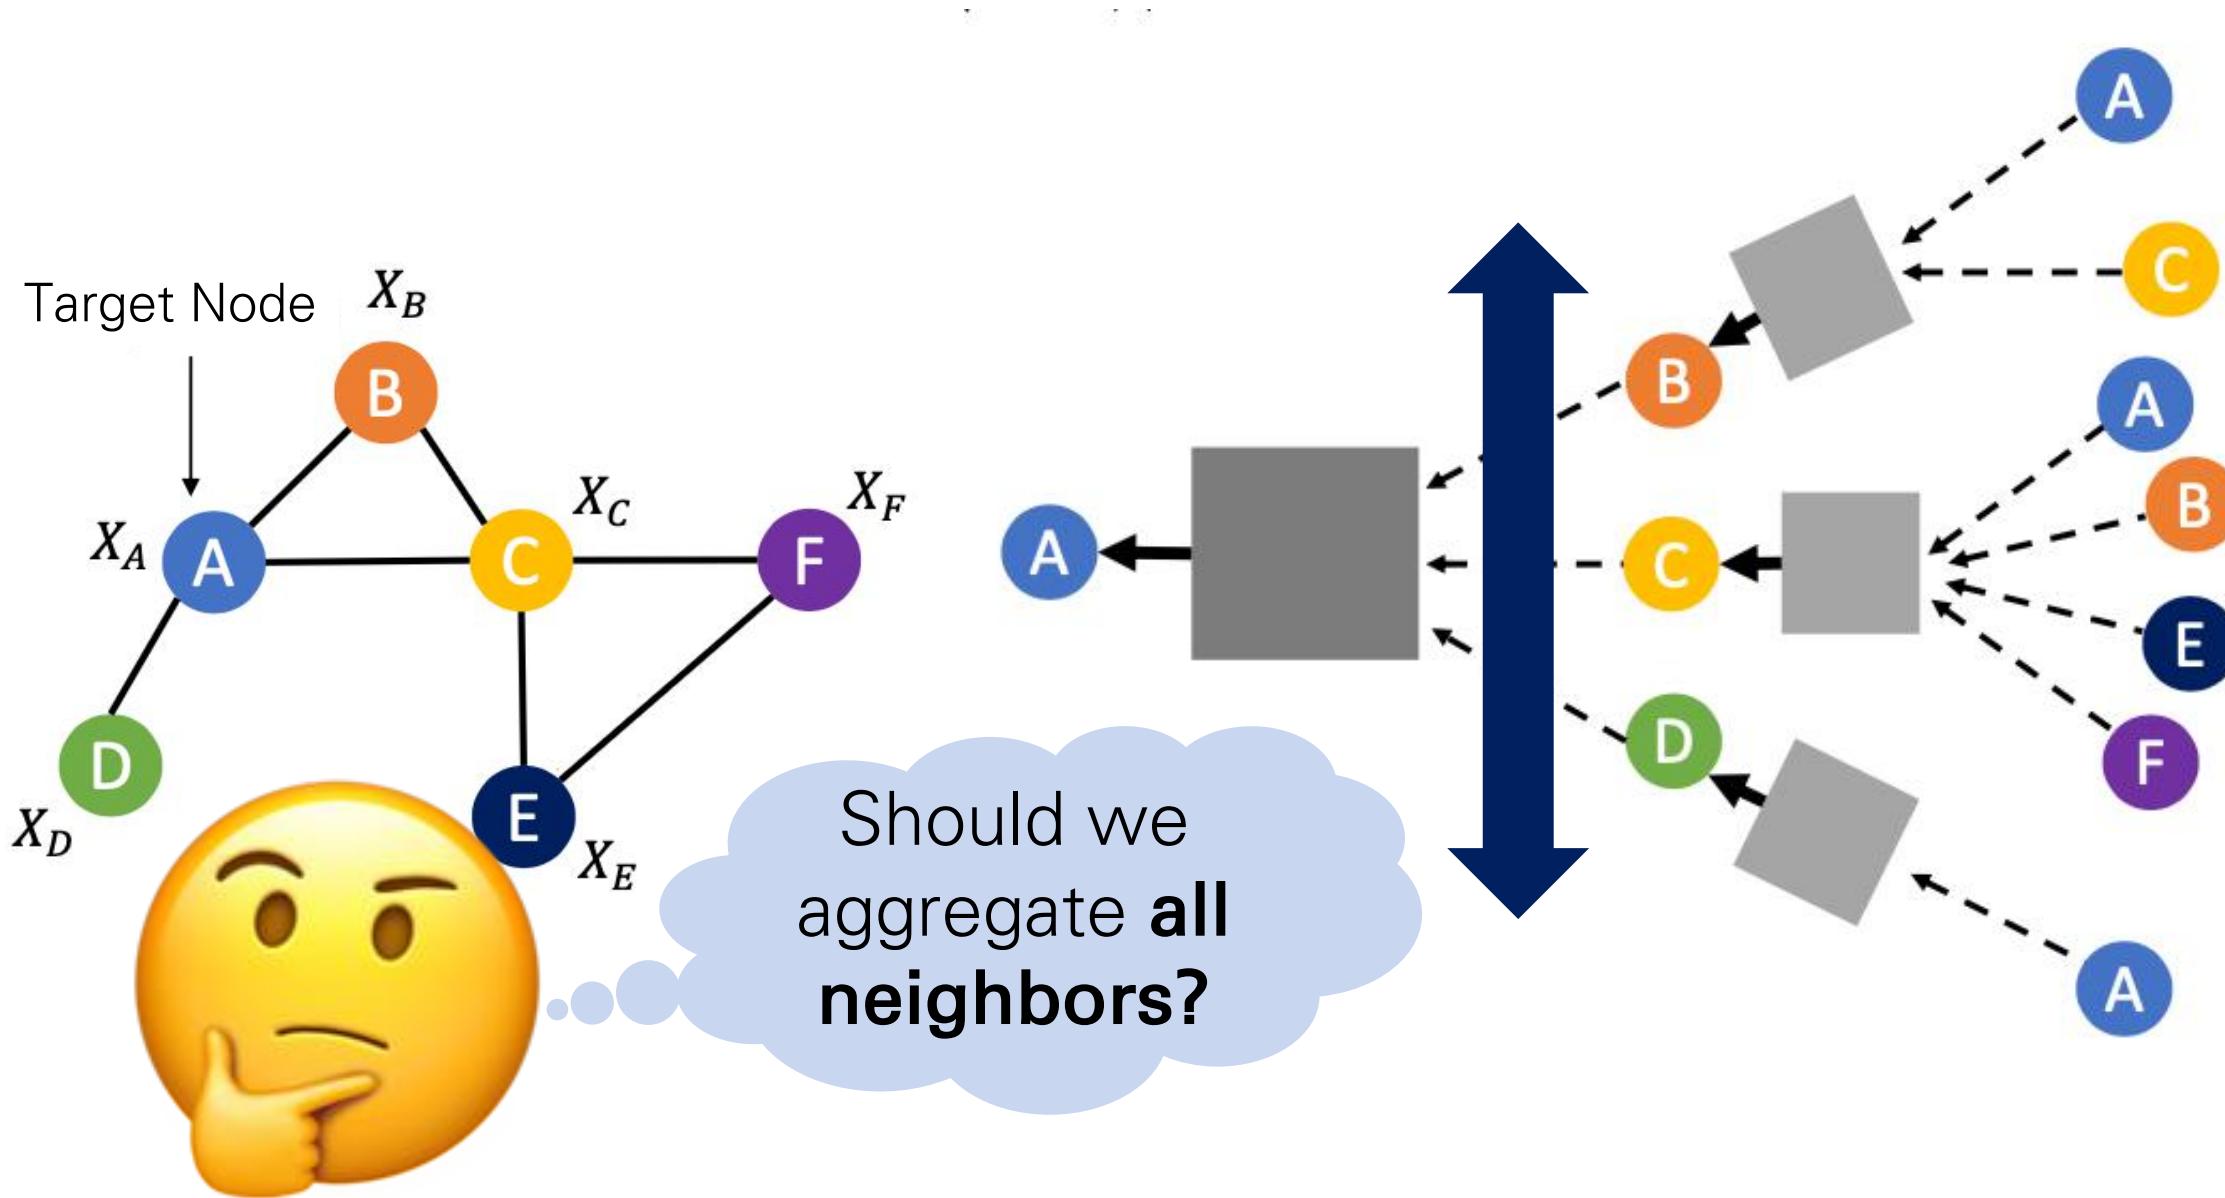
- Predict **properties of a molecule (graph)** where nodes are atoms and edges are chemical bonds

More on aggregation and Transformation operations

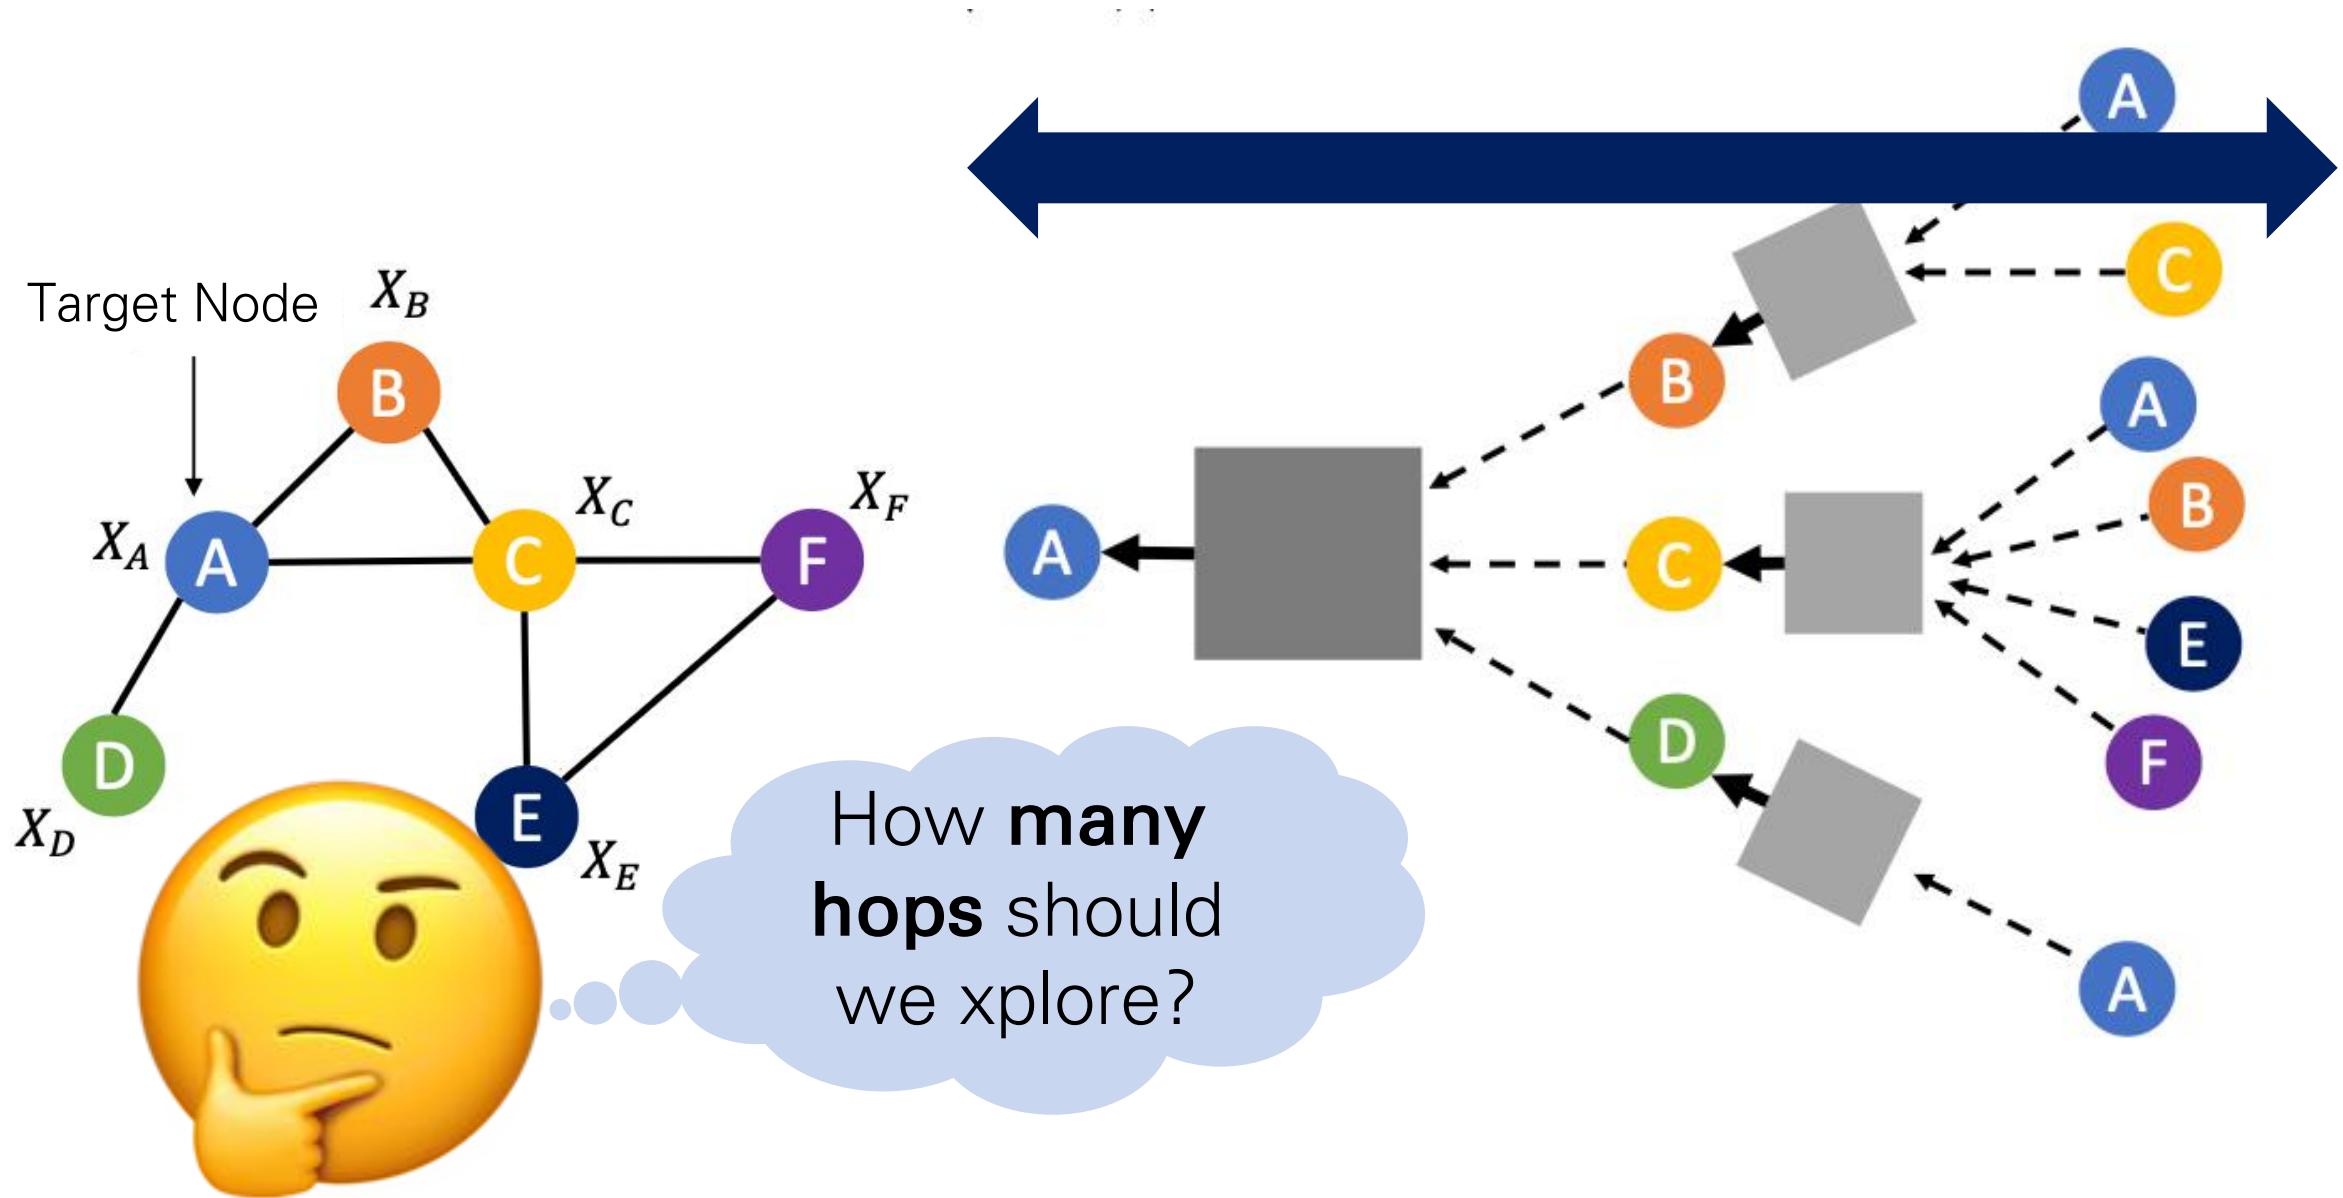
Graph Neural Networks – Width



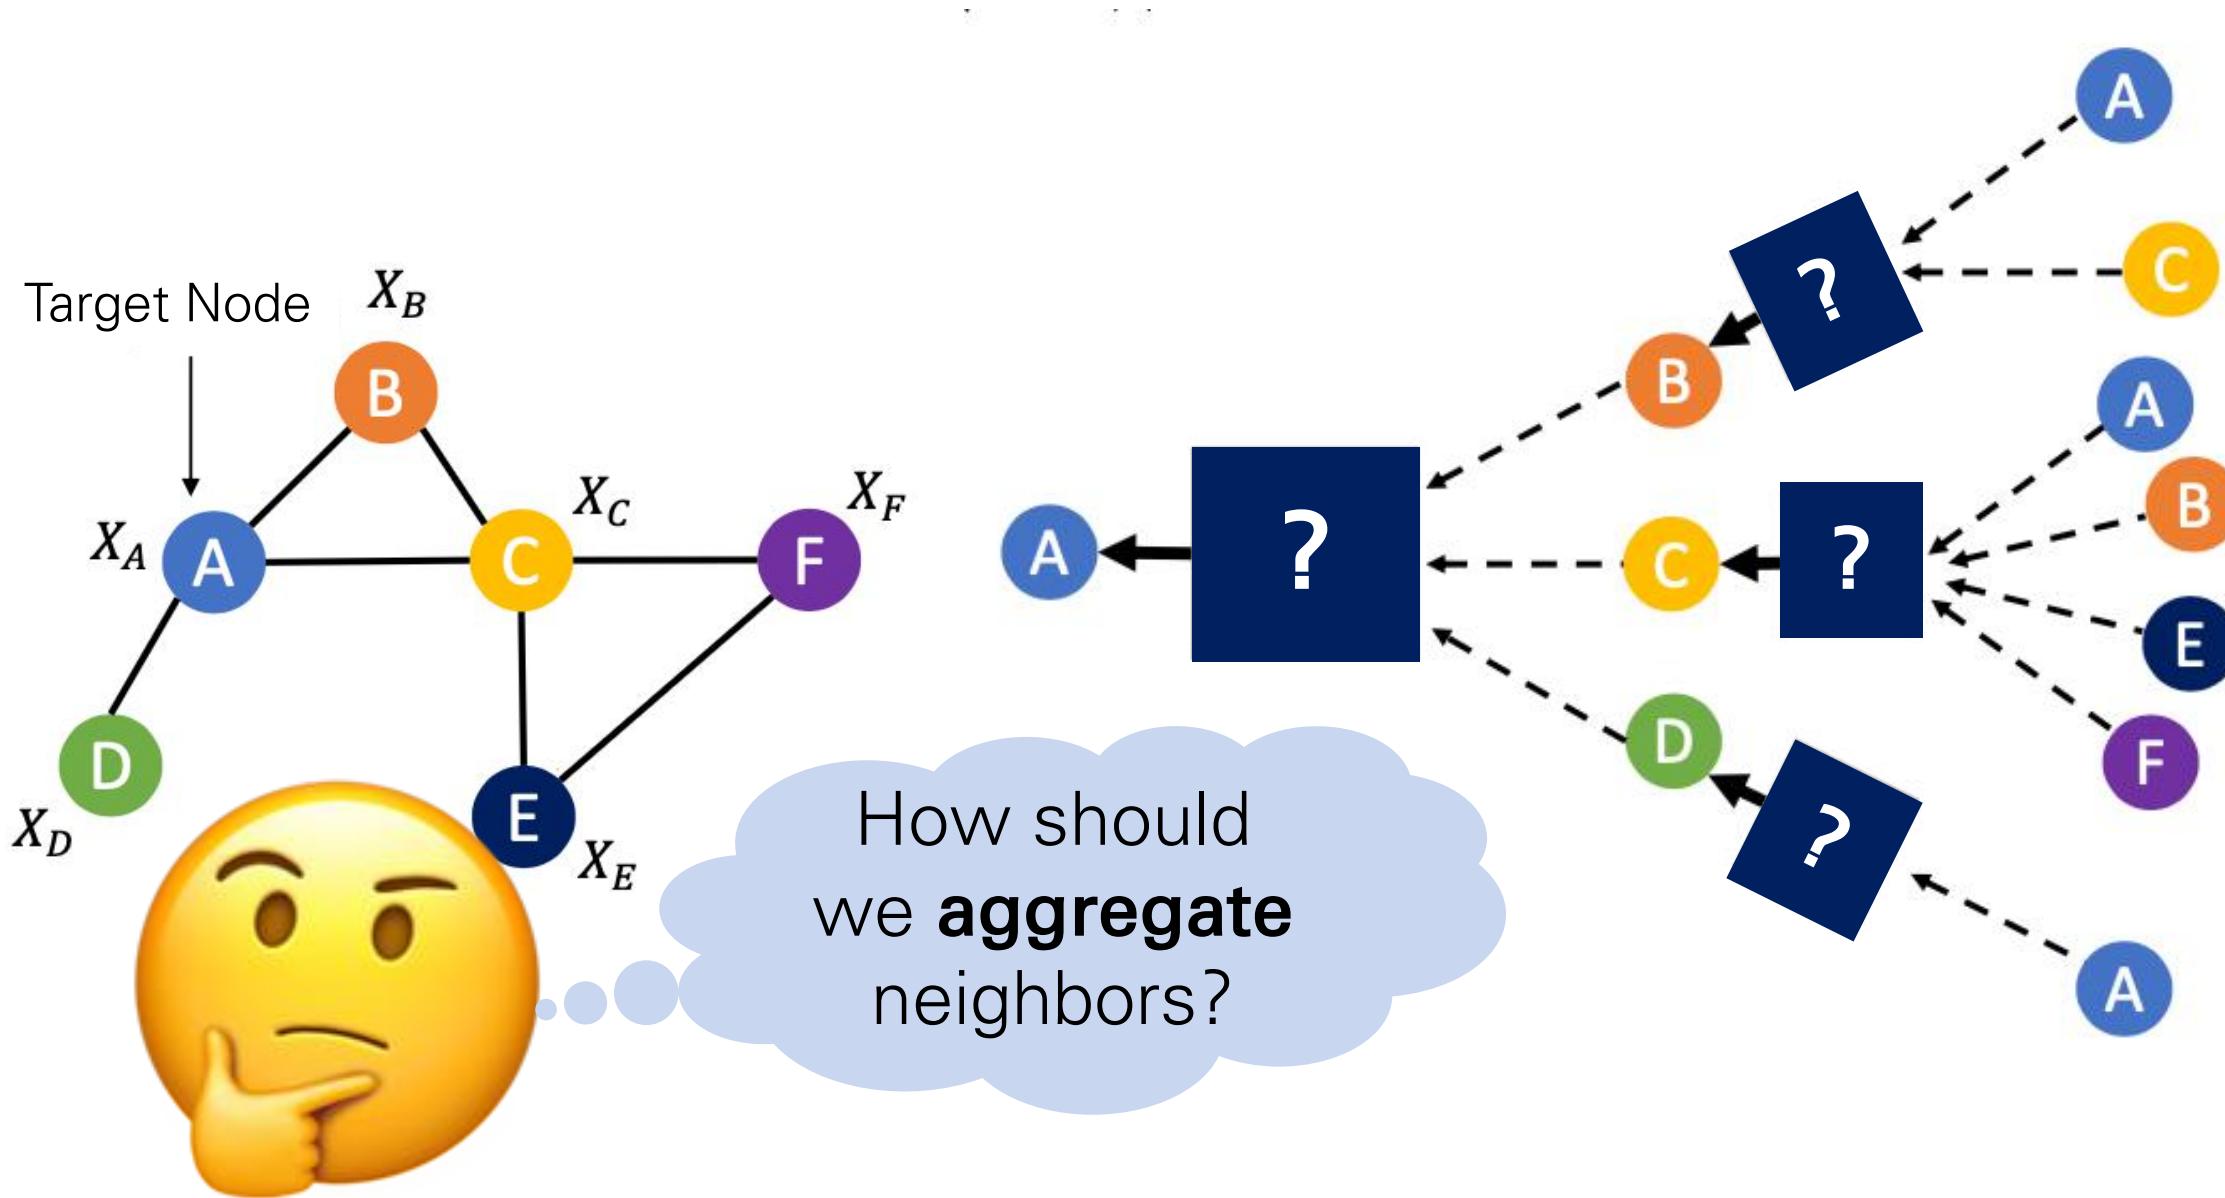
Graph Neural Networks (GNNs)



Graph Neural Networks (GNNs)

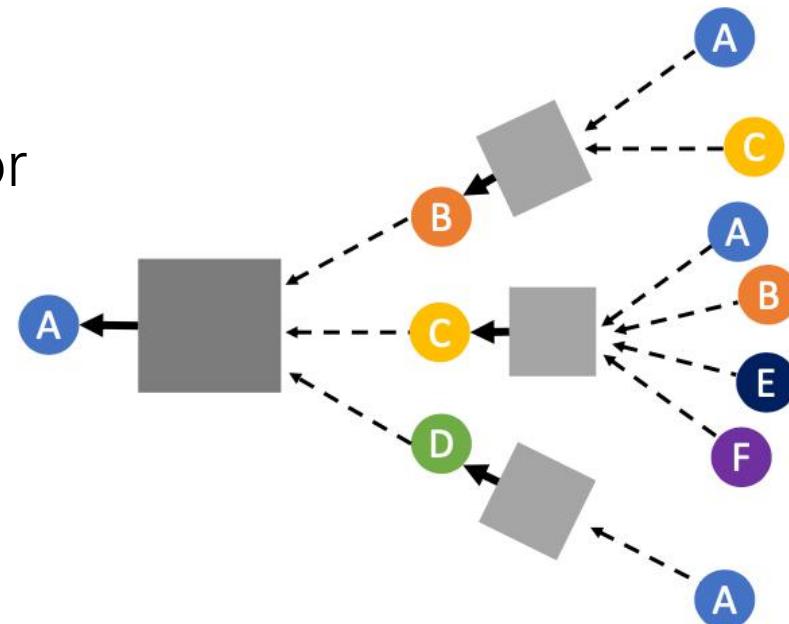


Graph Neural Networks (GNNs)



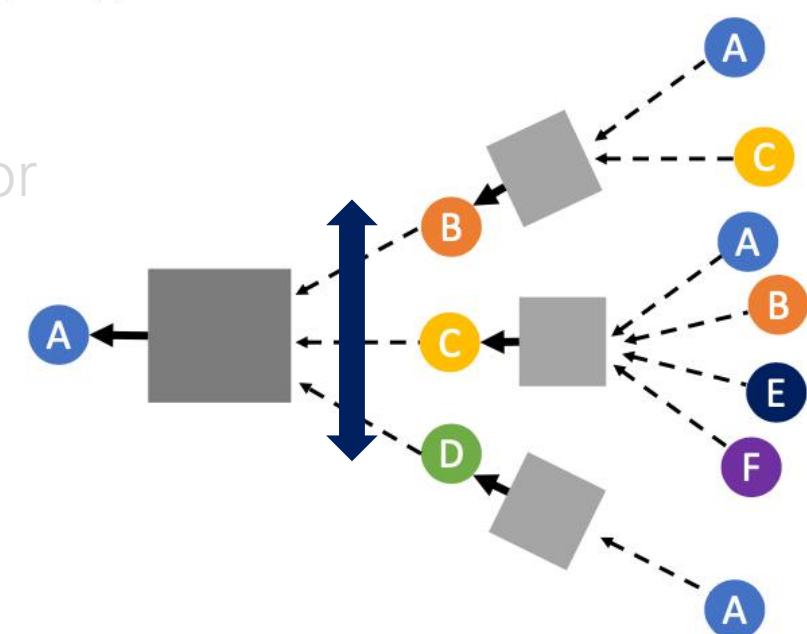
Graph Neural Network Architectures

- Width
 - Which neighbors should we aggregate messages from?
- Depth
 - How many hops should we check?
- Aggregation
 - How should we aggregate messages from neighbor



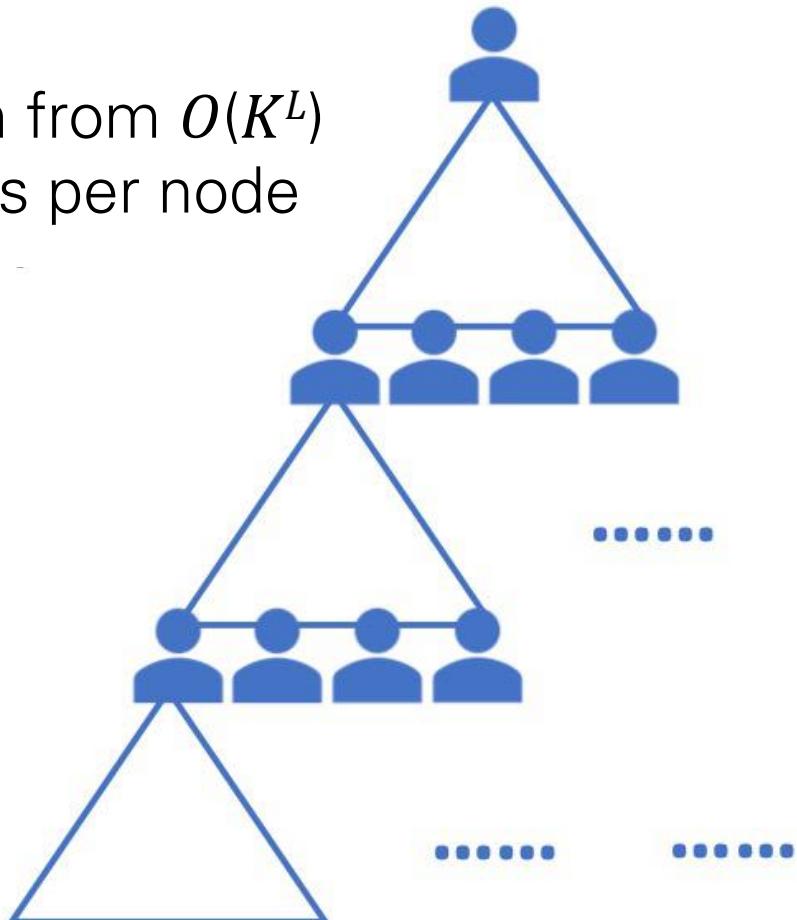
Graph Neural Network Architectures

- Width
 - Which neighbors should we aggregate messages from?
- Depth
 - How many hops should we check?
- Aggregation
 - How should we aggregate messages from neighbor



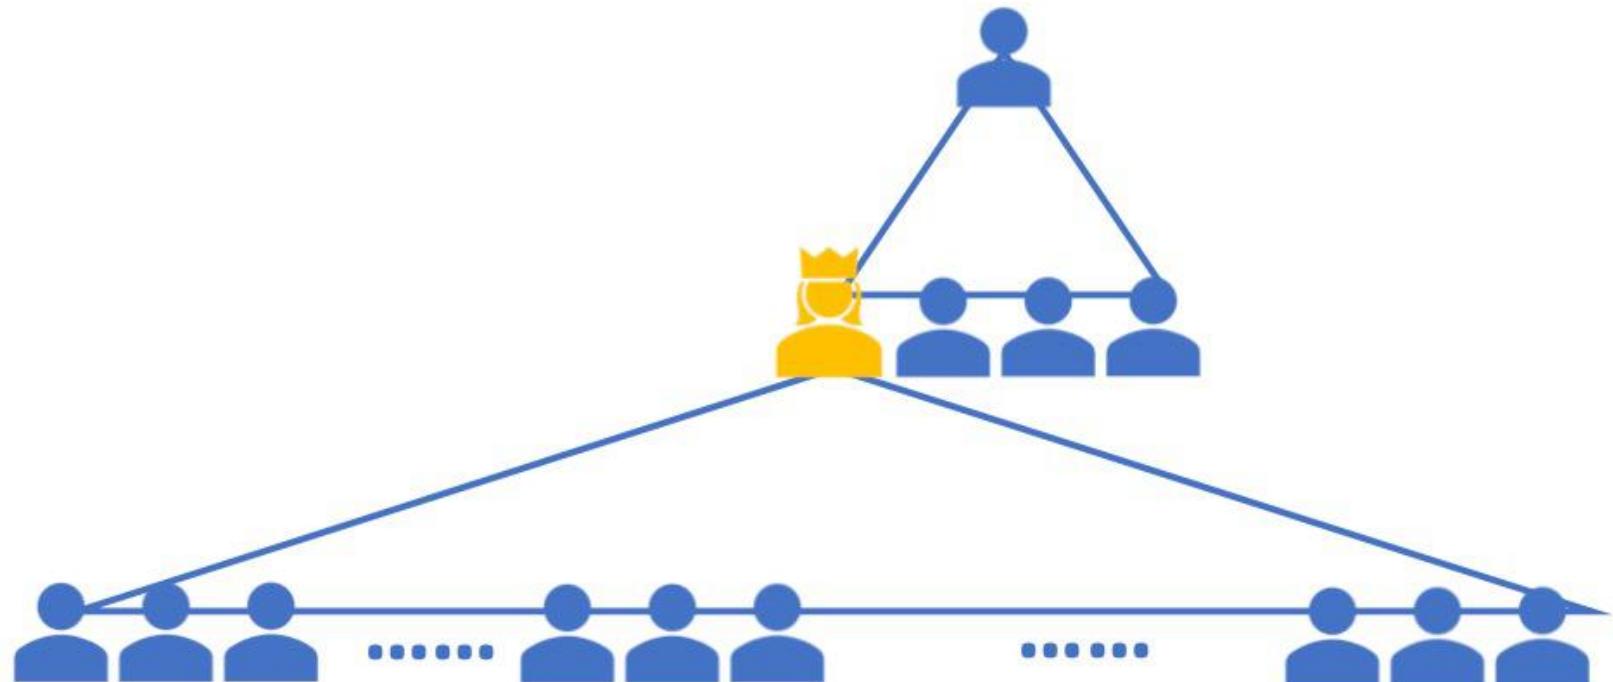
Aggregation Width in GNNs

- If we aggregate all neighbors, GNNs have scalability issues
- Neighbor explosion
 - In L -layer GNNs, one node aggregates information from $O(K^L)$ nodes where K is the average number of neighbors per node



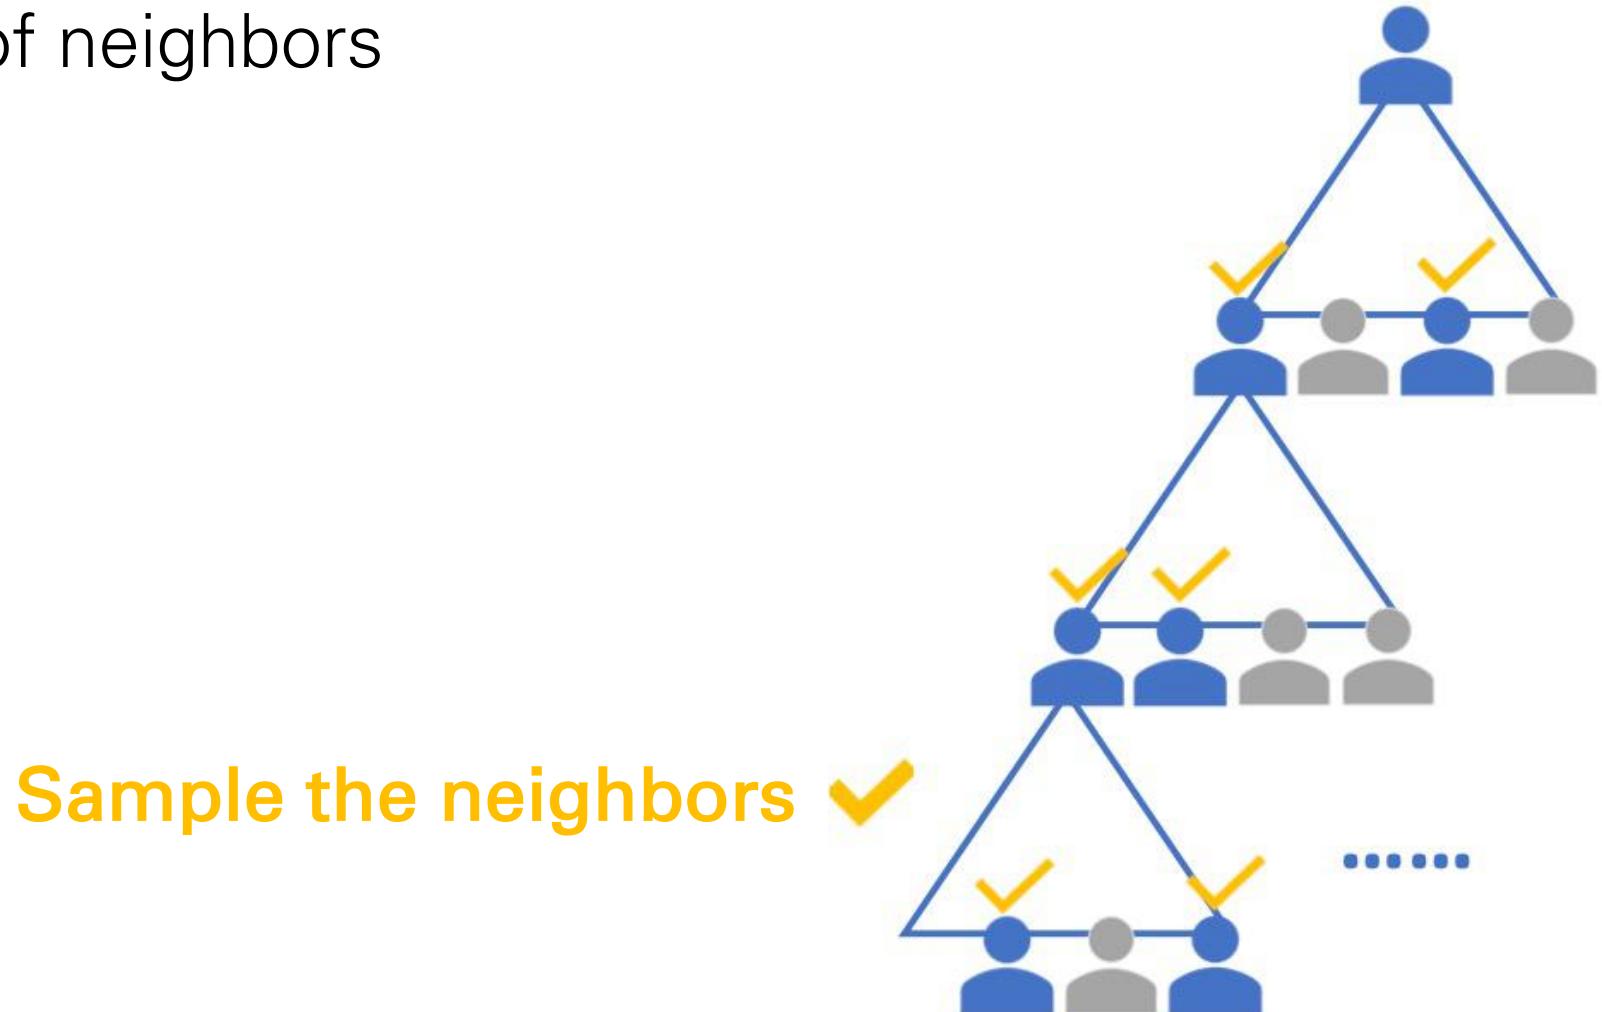
Aggregation Width in GNNs

- If we aggregate all neighbors, GNNs have scalability issues
- Neighbor explosion
 - Hub nodes who are connected to a huge number of nodes



Aggregation Width in GNNs

- Limit the neighborhood expansion by **sampling** a fixed number of neighbors



Aggregation Width in GNNs

- Random sampling
 - Assign **same** sampling probabilities to all neighbors
 - GraphSage^[4]
- Importance sampling
 - Assign **different** sampling probabilities to all neighbors
 - FastGCN^[5], LADIES^[6], AS-GCN^[7], GCN-BS^[8], PASS^[9]

[4] Will Hamilton, et al. "Inductive representation learning on large graphs"

[5] Jie Chen, et al. "Fastgcn: fast learning with graph convolutional networks via importance sampling"

[6] Difan Zou, et al. "Layer-Dependent Importance Sampling for Training Deep and Large Graph Convolutional Networks"

[7] Wenbing Huang, et al. "Adaptive sampling towards fast graph representation learning"

[8] Ziqi Liu, et al. "Bandit Samplers for Training Graph Neural Networks"

[9] Minji Yoon, et al. "Performance-Adaptive Sampling Strategy Towards Fast and Accurate Graph Neural Networks"

Aggregation Width in GNNs

Importance sampling

: assign higher sampling probabilities to neighbors who

- **Minimize variance in sampling**

- FastGCN^[5], LADIES^[6], AS-GCN^[7], GCN-BS^[8]

- **Maximize GNN performance**

- PASS^[9]

[4] Will Hamilton, et al. "Inductive representation learning on large graphs"

[5] Jie Chen, et al. "Fastgcn: fast learning with graph convolutional networks via importance sampling"

[6] Difan Zou, et al. "Layer-Dependent Importance Sampling for Training Deep and Large Graph Convolutional Networks"

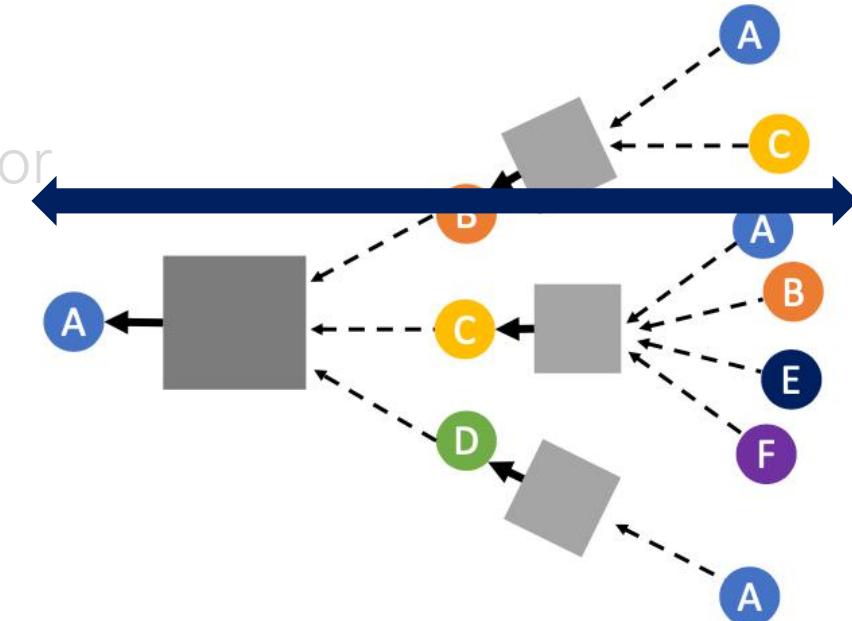
[7] Wenbing Huang, et al. "Adaptive sampling towards fast graph representation learning"

[8] Ziqi Liu, et al. "Bandit Samplers for Training Graph Neural Networks"

[9] Minji Yoon, et al. "Performance-Adaptive Sampling Strategy Towards Fast and Accurate Graph Neural Networks"

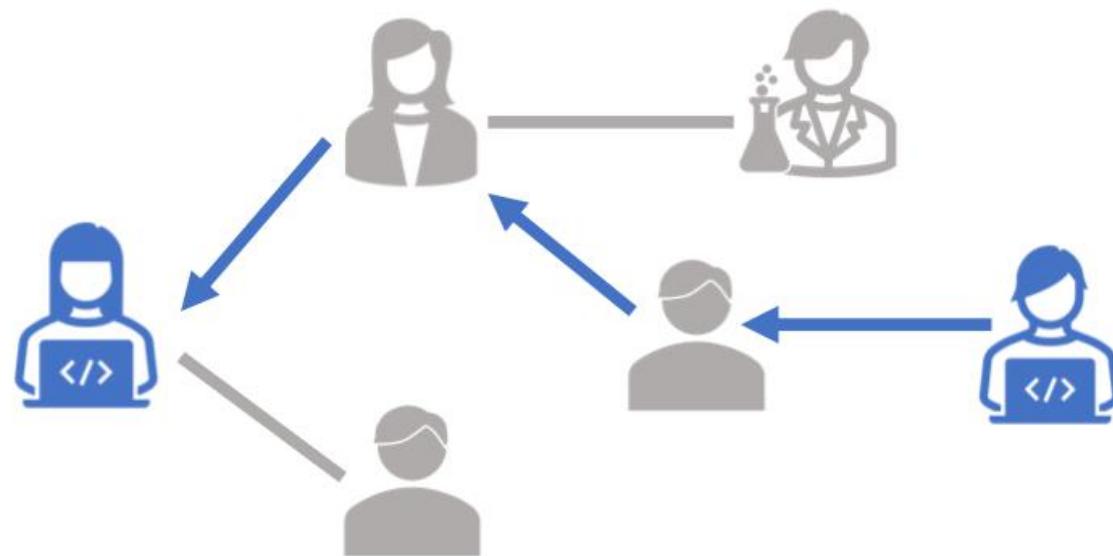
Graph Neural Network Architectures

- Width
 - Which neighbors should we aggregate messages from?
- Depth
 - How many hops should we check?
- Aggregation
 - How should we aggregate messages from neighbor



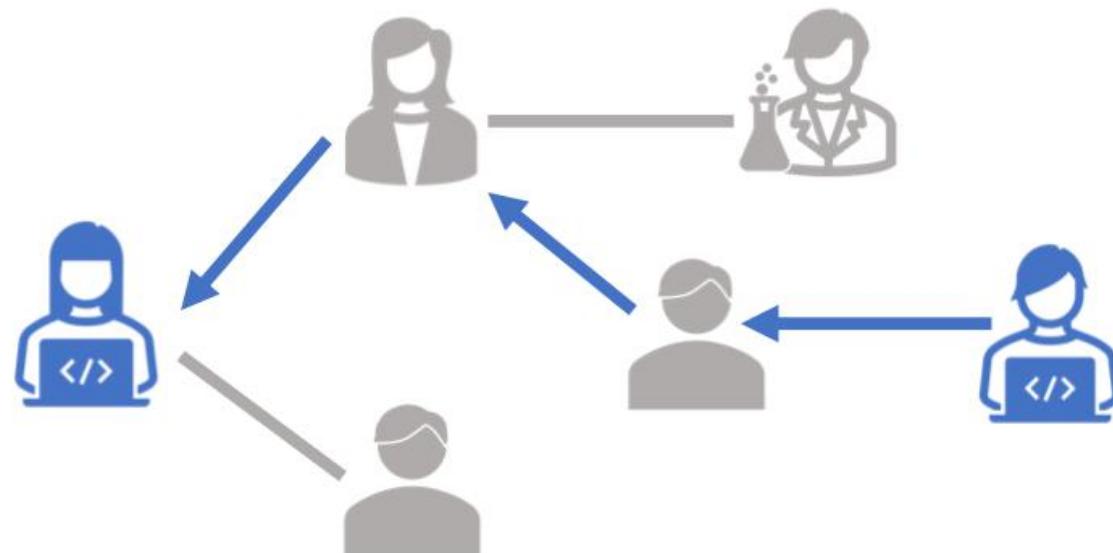
Aggregation Depth in GNNs

- Informative neighbors could be indirectly connected with a target node



Aggregation Depth in GNNs

- Informative neighbors could be indirectly connected with a target node
- Can't we just look multiple hops away from the target node?



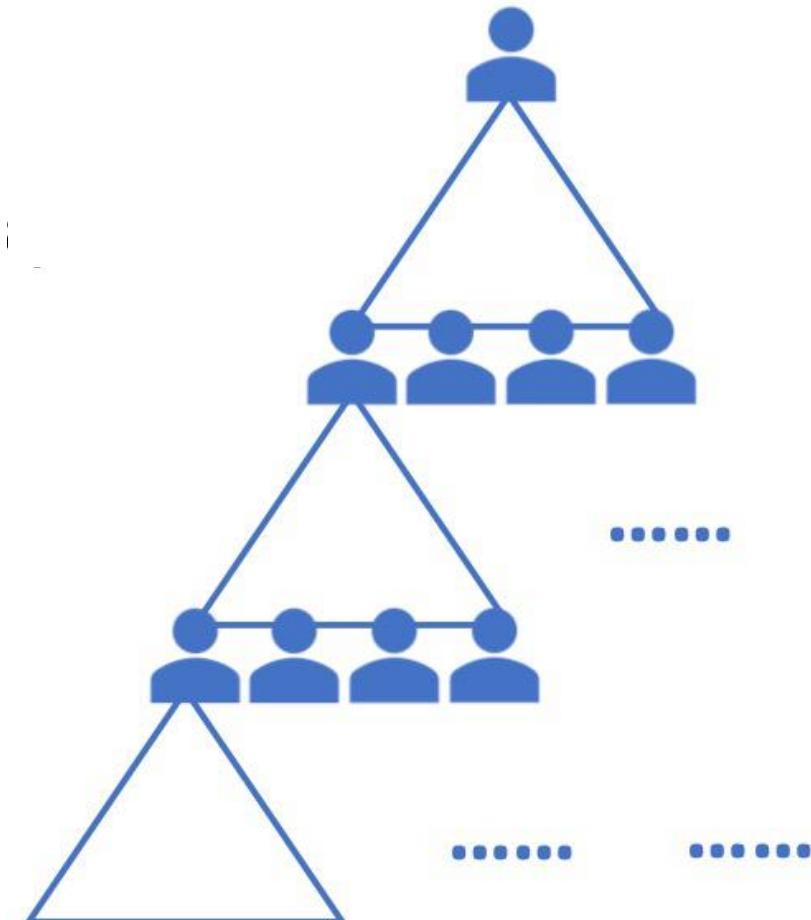
Aggregation Depth in GNNs

- 2-layer or 3-layer GNNs are commonly used in real worlds



Aggregation Depth in GNNs

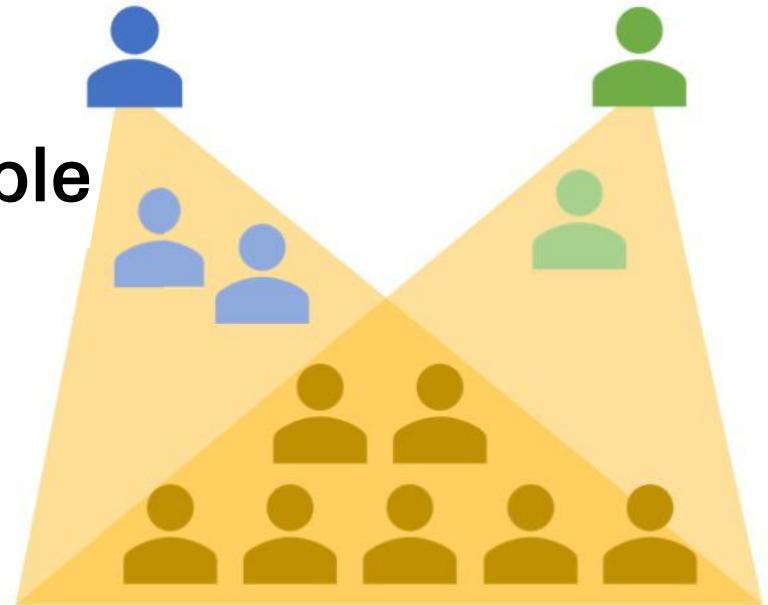
- When we increase the depth L more than this, GNNs face neighbor explosion $O(K^L)$
 - Over-smoothing
 - Over-squashing



Aggregation Depth in GNNs

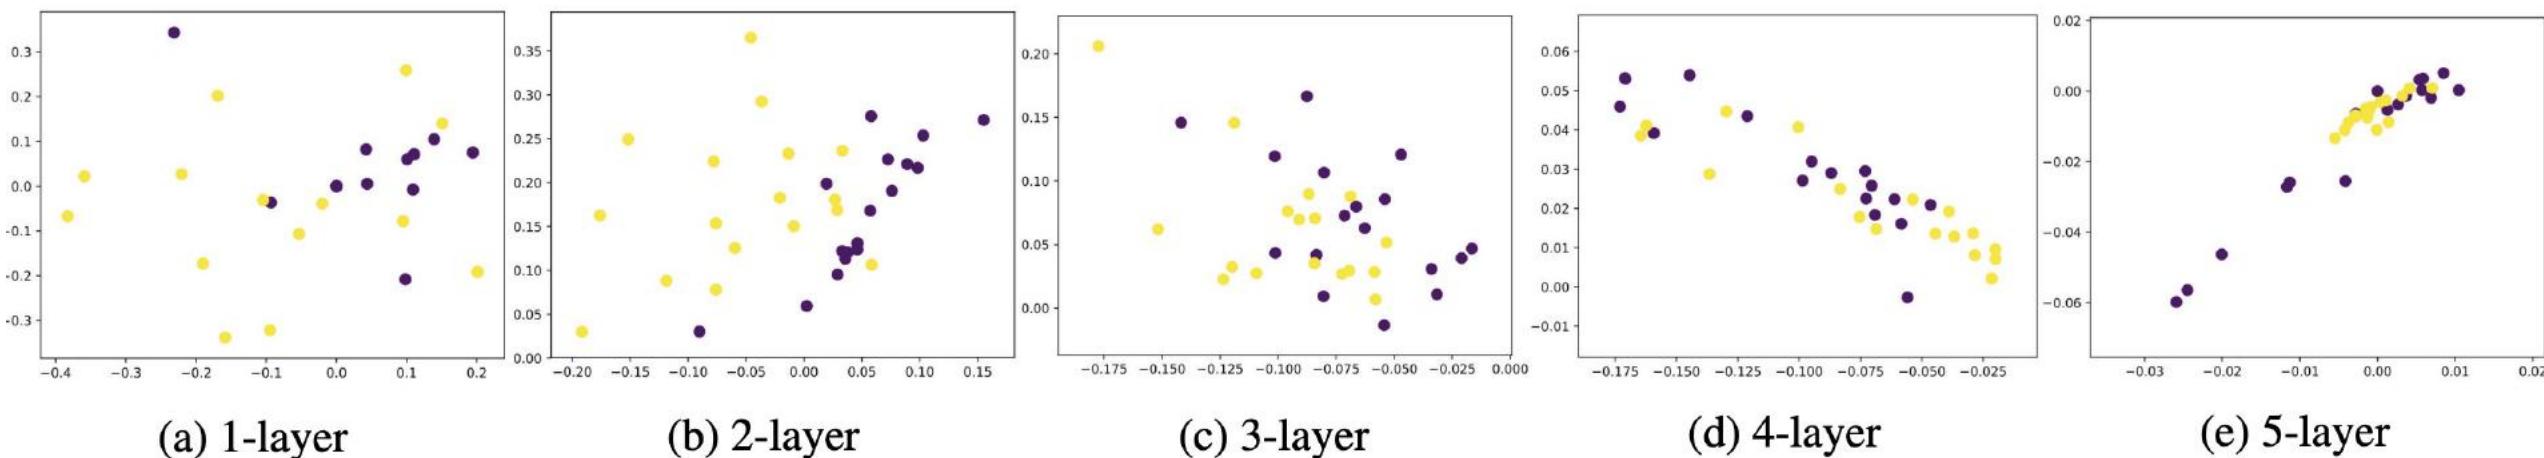
Over-smoothing^[10]

- When GNNs become deep, nodes share many neighbors
- Node embeddings become **indistinguishable**



Aggregation Depth in GNNs

- Over-smoothing^[10]
- Node embeddings of Zachary's karate club network with GNNs



[10] Qimai Li, et al. "Deeper Insights into Graph Convolutional Networks for Semi-Supervised Learning"

Aggregation Depth in GNNs

Mitigate over-smoothing

PairNorm^[11]

- Keep total pairwise squared distance (TPSD) **constant** across layers
- Push away pairs that are not connected

$$\text{TPSD}(\dot{\mathcal{X}}) = \sum_{(i,j) \in \mathcal{E}} \|\dot{x}_i - \dot{x}_j\|_2^2 + \sum_{(i,j) \notin \mathcal{E}} \|\dot{x}_i - \dot{x}_j\|_2^2 = c$$

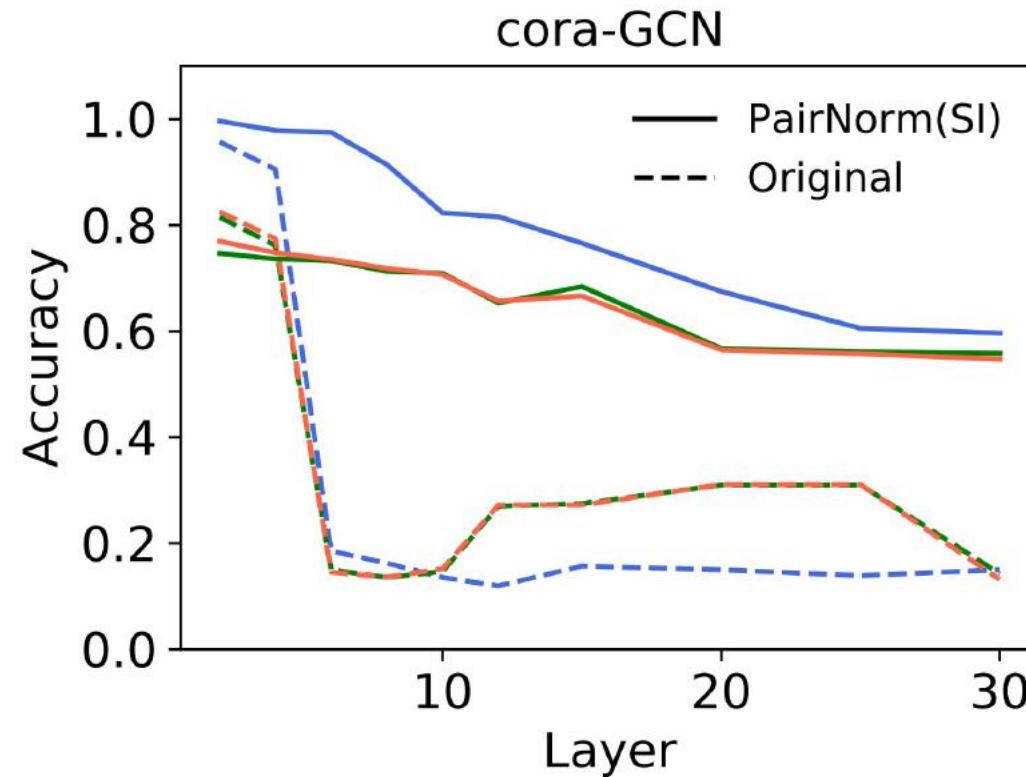
Connected pairs Disconnected pairs

The diagram illustrates the decomposition of the Total Pairwise Squared Distance (TPSD) into two components. The first component, labeled 'Connected pairs', is represented by a red dotted rectangle enclosing the sum of squared distances between nodes that are directly connected in the graph. The second component, labeled 'Disconnected pairs', is represented by a blue dotted rectangle enclosing the sum of squared distances between nodes that are not connected. The total TPSD is the sum of these two components, which is set to a constant value c .

Aggregation Depth in GNNs

Mitigate over-smoothing

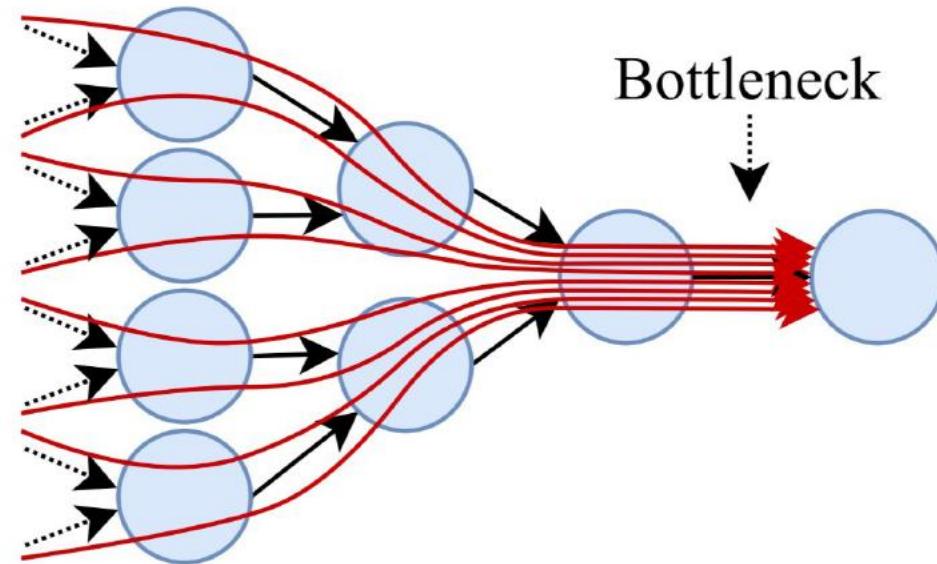
PairNorm^[11]



Aggregation Depth in GNNs

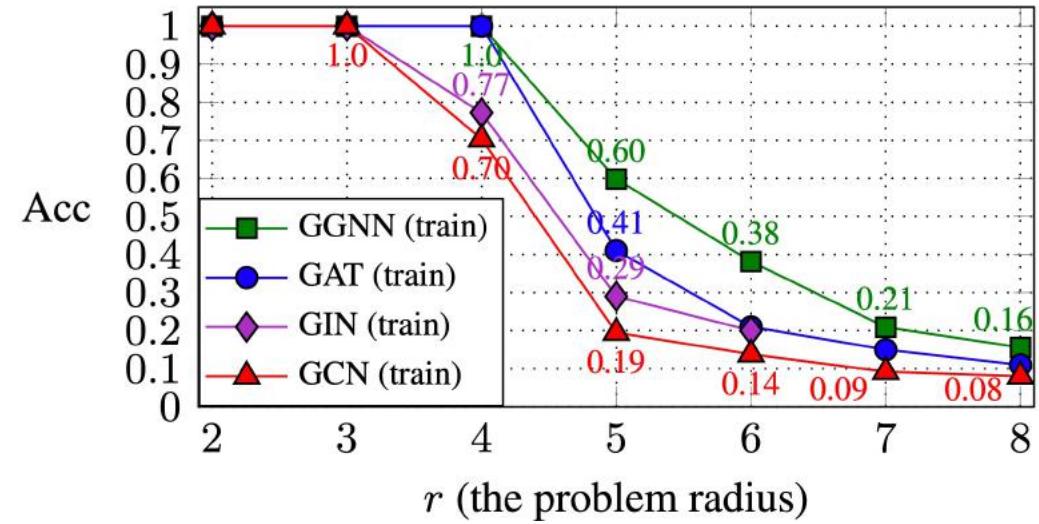
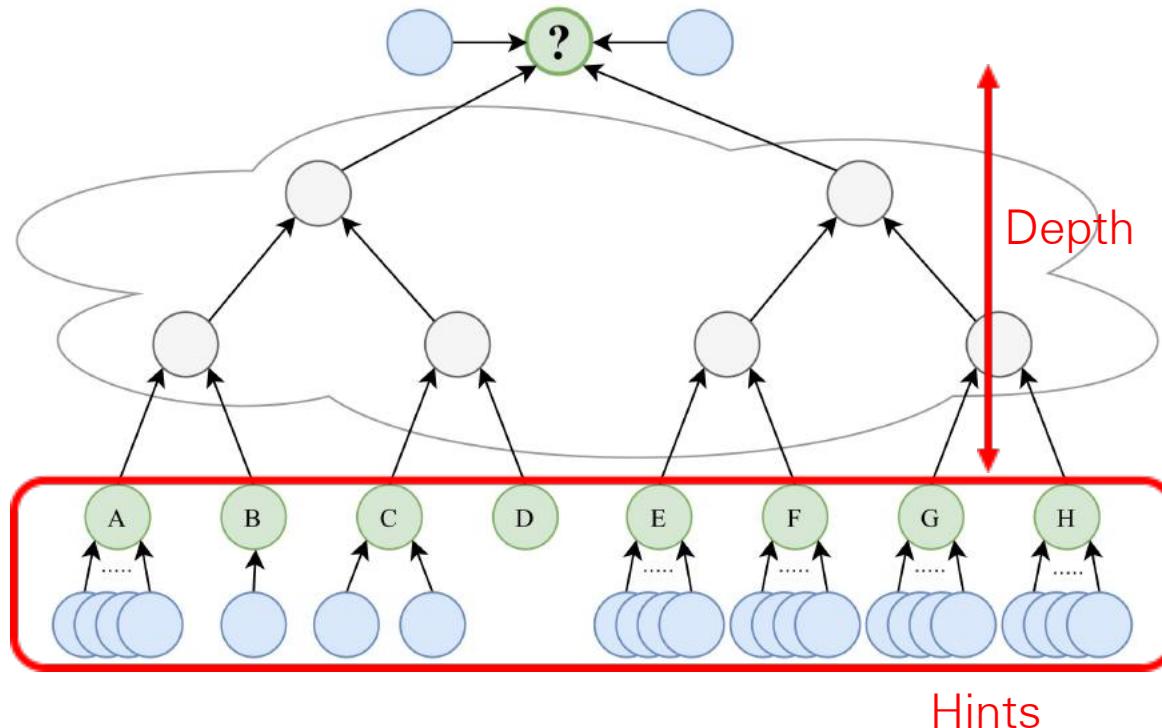
Over-squashing^[12]

- A node's exponentially-growing neighborhood is compressed into a fixed-size vector



Aggregation Depth in GNNs

Over-squashing^[12]



Aggregation Depth in GNNs

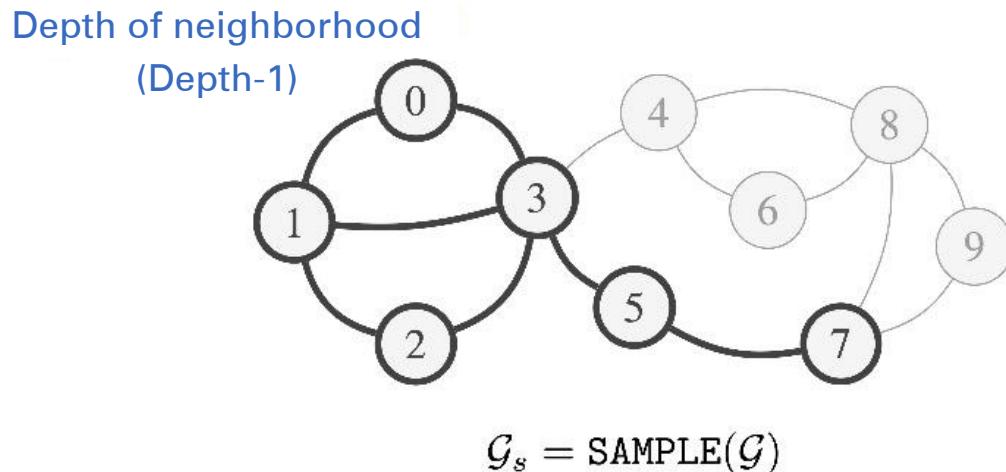
Decoupling the two concepts of depths in GNNs^[13]

- **Depth-1:** neighborhood that each node aggregates information from
- **Depth-2:** number of layers in GNNs

Aggregation Depth in GNNs

Decoupling the two concepts of depths in GNNs^[13]

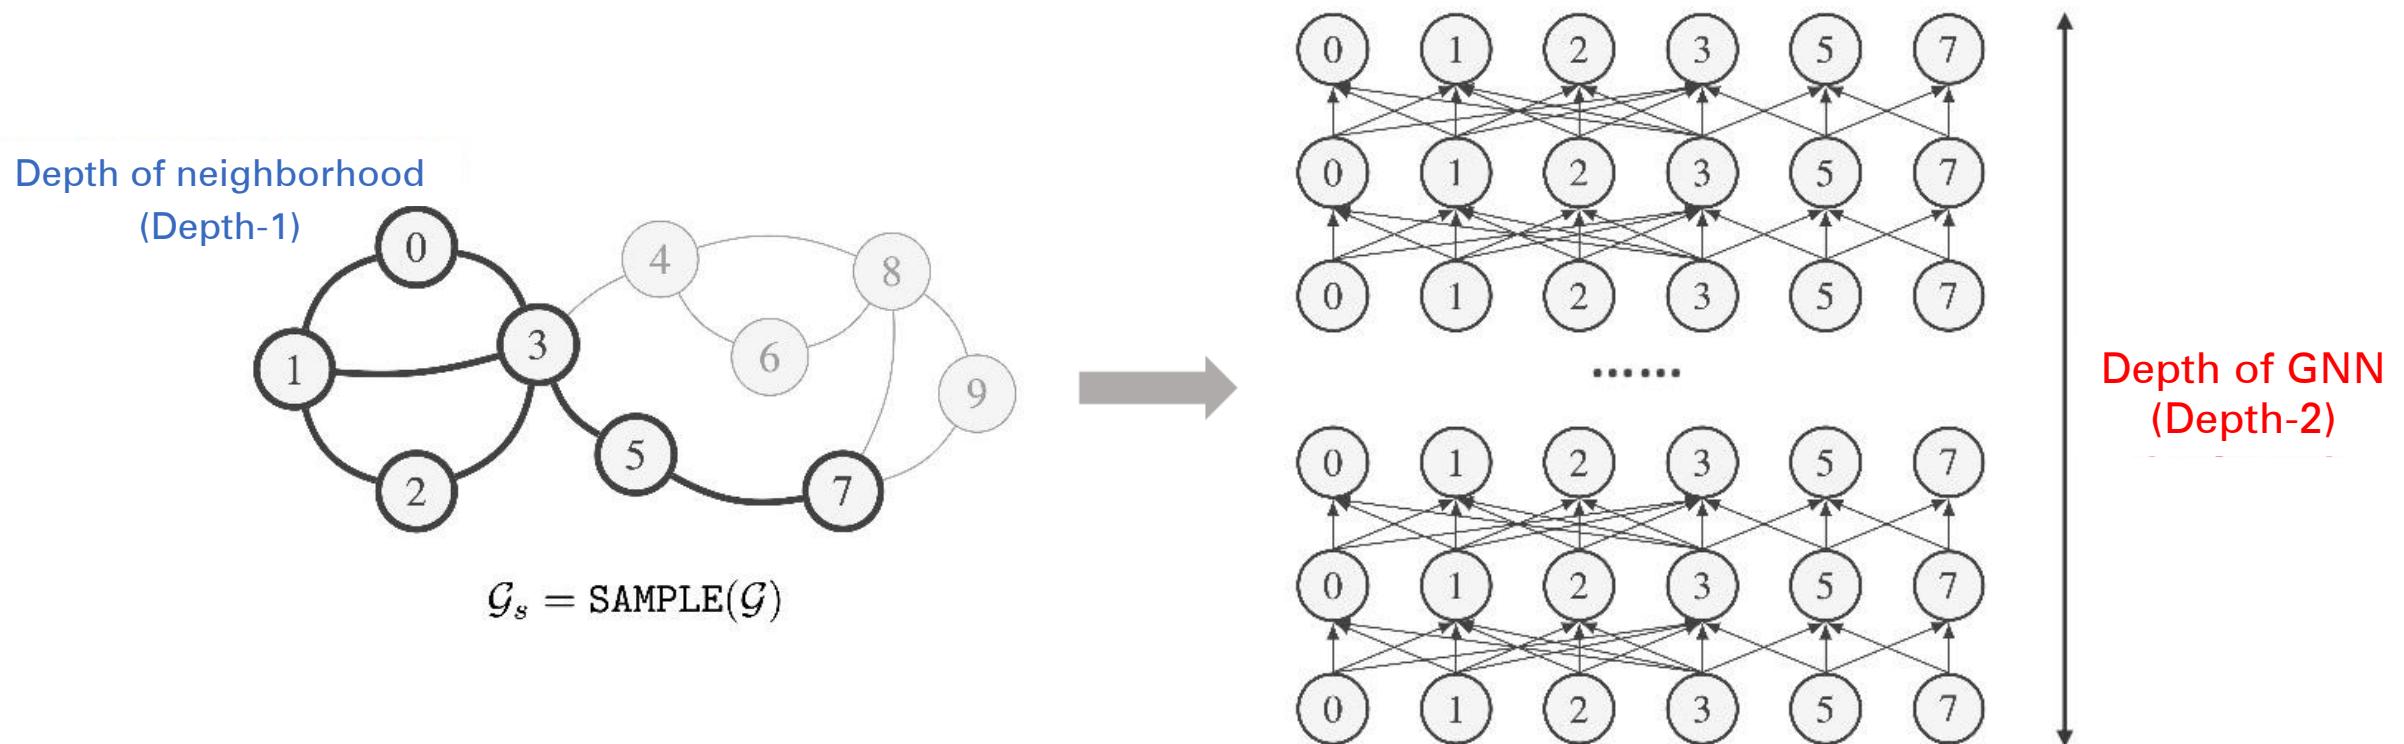
- Depth-1: neighborhood that each node aggregates information from
- Depth-2: number of layers in GNNs



Aggregation Depth in GNNs

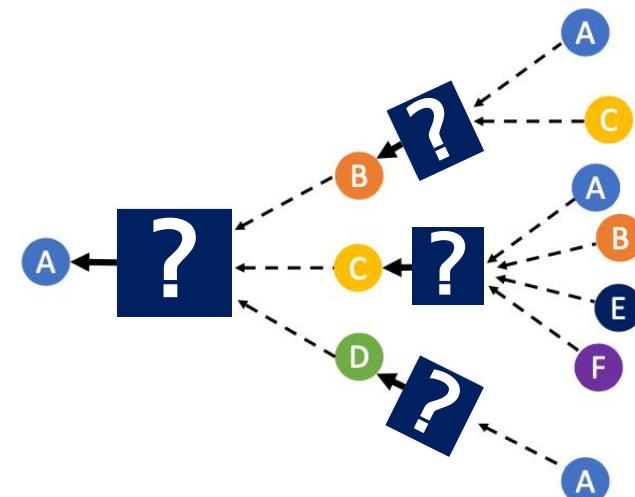
Decoupling the two concepts of depths in GNNs^[13]

- Depth-1: neighborhood that each node aggregates information from
- Depth-2: number of layers in GNNs



Graph Neural Network Architectures

- Width
 - Which neighbors should we aggregate messages from?
- Depth
 - How many hops should we check?
- Aggregation
 - How should we aggregate messages from neighbor



Aggregation strategy in GNNs

In each layer l :

Aggregate over neighbors

$$m_v^{(l-1)} = \boxed{f}^{(l)} \left(h_v^{(l-1)}, \{h_u^{(l-1)} : u \in \mathcal{N}(v)\} \right)$$

Transform messages

$$h_v^{(l)} = \boxed{g}^{(l)}(m_v^{(l-1)})$$

Aggregation strategy in GNNs

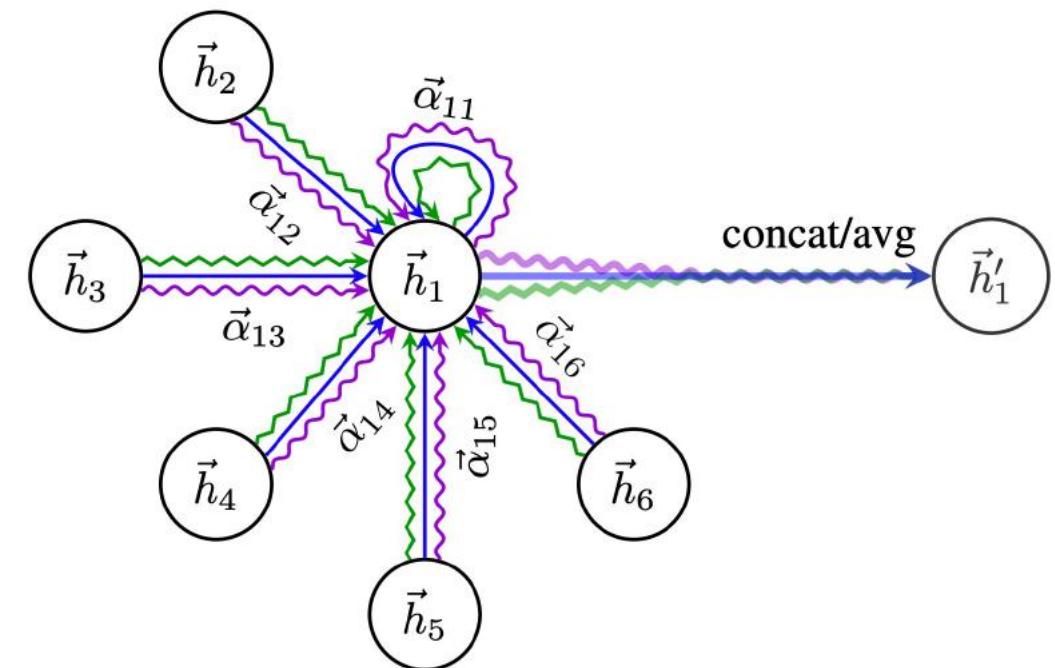
- GCN^[1]
 - Average embeddings of neighboring nodes

[1] Kipf, Thomas N., et al. "Semi-supervised classification with graph convolutional networks."

Aggregation strategy in GNNs

- GAT^[14]
 - Different weights to different nodes in a neighborhood
 - Multi-head attention

$$\alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{a}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{a}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_k] \right) \right)}$$



Aggregation strategy in GNNs

In each layer l :

Aggregate over neighbors

$$m_v^{(l-1)} = \boxed{f^{(l)}} \left(h_v^{(l-1)}, \{h_u^{(l-1)} : u \in \mathcal{N}(v)\} \right)$$

Core part of GNNs

Transform messages

$$h_v^{(l)} = \boxed{g^{(l)}}(m_v^{(l-1)})$$

Any neural network module can fit in
1-layer MLP is commonly used

Aggregation strategy in GNNs

Power of **GNNs**

=

Power of **aggregation strategies**

Aggregation strategy in GNNs

- By measuring the power of GNNs, we can find the best aggregation strategy!!



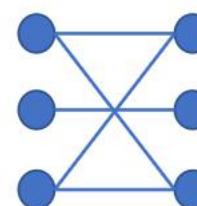
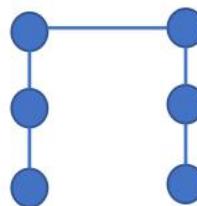
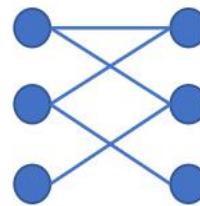
Aggregation strategy in GNNs

- By measuring the power of GNNs, we can find the best aggregation strategy!!
- But.. what is the power of GNNs and how can we measure it?



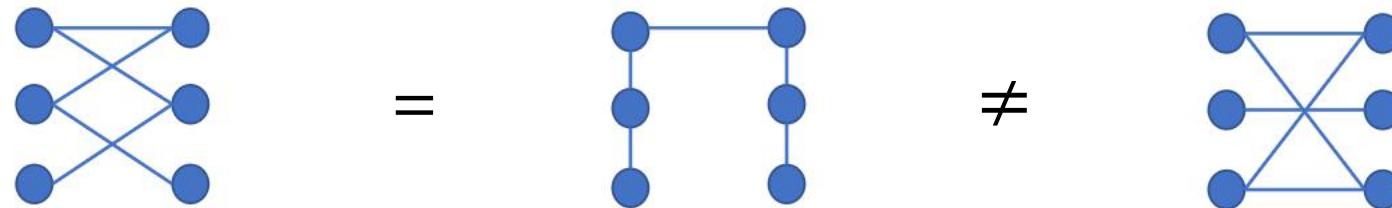
Aggregation strategy in GNNs

- How powerful are Graph Neural Networks?^[2]
- Metric
 - Graph-level prediction task
 - Can a GNN model distinguish two non-isomorphic graphs?



Aggregation strategy in GNNs

- How powerful are Graph Neural Networks?^[2]
- Metric
 - Graph-level prediction task
 - Can a GNN model distinguish two non-isomorphic graphs?



Aggregation strategy in GNNs

- How powerful are Graph Neural Networks?^[2]
 - Any aggregation-based GNN is at most as powerful as the **WL test**^[15]
 - Maximum power = aggregation strategy is injective

Weisfeiler-Lehman (WL)
graph isomorphism test

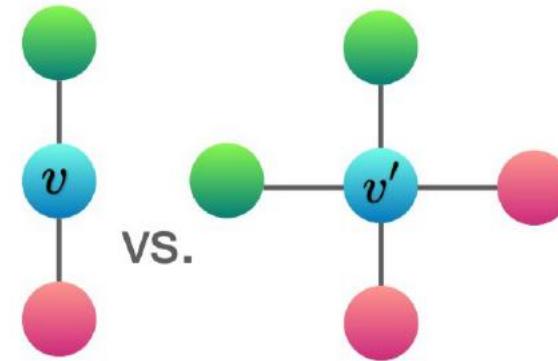
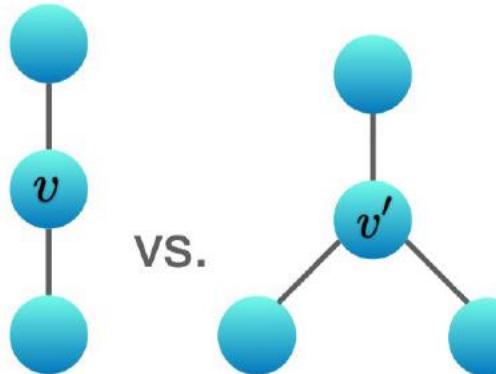
$$\mathbf{f}(\mathbf{x}_1) = \mathbf{f}(\mathbf{x}_2) \Rightarrow \mathbf{x}_1 = \mathbf{x}_2$$


[2] Keyulu Xu., et al. "How Powerful are Graph Neural Networks?"

[15] Boris Weisfeiler and AA Leman. "A reduction of a graph to a canonical form and an algebra arising during this reduction"

Aggregation strategy in GNNs

- How powerful are Graph Neural Networks?^[2]
 - Any aggregation-based GNN is at most as powerful as the **WL test**^[15]
 - Maximum power = aggregation strategy is injective
 - (ex) summation



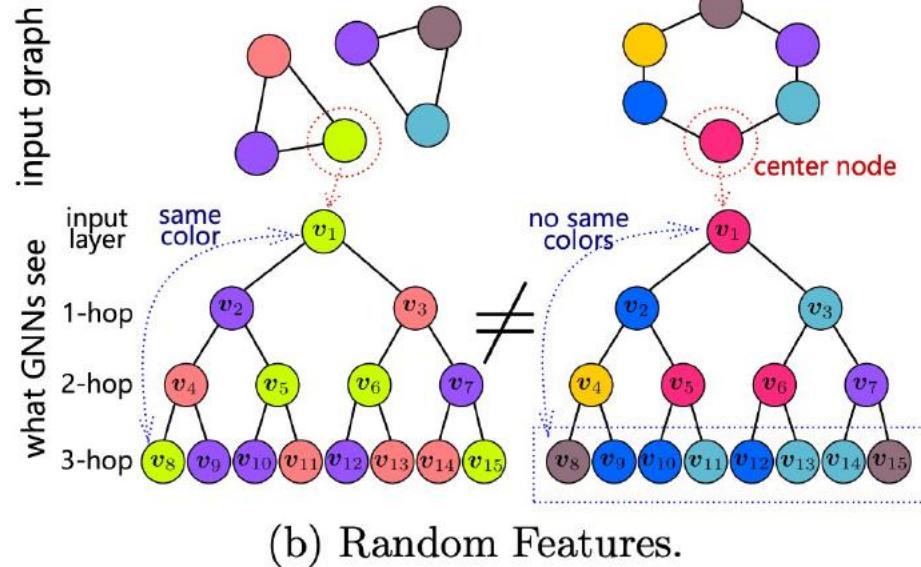
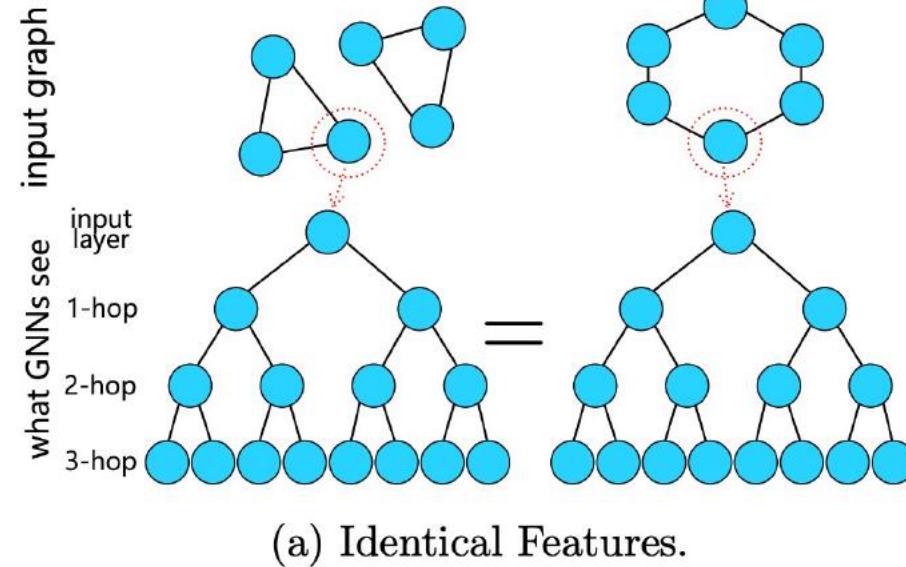
Mean and Max both fail, while Sum can distinguish them!!

Aggregation strategy in GNNs

- Can we make more powerful GNNs?
 - Very active area, with many open problems

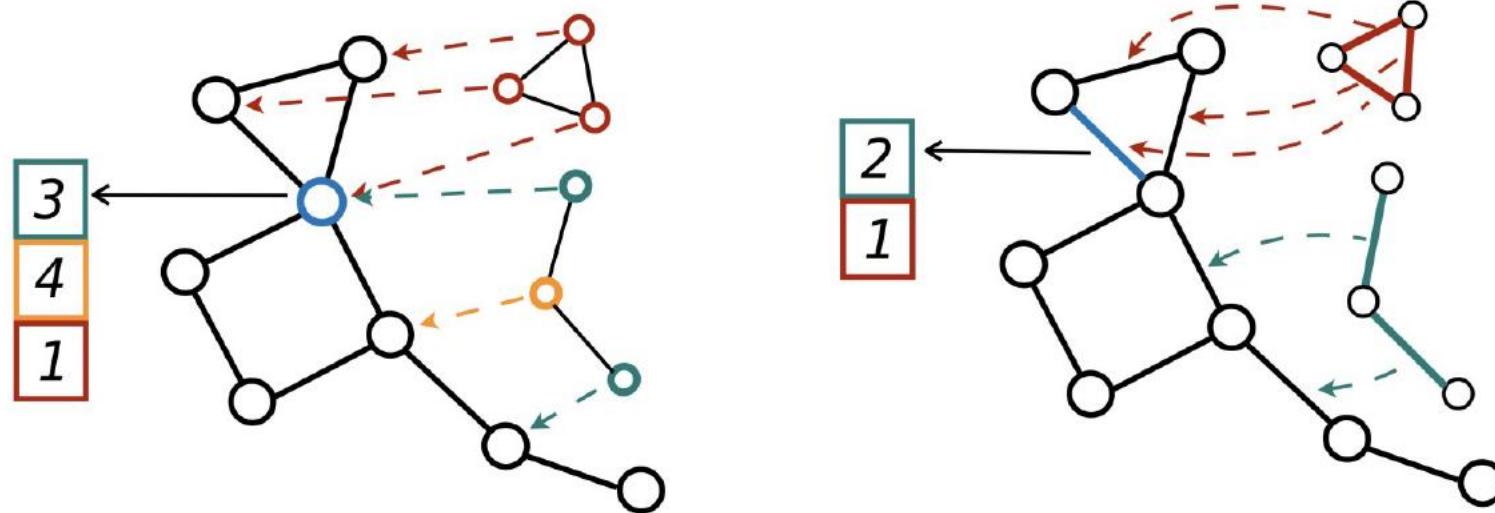
Aggregation strategy in GNNs

- Can we make more powerful GNNs?
- Augment nodes with randomized/positional features^[16]



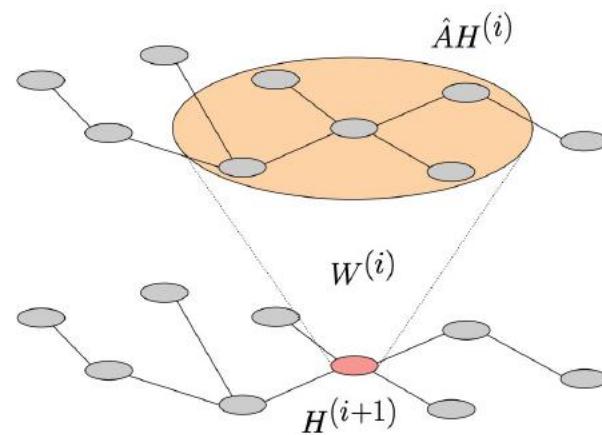
Aggregation strategy in GNNs

- Can we make more powerful GNNs?
- Augment nodes with randomized/positional features^[16]



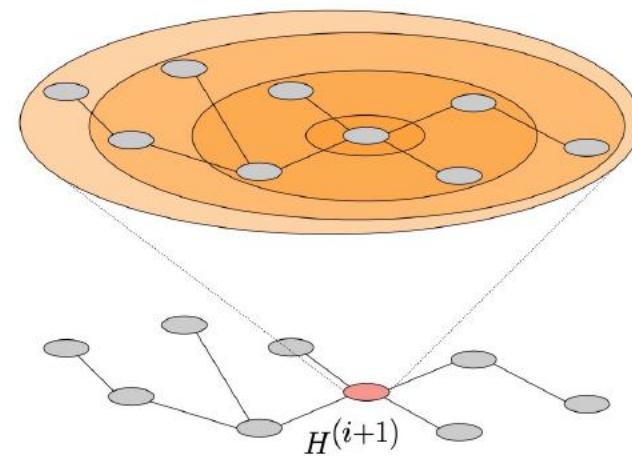
Aggregation strategy in GNNs

- Can we make more powerful GNNs?
- Directly aggregates k-hop information by using adjacency matrix powers^[18]



$$H^{(i+1)} = \sigma(\hat{A}H^{(i)}W^{(i)})$$

(a) Traditional graph convolution.



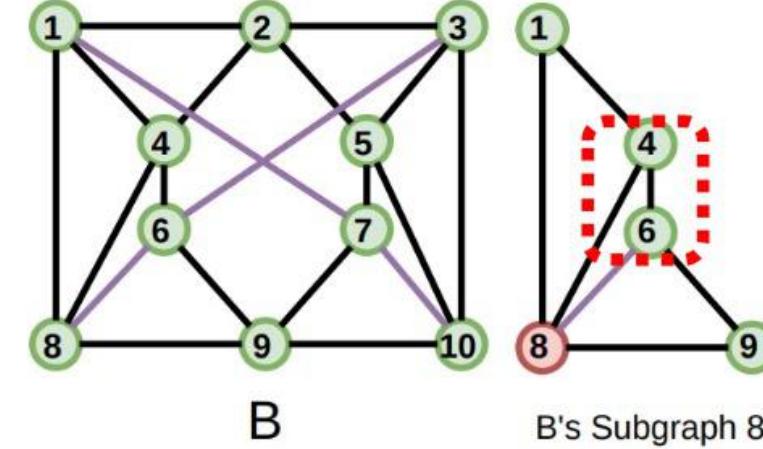
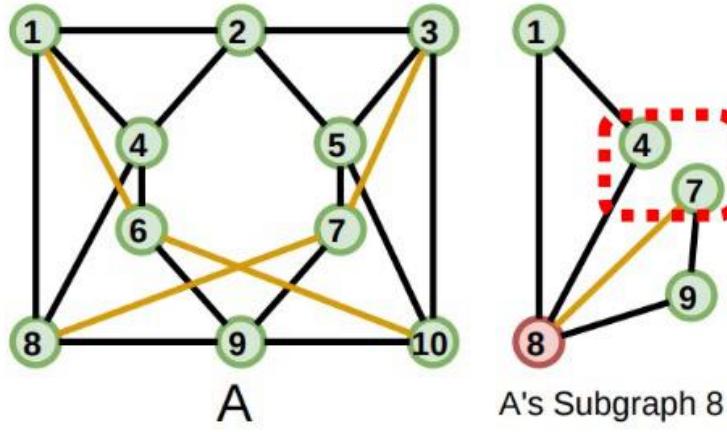
$$H^{(i+1)} = \sigma\left(\hat{A}^0 H^{(i)} W_0^{(i)} \mid \hat{A}^1 H^{(i)} W_1^{(i)} \mid \dots\right)$$

(b) Our mixed feature model.

[18] Sami Abu-El-Haija, et al. "MixHop: Higher-Order Graph Convolutional Architectures via Sparsified Neighborhood Mixing"

Aggregation strategy in GNNs

- Can we make more powerful GNNs?
- Extending local aggregation in GNNs from star patterns to general subgraph patterns^[19]



Aggregation strategy in GNNs

- [20] proves that there isn't a clear single “winner” aggregator

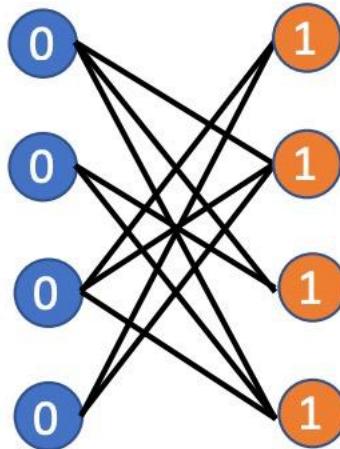
Theorem 1 (Number of aggregators needed). *In order to discriminate between multisets of size n whose underlying set is \mathbb{R} , at least n aggregators are needed.*

Aggregation strategy in GNNs

- Homophily assumption
 - Connected nodes are similar/related/informative

Aggregation strategy in GNNs

- Homophily assumption
 - Connected nodes are similar/related/informative
- How can we deal with heterophilous networks?^[21,22]
 - Connected nodes have different class labels and dissimilar features

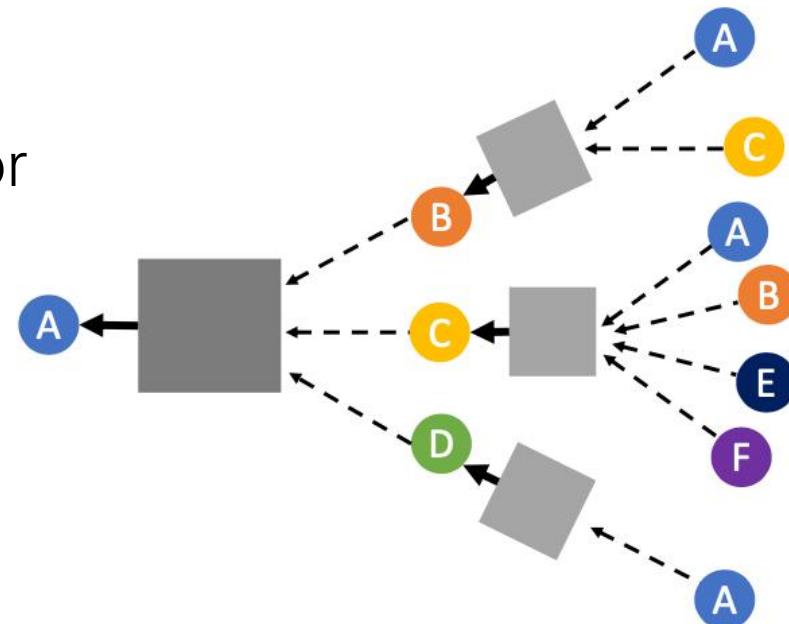


[21] Jiong Zhu., et al. "Beyond Homophily in Graph Neural Networks: Current Limitations and Effective Designs"

[22] Yao Ma, et al. "IS HOMOPHILY A NECESSITY FOR GRAPH NEURAL NETWORKS?"

Graph Neural Network Architectures

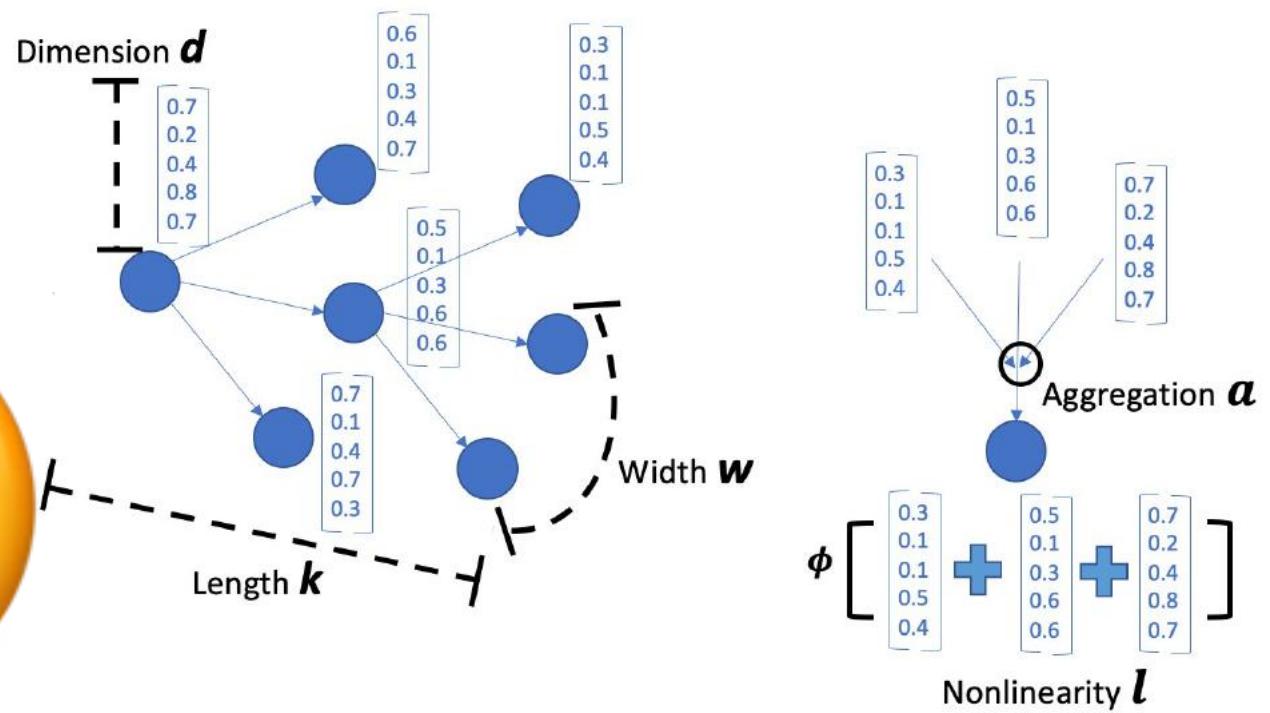
- Width
 - Which neighbors should we aggregate messages from?
- Depth
 - How many hops should we check?
- Aggregation
 - How should we aggregate messages from neighbor



Neural Architecture Search for GNNs

- Which width, depth, and aggregation strategy are proper for a given graph and task?

Width?
Depth?
Aggregation?



Neural Architecture Search for GNNs

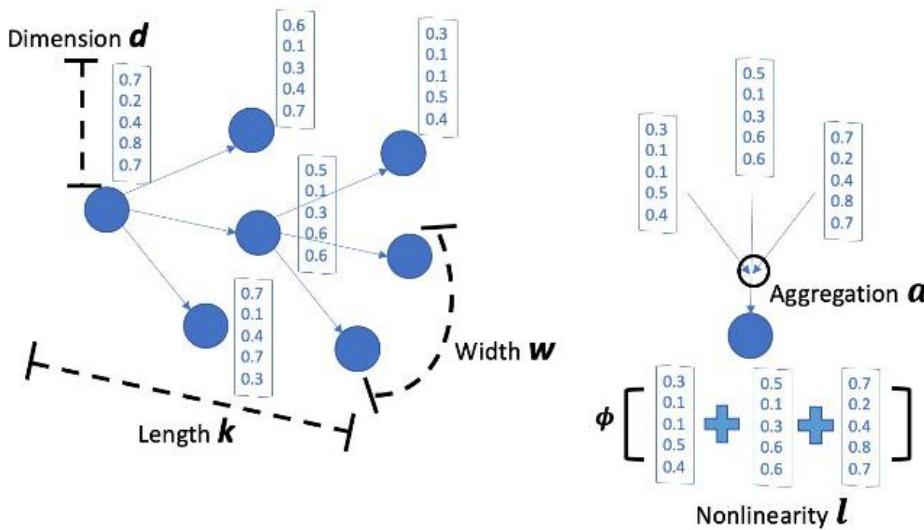
- Finding proper width, depth, and aggregation strategy for a given graph and task **automatically**^[1,2,3]

Here is the GNN you
requested

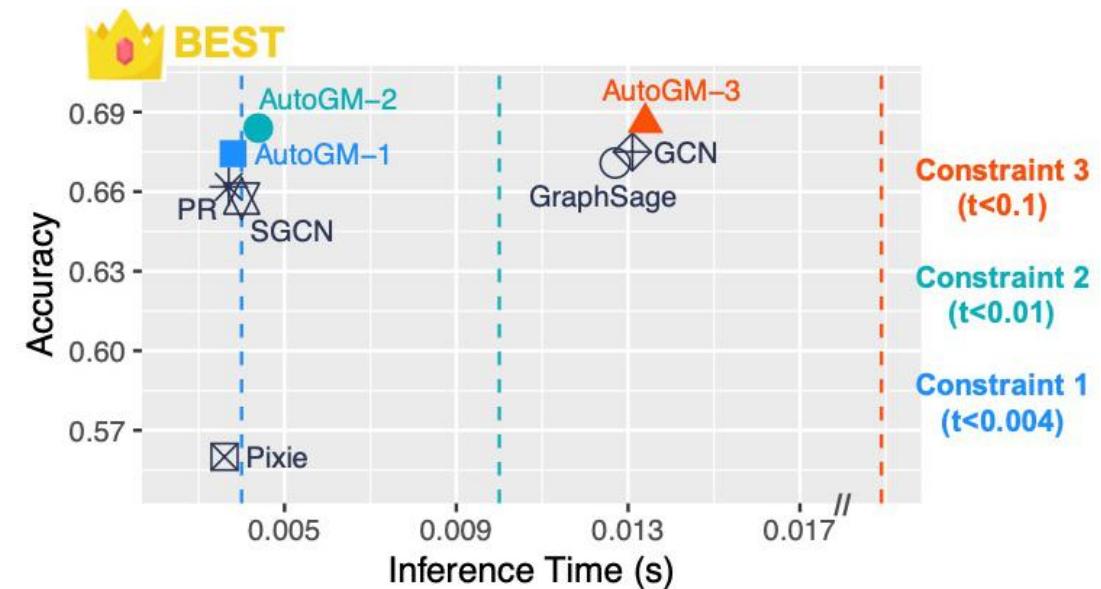


Neural Architecture Search for GNNs

- AutoGM[23]



Step 1: define a hyperparameter space



Step 2: explore the space efficiently

How to train GNNs?

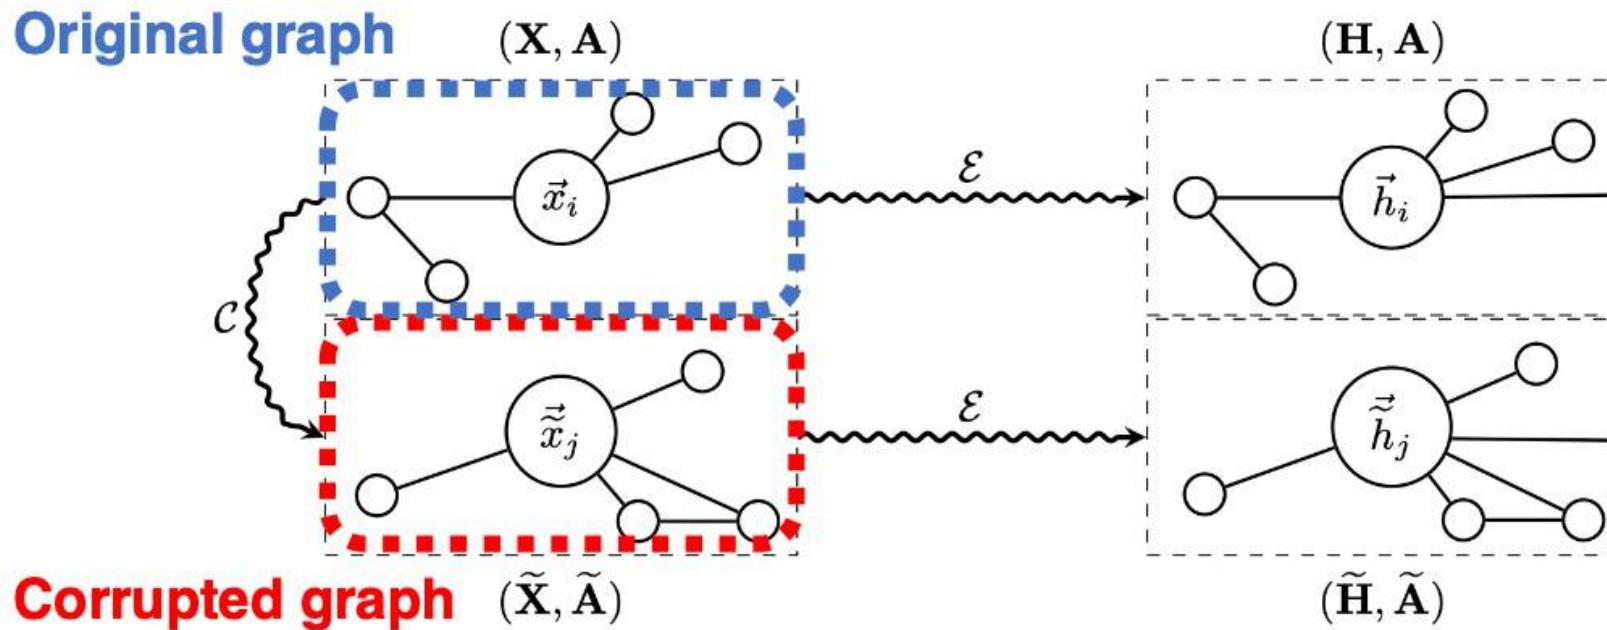
How to train GNNs

- **Semi-supervised learning**
 - Input node features are given for all nodes in a graph
 - Only a subset of nodes have labels

How to train GNNs

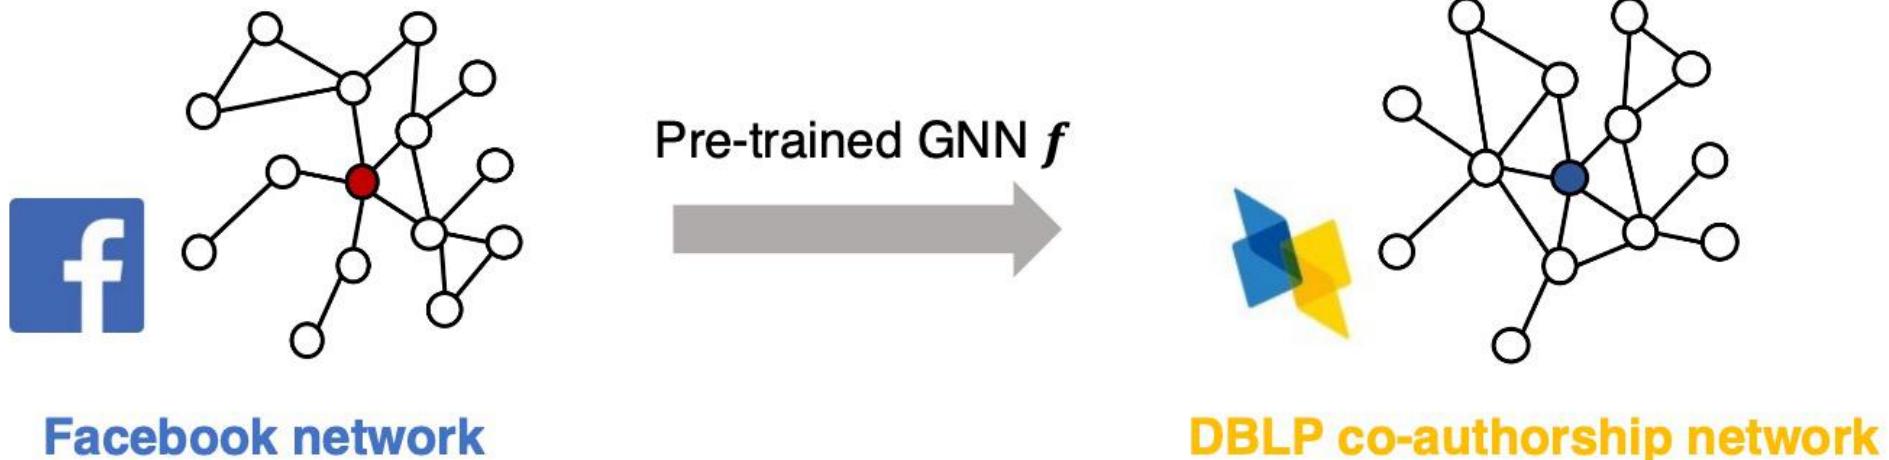
- Unsupervised learning^[26]

- Contrastive learning



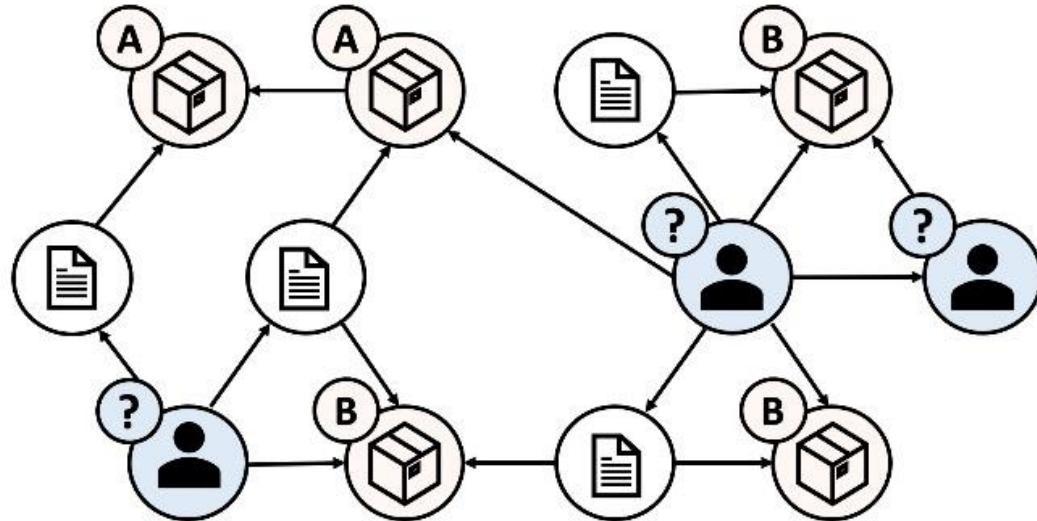
How to train GNNs

- Transfer learning
 - Transfer a pre-trained GNN model between graphs[27]



How to train GNNs

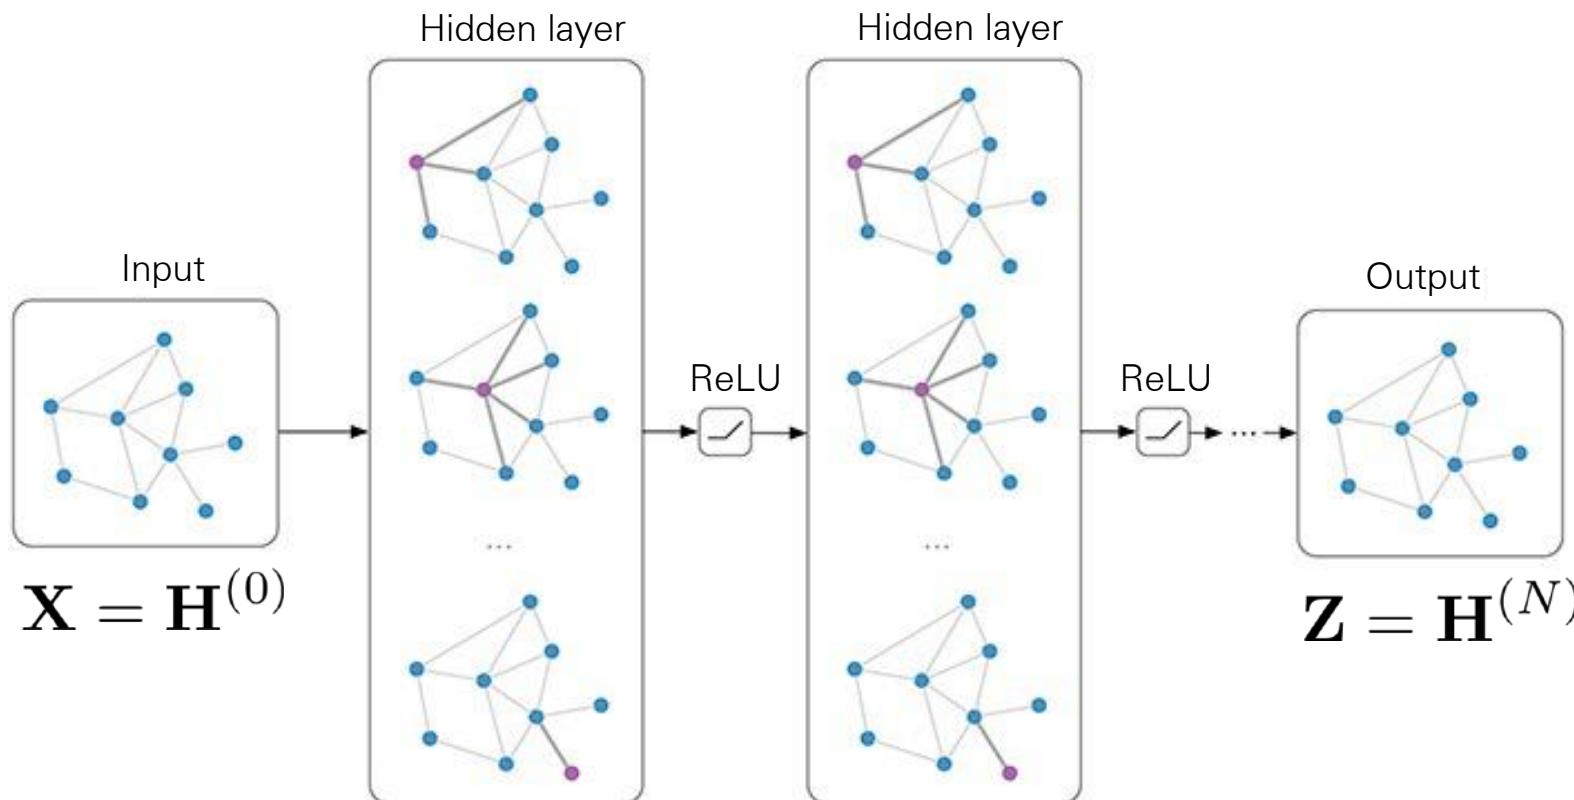
- Transfer learning
 - Transfer between different node types across a **heterogeneous graph**[28]



Applications to “classical” network problems

One fits all: Classification and link prediction with GNNs/GCNs

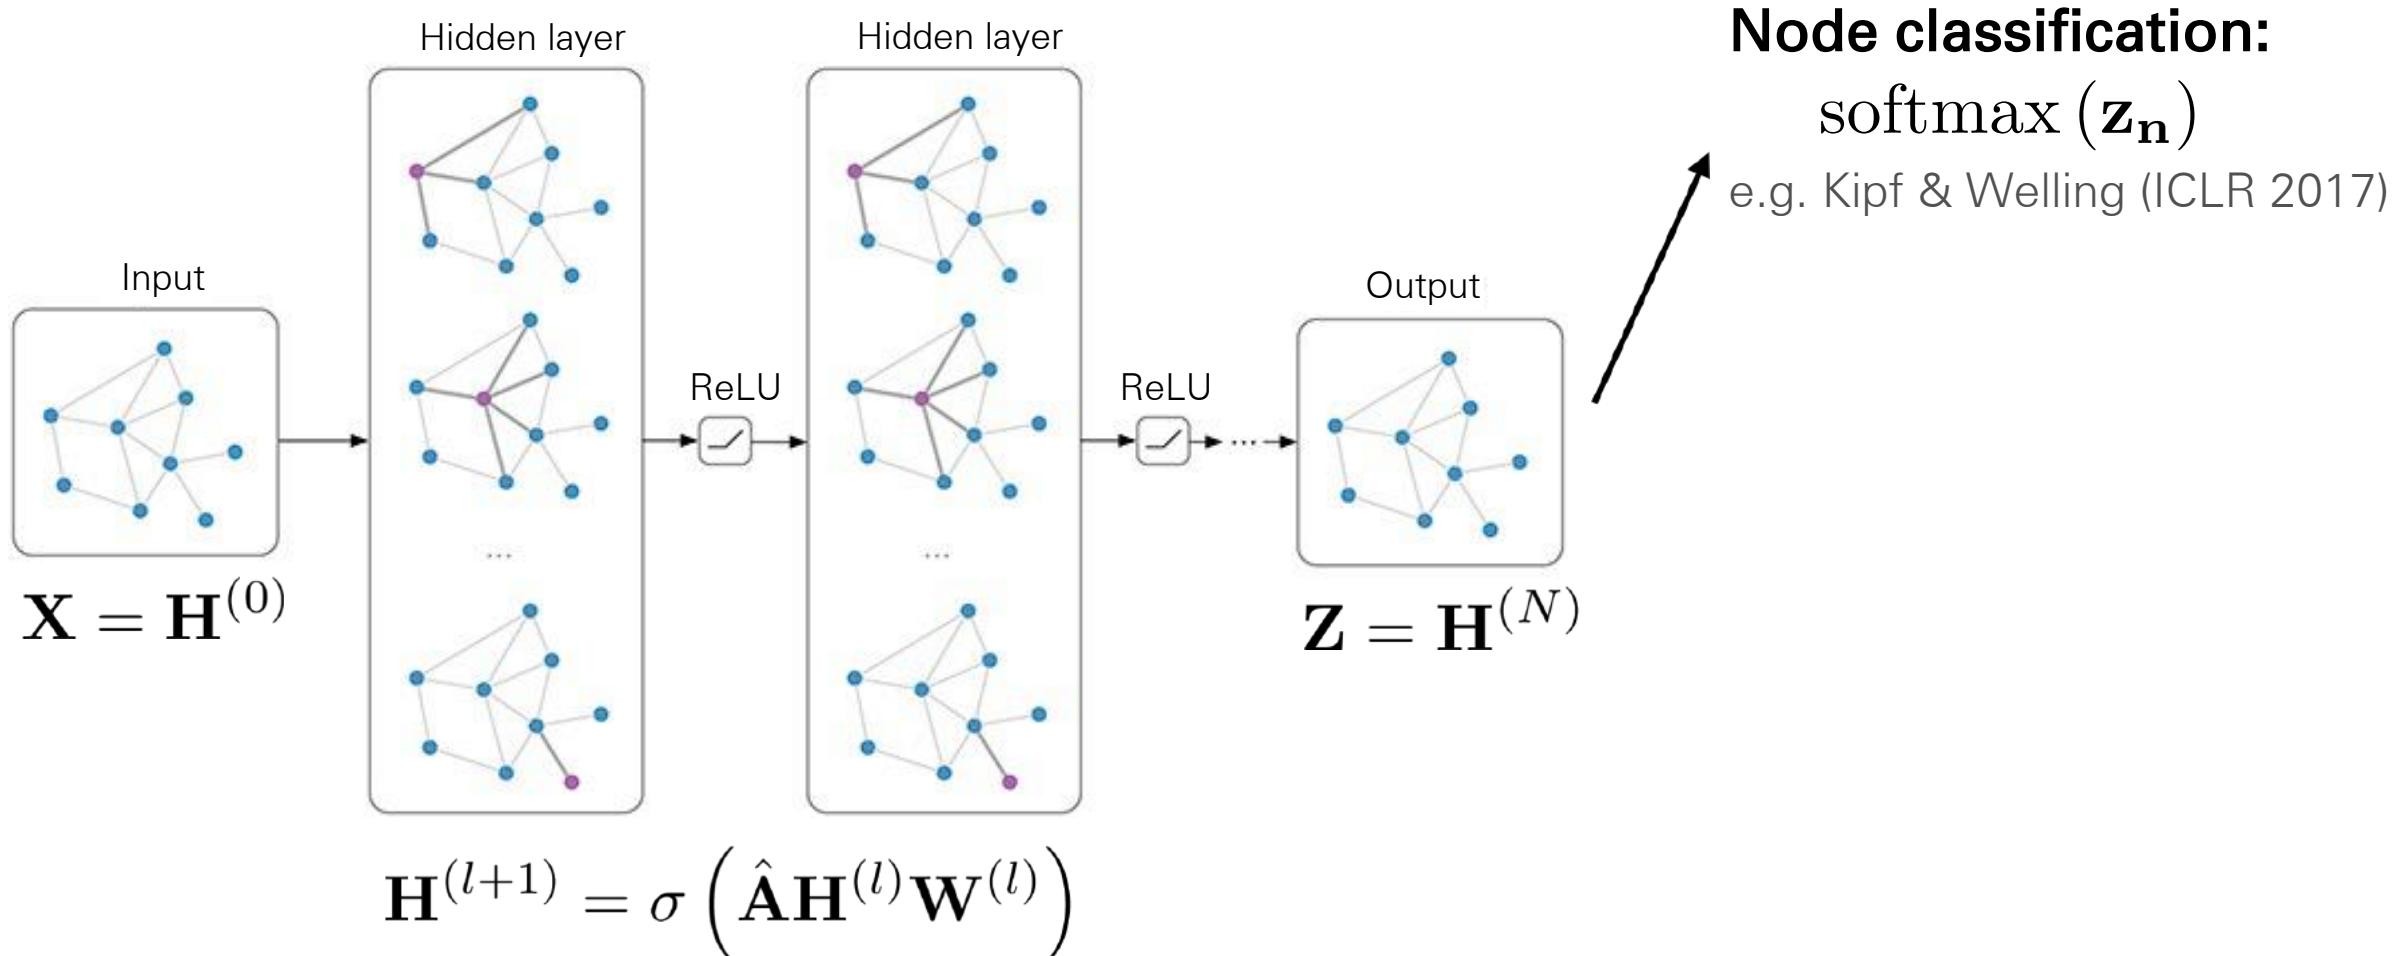
Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



$$\mathbf{H}^{(l+1)} = \sigma(\hat{\mathbf{A}}\mathbf{H}^{(l)}\mathbf{W}^{(l)})$$

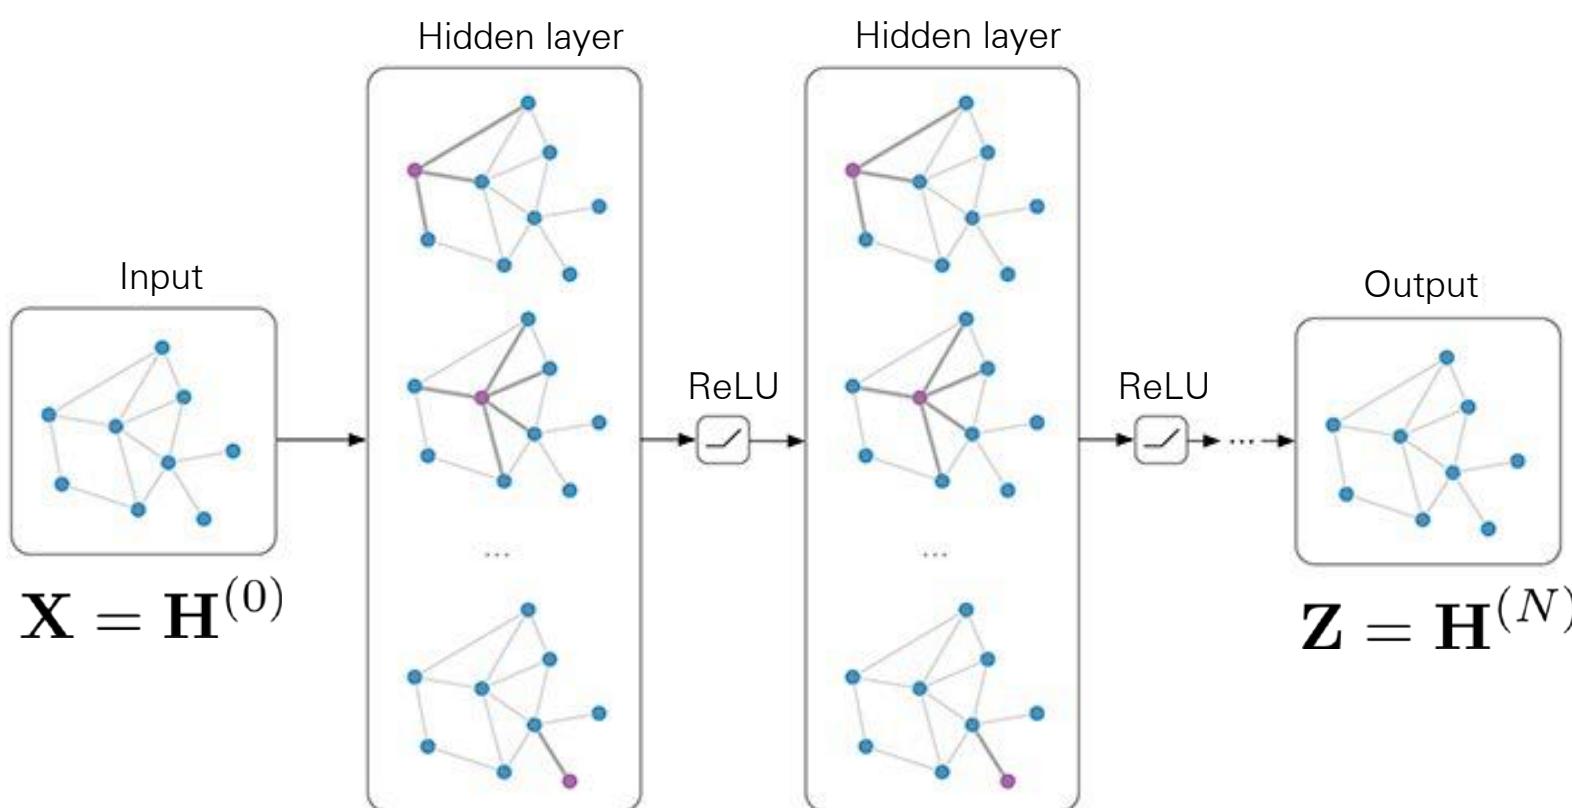
One fits all: Classification and link prediction with GNNs/GCNs

Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



One fits all: Classification and link prediction with GNNs/GCNs

Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



$$\mathbf{H}^{(l+1)} = \sigma \left(\hat{\mathbf{A}} \mathbf{H}^{(l)} \mathbf{W}^{(l)} \right)$$

Node classification:

softmax (\mathbf{z}_n)

e.g. Kipf & Welling (ICLR 2017)

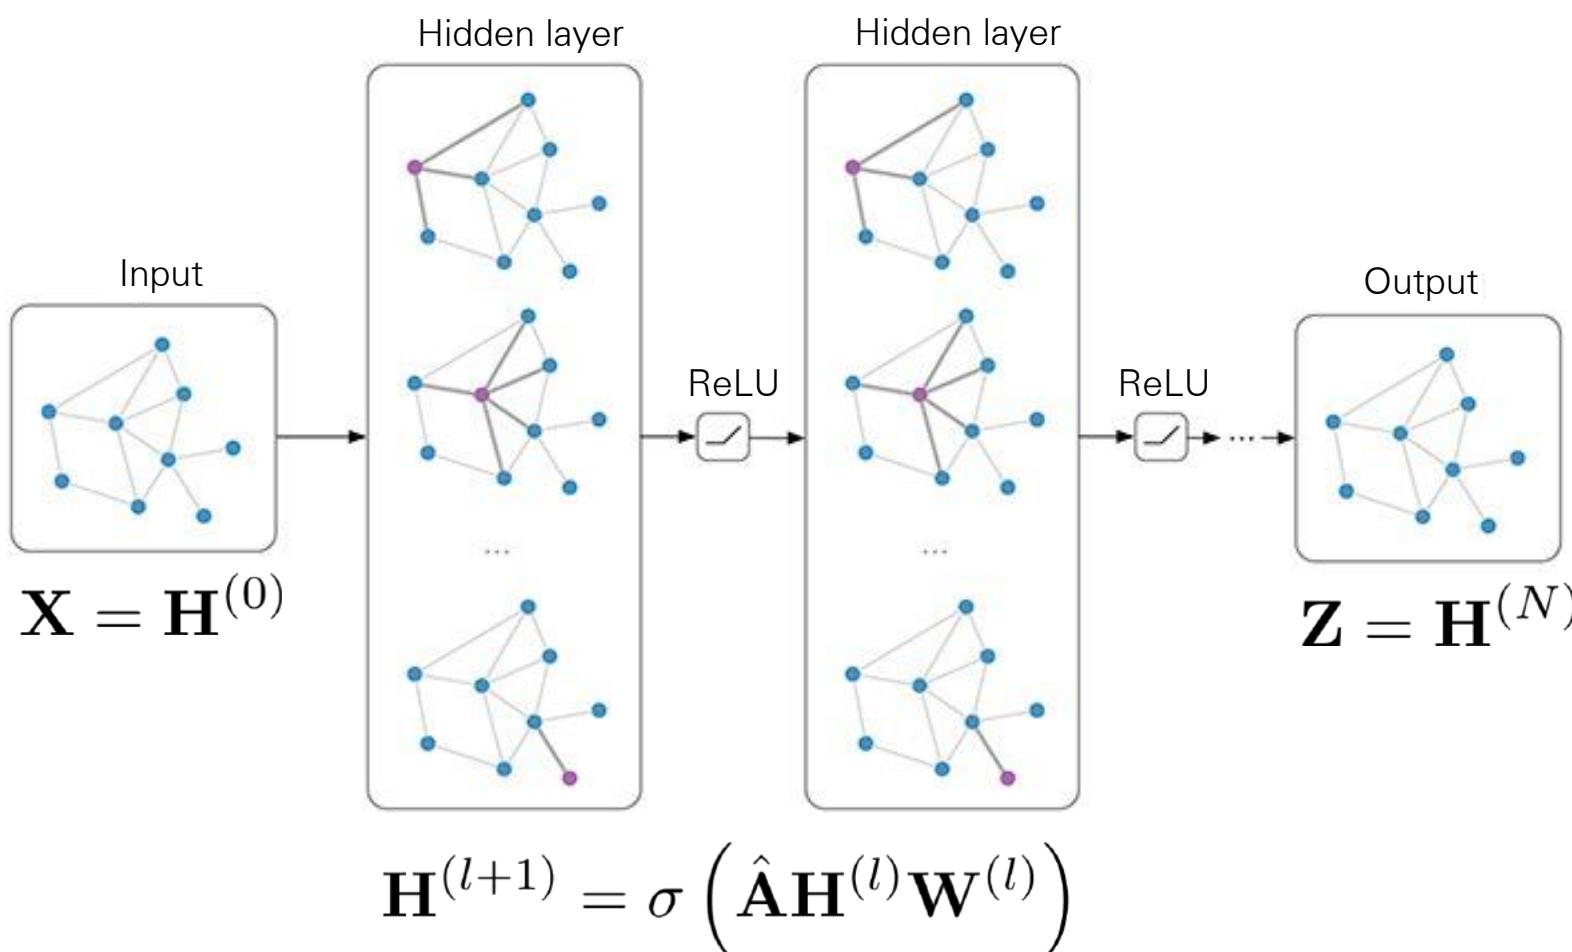
Graph classification:

softmax ($\sum_n \mathbf{z}_n$)

e.g. Duvenaud et al. (NIPS 2015)

One fits all: Classification and link prediction with GNNs/GCNs

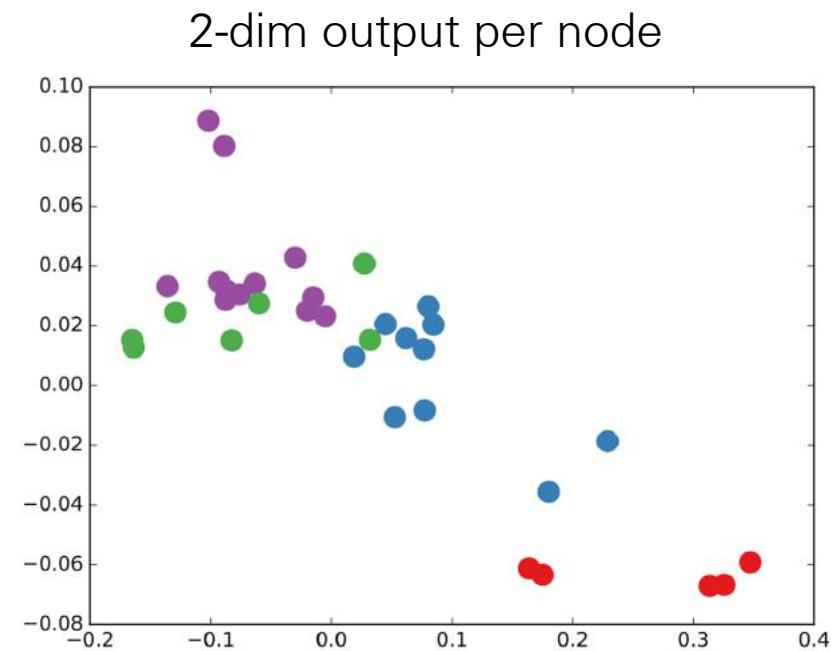
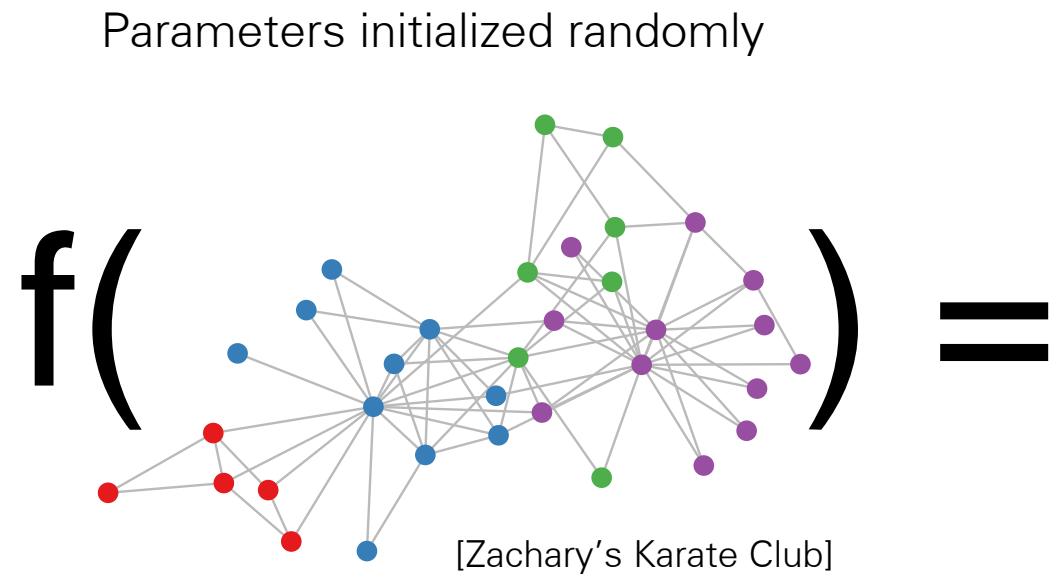
Input: Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times E}$, preprocessed adjacency matrix $\hat{\mathbf{A}}$



- Node classification:**
softmax (\mathbf{z}_n)
e.g. Kipf & Welling (ICLR 2017)
- Graph classification:**
softmax ($\sum_n \mathbf{z}_n$)
e.g. Duvenaud et al. (NIPS 2015)
- Link prediction:**
 $p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$
Kipf & Welling (NIPS BDL 2016)
“Graph Auto-Encoders”

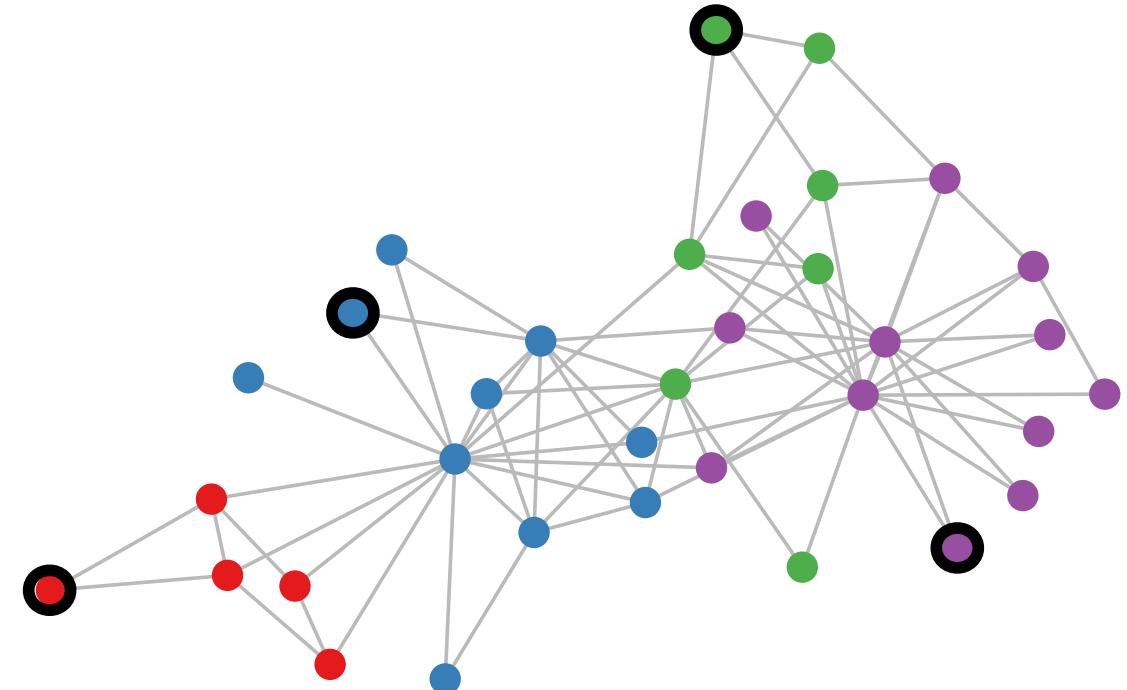
What do learned representations look like?

Forward pass through untrained 3-layer GCN model



Semi-supervised classification on graphs

- **Setting:**
Some nodes are labeled (black circle)
All other nodes are unlabeled
- **Task:**
Predict node label of unlabeled nodes

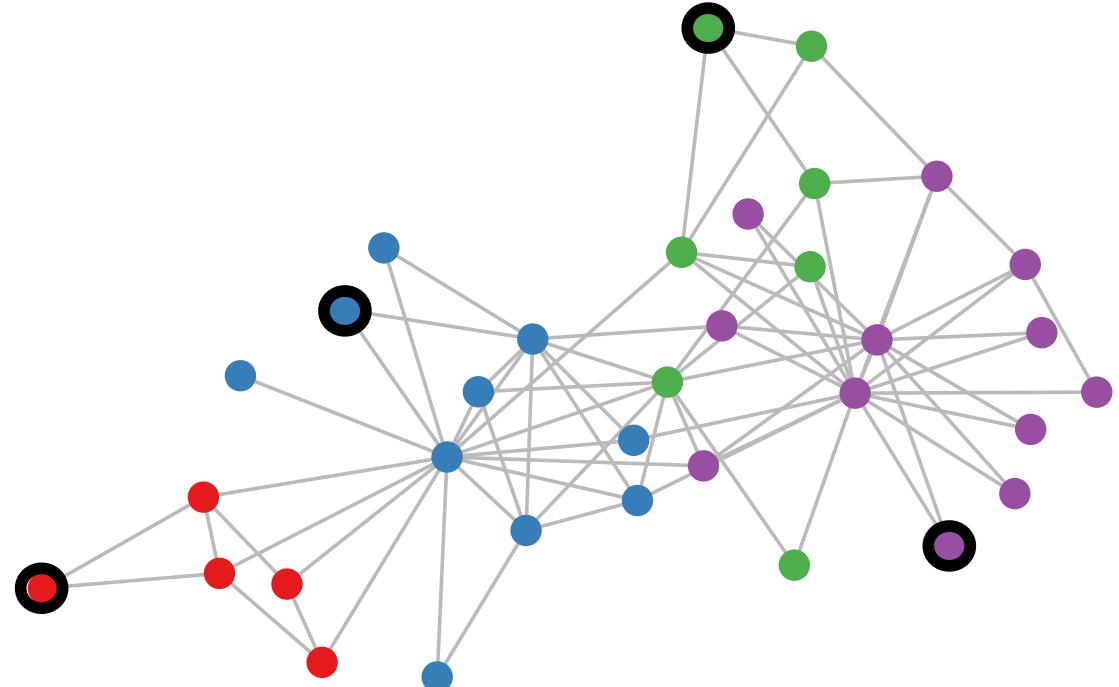


Semi-supervised classification on graphs

- **Setting:**
Some nodes are labeled (black circle)
All other nodes are unlabeled
- **Task:**
Predict node label of unlabeled nodes

Evaluate loss on labeled nodes only:

$$\mathcal{L} = - \sum_{l \in \mathcal{Y}_L} \sum_{f=1}^F Y_{lf} \ln Z_{lf}$$

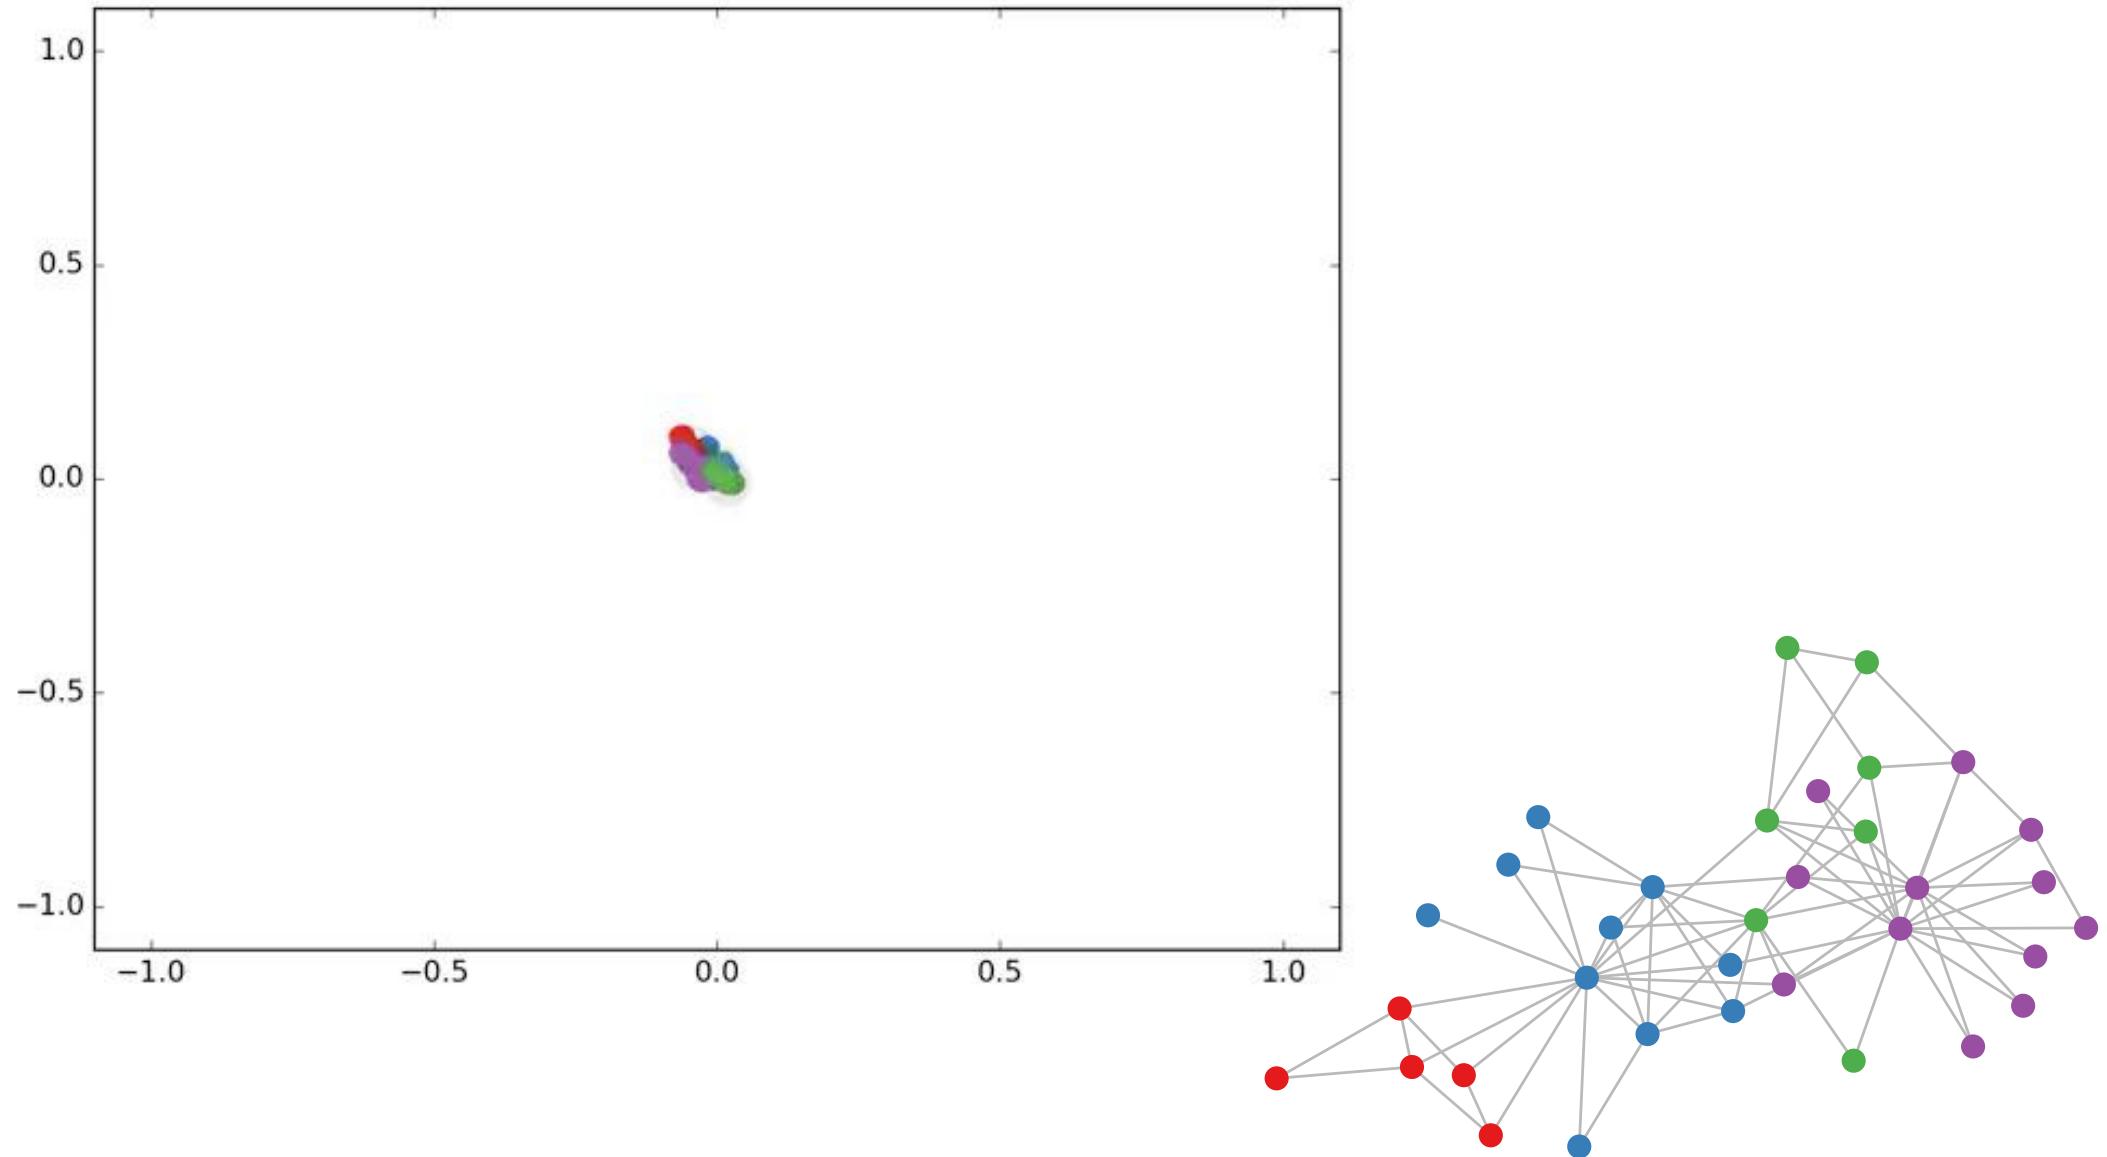


\mathcal{Y}_L set of labeled node indices

\mathbf{Y} label matrix

\mathbf{Z} GCN output (after softmax)

Toy example (semi-supervised learning)



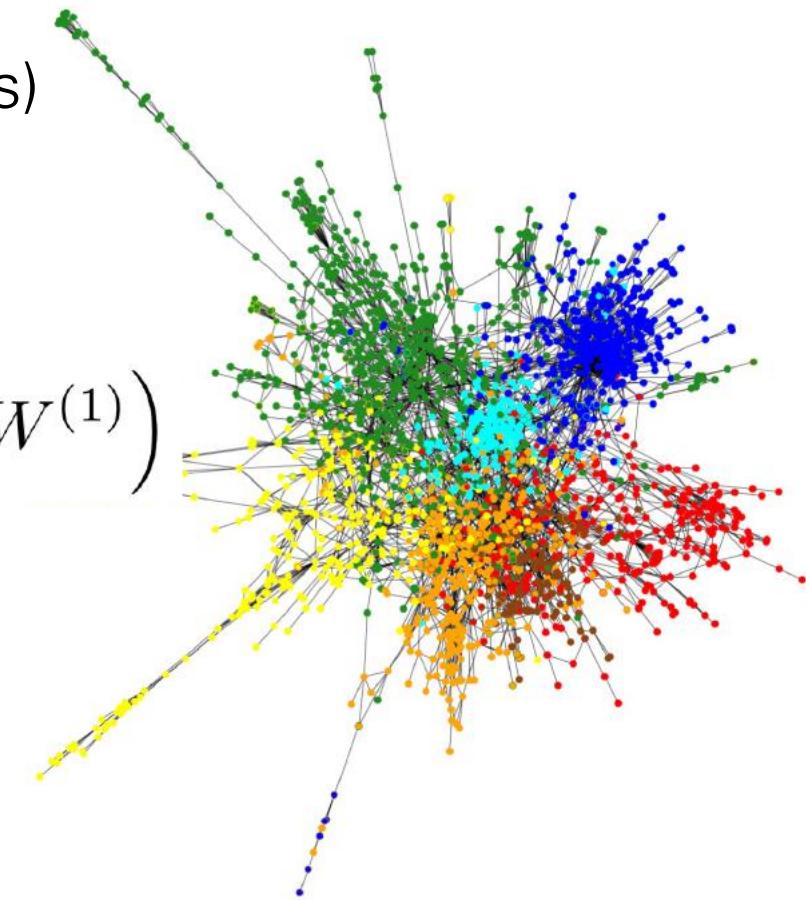
Application: Classification on citation networks

Input: Citation networks (nodes are papers, edges are citation links, optionally bag-of-words features on nodes)

Target: Paper category (e.g. stat.ML, cs.LG, ...)

Model: 2-layer GCN

$$Z = f(X, A) = \text{softmax}\left(\hat{A} \text{ReLU}\left(\hat{A}XW^{(0)}\right) W^{(1)}\right)$$



(Figure from: Bronstein, Bruna, LeCun, Szlam, Vandergheynst, 2016)

Application: Classification on citation networks

Input: Citation networks (nodes are papers, edges are citation links, optionally bag-of-words features on nodes)

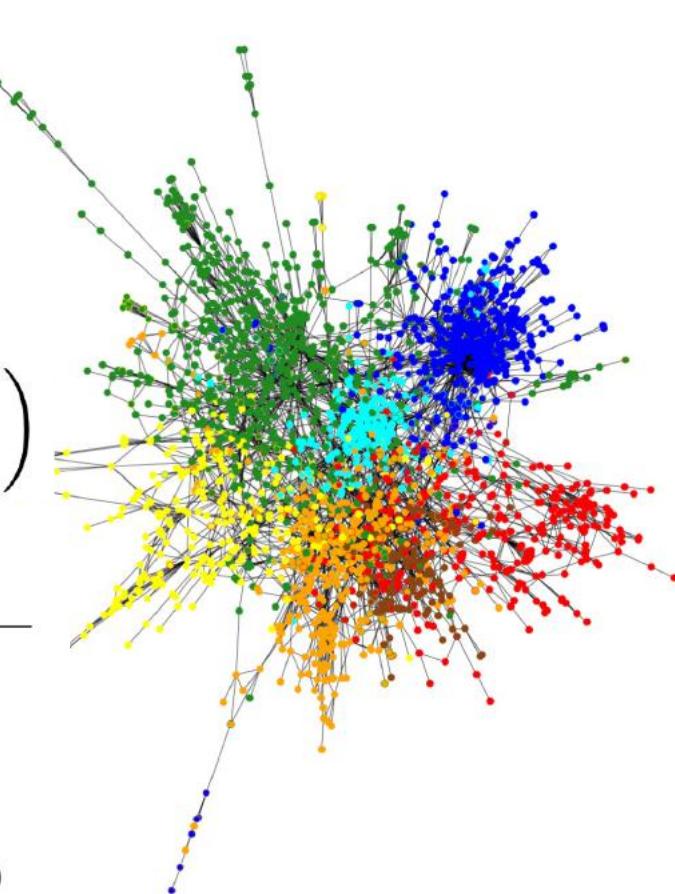
Target: Paper category (e.g. stat.ML, cs.LG, ...)

Model: 2-layer GCN

$$Z = f(X, A) = \text{softmax}\left(\hat{A} \text{ReLU}\left(\hat{A}XW^{(0)}\right) W^{(1)}\right)$$

Classification results (accuracy)

Method	Citeseer	Cora	Pubmed	NELL
ManiReg [3]	60.1	59.5	70.7	21.8
SemiEmb [24]	59.6	59.0	71.1	26.7
no input features	LP [27]	45.3	68.0	63.0
DeepWalk [18]	43.2	67.2	65.3	58.1
Planetoid* [25]	64.7 (26s)	75.7 (13s)	77.2 (25s)	61.9 (185s)
GCN (this paper)	70.3 (7s)	81.5 (4s)	79.0 (38s)	66.0 (48s)
GCN (rand. splits)	67.9 ± 0.5	80.1 ± 0.5	78.9 ± 0.7	58.4 ± 1.7



(Figure from: Bronstein, Bruna, LeCun, Szlam, Vandergheynst, 2016)

Still many open problems..

- And many more chances to do groundbreaking research
- ex) other graph formats
 - 3-dimensional graphs
 - Temporal graphs
 - ...

Next Lecture:

Pretraining Language Models