

PAPER LIST

Post-Training and RL

- **October 13:** [On the Interplay of Pre-Training, Mid-Training, and RL on Reasoning Language Models](#)
Charlie Zhang, Graham Neubig, Xiang Yue. ICML 2026.
<https://github.com/Interplay-LM-Reasoning/Interplay-LM-Reasoning>
- **October 14:** *PS Hour: Code walkthrough: Interplay-LM-Reasoning*
- **October 15:** [Front-Loading Reasoning: The Synergy between Pretraining and Post-Training Data](#)
Syeda Nahida Akter, Shrimai Prabhumoye, Eric Nyberg, Mostofa Patwary, Mohammad Shoeybi, Yejin Choi, Bryan Catanzaro. ICLR 2026.

Efficient Adaptation and Compression

- **October 20:** [Qwen-3D: A Generalist 3D Vision-Language Model for Spatial Understanding](#)
Lucy Lin, Ayush Jain, Yifan Liu, Katerina Fragkiadaki. ECCV 2026.
<https://qwen-3d.github.io>
- **October 21:** *PS Hour: Code walkthrough: Qwen-3D*
- **October 22:** [A Frame is Worth One Token: Efficient Generative World Modeling with Delta Tokens](#)
Tommie Kerssies, Gabriele Berton, Ju He, Qihang Yu, Wufei Ma, Daan de Geus, Gijs Dubbelman, Liang-Chieh Chen. CVPR 2026.
<https://github.com/amazon-far/deltatok>

Unified Multimodal Understanding and Generation

- **November 3:** [Lance: Unified Multimodal Modeling by Multi-Task Synergy](#)
YFengyi Fu, Mengqi Huang, Shaojin Wu, Yunsheng Jiang, Yufei Huo, Hao Li, Yinghang Song, Fei Ding, Jianzhu Guo, Qian He, Zheren Fu, Zhendong Mao, Yongdong Zhang. ArXiv preprint arXiv:2605.18678, May 2026.
<https://lance-project.github.io/>
- **November 4:** *PS Hour: Code walkthrough: Lance*
- **November 5:** [Beyond Language Modeling: An Exploration of Multimodal Pretraining](#)
Shengbang Tong, David Fan, John Nguyen, Ellis L Brown II, Gaoyue Zhou, et al. ICML 2026.
<https://beyond-llms.github.io>

Reasoning and Inference-time Computation

- **November 10:** [UniT: Unified Multimodal Chain-of-Thought Test-time Scaling](#)
Leon Liangyu Chen, Haoyu Ma, Zhipeng Fan, Ziqi Huang, Animesh Sinha et al. CVPR 2026.
- **November 11:** *No PS hour this week*
- **November 12:** [Rubrics as Rewards: Reinforcement Learning Beyond Verifiable Domains](#)
Anisha Gunjal, Anthony Wang, Elaine Lau, Vaskar Nath, Bing Liu, Sean Hendryx. ICLR 2026.

World Models and Physical Intelligence

- **November 17:** [PhyCo: Learning Controllable Physical Priors for Generative Motion](#)
Sriram Narayanan, Ziyu Jiang, Srinivasa Narasimhan, Manmohan Chandraker. CVPR 2026.
<https://phyco-video.github.io/>
- **November 18:** *PS Hour: Code walkthrough: PhyCo*
- **November 19:** [VideoWorld 2: Learning Transferable Knowledge from Real-world Videos](#)
Zhongwei Ren, Yunchao Wei, Xiao Yu, Guixun Luo, Yao Zhao, Bingyi Kang, Jiashi Feng, Xiaojie Jin. CVPR 2026.
<https://maverickren.github.io/VideoWorld2.github.io/>

Long-term Memory and Long-context Models

- **November 24:** [WorldMM: Dynamic Multimodal Memory Agent for Long Video Reasoning](#)
Woongyeong Yeo, Kangsan Kim, Jaehong Yoon, Sung Ju Hwang. CVPR 2026.
<https://worldmm.github.io>
- **November 25:** *PS Hour: Code walkthrough: WorldMM*
- **November 26:** [Attention-based Vision-Language Memory for Spatial Reasoning](#)
Zuntao Liu, Yi Du, Taimeng Fu, Shaoshu Su, Cherie Ho, Chen Wang. ECCV 2026.
<https://sairlab.org/vlm2>

Tool-using Agents

- **December 15:** [Don't Show Pixels, Show Cues: Unlocking Visual Tool Reasoning in Language Models via Perception Program](#)
Muhammad Kamran Janjua, Hugo Silva, Di Niu, Bahador Rashidi. CVPR 2026.
<https://github.com/AISmartPerception/perception-programs>
- **December 16:** *PS Hour: Code walkthrough: Perception Programs*
- **December 17:** [Diversity Over Frequency: Rethinking Tool Use in Visual Chain-of-Thought Agents](#)
Dong-Hee Kim, Reuben Tan, Donghyun Kim. ICML 2026.
<https://scaffolded-exploration.github.io/>

Mechanistic Interpretability and Model Steering

- **December 22:** [Weight-sparse transformers have interpretable circuits](#)
Leo Gao, Achyuta Rajaram, Jacob Coxon, Soham V. Govande, Bowen Baker, Daniel P. Mossing. ICML 2026.
https://github.com/openai/circuit_sparsity/
- **December 23:** *PS Hour: Code walkthrough: Circuit Sparsity*
- **December 24:** [Steer Like the LLM: Activation Steering that Mimics Prompting](#)
Geert Heyman, Frederik Vandeputte. ICML 2026.
<https://github.com/Nokia-Bell-Labs/steer-like-the-llm>

Evaluation and Benchmark Design

- **January 5:** [MMBench-Live: A Continuously Evolving Benchmark for Multimodal Models](#)
Yuanzhi Liu, Shousheng Zhao, Bo Zhou, Kongming Liang, Zhanyu Ma. ICML 2026.
<https://github.com/PRIS-CV/MMBench-Live>
- **January 6:** *PS Hour: Code walkthrough: MMBench-Live*
- **January 7:** [VGI-Bench: Probing Visual Intelligence in Video Generation Models](#)
JXuan He, Cong Wei, Yuhao Cheng, Linrui Ma, Yuxuan Zhang. EMNLP 2026.
<https://hexuan21.github.io/VGI-Bench/>