

Towards Visual Intelligence

Aykut Erdem

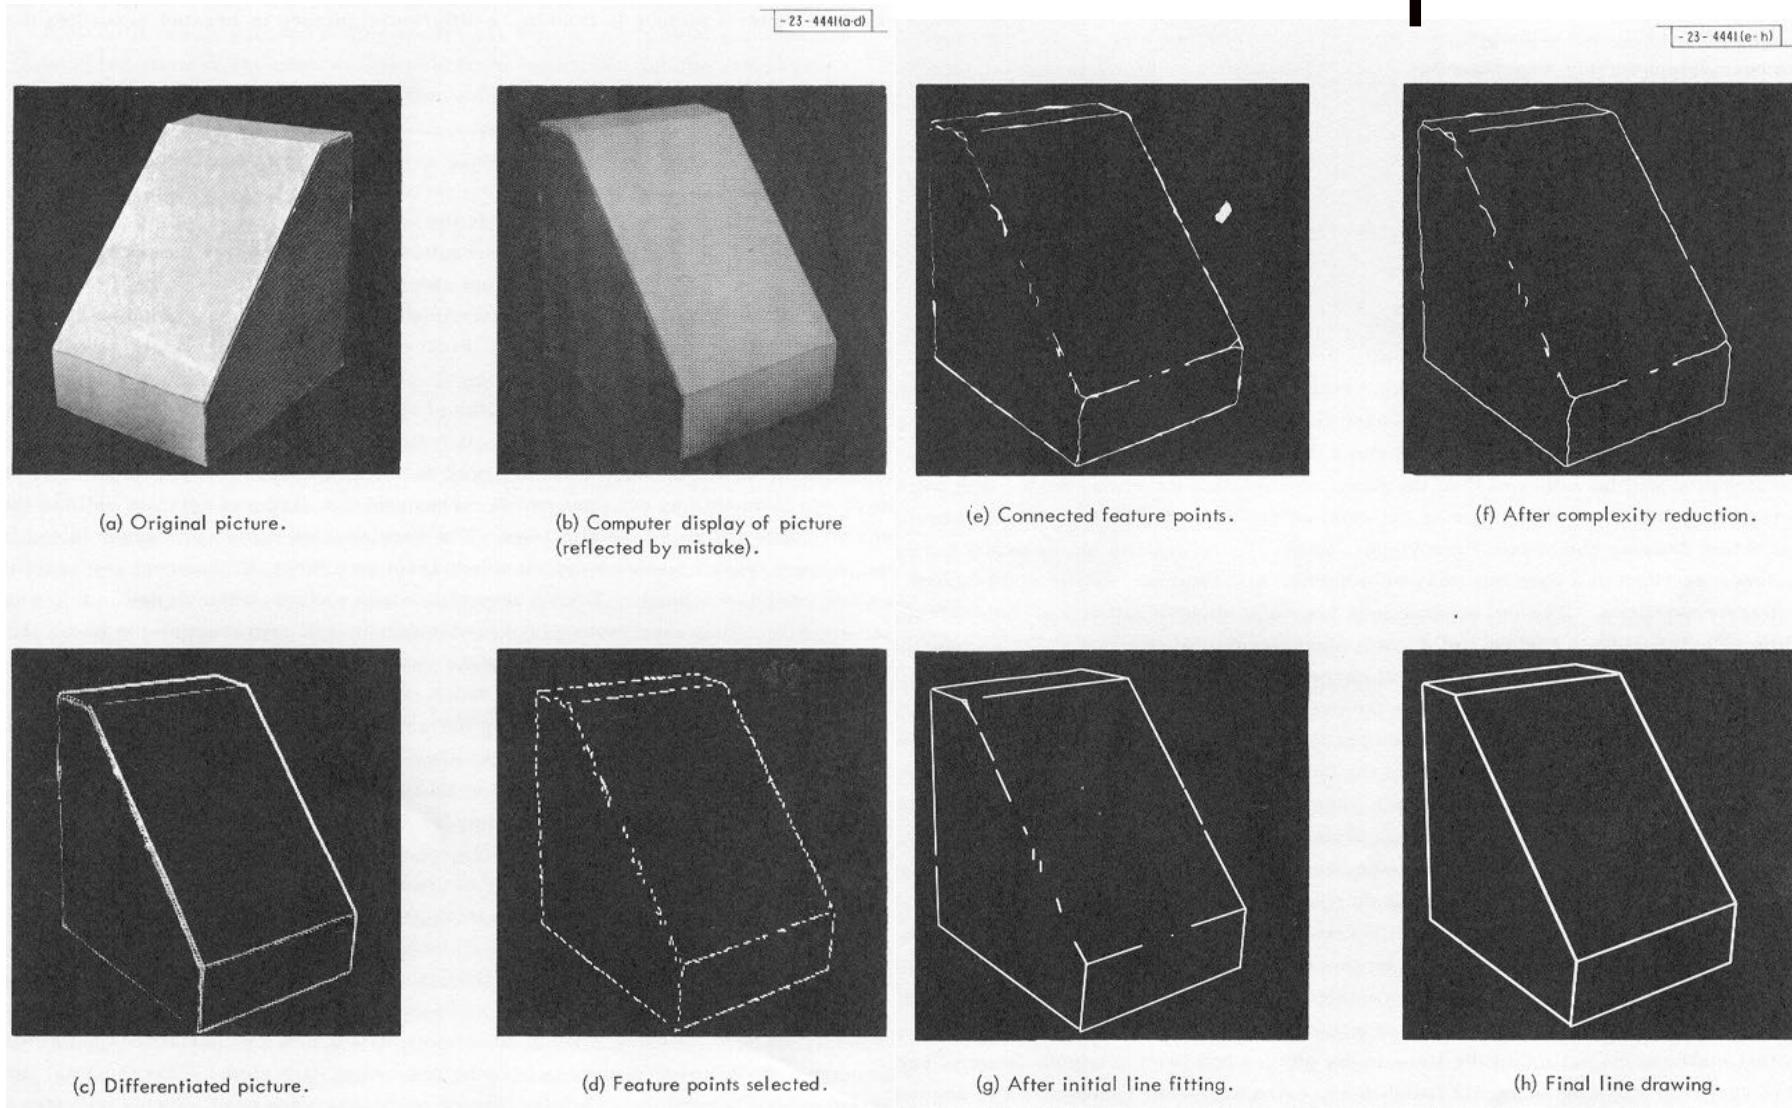
The Purpose of Vision

"What does it mean, to see? The plain man's answer (and Aristotle's too) would be, **to know what is where by looking.** In other words, vision is the process of discovering from images what is present in the world, and where it is."

[Marr, 1982]

Image credit: The Sense of Sight (Annie Louisa Swynnerton, 1895) 2

The First PhD Thesis on Computer Vision



- Machine perception of three-dimensional solids [Roberts 1963]

The Summer Vision Project

General goals:

FIGURE-GROUND.

divide a vidisector picture into regions such as likely objects, likely background areas and chaos

REGION DESCRIPTION.

analysis of shape and surface properties

OBJECT IDENTIFICATION.

name objects by matching them with a vocabulary of known objects

MASSACHUSETTS INSTITUTE OF TECHNOLOGY
PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.

July 7, 1966

THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

Why does vision appear easy to humans?

- Our brains are specialized to do vision.
- ~50% of the cortex in a human brain is devoted for visual processing
(cf. motor control ~20-30%, language ~10-20%)

Visual perception*:

540,000,000 years of data

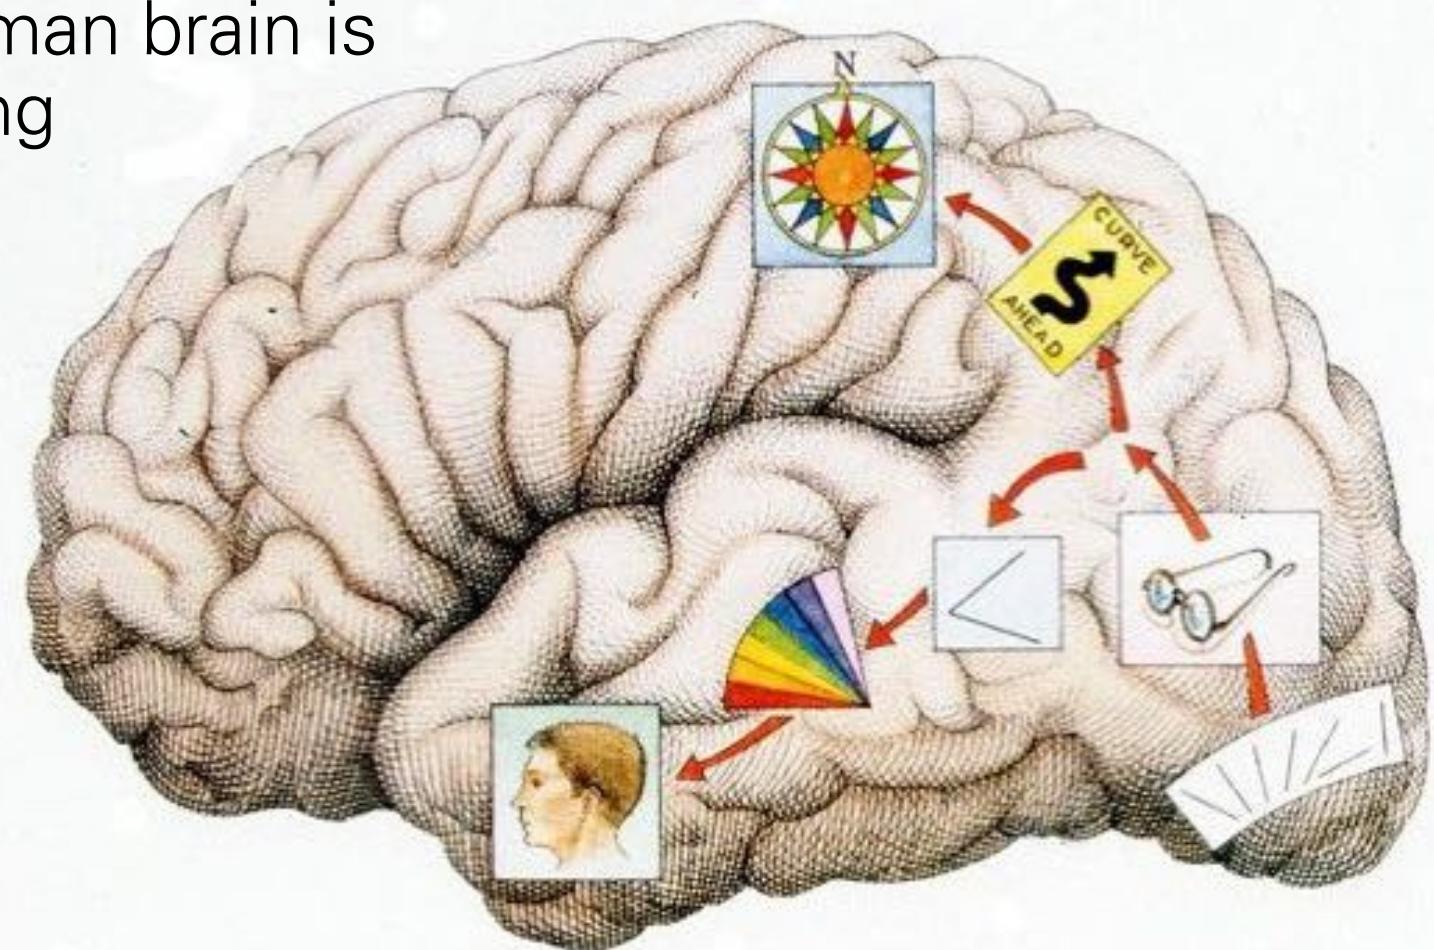
Bipedal movement:

230,000,000 years of data

Abstract thought:

100,000 years of data

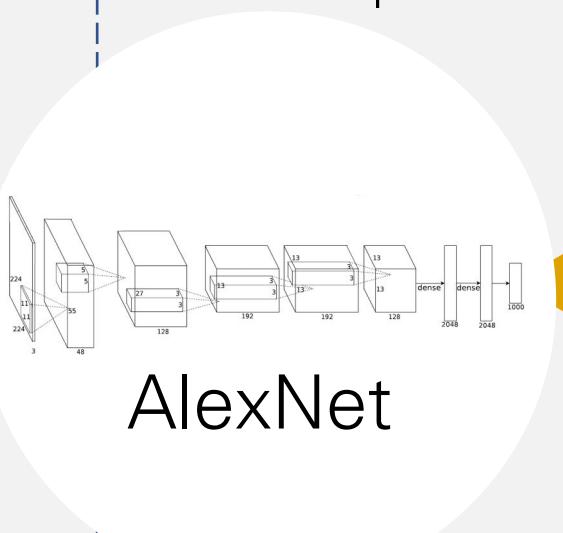
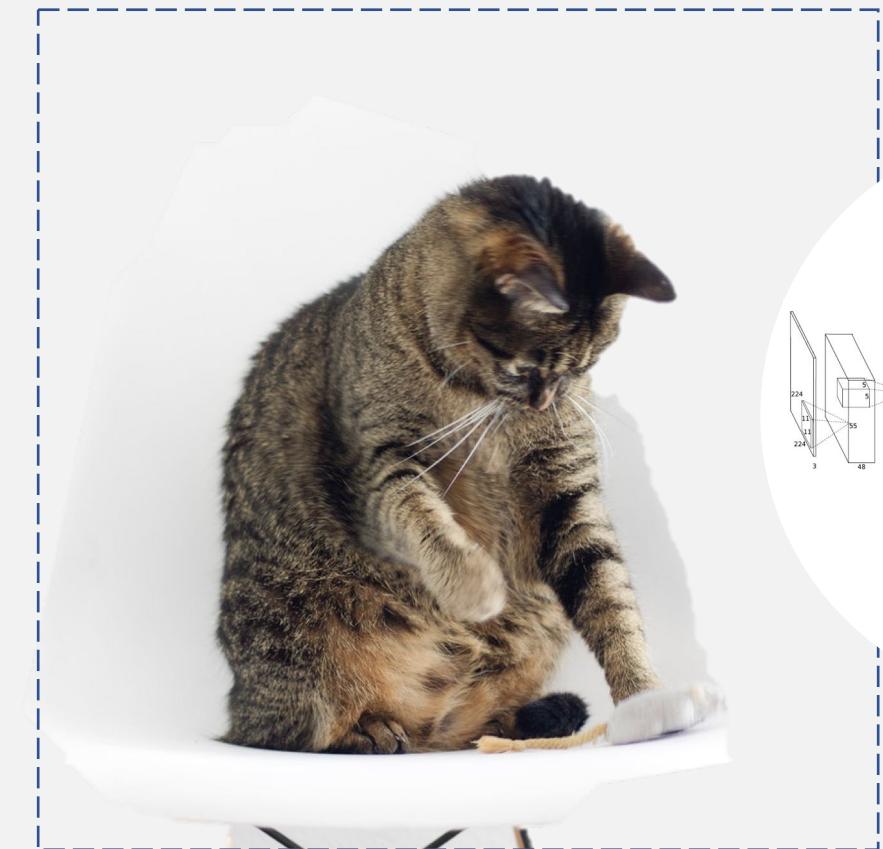
*Color vision



Fast Forward to 2012

IMAGENET Large Scale Visual Recognition Challenge (ILSVRC)

- **1.2M** training images, **1K** categories
- Measure top-5 classification error



The success of AlexNet, a deep convolutional network (CNN)

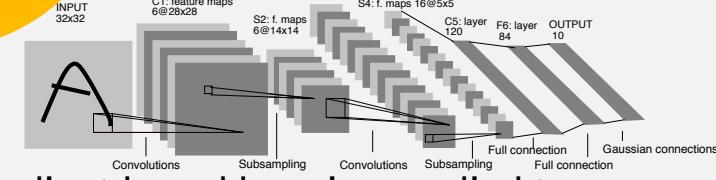
- 7 hidden layers (not counting some max pooling layers)
- 60M parameters



2012 Teams	%Error
Supervision (Toronto)	15.3
ISI (Tokyo)	26.1
VGG (Oxford)	26.9
XRCE/INRIA	27.0
UvA (Amsterdam)	29.6
INRIA/LEAR	33.4

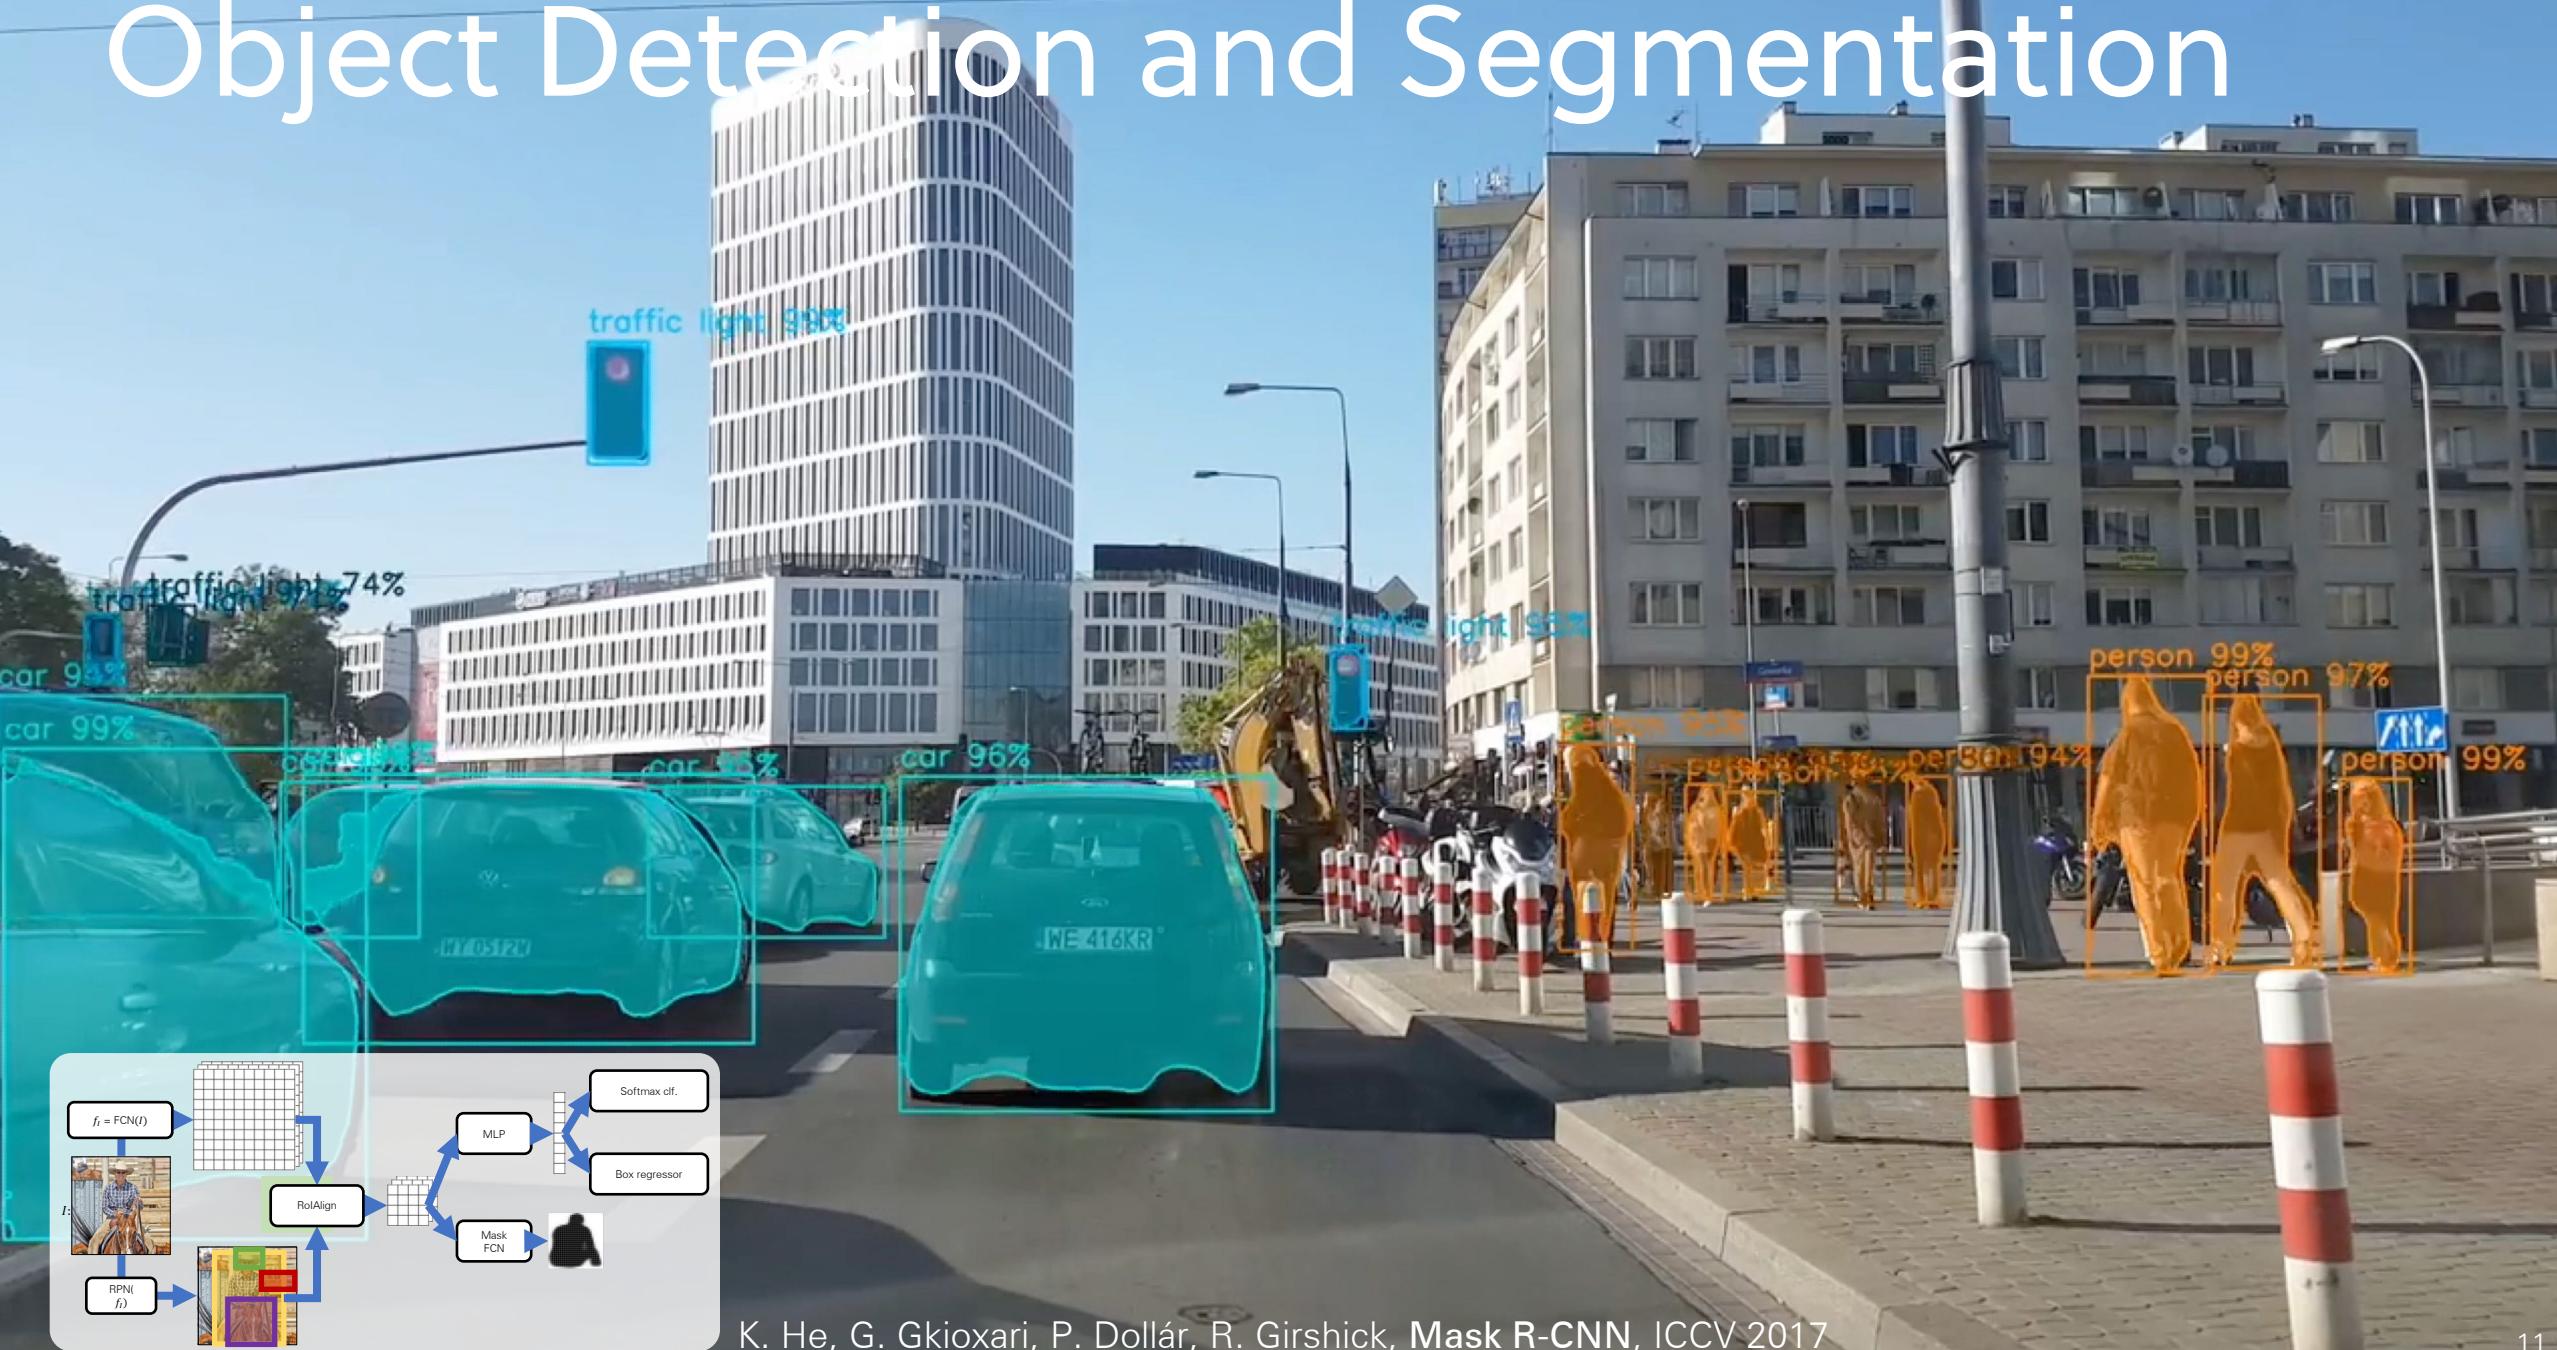
CNN based, non-CNN based

CNNs are biologically inspired by oriented cells in the visual cortex



- Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. **Gradient-based learning applied to document recognition**. Proceedings of the IEEE. 86 (11): 2278–2324, 1998.
- A. Krizhevsky, I. Sutskever, G.E. Hinton **ImageNet Classification with Deep Convolutional Neural Networks**. NIPS 2012

Object Detection and Segmentation



11.4 fps

Pose Estimation

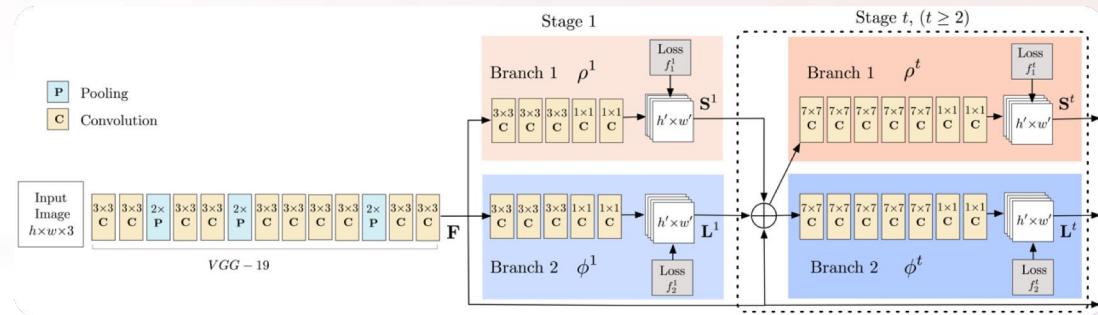


Photo Style Transfer



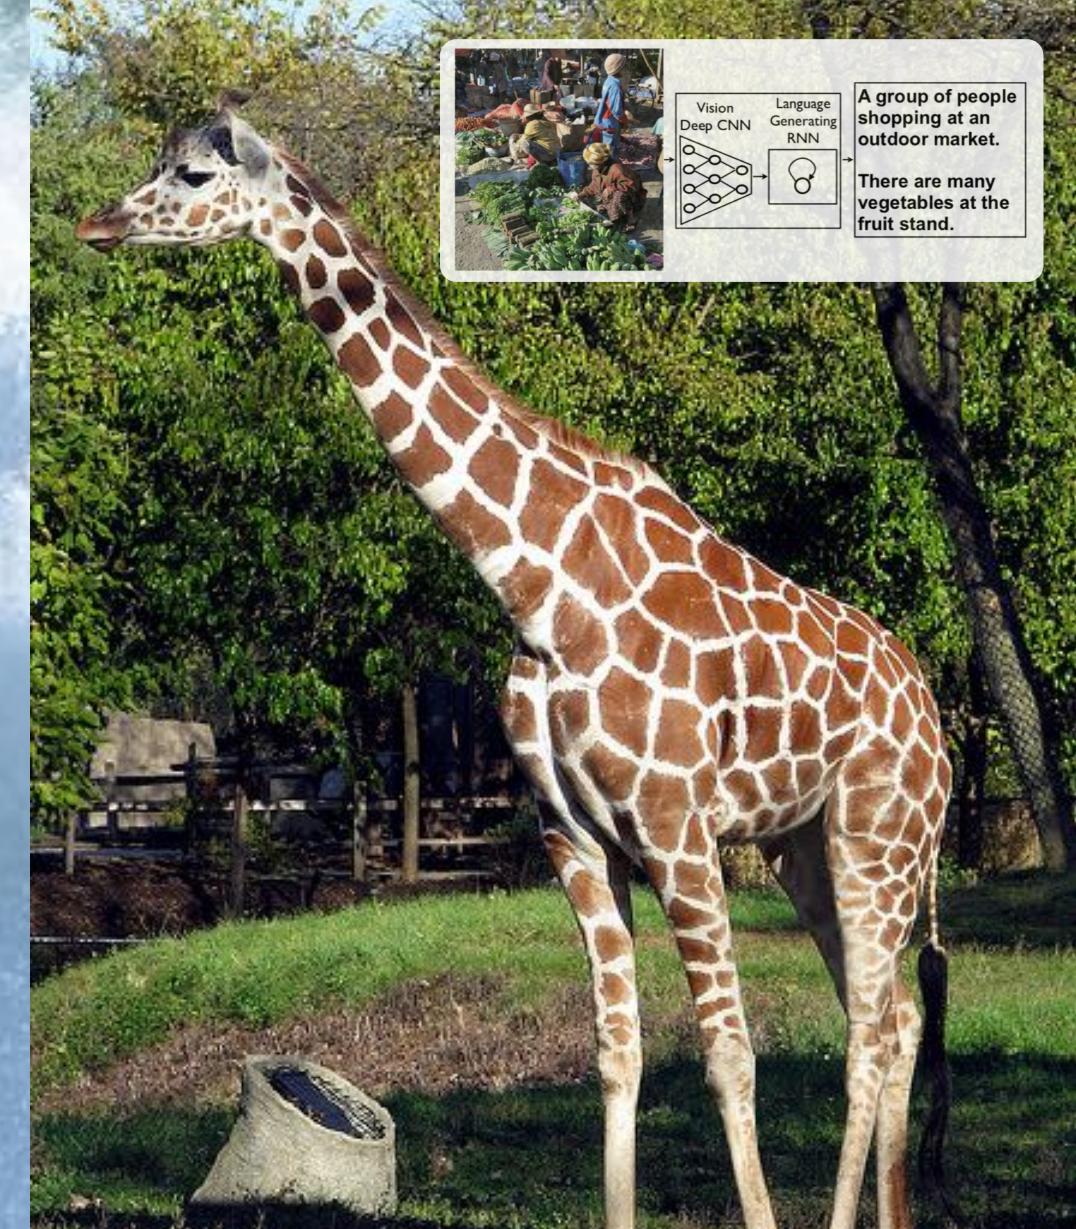
Photo Style Transfer



Image Captioning



A man riding a wave on a surfboard in the water.



A giraffe standing in the grass next to a tree.

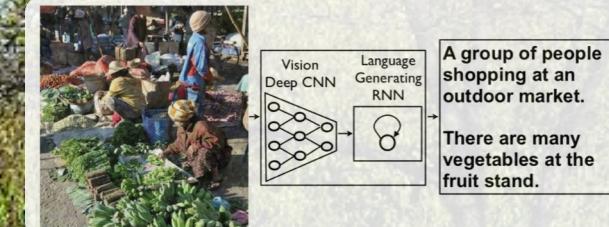
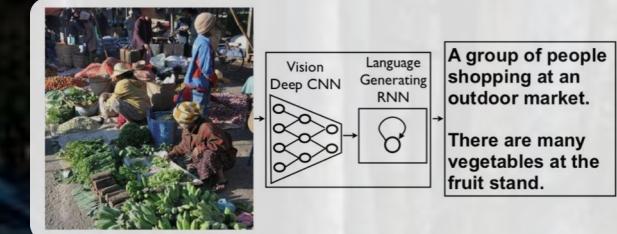


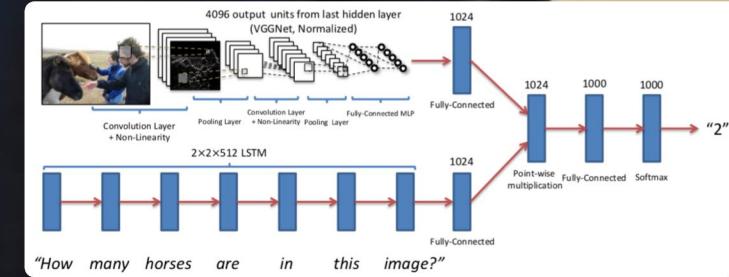
Image Captioning



Yarış pistinde virajı almakta olan bir yarış arabaşı



Visual Question Answering

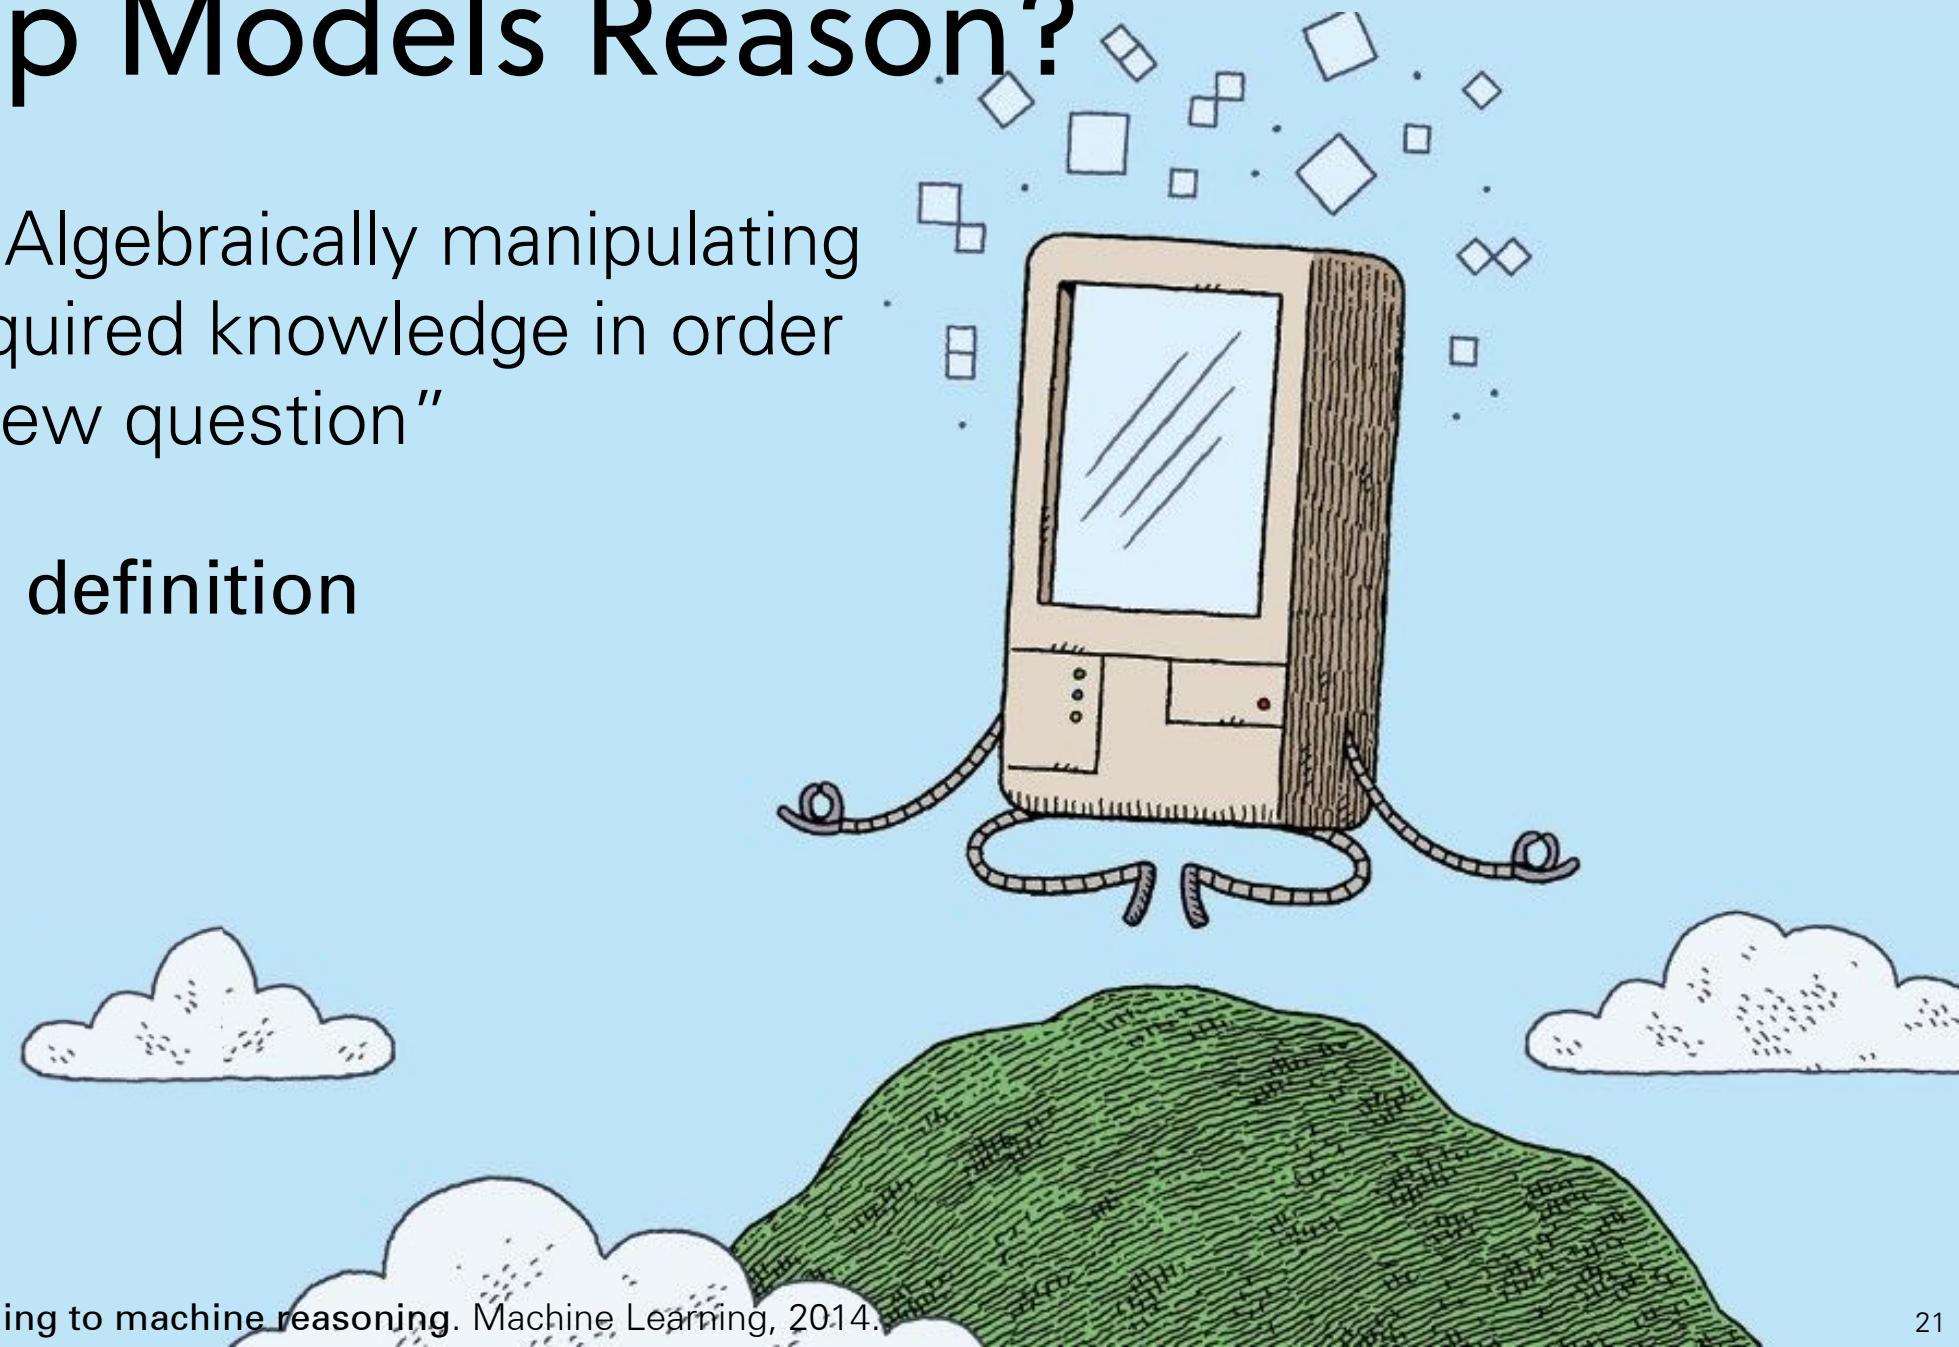


Question: What is the girl reaching into?

Answer: apples

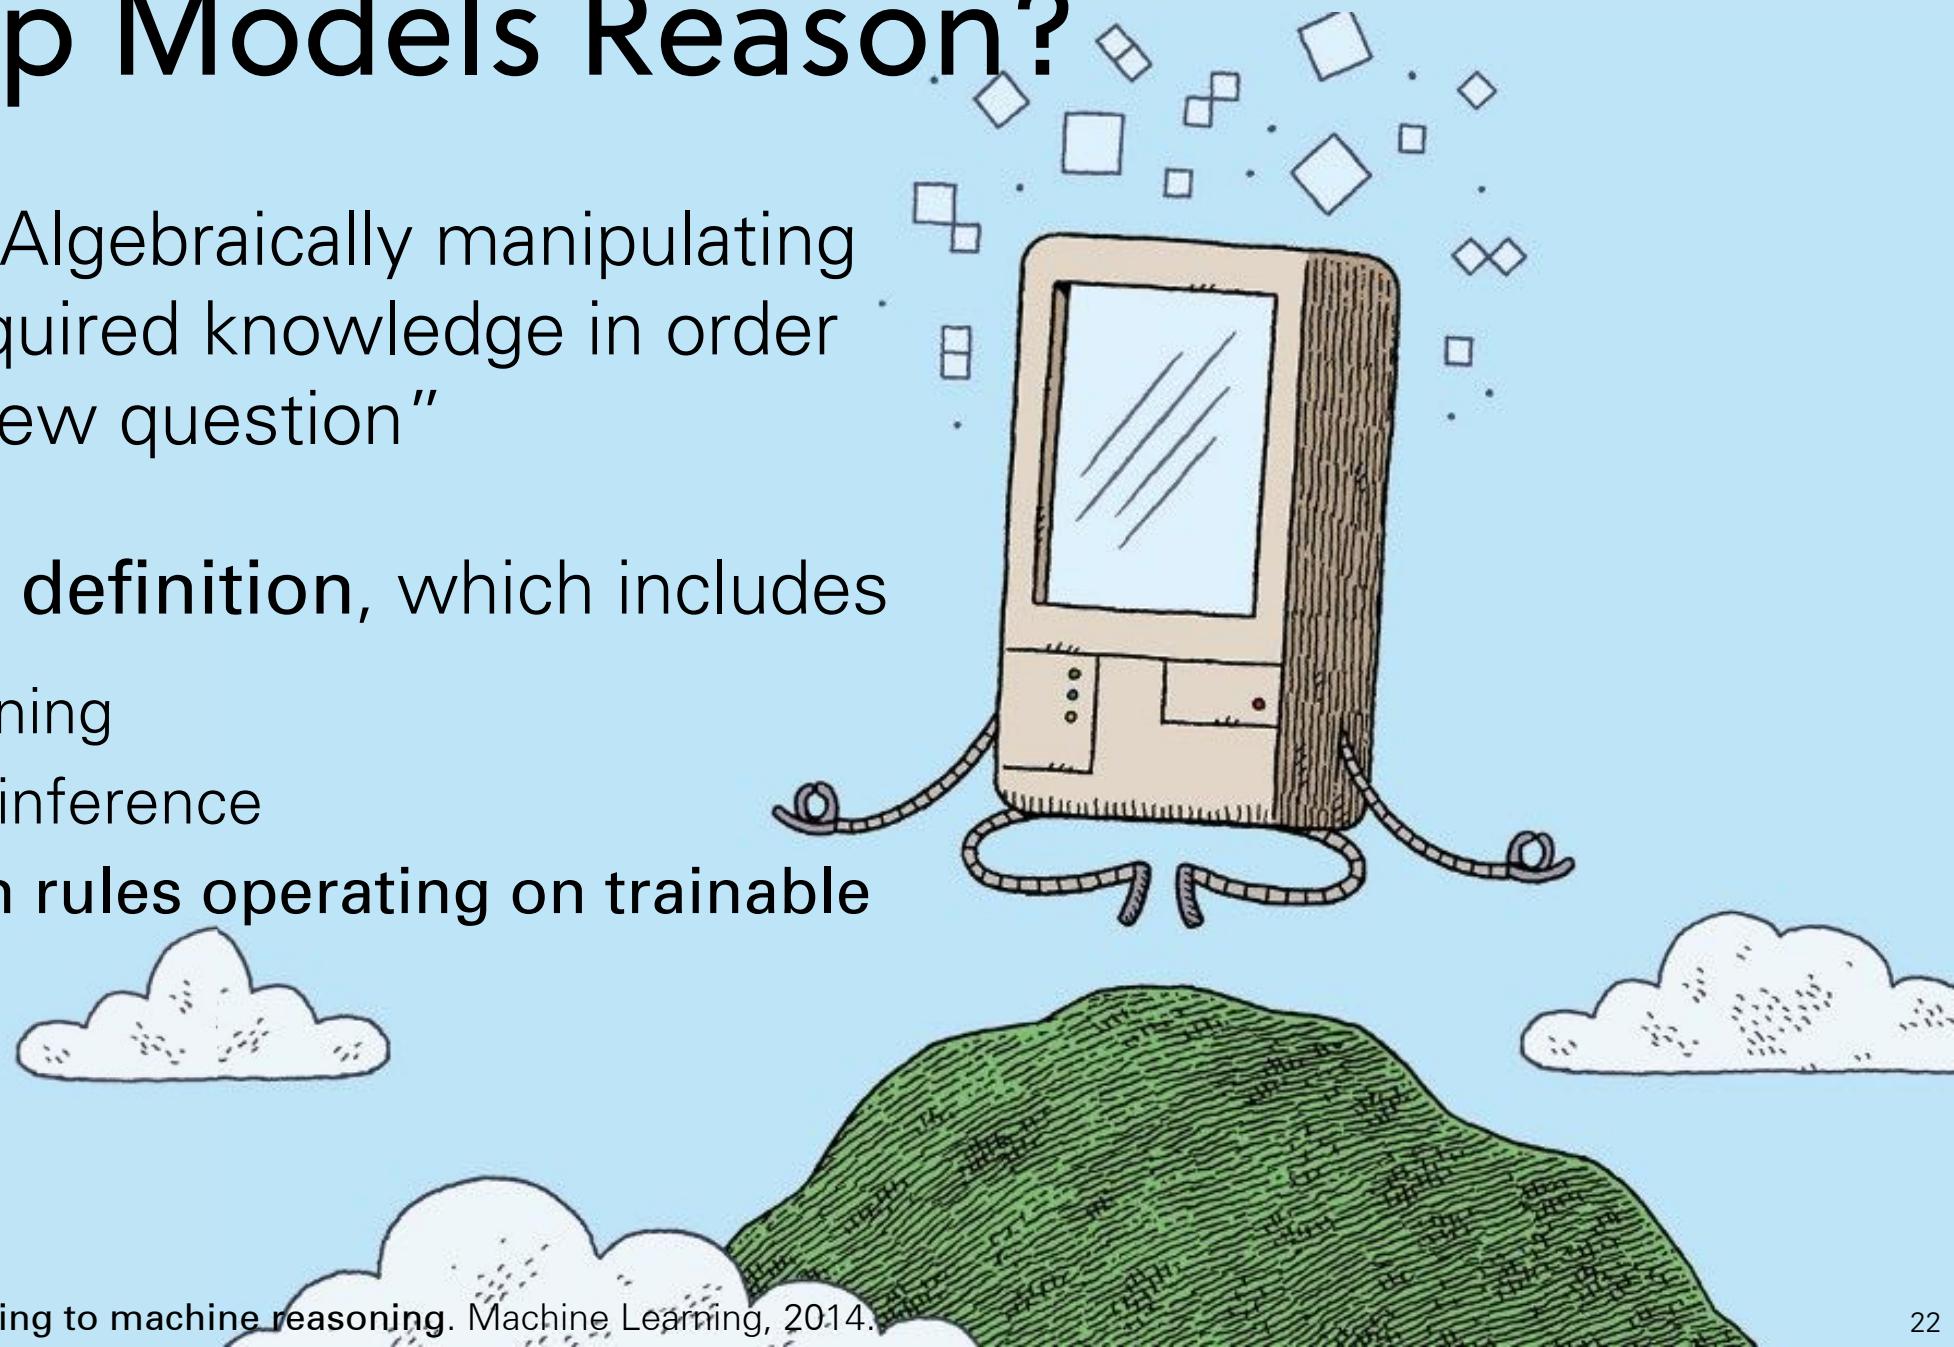
Can Deep Models Reason?

- Reasoning: “Algebraically manipulating previously acquired knowledge in order to answer a new question”
- A very broad definition



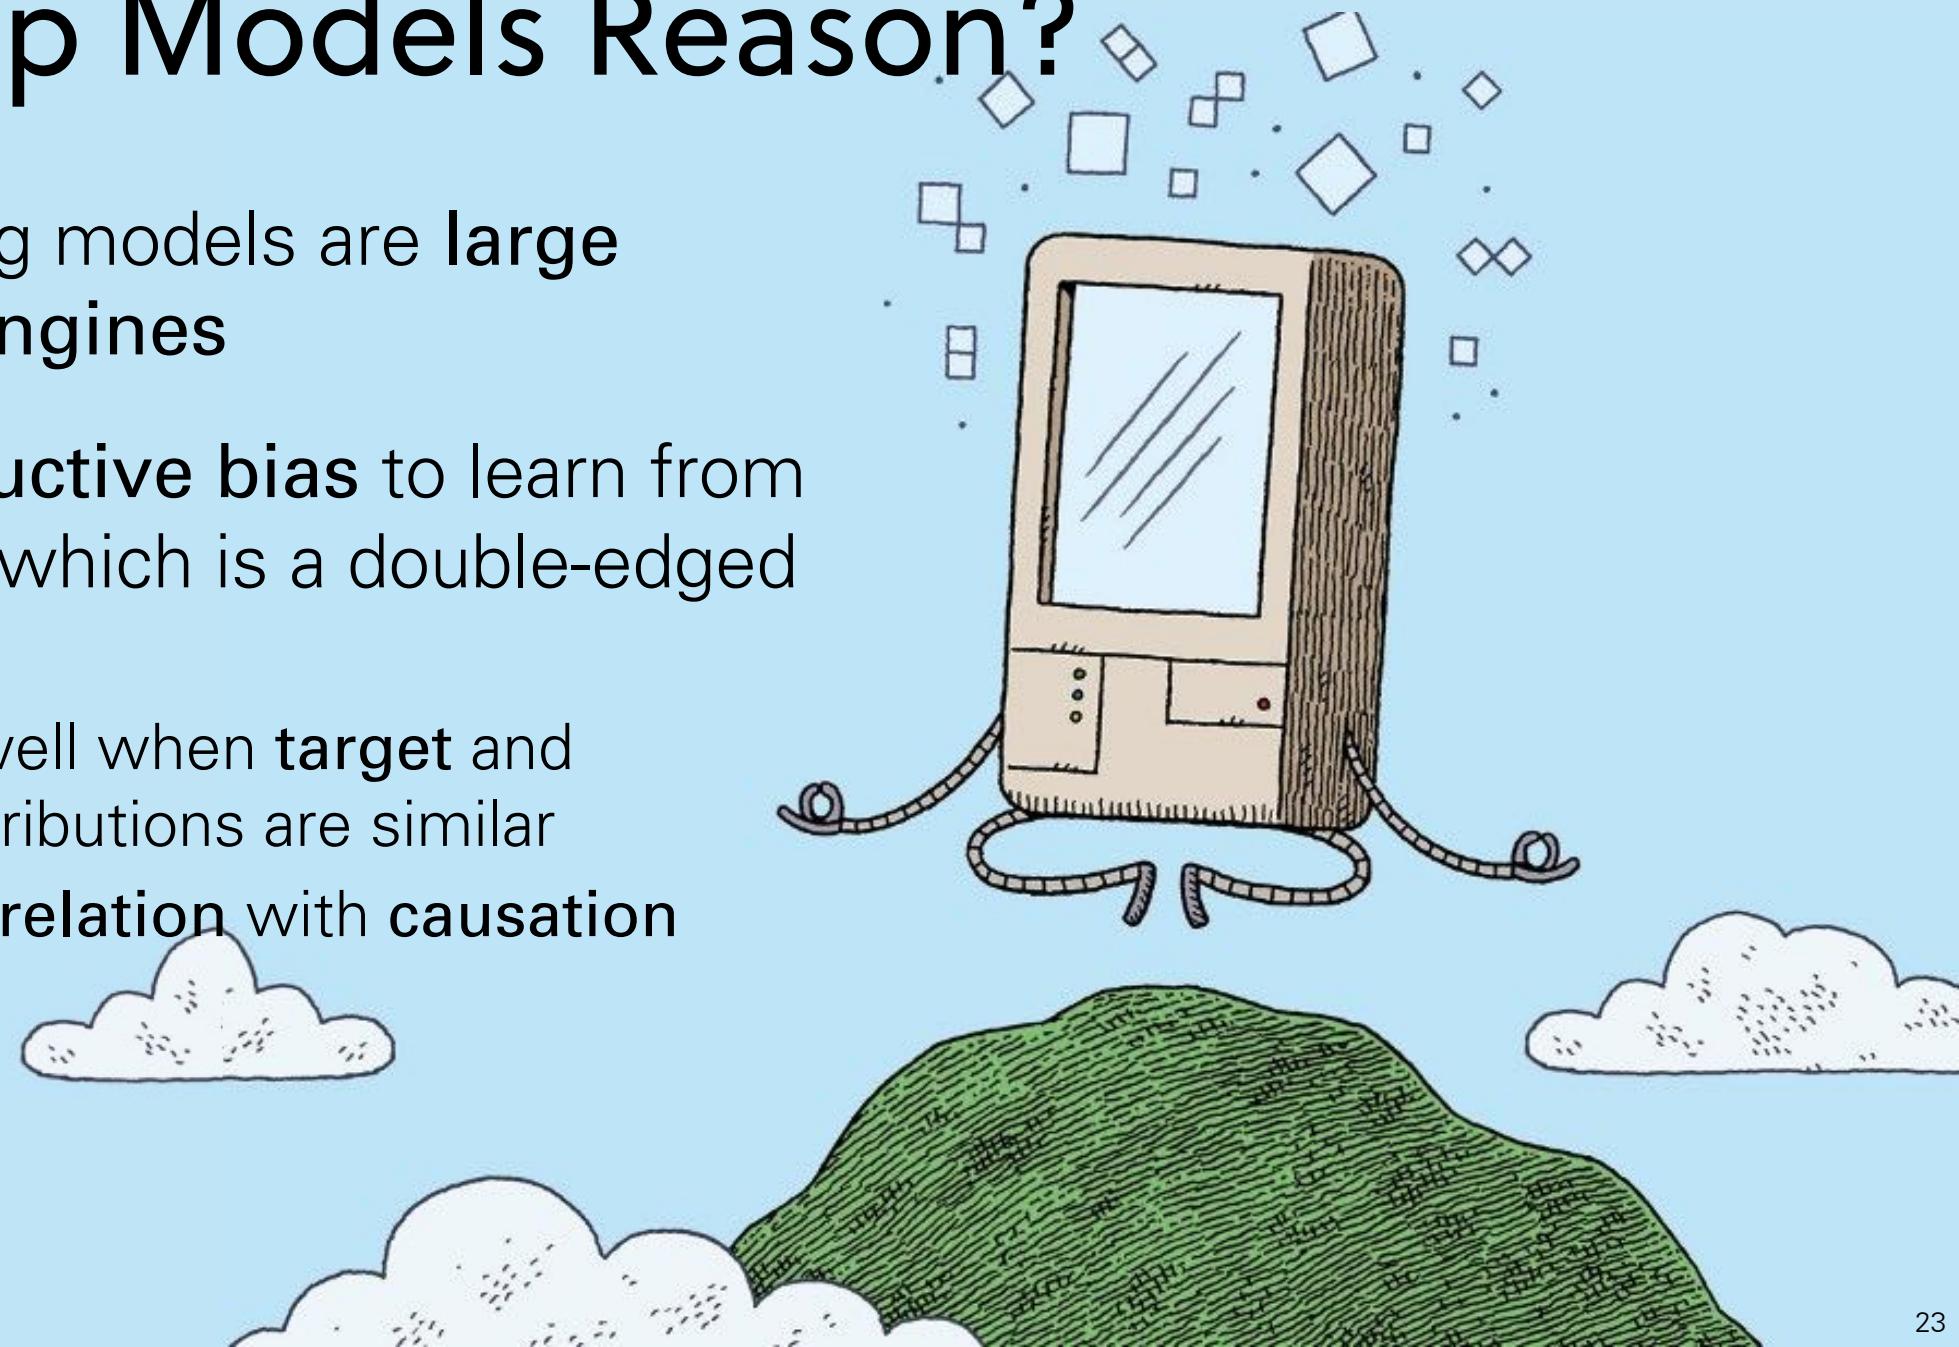
Can Deep Models Reason?

- **Reasoning:** “Algebraically manipulating previously acquired knowledge in order to answer a new question”
- A **very broad definition**, which includes
 - logical reasoning
 - probabilistic inference
 - composition rules operating on trainable modules



Can Deep Models Reason?

- Deep Learning models are **large correlation engines**
- They use **inductive bias** to learn from training data, which is a double-edged sword
 - Generalize well when **target** and **training** distributions are similar
 - Confuse **correlation** with **causation**

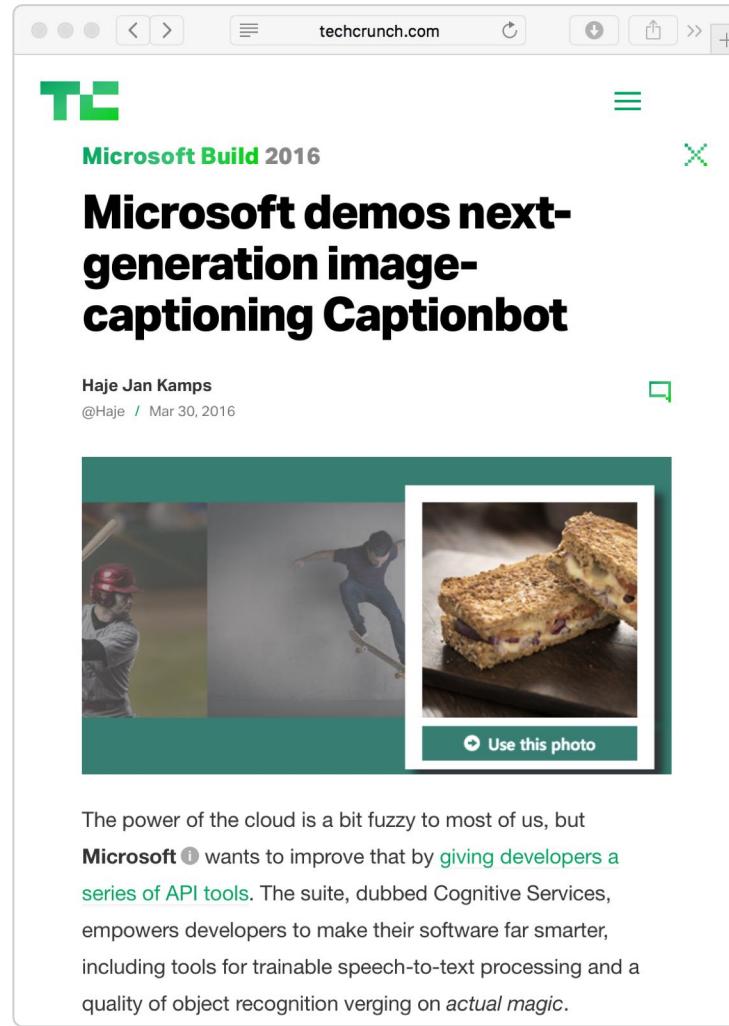


Take 2: Image Captioning



picdescbot @picdescbot · Feb 19

a herd of sheep grazing on a lush green field

A screenshot of a TechCrunch article page. The header reads "Microsoft Build 2016" and the main title is "Microsoft demos next-generation image-captioning Captionbot". Below the title, it says "Haje Jan Kamps" and "Mar 30, 2016". There are three thumbnail images at the bottom: a baseball player, a skateboarder, and a sandwich. A button labeled "Use this photo" is visible. The main text discusses Microsoft's Cognitive Services and their API tools.

6



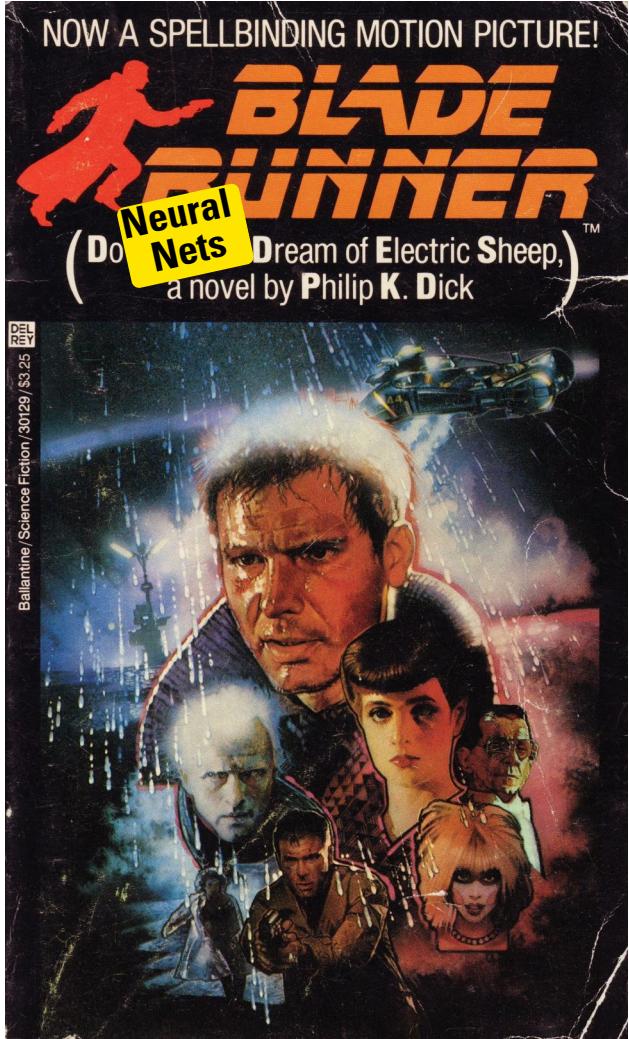
3



18



Take 2: Image Captioning



picdescbot @picdescbot · Feb 19

a herd of sheep grazing on a lush green field



6



3



18



[picdescbot](#) @picdescbot · Mar 8
a yellow and orange flowers in a field



8

41

103

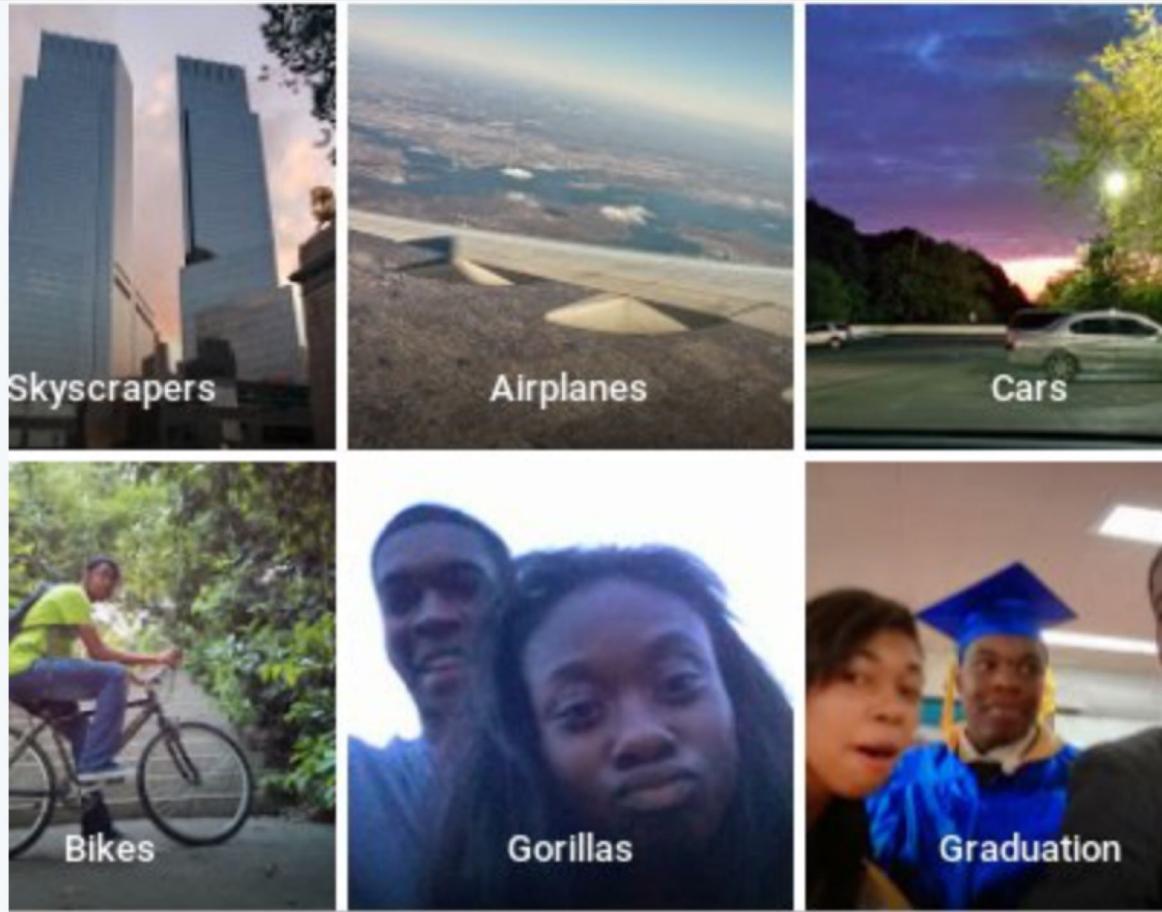
✉



▼

Suda yüzmekte olan bir köpek.





jackyalciné is about 40% into the IndieWeb.

@jackyalcine



Google Photos, y'all fucked up. My friend's not a gorilla.

4:22 AM - Jun 29, 2015

Heart 2,280 Comment 3,592 people are talking about this

Support The Guardian

The Guardian

News | Opinion | Sport | Culture | Lifestyle

World | UK | Science | Cities | Global development

i

Google says sorry for racist auto-tag in photo app

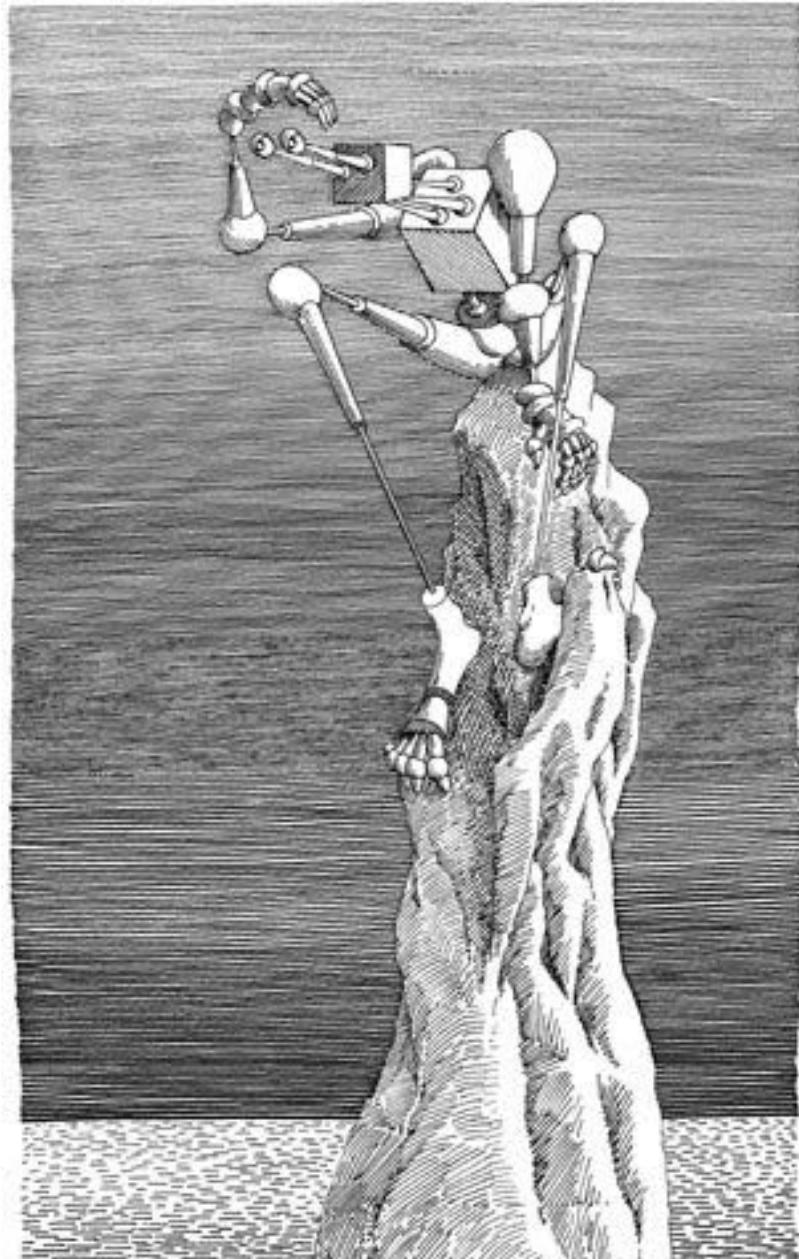
- Google Photos labelled a picture of two black people as 'gorillas'
- Google Maps and Flickr have also suffered from race-related problems

Jana Kasperkevic in New York

Twitter @kasperka Email
Wed 1 Jul 2015 18.52 BST

Looking Forward

- Intelligence is not just about **Pattern Recognition**
- Learning is the process of **modeling the world...**
 - **explaining and understanding** what we see
 - **imagining** things we could see but haven't yet.
 - **problem solving and planning** actions to make things real.
 - **building new models** as we learn more about the world.
 - **sharing our models**, communicating to others, understanding their models, and learning from them and with them.



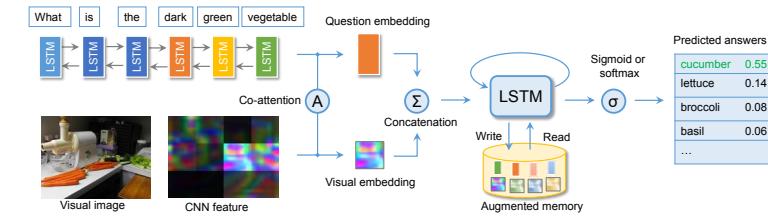
Explaining and understanding what we see



Q: What fruit is showing in this picture?



A: Bananas



Visual Reasoning

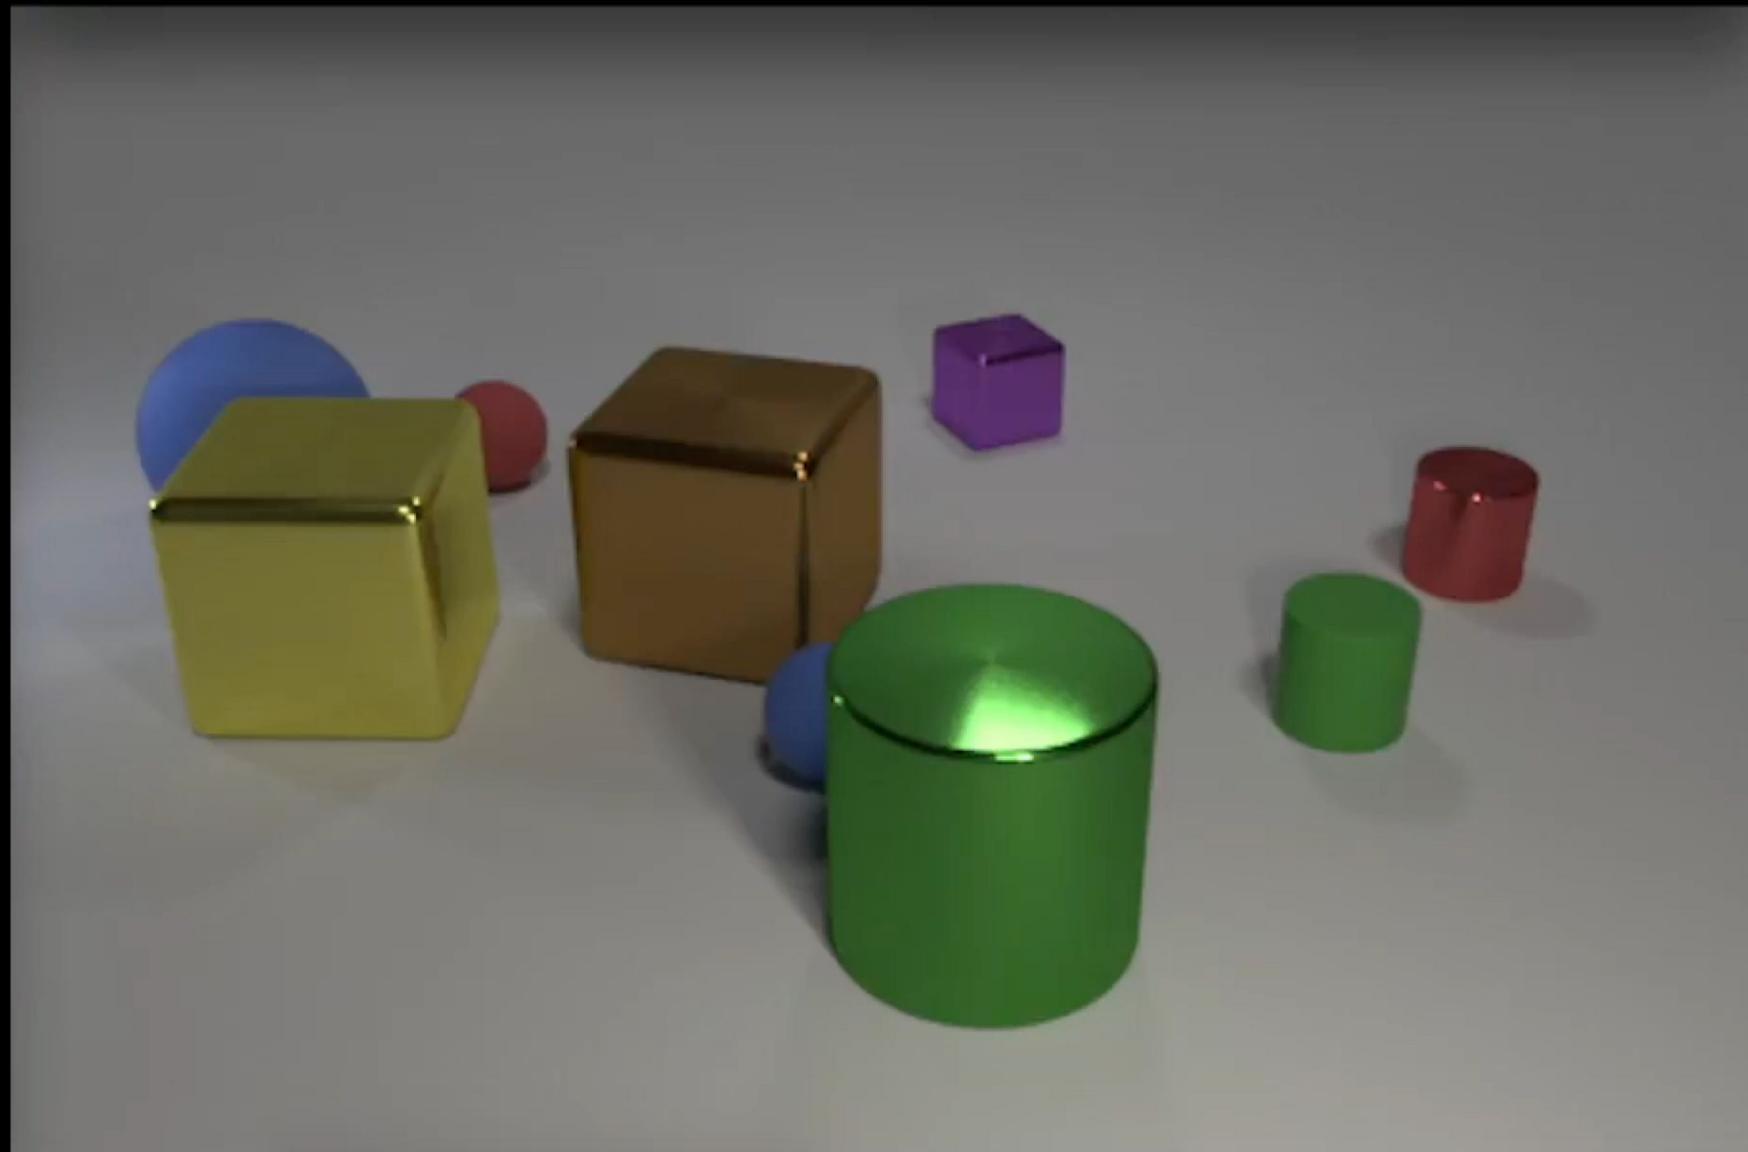
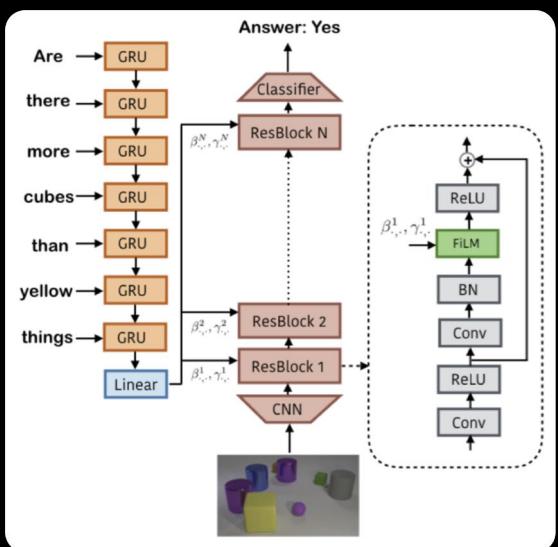


How many objects are
either small cylinders
or red things?

Answer: 5

Visual Reasoning

Ask me something!
>>> H

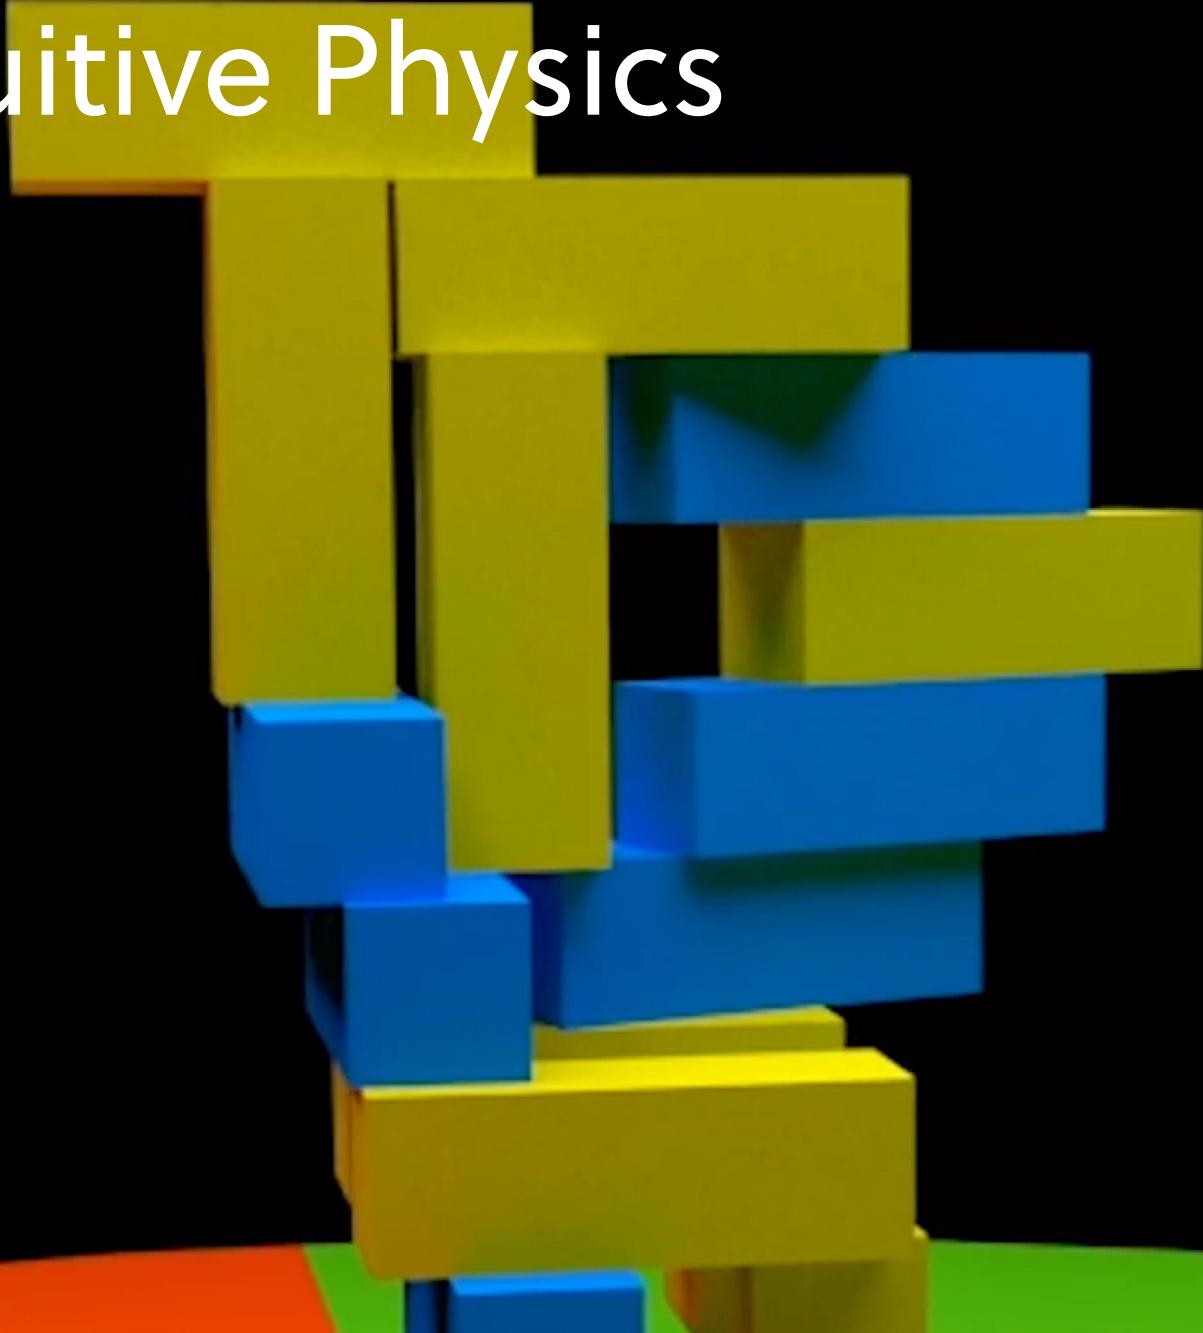


Intuitive Physics

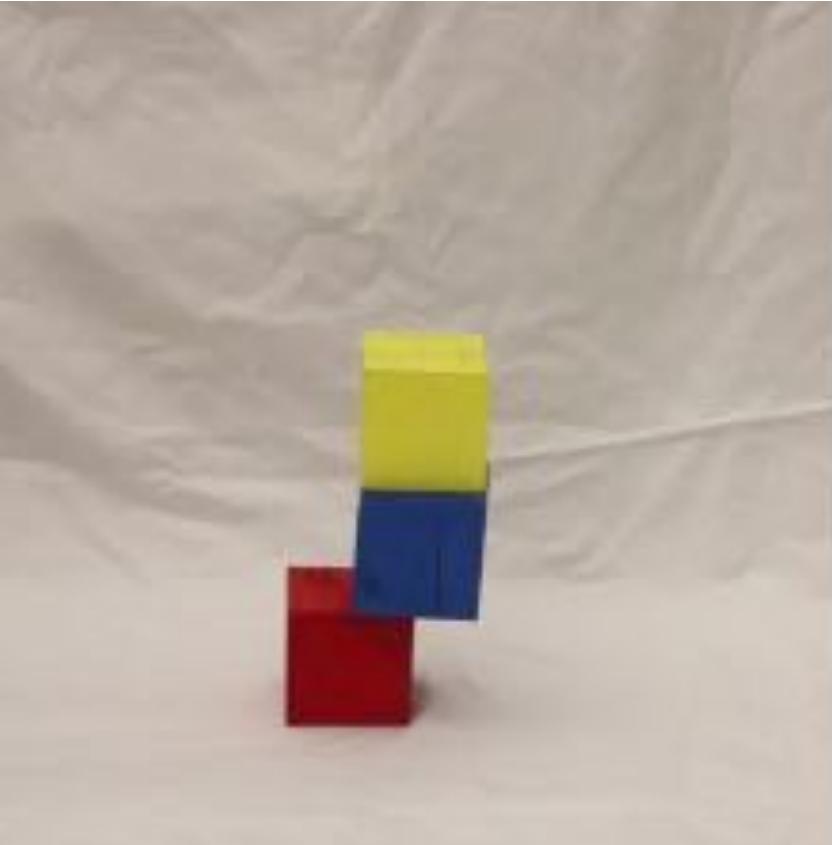
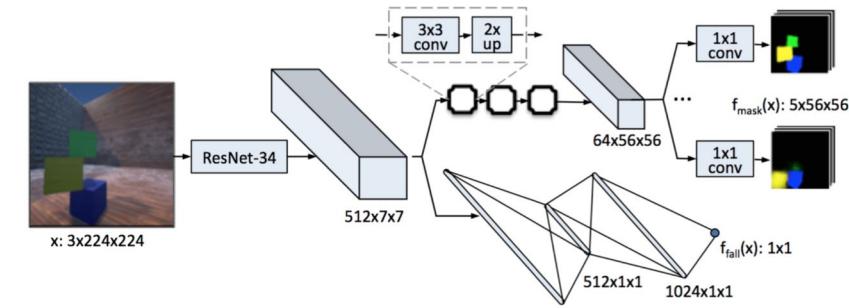
- Common-sense understanding of how the world operates at a physical level
- Helps us to perceive, understand and act with our environment



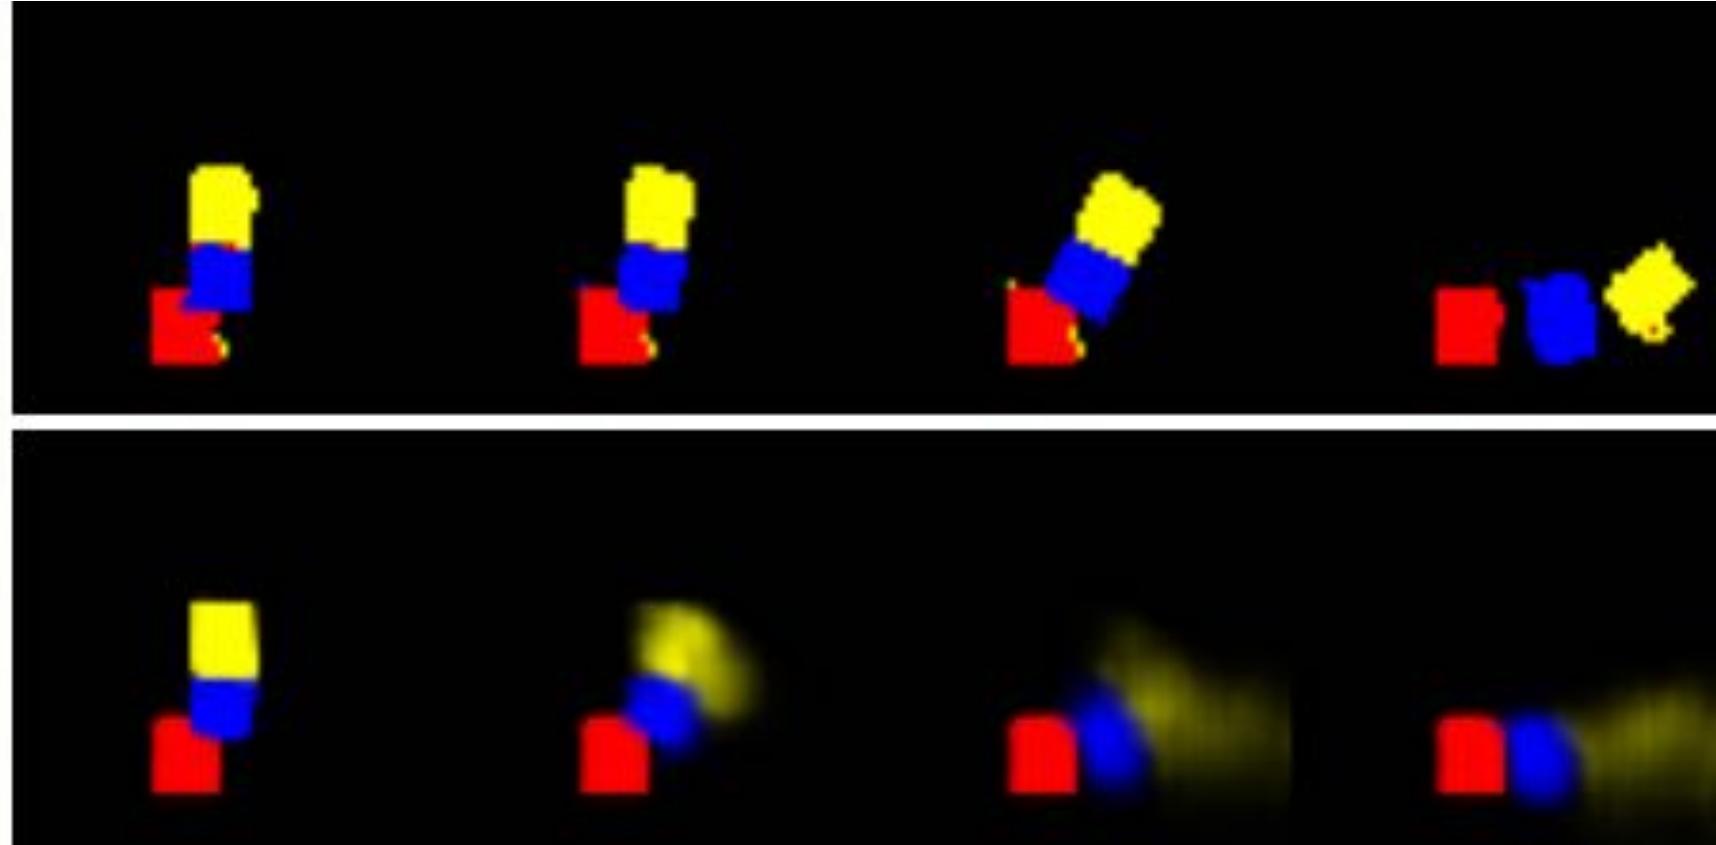
Intuitive Physics



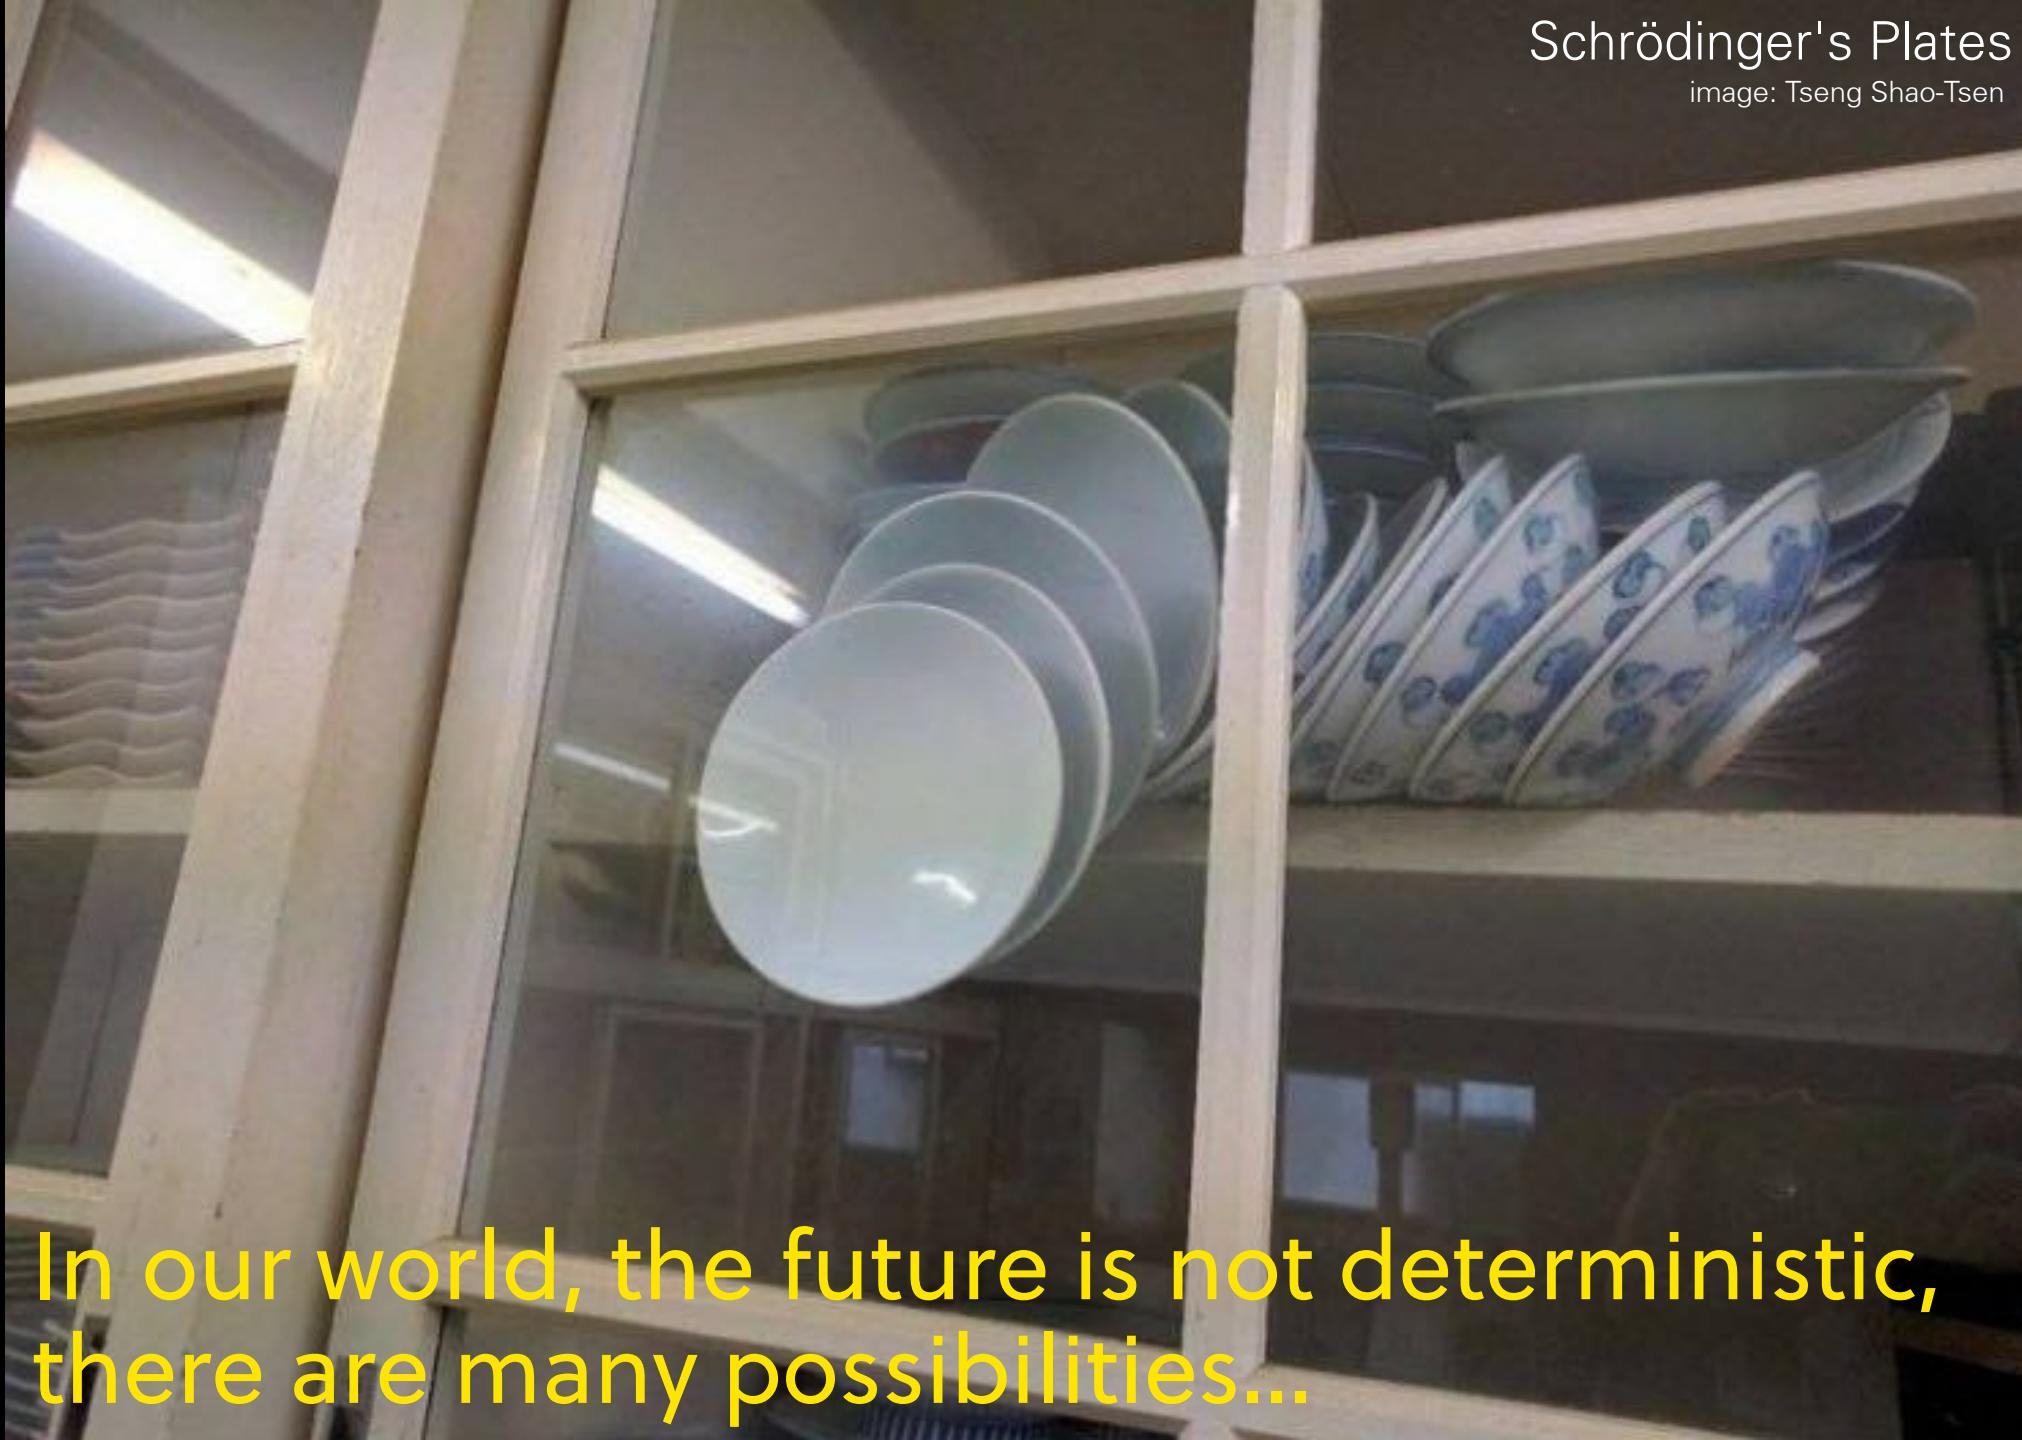
Intuitive Physics



Initial frame



PhysNet predictions of the future



In our world, the future is not deterministic,
there are many possibilities...

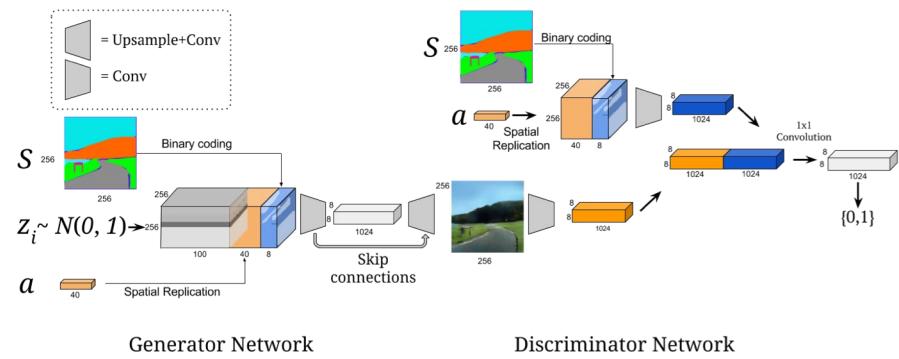
Imagining Things



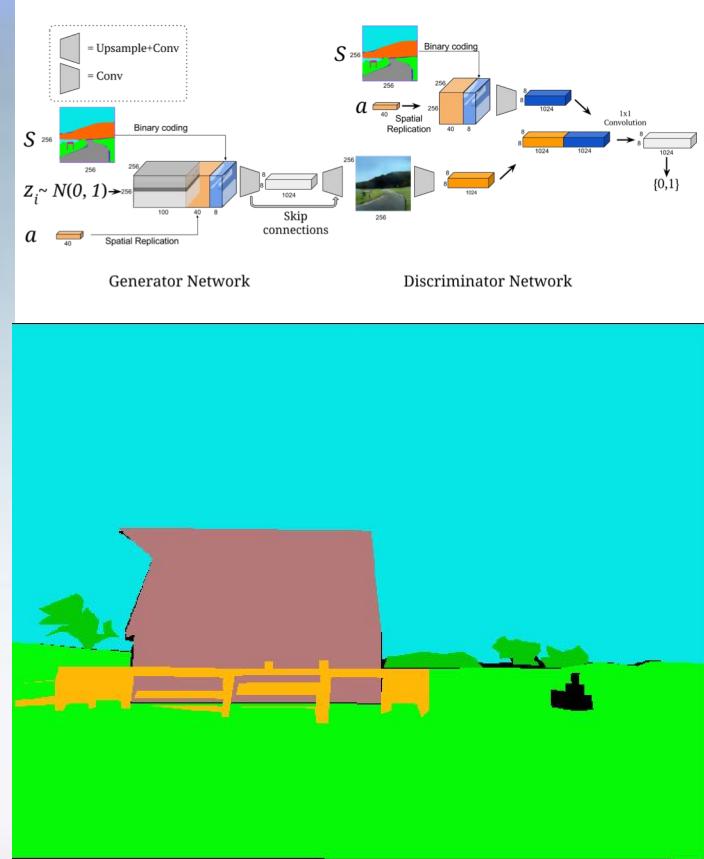
“Maybe in our world
there lives a happy little
tree over there.”

— Bob Ross (*The Joy of Painting*)

Imagining Things

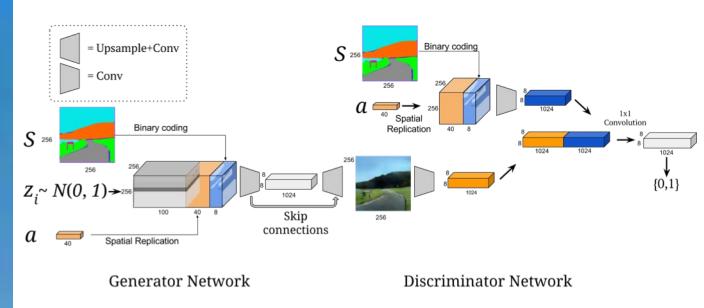


Imagining Things



Semantic Layout

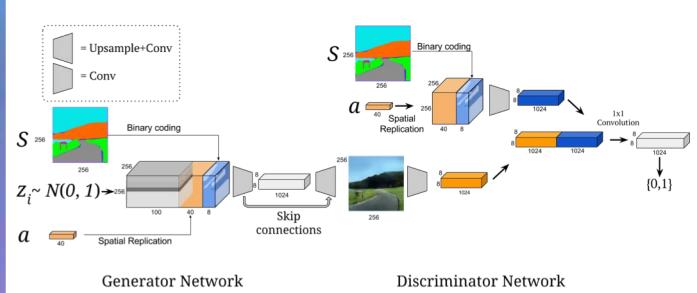
Imagining Things



Semantic Layout

Clear sky + flowers

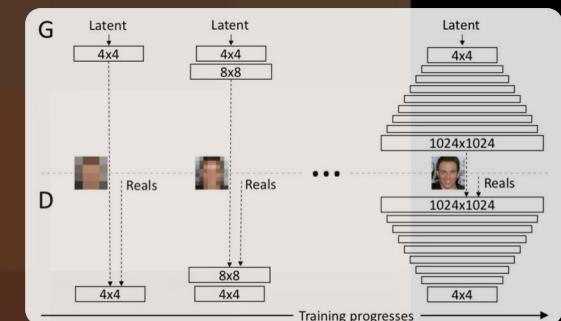
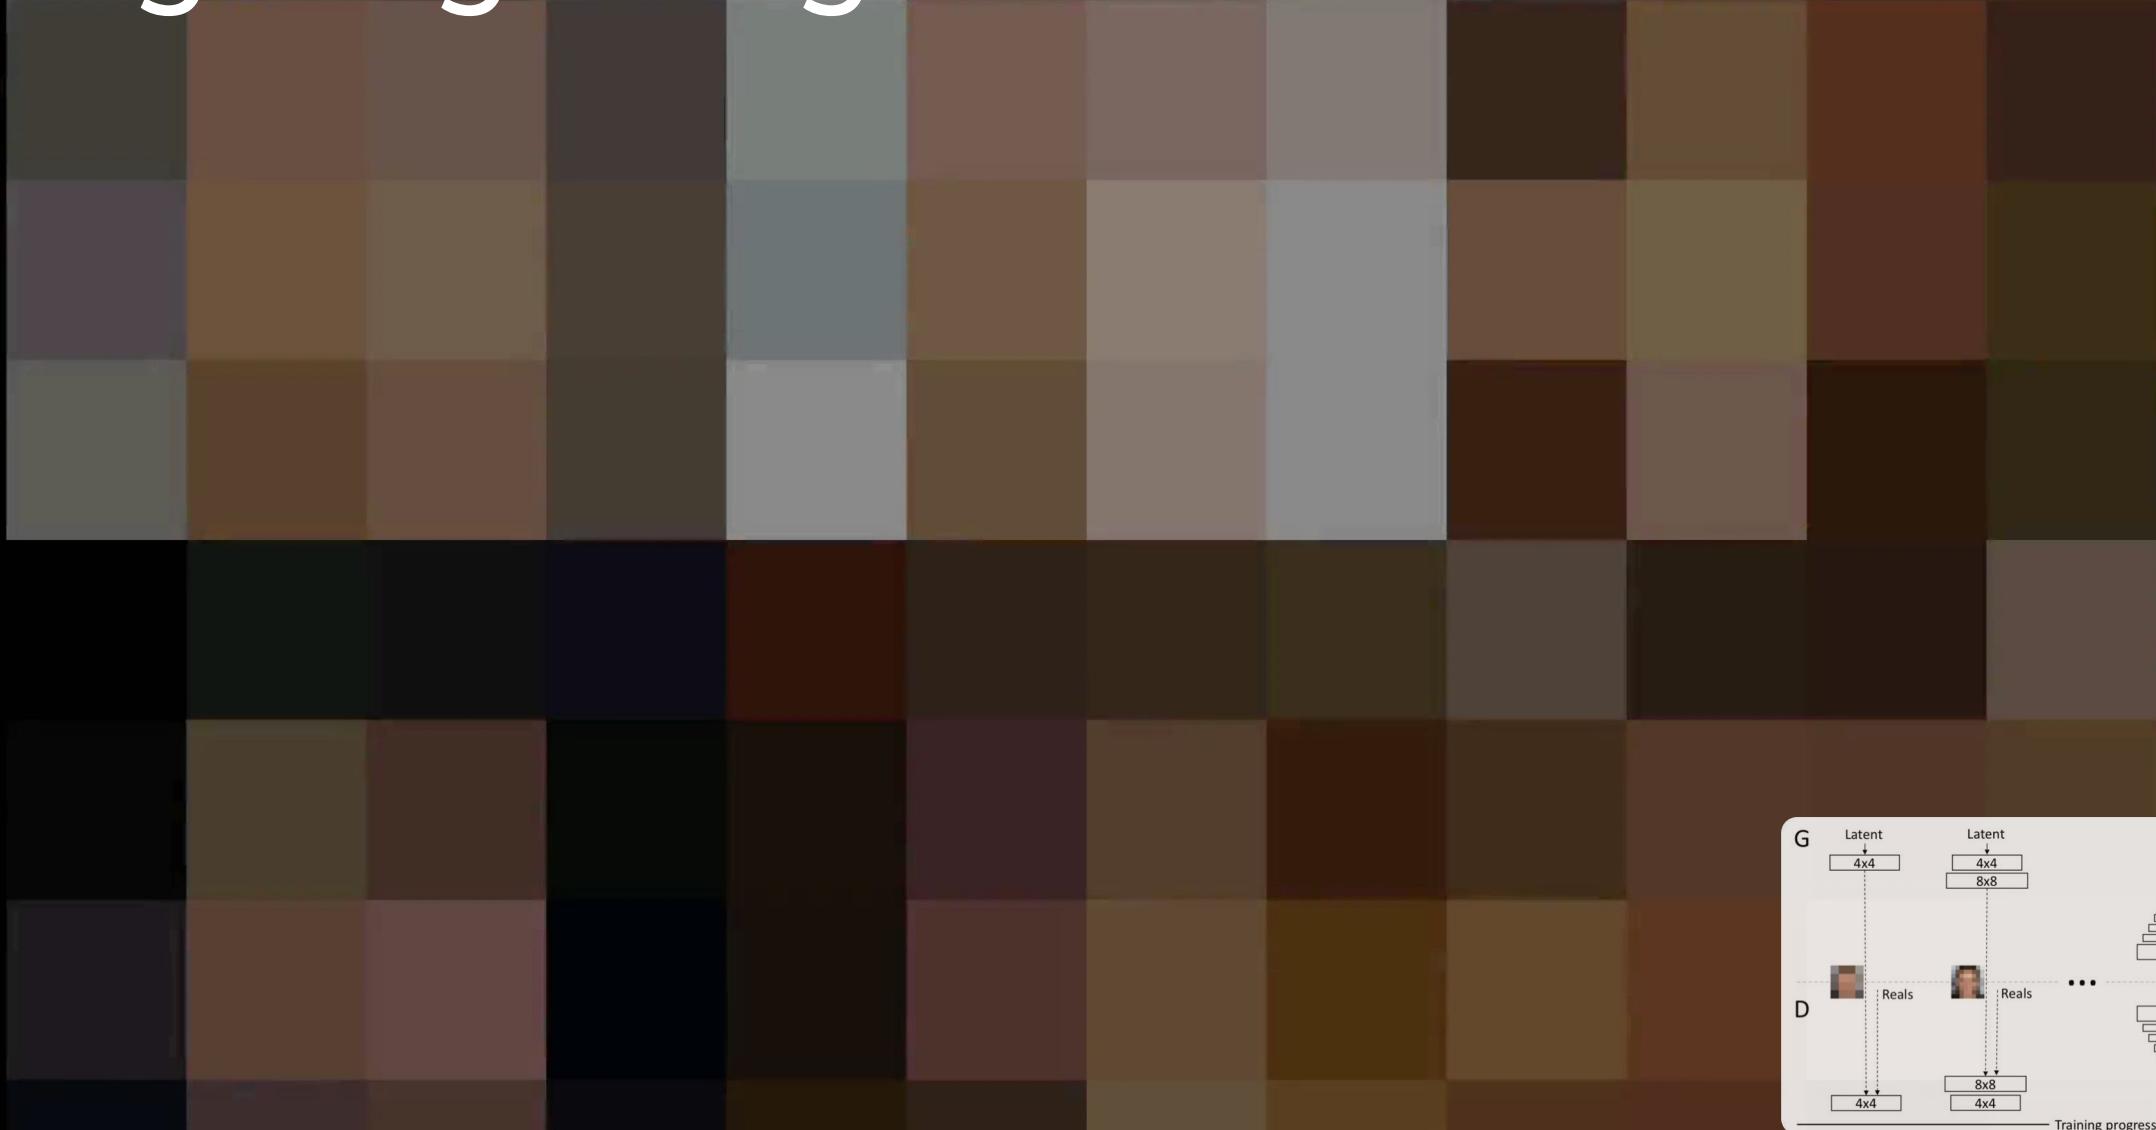
Imagining Things



Semantic Layout

Sunset

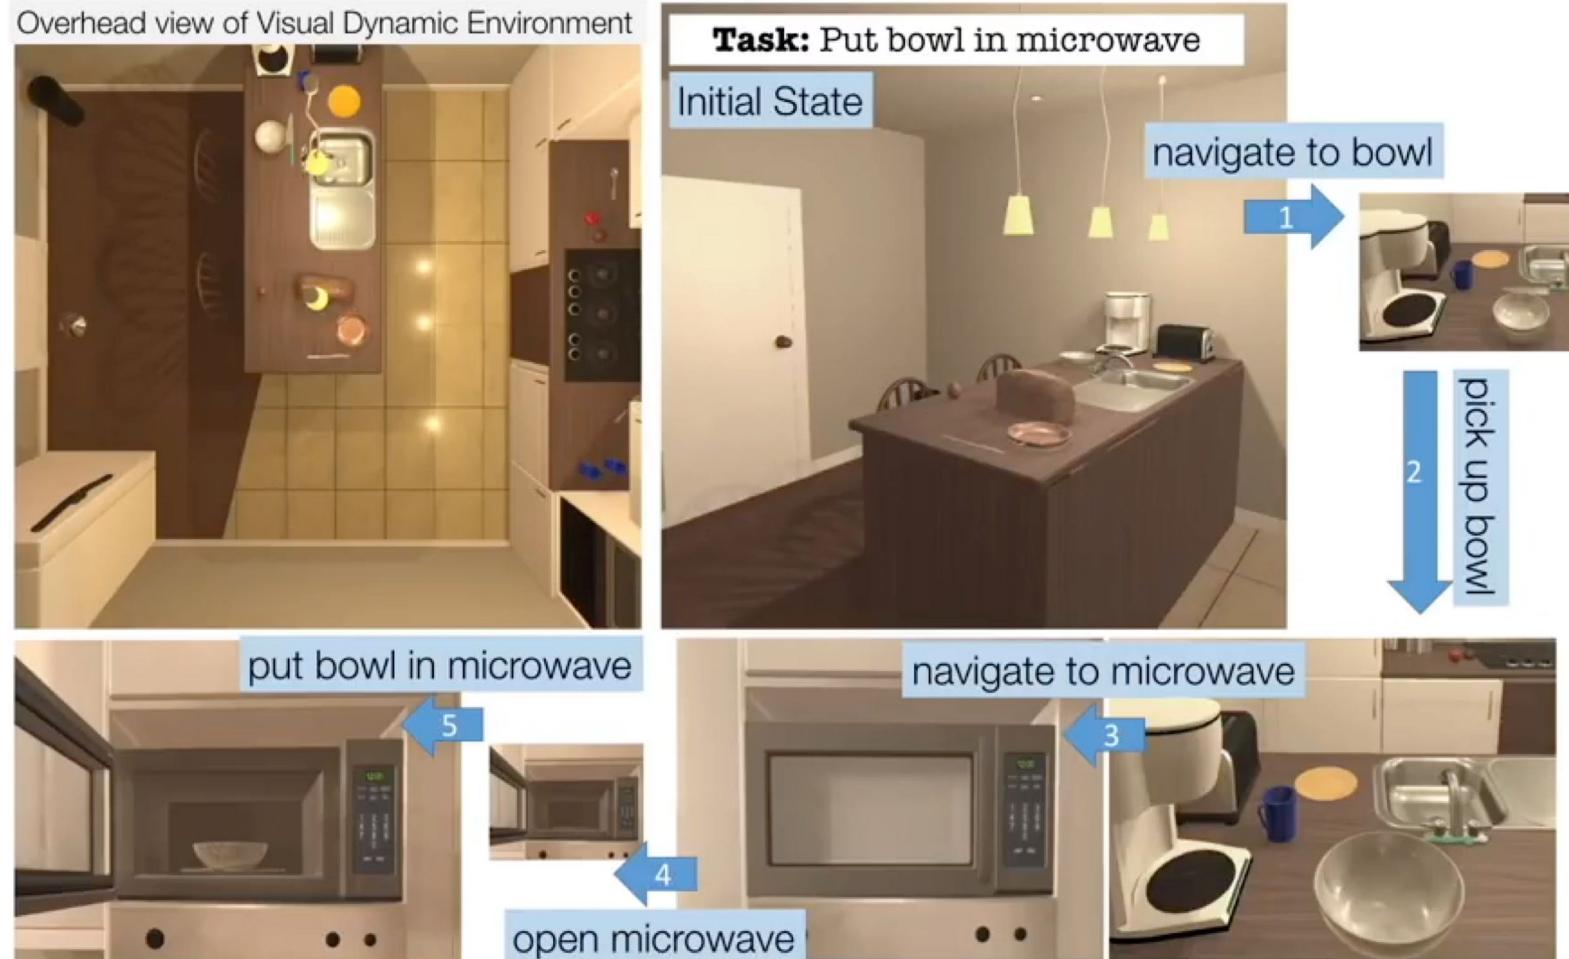
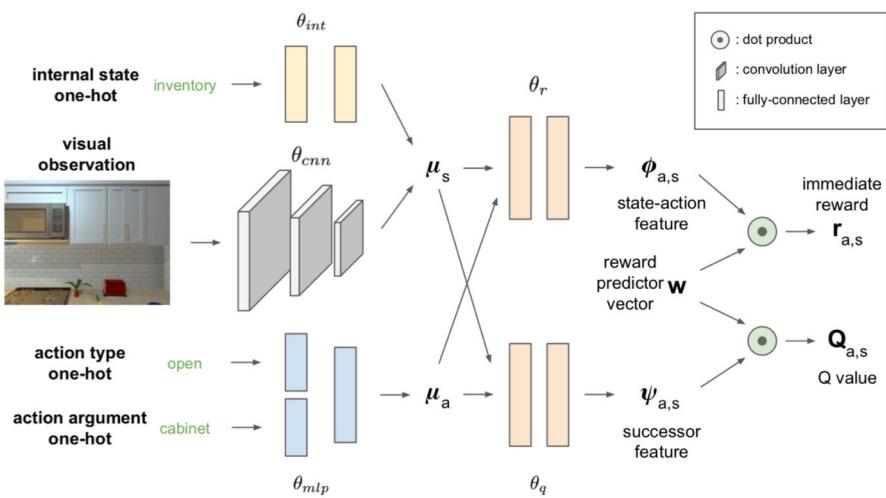
Imagining Things



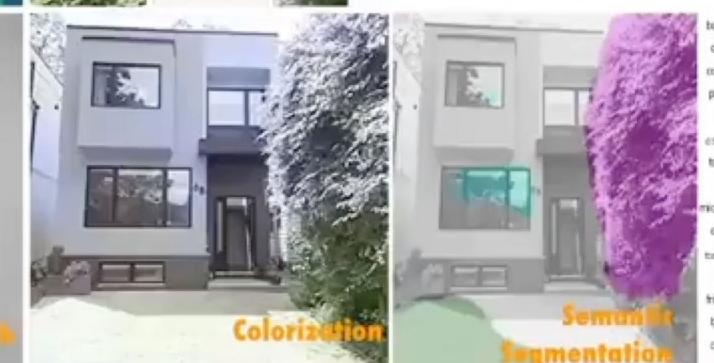
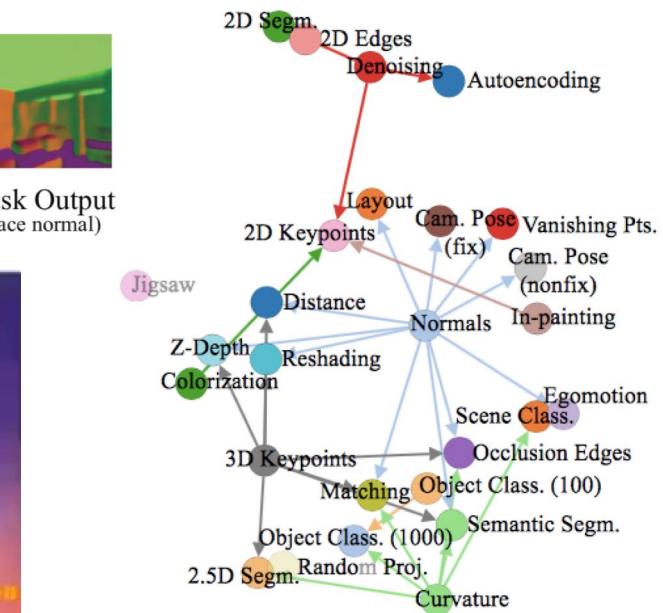
4x4

T.Karras, T.Aila, S.Laine and J, Lehtinen, "Progressive Growing of GANs for Improved Quality, Stability, and Variation", ICLR 2018

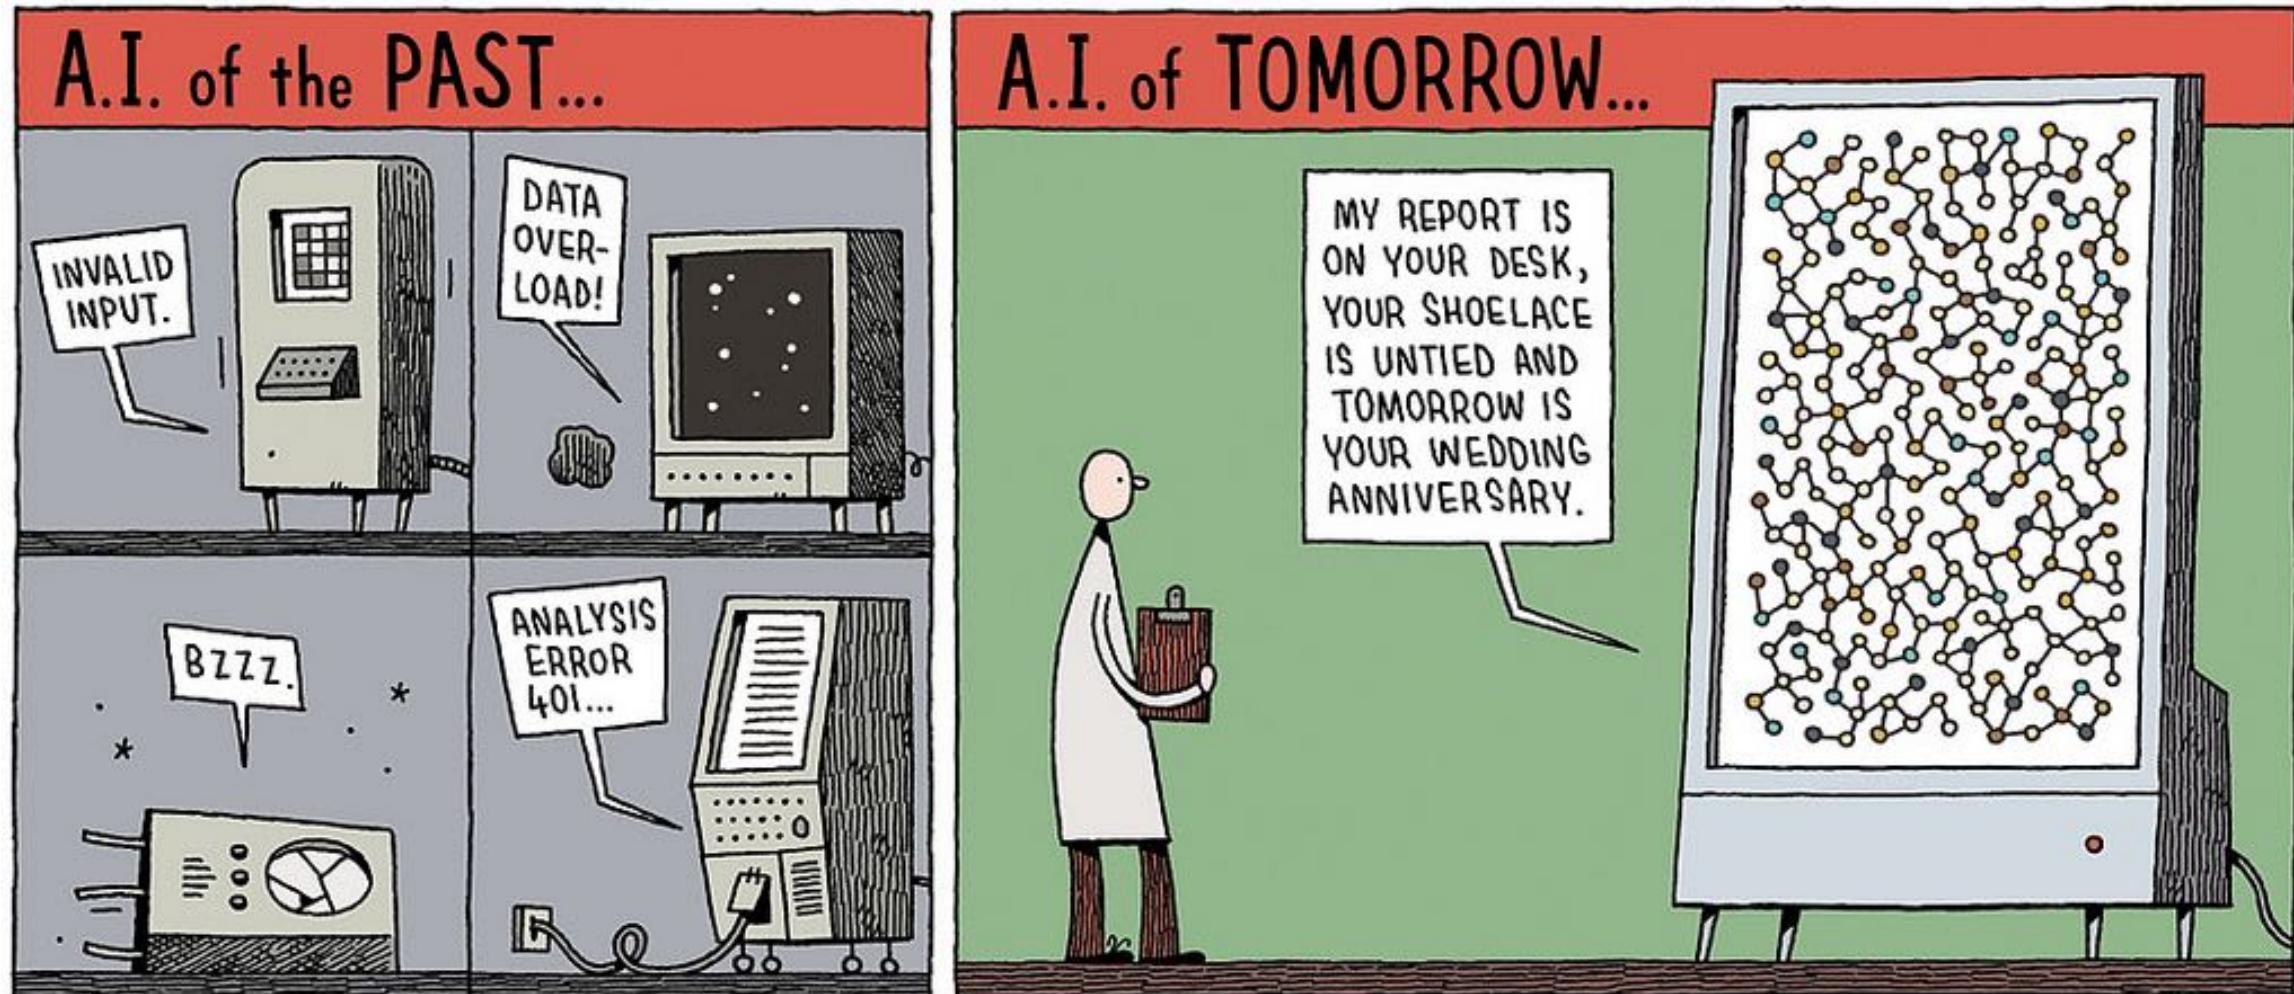
Planning



Transfer Learning



Can deep models reason?



- We are not there yet! But we can see real progress soon..